Maximal divergence of the leximin and leximax core selections in exact games

Maxwell Flanagan

August 21, 2026

Abstract

A singleton-normalized exact transferable-utility game carries two canonical egalitarian core selections: the leximin allocation Lmin and the leximax-min allocation Lmax. They coincide on every exact game with at most four players. We measure their separation by $\rho(v) = \|\operatorname{Lmin}(v) - \operatorname{Lmax}(v)\|_1/v(N)$ and let $s(n)$ be its supremum on n-player exact games. We prove $s(5) = 1/2$, $s(6) = 2/3$, and $\sup_n s(n) = 2$. An explicit two-parameter family $G(p,k)$ satisfies

$$\rho(G(p,k)) = \frac{2(p(k-1)-1)}{k(p+1)}, \qquad 2 - \rho(G(p,k)) = \frac{2n}{k(p+1)}, \quad n = p+1+k,$$

and yields $s(n) \geq 2 - 8n/(n^2 - (n \bmod 2))$. A common family of tight caps drives both selections: leximin treats each cap as a ceiling and concentrates the residual on one player, whereas leximax-min treats it as a budget and spreads the residual, making the two supports asymptotically disjoint. On the two-generator class we also give a dual upper-cap description of the core and a one-coalition certificate for the leximin floor; the mirror leximax-min ceiling statement fails.

The fixed-dimensional upper bounds are computer-assisted. At five players, two exact pipelines exhaust all configurations and use independently formulated Lmin-optimality components; they prove the sharp value and the $1/3$ bound in the leximin $(4,1)$ regime. For a separate sparse-lift question, we prove a sharp dimension threshold: every five-player exact game with distinct selectors and a two-player Lmax top has a faithful tight-pair lift for every fixed complete staged Lmin certificate, whereas an exact six-player game defeats every such certificate and every qualifying extreme top cover. The positive half is a 3840-cell exact-cone theorem replayed through two algebraically separate certificate encodings; the negative half has an exact primary replay and a standalone rational replay. At six players, singleton exactness and lexicographic optimality map a strict counterexample into one of fifteen sign-labelled blocker systems. Exact proof objects certify all fifteen infeasible; the largest systems reduce to nineteen role-level exchange schemas and three negative-cut templates. The human counterexample-to-relaxation implication, finite certificates, and novelty determination are kept as separate review obligations. For the dimension-free upper-bound program, we prove the conjectured atomic overlap inequality on every weighted four-role quotient satisfying the canonical unit-role normalization. A connected Lorenz-frontier argument confines the two selectors to one pair-box edge. The quantitative proof is human for atomic weight at least ten and uses a complete quantified linear-real audit of 1,152 cells for the six residual weight rows at weights six through nine. This quotient theorem makes common selector modules inductively contractible in the general carrier program. For strict-loss carrier cycles, proper tight packets exist and their cardinality base is automatically target-sized; comparison–residual cancellation reduces the remaining direct arm to scalar nongainer-block selection. A new exact seven-player control refutes selection of the full maximum-overlap root block and forces a labelled tight-support enlargement, locating the remaining gap in staged transport congestion or a marked-facet alternative.

Keywords: cooperative games; exact games; core; leximin; leximax; Lorenz order; computer-assisted proof.
MSC 2020: 91A12; 90C05; 52B12.

1 Introduction

Let (N, v) be a transferable-utility game on $N = \{1, \dots, n\}$ with

$$\mathrm{Core}(v) = \{x \in \mathbb{R}^N : x(N) = v(N),\ x(S) \geq v(S) \text{ for all } S \subseteq N\},$$

where $x(S) = \sum_{i \in S} x_i$. The game is *exact* if every coalition inequality is attained by some core allocation, and *singleton-normalized* if $v(\{i\}) = 0$ for every i; on that domain every core allocation is nonnegative. Exact games are the natural home for egalitarian core selections, and two such selections are canonical. The *leximin* allocation $\mathrm{Lmin}(v)$ maximizes the increasing rearrangement of x lexicographically over the core; the *leximax-min* allocation $\mathrm{Lmax}(v)$ minimizes the decreasing rearrangement lexicographically over the core. Both are single-valued on a nonempty compact convex set. The first raises the poorest as far as possible, the second lowers the richest as far as possible.

On small games the two agree. The companion paper (Flanagan, 2026) proves that every singleton-normalized exact game on at most four players is an exact partition game at Lmin; by the established implication from exact partition to Lorenz domination (Llerena and Mauri, 2017), Lmin then Lorenz-dominates every core allocation, and a Lorenz-greatest core point is simultaneously leximin- and leximax-optimal[1], so $\mathrm{Lmin}(v) = \mathrm{Lmax}(v)$ for every singleton-normalized exact game with $n \leq 4$. Our complete enumerations at $n \leq 4$ reconfirm this rigidity as a control (Section 5). The adjacent large-core chain — exact games on at most four players have large cores (Biswas et al., 1999), and on large-core games the procedural egalitarian solution coincides with Lmax (Dietzenbacher, 2021; Arin et al., 2003) — pins only the Lmax side: it does not deliver the coincidence, and indeed a five-player exact game with a large core has Lmin $\neq$ Lmax: the lower envelope of $(2, 0, 12, 6, 0)$, $(0, 0, 0, 0, 20)$ and $(0, 2, 10, 4, 4)$, with $\| \mathrm{Lmin} - \mathrm{Lmax} \|_1 = v(N)/5$ (verified exactly; archived with the repository). The companion's lower-envelope realization technique is used again here.

The pair itself has recent precedent. Vazirani (2024) and Gangam et al. (2024) study leximin and leximax core imputations of assignment, flow and matching games, observe on examples that the two selections differ, and describe their opposition in exactly the direction this paper quantifies; the leximax selection appears earlier as the leximax solution of neighbour games (Klijn et al., 2003) and as the lexmax solution in the egalitarian-set literature (Llerena and Mauri, 2016), besides its large-core characterization (Arin et al., 2003). The candidate contribution here is the quantitative worst case: the exact domain, the supremum 2 with its support-disjointness characterization, the closed-form extremal family, and the direction result that the *concentrating* selection is the leximin one. A targeted review of the principal egalitarian-core, antiduality, exact-partition, large-core, and exact-cone sources found no prior fixed-dimensional separation bound (Arin et al., 2008; Dietzenbacher and Yanovskaya, 2020; Studený and Kratochvíl, 2022). The closest adjacent base-polyhedron theory studies one decreasingly minimal or lexicographically optimal base and its canonical partitions (Frank and Murota, 2022). More sharply, Schoot Uiterkamp (2024) characterizes compact convex sets having a least translated and rescaled majorized point for every translation

[1] With equal totals, Lorenz domination gives $\sum_{i \leq k} x_i^{\uparrow} \geq \sum_{i \leq k} y_i^{\uparrow}$ for every k; induction on k yields $x^{\uparrow} \geq_{\mathrm{lex}} y^{\uparrow}$, and complementing partial sums gives the leximax-min half. See Dutta and Ray (1989) for the Lorenz-dominant core element on convex games.

as base-polyhedra, and derives the corresponding characterization of convex TU games. That theorem concerns a least-majorized point after every strategic translation, not the separation of two opposing selectors in one arbitrary exact-game core. The exact and totally balanced cone papers classify facets in game-value space (Kroupa and Studený, 2019; Studený and Kratochvíl, 2022). This is not a hidden change of vocabulary: if an exact game's core is directly the base-polyhedron $B(f)$ of a submodular f, exactness gives $v(S) = \min_{x \in B(f)} x(S) = f(N) - f(N \setminus S)$, so v is convex. Thus the direct base-polyhedron transfer covers the convex subdomain, not arbitrary exact games. On the special two-bound domain, where the core is a box intersected with the efficiency hyperplane, the egalitarian core is a singleton given by a clipped- level formula (Gong et al., 2022); hence the opposing Lorenz selections coincide there. The later one-bound class is a reduced-game-stable subclass of the same coincidence domain (Gong et al., 2024). Exact-core-bound reductions now extend to all balanced games, but the resulting coordinate-bound equality values can lie outside the core already with three players (Gong and Dietzenbacher, 2026). Thus neither checked formulation states or implies a distance theorem for the two opposing selections on arbitrary exact cores. This negative search is not a novelty proof, and the claim remains subject to specialist review. The Lorenz-monotonic allocation-scheme framework of Izquierdo et al. (2024) is another close terminological neighbor: its allocations are indexed across subgames and its exactness strengthening characterizes convexity, rather than bounding the two selections within one exact core. Likewise, the maximal population-monotonic egalitarian domain of Dietzenbacher and Doğan (2025) characterizes a Dutta–Ray allocation scheme across populations rather than the separation of the two selections in one core.

The question of this paper is what happens after the rigidity breaks. Write

$$\rho(v) = \frac{\|\operatorname{Lmin}(v) - \operatorname{Lmax}(v)\|_1}{v(N)}, \qquad s(n) = \sup\{\rho(v) : \ v \text{ exact on } n \text{ players}, \ v(N) > 0\},$$

the supremum running over singleton-normalized games. Since $\operatorname{Lmin}(v)$ and $\operatorname{Lmax}(v)$ are nonnegative with the same total $v(N)$, we have $\rho \le 2$ always, and $\rho = 2$ would mean disjoint supports. The rigidity says $s(n) = 0$ for $n \le 4$. How large can $s(n)$ get?

Results. The answer is: as large as arithmetically possible.

- $\sup_n s(n) = 2$, and at fixed n, $s(n) \ge 2 - 8n/(n^2 - (n \bmod 2))$, which is $2 - 8/n$ for even n (Theorem 18); in particular $s(n) \ge 1$ from $n = 8$ and $s(n) > 1$ from $n = 9$; the bound is the full exact-class value at $n = 6$, and is attained on the two-generator class at $n = 7, 8$ by complete enumeration.
- The witnesses form a two-parameter family $G(p, k)$ whose two allocations are computed in closed form (Theorem 14), with the deficit identity $2 - \rho(G(p, k)) = 2n/(k(p + 1))$ (Corollary 16).
- On the two-generator class the core has a dual description as a system of upper caps (Theorem 7), from which the leximin floor inherits a one-coalition certificate (Theorem 8); the mirror statement for the leximax-min ceiling is false, and we give the smallest counterexample (Observation 11).
- $n = 5$ is anomalous: the family formula gives $1/3$, a game analysed by hand attains $1/2$ (Theorem 22), and the value is the supremum — $s(5) = 1/2$, with $\rho \le 1/3$ in the $(4, 1)$ leximin regime (Theorems 34 and 33, by configuration-complete finite check, twice over) — while the $n = 5$ extremal class is qualitatively unlike the $n \ge 6$ one.
- At $n = 6$ the family lower bound is sharp: $s(6) = 2/3$ (Theorem 36). The upper bound is a table-free exact finite proof using singleton exactness witnesses and lexicographic exchange blockers; all fifteen sign-count cases have independently checked CPC proof objects.

- In the weighted four-role quotient program, the conjectured atomic overlap bound holds for every exact nonnegative quotient satisfying the unit-role singleton hypothesis (Theorem 43). A human adjacent-gap theorem closes all atomic weights $n \geq 10$; a stratum-free quantified linear-real audit closes the six residual rows at $n = 6, 7, 8, 9$. Weighted common-face rigidity and a connected positive-Lorenz continuation show that the selectors always share two pair endpoints. This completes the quotient theorem, not the full fixed-n theorem. Strong induction now discharges common carrier modules, and strict-loss cycles automatically have proper packets with target-sized cardinality base. An exact role-irreducible seven-player control refutes the full-root-block scalar shortcut while preserving the target: its root budget is -1, but a labelled tight-support enlargement closes at base $5/6$. The remaining arm is selected-support congestion or a marked-facet alternative.

The mechanism, in one sentence. In $G(p, k)$ a single family of tight caps forces $p+1$ players to share only p; leximin reads each cap as a ceiling, so the floor sticks at $p/(p+1)$ and the entire residual piles onto one player, while leximax-min reads the same cap as a budget, loading half the players up to $p(k-1)/k$ and thereby squeezing the capped players down to $1/k$ — so the two supports become asymptotically disjoint and $\rho \to 2$. It is worth stating the direction explicitly, because it inverts the usual reading: the *concentrating* selection is the leximin one. In $G(p, k)$ the leximin allocation gives one player a share $p(k-1)/((p+1)k) \to 1$ of $v(N)$ while everyone else sits just below 1; it is nevertheless leximin-optimal, because nothing better is feasible for the poorest. Both selections are driven into extreme shapes by the same constraint, in opposite directions.

What is not proved. An upper bound matching Theorem 18 at fixed n is open, and Section 6 shows that one natural route to it is closed. Section 7 separates the proved statements from the enumerated and the conjectured ones, and lists the open problems.

2 Preliminaries

Throughout, $N = \{1, \dots, n\}$, all games are singleton-normalized, and $T := v(N) > 0$. Vectors are indexed by N, and $\operatorname{supp}(x) = \{i : x_i \neq 0\}$.

2.1 Exact games as lower envelopes

The following generation fact is the only construction used in this paper. It is standard; see Csóka et al. (2011) and Studený and Kratochvíl (2022) for the surrounding theory of the exact cone, and Flanagan (2026) for the same statement in the companion.

Lemma 1 (lower-envelope generation)**.** *Let $g_1, \dots, g_r \in \mathbb{R}^N_+$ have a common total $g_1(N) = \cdots = g_r(N) = T$, and set $v(S) = \min_j g_j(S)$. Then v is exact, $v(N) = T$, every g_j lies in* $\operatorname{Core}(v)$*, and v is singleton-normalized as soon as every coordinate vanishes in some g_j. Conversely, every exact singleton-normalized game is of this form: one may take $g_1, \dots, g_r$ to be the vertices of* $\operatorname{Core}(v)$*.*

Proof. Each g_j satisfies $g_j(S) \geq v(S)$ by definition of the minimum and $g_j(N) = T = v(N)$, so $g_j \in \operatorname{Core}(v)$; and the inequality at S is attained by whichever g_j realizes the minimum, so v is exact. The singleton statement is immediate from $v(\{i\}) = \min_j (g_j)_i$. Conversely, if v is exact then $v(S) = \min_{x \in \operatorname{Core}(v)} x(S)$ for every S, and the minimum of a linear functional over the polytope $\operatorname{Core}(v)$ is attained at a vertex; the vertices are nonnegative because v is singleton-normalized, and all have total T. □

We call a game generated as in Lemma 1 by r allocations an *r-generator* game. Part of what follows is specific to $r = 2$.

2.2 The two selections

For $x \in \mathbb{R}^N_+$ let $x^\uparrow$ be its increasing and $x^\downarrow$ its decreasing rearrangement.

Definition 2. $\mathrm{Lmin}(v)$ *is the allocation in* $\mathrm{Core}(v)$ *whose increasing rearrangement is lexicographically largest;* $\mathrm{Lmax}(v)$ *is the allocation in* $\mathrm{Core}(v)$ *whose decreasing rearrangement is lexicographically smallest.*

Both exist and are unique on a nonempty compact convex set: existence is compactness applied stage by stage, and uniqueness follows by comparing two putative optima with their midpoint, which is feasible by convexity and strictly better at the first stage where they differ. We write $\mu(v) = \min_i \mathrm{Lmin}(v)_i$ for the *leximin floor* and $M(v) = \max_i \mathrm{Lmax}(v)_i$ for the *leximax-min ceiling*; equivalently $\mu = \max_{x \in \mathrm{Core}} \min_i x_i$ and $M = \min_{x \in \mathrm{Core}} \max_i x_i$.

2.3 Normalization and the ratio

Two conventions have to be stated honestly, because neither selection is translation-equivariant: both are applied to the raw payoff vector x, not to $x - (v(\{i\}))_{i \in N}$, and on a game that is not singleton-normalized these are different objects — the leximax-min allocation in particular changes under the shift. We therefore fix the domain once and for all: all games are singleton-normalized, and both selections are computed on the normalized game. No claim below asserts invariance of Lmin or Lmax under adding singleton values.

Scaling, by contrast, is harmless: $\mathrm{Core}(cv) = c\,\mathrm{Core}(v)$ for $c > 0$ and both selections are positively homogeneous, so $\mathrm{Lmin}(cv) = c\,\mathrm{Lmin}(v)$ and $\mathrm{Lmax}(cv) = c\,\mathrm{Lmax}(v)$. Hence $\rho(v) = \|\,\mathrm{Lmin}(v) - \mathrm{Lmax}(v)\|_1 / v(N)$ is scale-invariant, which is what makes it comparable across games of different size. We record the trivial ceiling and the form of the deficit.

Lemma 3 (ceiling and overlap)**.** *For any singleton-normalized exact v with $T = v(N) > 0$,*

$$\rho(v) = 2\Big(1 - \frac{1}{T}\sum_{i \in N} \min\big(\mathrm{Lmin}(v)_i, \mathrm{Lmax}(v)_i\big)\Big) \;\le\; 2,$$

with $\rho(v) = 2$ if and only if $\mathrm{Lmin}(v)$ *and* $\mathrm{Lmax}(v)$ *have disjoint supports.*

Proof. For nonnegative x, y with $x(N) = y(N) = T$ one has $|x_i - y_i| = x_i + y_i - 2\min(x_i, y_i)$, and summing gives $\|x - y\|_1 = 2T - 2\sum_i \min(x_i, y_i)$. The overlap term is nonnegative and vanishes exactly when no coordinate is positive in both. □

Lemma 4 (zero-floor null-player descent)**.** *Let v be singleton-normalized and exact. If*

$$\min_i \mathrm{Lmin}(v)_i = 0,$$

then some player r is null. If v' is the restriction to $N \setminus \{r\}$, then

$$\mathrm{Core}(v) = \mathrm{Core}(v') \times \{0_r\}, \quad \mathrm{Lmin}(v) = (\mathrm{Lmin}(v'), 0_r), \quad \mathrm{Lmax}(v) = (\mathrm{Lmax}(v'), 0_r).$$

In particular, deletion preserves total worth, overlap, and ρ.

Proof. The first Lmin stage is $\max_{y\in\mathrm{Core}(v)} \min_i y_i$. If its value is zero, no core allocation is positive in every coordinate. If each player i were positive at some $y^i \in \mathrm{Core}(v)$, then the average of the finitely many y^i would be a core allocation positive in every coordinate, a contradiction. Hence some fixed r satisfies $y_r = 0$ throughout the core. Exactness gives, for $S \subseteq N \setminus \{r\}$,

$$v(S \cup \{r\}) = \min_{y\in\mathrm{Core}(v)} y(S \cup \{r\}) = \min_{y\in\mathrm{Core}(v)} y(S) = v(S),$$

so r is null.

Since $v(N \setminus \{r\}) = v(N)$, efficiency and nonnegativity force the r coordinate of every core point to zero. Restriction maps the core into $\mathrm{Core}(v')$. Conversely, extending any point of $\mathrm{Core}(v')$ by zero checks all coalitions containing r by nullity, so it lies in $\mathrm{Core}(v)$. This proves the product identity. Deleting a coordinate fixed throughout the feasible set leaves both lexicographic comparisons unchanged, proving the selection and overlap assertions. □

Put $c_n = n/\lfloor n^2/4 \rfloor$. These constants strictly decrease: for $n = 2k$,

$$c_{n-1} - c_n = \frac{2k-1}{k(k-1)} - \frac{2}{k} > 0,$$

and for $n = 2k+1$,

$$c_{n-1} - c_n = \frac{2}{k} - \frac{2k+1}{k(k+1)} > 0.$$

Consequently the zero-floor branch in dimension n follows inductively from the hybrid-overlap theorem in dimension $n-1$. The five-player overlap floor $3/4$ proved below exceeds $c_6 = 2/3$, and Theorem 36 therefore closes every zero-floor seven-player game. More generally, a minimal-dimensional counterexample to the proposed all-n bound must have strictly positive Lmin floor and no null player. This elementary descent is not claimed as a novel null-player result; its role is to remove the entire zero-floor branch from the remaining packet argument.

So $\rho \to 2$ is precisely asymptotic support disjointness of the two egalitarian selections, and the whole content of the paper is that this can be forced.

3 The two-generator class

Throughout this section v is the lower envelope of exactly two nonnegative allocations g, h with common total T, so that $v(S) = \min(g(S), h(S))$ for every S (including $S = N$, where both sides equal T).

Lemma 5 (support disjointness)**.** *If v is singleton-normalized then $A := \mathrm{supp}(g)$ and $B := \mathrm{supp}(h)$ are disjoint, and every player outside $A \cup B$ is null, receiving 0 in every core allocation. Deleting null players changes neither* Lmin*,* Lmax*, nor T, so we may and do assume $A \cup B = N$ and $A \cap B = \emptyset$.*

Proof. $v(\{i\}) = \min(g_i, h_i) = 0$ forces $g_i = 0$ or $h_i = 0$, which is exactly $A \cap B = \emptyset$. If $i \notin A \cup B$ then $g_i = h_i = 0$, so $v(N \setminus \{i\}) = \min\big(g(N \setminus \{i\}), h(N \setminus \{i\})\big) = T$; any $x \in \mathrm{Core}(v)$ has $x(N \setminus \{i\}) \geq T$ and $x(N) = T$ with $x \geq 0$, hence $x_i = 0$. Both selections assign a null player 0, so deletion is harmless. □

Definition 6. *For $S \subseteq N$ put $f(S) := \max\big(g(S), h(S)\big) = \max\big(g(S \cap A), h(S \cap B)\big)$.*

3.1 A dual description of the core

Theorem 7 (two-generator core duality)**.**

$$\mathrm{Core}(v) = \big\{x \geq 0 :\ x(N) = T \text{ and } x(S) \leq f(S) \text{ for all } S \subseteq N\big\}.$$

Proof. For x with $x(N) = T$ and any S, the inequality $x(S) \geq v(S)$ is equivalent to $x(N \setminus S) \leq T - v(S)$. It therefore suffices to prove

$$T - v(S) = f(N \setminus S) \qquad \text{for every } S \subseteq N. \tag{1}$$

Since g and h are modular with $g(N) = h(N) = T$,

$$T - v(S) = T - \min\big(g(S), h(S)\big) = \max\big(T - g(S),\, T - h(S)\big) = \max\big(g(N \setminus S),\, h(N \setminus S)\big) = f(N \setminus S).$$

Replacing S by $N \setminus S$ in (1) turns the family $\{x(S) \geq v(S)\}$ into the family $\{x(S) \leq f(S)\}$. Finally $x \geq 0$ is exactly the family of singleton inequalities $x(\{i\}) \geq v(\{i\}) = 0$. □

Identity (1) uses only modularity and the common total; the disjointness of Lemma 5 is not needed for it, and enters only to normalize the presentation of f. Note also that $S = N$ gives $f(N) = T$, so the cap at N is implied by efficiency, and $S = \emptyset$ gives $f(\emptyset) = 0$.

In fact the duality is fully general: for *any* singleton-normalized game, setting $f(S) := T - v(N \setminus S)$, the same complementation involution gives $\mathrm{Core}(v) = \{x \geq 0 : x(N) = T,\ x(S) \leq f(S)\ \forall S\}$. The two-generator content of Theorem 7 is the computation $f = \max(g, h)$. The general form is used in Section 5.4.

3.2 The leximin floor

Theorem 8 (leximin floor: the easy direction)**.** *Let*

$$\lambda := \min_{\emptyset \neq S \subseteq N} \frac{f(S)}{|S|} = \min_{(A', B') \neq (\emptyset, \emptyset)} \frac{\max\big(g(A'), h(B')\big)}{|A'| + |B'|},$$

the second minimum running over $A' \subseteq A$, $B' \subseteq B$. Then $\mu \leq \lambda$.

Proof. $\mathrm{Lmin}(v)$ is a core allocation, so by Theorem 7, for every nonempty S, $|S|\,\mu \leq \mathrm{Lmin}(v)(S) \leq f(S)$, i.e. $\mu \leq f(S)/|S|$. Take the minimum over S. □

The bound is a *one-coalition certificate*: a single tightest cap already pins the floor from above. Its natural habitat is polymatroid theory: when f is monotone submodular, the polytope of Theorem 7 is a polymatroid base polytope, the leximin floor *equals* λ, and the decreasingly-minimal and increasingly-maximal points coincide — the theory of lexicographically optimal bases (Fujishige, 1980), where $\rho = 0$ identically. The two-generator $f = \max(g, h)$ is not submodular in general; the conjecture below asks whether the floor identity survives the loss of submodularity, and the rest of the paper measures exactly how far the dec-min/inc-max coincidence fails. Empirically the bound is exact.

Conjecture 9 (leximin floor formula)**.** *$\mu = \lambda$ for every singleton-normalized two-generator exact game.*

Evidence and reduction. Conjecture 9 holds in 586 of 586 random two-generator games at $n = 4, 5, 6, 7$ (sample sizes 120, 218, 165, 83), with all arithmetic exact. What remains to prove is a clean linear-programming statement. Equality means that the system

$$y \geq 0, \qquad y(N) = T - n\lambda, \qquad y(S) \leq f'(S) := f(S) - \lambda|S|$$

is feasible; indeed, given such y, the allocation $x = y + \lambda\mathbf{1}$ satisfies $x(N) = T$, $x \geq \lambda\mathbf{1} \geq 0$ and $x(S) \leq f(S)$, hence lies in the core with floor at least λ. Note $f' \geq 0$ by definition of λ. By linear-programming duality,

$$\max\{y(N) :\ y \geq 0,\ y(S) \leq f'(S)\ \forall S\} = \min\Big\{\textstyle\sum_S w_S f'(S) :\ w \geq 0,\ \sum_{S\ni i} w_S \geq 1\ \forall i\Big\},$$

the minimum running over fractional covers. Taking $w = \mathbf{1}_{\{S=N\}}$ shows the minimum is at most $f'(N) = T - n\lambda$, so Conjecture 9 is exactly the assertion

$$\sum_S w_S f'(S) \ \geq\ f'(N) \qquad \text{for every fractional cover } w. \tag{2}$$

Proposition 10 (balanced covers)**.** *Inequality (2) holds for every exact (balanced) cover, i.e. whenever* $\sum_{S\ni i} w_S = 1$ *for all* i.

Proof. Balancedness gives $\sum_S w_S|S| = \sum_i \sum_{S\ni i} w_S = n$, so, using $f \geq g$ and the modularity of g,

$$\sum_S w_S f'(S) = \sum_S w_S f(S) - \lambda n \ \geq\ \sum_S w_S g(S) - \lambda n = \sum_i g_i \sum_{S\ni i} w_S - \lambda n = g(N) - \lambda n = f'(N),$$

the last step because $g(N) = T$ and $f'(N) = T - \lambda n$. The same computation works with h in place of g. □

What is missing is the case of covers that strictly over-cover some players. Writing $e_i := \sum_{S\ni i} w_S - 1 \geq 0$, the same computation gives

$$\sum_S w_S f'(S) \ \geq\ f'(N) + \max\Big(\sum_{i\in A} g_i e_i,\ \sum_{i\in B} h_i e_i\Big) - \lambda \sum_{i\in N} e_i,$$

and since $\lambda \leq f(\{i\})$, which equals g_i for $i \in A$ and h_i for $i \in B$, each inner sum dominates λ times its *own* block's excess but not the other block's. Closing that gap, or constructing the attaining core allocation directly, is the remaining work.

3.3 The leximax-min ceiling is not dual

The mirror-image statement is an inequality but not an identity, and this is the first place where the two selections part company.

Observation 11 (M-side asymmetry)**.** *For every core allocation, and in particular for* $\mathrm{Lmax}(v)$,

$$M \ \geq\ \max_{\substack{P\subseteq A,\ Q\subseteq B\\ (P,Q)\neq(\emptyset,\emptyset)}} \frac{\min\big(g(P), h(Q)\big)}{|P| + |Q|}. \tag{3}$$

Inequality (3) is not an identity: it fails strictly already at $n = 6$.

Proof of the inequality. Let $x \in \mathrm{Core}(v)$ and put $S = P \cup Q$. Then $g(S) = g(P)$ and $h(S) = h(Q)$ because g vanishes on B and h on A, so

$$\max_i x_i \ \geq\ \frac{x(S)}{|S|} \ \geq\ \frac{v(S)}{|S|} = \frac{\min\big(g(P), h(Q)\big)}{|P| + |Q|}.$$

Apply this to $x = \mathrm{Lmax}(v)$. Equivalently, (3) is Theorem 7 read on complements. □

How to read (3). The maximum runs over pairs (P, Q) other than $(\emptyset, \emptyset)$, with the convention $g(\emptyset) = h(\emptyset) = 0$. There is no infinity convention: a pair with an empty part contributes $\min(\cdot, 0) = 0$ and is simply inactive, so only pairs with *both* parts nonempty can bind. This is visible in the proof, where an empty part makes $v(S)$ the minimum of something and 0. Reading an empty part as $+\infty$ instead would reduce (3) to a maximum over singletons and give a different — and wrong — value; the misreading is easy enough that we spell the convention out.

The smallest counterexample. Scanning the complete two-generator enumeration in increasing T, no counterexample to equality in (3) exists at $n \le 5$ for any T, nor at $n = 6$ for $T \le 4$. At $n = 6$, $T = 5$:

$$g = (3, 2, 0, 0, 0, 0), \qquad h = (0, 0, 2, 1, 1, 1),$$
$$\mathrm{Lmin}(v) = \left(1, \tfrac{2}{3}, \tfrac{4}{3}, \tfrac{2}{3}, \tfrac{2}{3}, \tfrac{2}{3}\right), \qquad \mathrm{Lmax}(v) = \left(\tfrac{6}{5}, \tfrac{4}{5}, \tfrac{6}{5}, \tfrac{3}{5}, \tfrac{3}{5}, \tfrac{3}{5}\right),$$

so $M = 6/5$, whereas the right-hand side of (3) equals 1, attained for instance at $P = \{1\}, Q = \{3\}$ with $\min(3, 2)/2 = 1$. The gap is $1/5$. One further counterexample class occurs at $n = 6$, $T = 5$: $g = (3, 1, 1)$, $h = (2, 2, 1)$ on the respective supports (its block-swapped relabelling is the same game up to permutation), again with $M = 6/5$ against a bound of 1. Over the same random samples used for Conjecture 9, (3) is tight in 120/120 games at $n = 4$, 218/218 at $n = 5$, 160/165 at $n = 6$ and 74/83 at $n = 7$.

Remark 12 (diagnosis)**.** *Theorem 8 bounds the floor using only the upper caps $x(S) \le f(S)$, and a single tightest cap can be saturated on its own, which is why one coalition certifies the optimum. Bounding the ceiling instead requires the lower bounds $x(N \setminus S) \ge T - f(S)$, and* minimizing *a maximum must respect several of them simultaneously; no single one certifies the optimum. This is a genuine asymmetry between the two selections, and it is the reason* Lmax *is the harder object throughout.*

4 The master family

Definition 13. *For integers $p \ge 1$ and $k \ge 2$ put $n = p + 1 + k$ and $T = pk$, and let $G(p, k)$ be the lower envelope of*

$$g = \big(p(k-1),\ \underbrace{1, \dots, 1}_{p},\ \underbrace{0, \dots, 0}_{k}\big), \qquad h = \big(\underbrace{0, \dots, 0}_{p+1},\ \underbrace{p, \dots, p}_{k}\big).$$

We index the players $0, 1, \dots, p + k$ and call player 0 the head*, $O := \{1, \dots, p\}$ the* ones*, and $B := \{p + 1, \dots, p + k\}$; so $A = \{0\} \cup O$ has $p + 1$ elements and $|B| = k$.*

Both generators are nonnegative with total $p(k-1)+p = pk = T$, their supports are disjoint, and every coordinate vanishes in one of them; by Lemma 1, $G(p, k)$ is exact and singleton-normalized.

Theorem 14 (closed forms)**.** *For all $p \ge 1$, $k \ge 2$,*

$$\mathrm{Lmin}\big(G(p,k)\big) = \left(\tfrac{p^2(k-1)}{p+1},\ \left[\tfrac{p}{p+1}\right]^{p+k}\right),$$
$$\mathrm{Lmax}\big(G(p,k)\big) = \left(\tfrac{p(k-1)}{k},\ \left[\tfrac{1}{k}\right]^{p},\ \left[\tfrac{p(k-1)}{k}\right]^{k}\right),$$
$$\big\|\mathrm{Lmin} - \mathrm{Lmax}\big\|_1 = \frac{2p\big(p(k-1)-1\big)}{p+1}, \qquad \rho\big(G(p,k)\big) = \frac{2\big(p(k-1)-1\big)}{k(p+1)},$$

where the exponents indicate the number of coordinates, listed in the order head, ones, B.

The proof occupies Sections 4.1–4.6: a key cap, the two optimizations, a membership check, and the arithmetic of the gap.

4.1 The key cap

For each $b \in B$,

$$f\big(O \cup \{b\}\big) = \max\big(g(O), h(\{b\})\big) = \max(p, p) = p,$$

so by Theorem 7 every core allocation x satisfies

$$x_1 + \cdots + x_p + x_b \;\le\; p \qquad \text{for every } b \in B. \tag{KC}$$

That single family of $(p+1)$-player caps drives everything below. Note that the two generators *tie* at these coalitions: the maximum defining f is attained on both sides.

4.2 Leximin

Let $x \in \mathrm{Core}(G(p,k))$ and $m := \min_i x_i$. Applying (KC) with any fixed b gives $(p+1)m \le p$, so

$$m \le \frac{p}{p+1}, \qquad \text{hence} \qquad \mu \le \frac{p}{p+1}.$$

Suppose $m = p/(p+1)$. Then in (KC) the $p+1$ numbers $x_1, \ldots, x_p, x_b$ are each at least $p/(p+1)$ and sum to at most $p = (p+1)\cdot p/(p+1)$, so all of them equal $p/(p+1)$. This holds for *every* $b \in B$, hence

$$x_1 = \cdots = x_p = x_{p+1} = \cdots = x_{p+k} = \frac{p}{p+1},$$

and efficiency forces the head to

$$x_0 = T - (p+k)\frac{p}{p+1} = pk - \frac{p(p+k)}{p+1} = \frac{p^2(k-1)}{p+1},$$

using $k - (p+k)/(p+1) = (kp+k-p-k)/(p+1) = p(k-1)/(p+1)$. Thus the max-min allocation is *unique*, and a unique max-min allocation is a fortiori the leximin allocation. Since $p(k-1) \ge 1$, the head's coordinate is at least $p/(p+1)$, so the floor of this allocation is indeed $p/(p+1)$. It remains to check that this point is in the core, which is done in Section 4.4.

4.3 Leximax-min

Let $x \in \mathrm{Core}(G(p,k))$, $M := \max_i x_i$ and $s := x_1 + \cdots + x_p$. By (KC), $x_b \le p - s$ for every $b \in B$, and trivially $x_b \le M$ and $x_0 \le M$. Efficiency gives

$$pk = x_0 + s + \sum_{b \in B} x_b \;\le\; M + s + k\min(M,\, p-s). \tag{4}$$

Case A: $M \le p - s$. Then (4) reads $pk \le (k+1)M + s$. The case hypothesis is $s \le p - M$, so $pk \le (k+1)M + p - M = kM + p$, i.e. $M \ge p(k-1)/k$.

Case B: $M > p - s$. Then (4) reads $pk \le M + s + k(p-s)$, i.e. $(k-1)s \le M$, so $s \le M/(k-1)$. The case hypothesis then gives $M > p - s \ge p - M/(k-1)$, i.e. $Mk/(k-1) > p$, i.e. $M > p(k-1)/k$.

In both cases $M \ge p(k-1)/k$, so the leximax-min ceiling is at least $p(k-1)/k$. Equality can occur only in Case A with every inequality used tight, that is with $x_0 = M$, with $x_b = M$ for all $b \in B$, and with

$$s = pk - (k+1)M = pk - (k+1)\frac{p(k-1)}{k} = \frac{p\big(k^2 - (k+1)(k-1)\big)}{k} = \frac{p}{k}.$$

(Consistency: $p - s = p - p/k = p(k-1)/k = M$, so Case A holds with equality, as it must.) Thus at the min-max level the head and all of B are pinned at $p(k-1)/k$ and the ones share exactly p/k; since $p/k \le p(k-1)/k$ for $k \ge 2$, the ones occupy the last p positions of the decreasing rearrangement. The next leximax stage minimizes their maximum subject to their sum being p/k, whose unique optimum is $1/k$ each. Hence Lmax is as claimed, again modulo membership.

4.4 Membership

This step is load-bearing: the two optimizations above determine the optima only conditionally on the described allocations being feasible.

Index as in Definition 13 and, for $S \subseteq N$, put

$$\varepsilon := [\,0 \in S\,], \qquad i := |S \cap O|, \qquad j := |S \cap B|,$$

so that $0 \le i \le p$, $0 \le j \le k$ and

$$g(S) = \varepsilon\, p(k-1) + i, \qquad h(S) = jp, \qquad f(S) = \max\big(\varepsilon\, p(k-1) + i,\ jp\big). \tag{5}$$

Lemma 15. *Both allocations of Theorem 14 lie in* $\mathrm{Core}(G(p,k))$ *for all* $p \ge 1$, $k \ge 2$. *Moreover, for both of them the tight coalitions are exactly*

$$\begin{array}{lll} \textit{(T1)} & S = O \cup \{b\},\ b \in B, & f(S) = \max(p,p) = p, \\ \textit{(T2)} & S = \{0\} \cup B',\ |B'| = k-1, & f(S) = \max\big(p(k-1),(k-1)p\big) = p(k-1), \\ \textit{(T3)} & S = N, & f(N) = T. \end{array}$$

Proof. The leximin allocation. Write $q := p/(p+1)$ and $H := p^2(k-1)/(p+1) = p(k-1)q$, so $x^* = (H, q, \dots, q)$ with $p+k$ copies of q. Efficiency:

$$H + (p+k)q = q\big(p(k-1) + p + k\big) = q(pk+k) = k(p+1)\frac{p}{p+1} = pk = T,$$

and $x^* > 0$. It remains to check $x^*(S) = \varepsilon H + (i+j)q \le f(S)$.

$\varepsilon = 0$, $j = 0$: S is a subset of O, and $x^*(S) = iq < i = f(S)$ for $i \ge 1$ because $q < 1$; for $i = 0$, $S = \emptyset$. Strict.

$\varepsilon = 0$, $j \ge 1$: bound $f(S) \ge jp$ and use $i \le p$:

$$x^*(S) = (i+j)\frac{p}{p+1} \le (p+j)\frac{p}{p+1} \le jp \iff p + j \le j(p+1) \iff p \le jp,$$

which holds since $j \ge 1$. Equality throughout requires $i = p$ and $j = 1$, i.e. $S = O \cup \{b\}$, which is (T1).

$\varepsilon = 1$, $j \le k-1$: put $G_0 := p(k-1) + i = g(S)$, so $x^*(S) = q(G_0 + j)$ and

$$x^*(S) \le G_0 \iff p(G_0 + j) \le (p+1)G_0 \iff jp \le G_0 = p(k-1) + i,$$

which holds because $jp \le (k-1)p$. Hence $x^*(S) \le g(S) \le f(S)$, with equality exactly when $j = k-1$ and $i = 0$, i.e. $S = \{0\} \cup B'$ with $|B'| = k-1$, which is (T2).

$\varepsilon = 1$, $j = k$: then $h(S) = kp = T$, so $f(S) \ge T \ge x^*(S)$, with equality only for $S = N$ since $x^* > 0$; this is (T3).

The leximax-min allocation. Write $L := p(k-1)/k$ and $r := 1/k$, so $y^* = (L, r, \dots, r, L, \dots, L)$ with p copies of r and k copies of L. Efficiency:

$$L + pr + kL = (k+1)L + \frac{p}{k} = \frac{p\big[(k+1)(k-1)+1\big]}{k} = \frac{pk^2}{k} = pk = T,$$

and $y^* > 0$. Now $y^*(S) = \varepsilon L + ir + jL$.

$\varepsilon = 0$, $j = 0$: $y^*(S) = i/k \le i = f(S)$ since $k \ge 2$, strictly for $i \ge 1$.

$\varepsilon = 0$, $j \ge 1$: comparing with the h-side,

$$\frac{i + jp(k-1)}{k} \le jp \iff i + jp(k-1) \le jpk \iff i \le jp,$$

which holds since $i \le p \le jp$, with equality exactly when $i = p$ and $j = 1$, i.e. (T1).

$\varepsilon = 1$, $j \le k-1$: $y^*(S) = (j+1)L + i/k$, and multiplying by $k > 0$,

$$y^*(S) \le p(k-1) + i \iff (j+1)p(k-1) + i \le kp(k-1) + ki \iff p(j+1-k) \le i,$$

after dividing by $k - 1 \ge 1$; the left-hand side is at most $0 \le i$, with equality exactly when $j = k-1$ and $i = 0$, i.e. (T2).

$\varepsilon = 1$, $j = k$: as before $f(S) \ge h(S) = kp = T \ge y^*(S)$, with equality only at $S = N$. □

4.5 The hinge: both selections saturate the same caps

Lemma 15 contains the expository point of the paper. The active constraints are the same for Lmin and Lmax: the two families (T1) and (T2), plus efficiency. Both are coalitions at which the two generators *tie*, which is the structural reason they are the active ones — there is no slack on either side to give. What differs is only how the two selections distribute mass *inside* a tight coalition. On $S = O \cup \{b\}$, which has total exactly p:

	each of the p ones	the single $b \in B$
Lmin :	$\frac{p}{p+1}$	$\frac{p}{p+1}$
Lmax :	$\frac{1}{k}$	$\frac{p(k-1)}{k}$

Leximin splits the budget p evenly over all $p+1$ members; leximax-min gives the ones the least it can and puts the whole remainder on the one member of B, because loading B is how it lowers the global maximum. The divergence is visible in this one coalition. Everything else is bookkeeping: leximin's even split leaves the residual $T - (p+k)p/(p+1)$ with nowhere to go but the head, while leximax-min's loading of every $b \in B$ up to $p(k-1)/k$ consumes almost all of T across B.

4.6 The gap

Proof of Theorem 14, conclusion. Put $D := p(k-1) - 1 \ge 0$. Coordinatewise,

$$\text{head:} \quad \frac{p^2(k-1)}{p+1} - \frac{p(k-1)}{k} = \frac{p(k-1)D}{k(p+1)},$$

$$\text{each of the } p \text{ ones:} \quad \frac{p}{p+1} - \frac{1}{k} = \frac{D}{k(p+1)},$$

$$\text{each of the } k \text{ members of } B\text{:} \quad \frac{p(k-1)}{k} - \frac{p}{p+1} = \frac{pD}{k(p+1)},$$

all three nonnegative. Hence

$$\| \text{Lmin} - \text{Lmax} \|_1 = \frac{p(k-1)D + pD + kpD}{k(p+1)} = \frac{D\big(p(k-1) + p + kp\big)}{k(p+1)} = \frac{2pkD}{k(p+1)} = \frac{2pD}{p+1},$$

and $\rho = 2pD/\big((p+1)pk\big) = 2D/\big(k(p+1)\big)$. □

Corollary 16 (deficit identity). $2-\rho\big(G(p,k)\big)=\dfrac{2n}{k(p+1)}$, *where* $n=p+1+k$.

Proof. $2-\dfrac{2\big(p(k-1)-1\big)}{k(p+1)}=\dfrac{2\big[k(p+1)-p(k-1)+1\big]}{k(p+1)}=\dfrac{2(kp+k-pk+p+1)}{k(p+1)}=\dfrac{2(k+p+1)}{k(p+1)}$. □

Remark 17 (the deficit is twice the overlap). *Corollary 16 also follows from Lemma 3. In* $G(p,k)$ *the coordinatewise minima of the two allocations are* $p(k-1)/k$ *at the head,* $1/k$ *at each one, and* $p/(p+1)$ *at each member of* B*, so*

$$\sum_i \min\big(\mathrm{Lmin}_i,\mathrm{Lmax}_i\big)=\frac{p(k-1)}{k}+\frac{p}{k}+\frac{kp}{p+1}=p+\frac{kp}{p+1}<p+k=n-1,$$

and $2\cdot\big(p+kp/(p+1)\big)/(pk)=2(p+1+k)/\big(k(p+1)\big)$. *The overlap between the two selections is thus bounded by* n *while* $T=pk$ *grows quadratically: this is the precise sense in which their supports become asymptotically disjoint.*

4.7 The supremum

Theorem 18 (value of the supremum). *For every* $n\ge 4$,

$$s(n)\;\ge\;2-\frac{8n}{n^2-(n \bmod 2)}=\begin{cases}2-\dfrac{8}{n}, & n \text{ even},\\[2mm] 2-\dfrac{8n}{n^2-1}, & n \text{ odd},\end{cases}$$

attained on the two-generator class by $G(p,k)$ *with the balanced split* $k=\lceil n/2\rceil$, $p+1=\lfloor n/2\rfloor$. *Consequently* $\sup_n s(n)=2$.

Proof. By Corollary 16 the deficit of $G(p,k)$ is $2n/(k(p+1))$ with $k+(p+1)=n$, $k\ge 2$ and $p+1\ge 2$. Minimizing the deficit means maximizing the product $k(p+1)$ of two positive integers with fixed sum n, whose maximum is $\lfloor n/2\rfloor\lceil n/2\rceil$, equal to $n^2/4$ for even n and $(n^2-1)/4$ for odd n, attained at the balanced split (for $n\ge 4$ this split satisfies the constraints $k\ge 2$, $p\ge 1$). Substituting gives the displayed bound. For the last claim, $\rho\le 2$ always by Lemma 3, and $2-8/n\to 2$. □

Table 1 lists the resulting values. The ratio reaches 1 at $n=8$ and passes it at $n=9$: from eight players on the two egalitarian core selections can differ by the entire worth of the grand coalition, and from nine players by more.

Remark 19 (the degenerate corner). $G(1,2)$ *has* $n=4$, $T=2$, $g=(1,1,0,0)$, $h=(0,0,1,1)$ *and* $\rho=2(1\cdot 1-1)/(2\cdot 2)=0$, *consistent with the four-player rigidity; the deficit identity reads* $2-0=2n/(k(p+1))=8/4=2$.

Remark 20 (there is nothing special about $|B|=4$). *An earlier version of this family fixed* $k=|B|=4$, *giving* $\rho=(3p-1)/\big(2(p+1)\big)\to 3/2$. *Corollary 16 shows that* k *is a free parameter and that the optimum is* $k\approx n/2$, *so the limit* $3/2$ *was an artifact of the fixed block size; the* $k=4$ *slice is already beaten from* $n=10$ *on (*$7/6$ *against* $6/5$*). The number* 4 *that appears as* $|\,\mathrm{supp}(\mathrm{Lmax})|$ *in the* $n=5$ *extremal class (Section 5) is a small-*n *phenomenon and not a feature of the asymptotics: in* $G(p,k)$ *every coordinate of* Lmax *is strictly positive, so* $|\,\mathrm{supp}(\mathrm{Lmax})|=n$ *and no player is starved at all.*

n	(p,k)	lower bound on $s(n)$	decimal
6	$(2,3)$	$2/3$	0.6667
7	$(2,4)$	$5/6$	0.8333
8	$(3,4)$	1	1.0000
9	$(3,5)$	$11/10$	1.1000
10	$(4,5)$	$6/5$	1.2000
11	$(4,6)$	$19/15$	1.2667
12	$(5,6)$	$4/3$	1.3333
16	$(7,8)$	$3/2$	1.5000
20	$(9,10)$	$8/5$	1.6000
100	$(49,50)$	$48/25$	1.9200

Table 1: The bound of Theorem 18 at the balanced split. All entries are exact; decimals are rounded.

n	totals	games	enumerated max ρ	Theorem 18
3	2–30	225	0	—
4	2–30	3,040	0	0
5	2–24	4,357	$1/2$	$1/3$
6	2–20	8,213	$2/3$	$2/3$
7	2–14	1,609	$5/6$	$5/6$
8	2–12	1,049	1	1

Table 2: Complete two-generator enumerations. The rows $n = 3, 4$ are controls: they reproduce the four-player rigidity. The last column is the value of the balanced-split family bound.

5 Fixed n: enumeration, certificates, and the $n = 5$ anomaly

5.1 Complete two-generator enumeration

Write $s_2(n)$ for the supremum of ρ over singleton-normalized two-generator exact games on n players. Table 2 reports complete enumerations, up to player permutation, generator swap and scaling, of all integer generator pairs with common total in the stated range. All arithmetic is exact rational.

Two readings must be kept apart.

What is complete. Because ρ is scale-invariant, every two-generator game with rational generators has a unique primitive integer representative, and enumeration to total $T_{\max}$ is exactly complete for the games whose primitive representative has total at most $T_{\max}$ — equivalently, for rational shapes of denominator at most $T_{\max}$. It says nothing about finer denominators or about irrational generators, so enumeration-to-T is a lower-bound method for $s_2(n)$: the entries of Table 2 are maxima over the enumerated range, not values of s_2. As a partial probe of the excluded region, shape hill-climbing at fine resolution was run from the known witnesses (resolution $1/60$ at $n = 5$; $1/56$ and $1/60$ at $n = 6$; $1/64$ at $n = 7$). None improved on the tabulated maximum; at $n = 7$ the climb from the $5/6$ witness stalled after 42 steps, and a second start from a base of $2/3$ climbed only to $3/4$.

What agrees. The enumerated maxima agree exactly with Theorem 18 at $n = 6, 7, 8$, so on the two-generator class that bound is attained there, and

$$s_2(n) \stackrel{?}{=} 2 - \frac{8n}{n^2 - (n \bmod 2)} \qquad (n \geq 6)$$

is a live candidate, currently confirmed at three values of n. At $n = 5$ the two disagree, and the enumeration is the correct one; see Section 5.3.

5.2 Searching the full exact class

The two-generator restriction is not innocuous. Adding a generator *lowers* the envelope and hence *grows* the core, so it is not automatic that two generators are extremal. The following observation lets one search the entire exact class rather than a bounded-generator slice.

Lemma 21 (core closing)**.** *Fix $T > 0$ and a vector $c = (c_S)$ indexed by the proper coalitions, and put $P(c) = \{x : x(N) = T,\ x(S) \geq c_S$ for all proper $S\}$. If $P(c) \neq \emptyset$, then the game v defined by $v(S) = \min_{x \in P(c)} x(S)$ for proper S and $v(N) = T$ is exact, and $\mathrm{Core}(v) = P(c)$.*

Proof. Every $x \in P(c)$ satisfies $x(S) \geq v(S)$ by definition of v and $x(N) = T$, so $P(c) \subseteq \mathrm{Core}(v)$; and each minimum is attained, so v is exact. Conversely, $v(S) \geq c_S$ for every proper S because every point of $P(c)$ satisfies $x(S) \geq c_S$; hence $x \in \mathrm{Core}(v)$ implies $x(S) \geq v(S) \geq c_S$ and $x(N) = T$, i.e. $x \in P(c)$. □

Perturbing c and re-closing therefore moves within the exact class and reaches cores with arbitrarily many vertices. Hill climbing in this "c-space", *seeded at the extremal two-generator game* of Table 2, produced no improvement at $n = 5$ (240 steps, stays at $1/2$), $n = 6$ (496 steps, stays at $2/3$) or $n = 7$ (301 steps, stays at $5/6$). Since the search is over all exact games, this certifies the two-generator optimum as a *local* optimum in the full exact class at those three values of n. Independently, from the known extremal at each of $n = 5, 6, 7, 8$, about 200 structured and random third generators at fixed common total were added: zero improvements.

One limitation must be stated. Unseeded c-space restarts reach only $1/3$ at $n = 6$, far below the known $2/3$; that variant of the search is under-powered for $n \geq 6$ and certifies nothing. Only the seeded runs support the local-optimality claim, and none of these runs bears on global optimality.

5.3 The $n = 5$ anomaly

At $n = 5$ the balanced-split family gives only $1/3$, but the truth is $1/2$, and the witness has a non-uniform h-block. The proof is by hand from two flat splits of the game; no linear program and no vertex enumeration is used.

Theorem 22 ($n = 5$)**.** *Let w be the lower envelope of $g^1 = (1, 0, 0, 1, 2)$ and $g^2 = (0, 1, 3, 0, 0)$. Then w is exact and singleton-normalized with $w(N) = 4$, and*

$$\mathrm{Lmin}(w) = \left(\tfrac{1}{2}, \tfrac{1}{2}, \tfrac{5}{4}, \tfrac{1}{2}, \tfrac{5}{4}\right), \qquad \mathrm{Lmax}(w) = (1, 0, 1, 1, 1), \qquad \| \mathrm{Lmin} - \mathrm{Lmax} \|_1 = 2, \qquad \rho(w) = \tfrac{1}{2}.$$

In particular $s(5) \geq 1/2 > 1/3$.

Proof. Exactness and normalization are immediate from Lemma 1, each coordinate vanishing in one generator. Two flat splits control the core: $w(\{1, 2\}) = 1$ and $w(\{3, 4, 5\}) = 3$ sum to $w(N) = 4$, so

$x_1 + x_2 = 1$ on all of Core(w), and likewise $w(\{2,4\}) = 1$ with $w(\{1,3,5\}) = 3$ gives $x_2 + x_4 = 1$. Putting $s := x_1$, every core point therefore satisfies

$$x_1 = x_4 = s, \qquad x_2 = 1 - s, \qquad x_3 + x_5 = 3 - s, \tag{6}$$

with $s \in [0,1]$ by nonnegativity of x_1 and x_2.

Leximin. By (6), $\min_i x_i \le \min(s, 1-s) \le 1/2$, with equality forcing $s = 1/2$ and hence $x_1 = x_2 = x_4 = 1/2$. Then $x_3 + x_5 = 5/2$, and the leximin continuation splits this evenly, giving $5/4$ each. The resulting allocation $x = (1/2, 1/2, 5/4, 1/2, 5/4)$ lies in the core, as direct substitution against both generators confirms. Hence Lmin(w) is as claimed.

Leximax-min. By (6), $\max_i x_i \ge (x_3 + x_5)/2 = (3-s)/2 \ge 1$ since $s \le 1$. Equality forces $s = 1$, so $x = (1, 0, x_3, 1, x_5)$ with $x_3 + x_5 = 2$; the ceiling $\max_i x_i = 1$ then forces $x_3 = x_5 = 1$. The unique candidate $(1,0,1,1,1)$ lies in the core, so Lmax(w) $= (1,0,1,1,1)$.

Finally $\| \text{Lmin} - \text{Lmax} \|_1 = \frac{1}{2} + \frac{1}{2} + \frac{1}{4} + \frac{1}{2} + \frac{1}{4} = 2$ and $\rho = 2/4 = 1/2$. □

Four independent methods agree on this game: the hand argument above, exact rational core-vertex enumeration, a staged exact solver over linear real arithmetic, and a from-scratch exact-rational two-phase simplex with Bland's rule.

The value $\frac{1}{2}$ is in fact the supremum: $s(5) = \frac{1}{2}$ (Theorem 34 below, by a configuration-complete exact finite check). Before that machinery, the searches already pointed there: roughly 4×10^4 distinct exact five-player games evaluated in exact rational arithmetic produced nothing above $1/2$ — the complete two-generator enumeration of Table 2, complete three-generator enumerations at totals 4–7, c-space hill climbing in the sense of Lemma 21, and fine-denominator shape climbing from the extremal witness — with the $n = 3, 4$ controls returning maximum 0 from the same pipeline.

Structure of the extremal class. The following are computational findings over 360 distinct profiled games at $n = 5$, of which 199 have $\rho > 1/3$ and 107 attain $\rho = 1/2$. Among the 107 extremal games, without exception:

- Lmin has stratum composition $(3,2)$, and Lmax has stratum composition $(1,4)$;
- Lmax has exactly one zero coordinate, so $|\text{supp}(\text{Lmax})| = 4$;
- normalized to $T = 4$ the pair is unique up to permutation: Lmin $\sim (\frac{1}{2}, \frac{1}{2}, \frac{1}{2}, \frac{5}{4}, \frac{5}{4})$ and Lmax $\sim (0,1,1,1,1)$ — that is, exactly the pair of Theorem 22.

Among the 199 games above $1/3$, the Lmin strata are $(3,2)$ in 196 cases and $(2,1,2)$ in 3; the pattern $(4,1)$ never occurs above $1/3$, and Lmax is never the uniform allocation.

One tempting invariant is refuted. Call S a *flat split* if $v(S) + v(N \setminus S) = v(N)$; then $x(S)$ is constant on the core, which is the affine seesaw visible in (6) — in the game of Theorem 22, $w(\{1,2\}) + w(\{3,4,5\}) = 1 + 3 = 4$. Only 76 of the 107 extremal games have any flat split, and 31 attain exactly $1/2$ with none, so the seesaw is a feature of the particular presentation in Theorem 22 and not an invariant of extremality. When a flat split does exist its sizes are always $\{2,3\}$, and the Lmax-starved player lies on the size-2 side in all 76 cases.

Two regimes. Theorem 18 gives only $1/3$ at $n = 5$ because the family $G(p,k)$ has a uniform h-block and the available splits $(k, p+1) \in \{(2,3),(3,2)\}$ are too coarse; the extremal five-player game has a non-uniform h-block. The crossover to the G-family shape happens at $n = 6$. The two regimes differ qualitatively and not just numerically: at $n = 5$ the extremal Lmax starves one player to 0 and has support of size 4, whereas in $G(p,k)$ the extremal Lmax is strictly positive everywhere. Explaining the crossover is open (Section 7).

5.4 One-coalition certificates and the regime change

Write, for any singleton-normalized exact game with the general dual bound $f(S) = T - v(N \setminus S)$,

$$\lambda := \min_{\emptyset \neq S} \frac{f(S)}{|S|}, \qquad \nu := \max_{\emptyset \neq S} \frac{v(S)}{|S|}, \qquad \mu := \min_i \mathrm{Lmin}(v)_i, \qquad M := \max_i \mathrm{Lmax}(v)_i.$$

The general duality gives the bracketing $\mu \leq \lambda$ and $M \geq \nu$ at once (for $\mu \leq \lambda$ the proof of Theorem 8 applies verbatim). On a two-generator game the right-hand side of (3) is exactly ν, since $v(P \cup Q) = \min(g(P), h(Q))$. Observation 11 says the second bracket can be strict from $n = 6$ on. Two facts sharpen this into a statement about the regime change at five players.

Proposition 23 (covering obstruction)**.** *Let $\tau_\nu := \{S : v(S) = \nu|S|\}$ be the ν-tight family. If τ_ν covers N and $\nu > T/n$, then no core point satisfies $x \leq \nu \mathbf{1}$, and hence $M > \nu$.*

Proof. Sum the core inequalities $x(S) \geq v(S) = \nu|S|$ over a cover with weights $w \geq 0$: with $c_i := \sum_{S \ni i} w_S \geq 1$ this gives $\sum_i c_i(x_i - \nu) \geq 0$. If $x \leq \nu \mathbf{1}$, every term is nonpositive, so $x_i = \nu$ for every i; then $T = \nu n$, contradicting $\nu > T/n$. □

Moreover no *balanced* collection can witness $M > \nu$: for balanced w, $\sum_S w_S v(S) \leq T$ because the core is nonempty, and the same computation then yields $\sum_S w_S\big(\nu|S| - v(S)\big) \geq \nu n - T$, which is what equality requires. So a failure of $M = \nu$ needs a strictly over-covering fractional cover: no balanced collection can witness one. A covering ν-tight family is one way this arises, and it is what occurs at the smallest counterexample of Observation 11: the four ν-tight sets $\{2,3\}, \{1,3,4\}, \{1,3,5\}, \{1,3,6\}$ cover N with per-player excess $(2,0,3,0,0,0)$, and player 2, lying only in $\{2,3\}$, blocks every balanced subfamily. At the five-player extremal game of Theorem 22, by contrast, $\nu = 1$ and the ν-tight family $\{3,5\}, \{1,3,5\}, \{3,4,5\}$ misses player 2, so Proposition 23 does not apply, and the dual minimum $h(N) = 1$ is attained at the balanced partition $\{2,4\} \mid \{1,3,5\}$. The covering test is sufficient for $M > \nu$ but not necessary: at six players, games with $M > \nu$ outnumber those with a covering ν-tight family.

Theorem 24 (one-coalition certificates at five players)**.** *For every singleton-normalized exact game on five players, $\mu = \lambda$ and $M = \nu$: both egalitarian levels are certified by a single coalition.*

Proof, by finite exact check. By the duality above, $M = \nu$ fails if and only if some fractional cover w has $\sum_S w_S h(S) < h(N)$, where $h(S) = \nu|S| - v(S) \geq 0$; likewise for μ with $f'(S) = f(S) - \lambda|S|$. Both objectives are nonnegative on the cover polyhedron, which is pointed with recession cone the nonnegative orthant, so each minimum is attained at a vertex; on five players there are 2287 vertices, each of support at most five. Fixing a vertex and a coalition attaining ν (resp. λ) makes every hypothesis, and the failure inequality, linear in v, and $v(N) = 1$ is free by homogeneity. Exactness is imposed through the 280 inequalities that the facet theorem of Studený and Kratochvíl (2022) assigns to the sixteen indecomposable min-semi-balanced systems on five players under all permutations; this description was validated in both directions against independent exactness tests. Each of the 1142 symmetry classes of (vertex, anchor) pairs, on each side, is infeasible in exact rational arithmetic. □

Two remarks calibrate the proof. First, deleting the 280 exact-cone inequalities makes 361 (resp. 339) of the over-covering classes feasible, so it is exactness — not the shape of the certificate — that forces the one-coalition property at five players. Second, the M-side is false from six players on (Observation 11), and the same reduction detects that failure. The theorem is therefore the formal content of the regime change of Section 5.3: at five players both selections are governed by

single bottleneck coalitions, while from six players the ceiling requires genuinely fractional covers — which is precisely where $s(n)$ begins its climb toward 2. One consequence localizes all the difficulty of $s(5)$.

Corollary 25 (floor from ceiling)**.** *For every singleton-normalized exact five-player game,* $\mu \geq \min_i \mathrm{Lmax}(v)_i$.

Proof. Lmax is a core point, so $f(S) = T - v(N \setminus S) \geq T - \mathrm{Lmax}(N \setminus S) = \mathrm{Lmax}(S)$ for every S; hence $\lambda \geq \min_S \mathrm{Lmax}(S)/|S| \geq \min_i \mathrm{Lmax}_i$, and $\mu = \lambda$ by Theorem 24. □

If Lmax has no zero coordinate, the leximin floor is therefore bounded below for free and elementary gap arithmetic bounds the ratio; every genuinely hard configuration for $s(5)$ has a *zero* in Lmax. The theorem does *not* by itself deliver the starvation bound needed for those configurations; that bound is a statement about a different functional and requires its own argument, which the next theorem supplies in the extremal one.

Theorem 26 (the extremal configuration is pinned)**.** *Let v be a singleton-normalized exact five-player game whose* Lmin *strata form the pattern* $(3, 2)$ *and whose* Lmax *strata form the pattern* $(1, 4)$ *with the singleton stratum at zero, at player j. Then*

$$f(\{j\}) = T - v(N \setminus \{j\}) \;\geq\; \tfrac{T}{8}.$$

Proof. Normalize $T = 1$ and suppose $f(\{j\}) < \frac{1}{8}$. The starved player j is Lmin-poor by Corollary 31 below (whose two-line proof uses only Lemma 30 and is independent of this theorem); write $P = \{j, p_1, p_2\}$ for the poor stratum and $R = \{r_1, r_2\}$ for the rich pair, with $x = \mathrm{Lmin}$ equal to μ on P and $\varrho = (1 - 3\mu)/2$ on R, and $z = \mathrm{Lmax}$ equal to 0 at j and $\frac{1}{4}$ elsewhere. Call S *f-tight* when $x(S) = f(S)$, equivalently when $N \setminus S$ is v-tight at x.

Step 0. $\lambda \leq f(\{j\}) < \frac{1}{8}$, and $\mu \leq \lambda$ holds for every game (the easy half of Theorem 24, i.e. Theorem 8's argument), so $\mu < \frac{1}{8}$; hence $\varrho > \frac{5}{16} > \frac{1}{4}$ and $c := 1 - 5\mu > 0$.

Step 1 (key lemma). Every f-tight S satisfies $|S \cap \{p_1, p_2\}| \leq |S \cap R|$. Indeed, $U = N \setminus S$ is v-tight, and $z \in \mathrm{Core}(v)$ gives, with $a = |U \cap P|$, $b = |U \cap R|$, $\varepsilon = [\, j \in U \,]$,

$$\mu a + \varrho b \;=\; x(U) = v(U) \;\leq\; z(U) \;=\; \tfrac{a+b-\varepsilon}{4}.$$

The claim is $b \leq a - \varepsilon$. A violation leaves four shapes, each contradictory: $a - \varepsilon = 0$ with $\varepsilon = 0$ gives $\varrho \leq \frac{1}{4}$; with $\varepsilon = 1$ it gives $\mu \leq b(\frac{1}{4} - \varrho) < 0$; $a - \varepsilon = 1$ with $\varepsilon = 1$ (so $b = 2$) gives $1 - \mu \leq \frac{3}{4}$, i.e. $\mu \geq \frac{1}{4}$; and $a - \varepsilon = 1$ with $\varepsilon = 0$, $b = 2$ — the coalition $U = \{p_i\} \cup R$ — gives $1 - 2\mu \leq \frac{3}{4}$, i.e. $\mu \geq \frac{1}{8}$, against the failure hypothesis. This last case is the only one that mentions $\frac{1}{8}$.

Step 2 (the pinning, dualized). That P is the Lmin bottom stratum means every core point $y \geq \mu\mathbf{1}$ has $y(P) = 3\mu$. Substituting $u = y - \mu\mathbf{1}$, this is $\max\{u(P) : u \geq 0,\ u(N) = c,\ u(S) \leq h(S)\} = 0$ with $h(S) := f(S) - \mu|S| \geq 0$. LP duality supplies α and $\beta \geq 0$ with $\alpha + B_k \geq [\, k \in P \,]$ for every player k, where $B_k = \sum_{S \ni k} \beta_S$, and $\alpha c + \sum_S \beta_S h(S) \leq 0$. Since $c > 0$ and $h \geq 0$, necessarily $\alpha \leq 0$; put $A = -\alpha \geq 0$, so $B_k \geq 1 + A$ on P and $B_k \geq A$ on R.

Step 3 (exact tightness). $x \in \mathrm{Core}(v)$ gives $h(S) \geq x(S) - \mu|S| = (\varrho - \mu)\,|S \cap R|$, so

$$\sum_S \beta_S h(S) \;\geq\; (\varrho - \mu)(B_{r_1} + B_{r_2}) \;\geq\; 2A(\varrho - \mu) \;=\; cA,$$

using $2(\varrho - \mu) = 1 - 5\mu = c$ *exactly*. Step 2 bounds the same sum by cA, so both inequalities are equalities: $B_{r_1} + B_{r_2} = 2A$, and every support set of β is f-tight.

Step 4 (counting). The key lemma now applies to every support set, so $B_{p_1} + B_{p_2} = \sum_S \beta_S |S \cap \{p_1, p_2\}| \leq \sum_S \beta_S |S \cap R| = 2A$, while Step 2 demands $B_{p_1} + B_{p_2} \geq 2 + 2A$. Contradiction. □

Remark 27 (sharpness, and the archived check)**.** *The constant is not inserted; it falls out. Exactly one line of the proof mentions* $\frac{1}{8}$ *— the coalition shape* $U = \{p_i\} \cup R$ *— and at* $\mu = T/8$ *that inequality becomes an equality and the key lemma fails at* $S = \{j, p_{i'}\}$*. This is realized: in the extremal game of Theorem 22 (which lies in the configuration with* $\mu = T/8$*) the* f*-tight sets* $\{1, 2\}$ *and* $\{2, 4\}$ *are precisely of that shape. The proof also consumes strictly less than the original finite check did: neither the hard direction of Theorem 24 nor the facet description appears — only core membership of the two selections, the floor bracketing* $\mu \le \lambda$*, Corollary 31, and the pinning. The original check (*10,974 *exact programs over the* 354 *pinning-dual vertices) remains archived as an independent confirmation.*

Corollary 28 (the pinned solution pair)**.** *In the configuration of Theorem 26, suppose in addition that the starved player is* Lmin*-poor and that some coalition attaining* ν *has three players. Then, up to relabelling,* $\mathrm{Lmin} = \left(\frac{T}{8}, \frac{T}{8}, \frac{T}{8}, \frac{5T}{16}, \frac{5T}{16}\right)$*,* $\mathrm{Lmax} = \left(0, \frac{T}{4}, \frac{T}{4}, \frac{T}{4}, \frac{T}{4}\right)$*, and* $\rho(v) = \frac{1}{2}$ *exactly.*

Proof. Normalize $T = 1$. The Lmax pattern gives $M = 1/4$ directly. Core membership of the Lmax vector gives, for $S \neq \{j\}$, $f(S) \ge \big(|S| - [\, j \in S \,]\big)/4 \ge |S|/8$, and Theorem 26 gives $f(\{j\}) \ge 1/8$, so $\lambda \ge 1/8$. By the M-side of Theorem 24, $\nu = M = 1/4$, so a three-player coalition S_0 attaining ν has $v(S_0) = 3/4$ and $f(N \setminus S_0) = 1/4$ on two players, so $\lambda \le 1/8$; hence $\lambda = 1/8$, and Theorem 24 gives $\mu = 1/8$. With the starved player Lmin-poor, $\mu = 1/8 \le 1/6$ gives $(1 - 3\mu)/2 \ge 1/4 = M$, so the coordinatewise gap telescopes to $\| \mathrm{Lmin} - \mathrm{Lmax} \|_1 = T - 4\mu = T/2$, and the two displayed allocations are the unique ones consistent with the strata, the floor and the totals. □

The first additional hypothesis is automatic: Corollary 31 below proves, in two lines, that the starved player is always Lmin-poor. The second — a three-player coalition attaining ν — holds in every configuration game in the combined harvests (53 of 53) but remains open in general, and is recorded in Section 7.

For the ratio *bound*, as opposed to the pinned pair, both hypotheses can in fact be dispensed with.

Theorem 29 (unconditional configuration bound)**.** *In the configuration of Theorem 26,* $\rho(v) \le \frac{1}{2}$.

Proof, by gap arithmetic. Normalize $T = 1$; the Lmax pattern gives $M = 1/4$ and the $(3, 2)$ pattern gives $\mathrm{Lmin} = (\mu, \mu, \mu, \varrho, \varrho)$ with $\varrho = (1 - 3\mu)/2$. The starved player is Lmin-poor by Corollary 31 below (a two-line consequence of Lemma 30, independent of this theorem), so the coordinatewise gap is $1 - 4\mu$ for $\mu \le 1/6$ and 2μ otherwise, and Theorem 26 gives $\mu \ge 1/8$, so both branches are at most $\frac{1}{2}$. (Historically the Lmin-rich branch, whose gap $3/2 - 6\mu$ would exceed $\frac{1}{2}$ for $\mu < 1/6$, was first closed by a finite check of the same design as Theorem 26 — the rich relabelling of the pinning polyhedron, 31×354 exact programs, all infeasible, with the drop-the-pinning control restoring feasibility. The corollary supersedes that check by showing the branch is vacuous; the check remains archived as an independent confirmation.) □

Two calibrations. First, the floor bound is not removable in favour of soft arguments: from Theorem 24 and core membership of both allocation vectors alone — everything that is game-free — the best obtainable bound on ρ at five players is exactly $\frac{3}{2}$, so every configuration needs a lower bound on the floor of the strength of Theorem 26. Second, the attempted refutation of the ν-triple hypothesis produced, once again, a certificate whose audited model is a genuine exact game *outside* the configuration (a $(1, 4)$-shaped Lmax with minimum $1/30$, ratio $5/12$) — the third certificate-not-realizable instance of this program, and the reason hypotheses are retired here only by finite checks, never by feasible certificates alone.

The configuration hypothesis is essential, not decorative: dropping the $(3,2)$ pinning makes the same programs satisfiable, and their models are genuine exact five-player games with $T-v(N\setminus\{i\}) = T/16$ for some player i — half the bound — but with Lmin = Lmax, hence no separation at all. The cases outside Corollary 28 — the Lmax stratum shapes $(1,2,2)$ and $(1,1,3)$, the $(4,1)$ Lmin boundary, and the Lmin-rich starved player — are closed in the next subsection, which settles the supremum in full.

5.5 The five-player supremum in full

Three localization lemmas come first; all are consequences of Theorem 24 together with core membership of the two allocation vectors, and none needs a computation.

Lemma 30 (anchor localization)**.** *Let v be a singleton-normalized exact five-player game, S_0 a coalition attaining ν and S_* a coalition attaining λ (both exist by Theorem 24). Then*

(a) $|S_0| \geq 2$, and $\mathrm{Lmax}_i = M$ for every $i \in S_0$; consequently the Lmax *top stratum contains at least two players;*

(b) $\mathrm{Lmin}_i = \mu$ for every $i \in S_*$, and $\mathrm{Lmax}(S_*) \leq \mu|S_*|$; consequently $\mu \geq \min_{i\in P} \mathrm{Lmax}_i$, where P is the Lmin *bottom stratum;*

(c) $2M + 3\mu \leq T$.

Proof. (a) If $S_0 = \{i\}$ then $M = \nu = v(\{i\}) = 0$, but $\mathrm{Lmax}(N) = T > 0$ forces $M \geq T/5$. For the second claim, $\mathrm{Lmax}(S_0) \geq v(S_0) = M|S_0|$ by core membership while $\mathrm{Lmax}_i \leq M$ throughout, so Lmax is constant M on S_0. (b) Core membership gives $\mathrm{Lmin}(S_*) \leq f(S_*) = \mu|S_*|$ while $\mathrm{Lmin}_i \geq \mu$ gives the reverse, so Lmin $= \mu$ on S_*; the same cap applied to Lmax gives $\mathrm{Lmax}(S_*) \leq \mu|S_*|$, whence $\mu \geq \mathrm{Lmax}(S_*)/|S_*| \geq \min_{i\in S_*} \mathrm{Lmax}_i \geq \min_{i\in P} \mathrm{Lmax}_i$. (c) By the floor side, $f(N\setminus S_0) \geq \mu(5-|S_0|)$, i.e. $M|S_0| + \mu(5-|S_0|) \leq T$; since $M \geq T/5 \geq \mu$ the left side is nondecreasing in $|S_0|$, and $|S_0| = 2$ is admissible by (a). □

Part (b) sharpens Corollary 25: the minimum ranges over the Lmin-*poor* players only. Combined with the trivial $\mu \leq T/5$ it closes, by two lines, the hypothesis that resisted the configuration analysis.

Corollary 31 (the starved player is Lmin-poor)**.** *In the configuration of Theorem 26 —* Lmax *strata $(1,4)$ with the singleton at zero, at player j — the player j lies in the* Lmin *bottom stratum.*

Proof. The stratum pattern forces $4M = T$. If $j \notin P$ then $\mathrm{Lmax}_i = M = T/4$ for every $i \in P$, so Lemma 30(b) gives $\mu \geq T/4$, contradicting $\mu \leq T/5$. □

In particular the Lmin-rich branch of Theorem 29's proof is vacuous, the first additional hypothesis of Corollary 28 is automatic, and the corresponding finite check (31×354 programs) is superseded by the two lines above.

The second ingredient is an exact certificate form for the two selections, replacing anchors and pinnings by linear identities.

Lemma 32 (staged Farkas certificates)**.** *Let v be a game with nonempty core, ϕ a core point with level sets $B_1 \prec \cdots \prec B_k$ in increasing order of payoff. Then $\phi = \mathrm{Lmin}(v)$ if and only if for every stage t and every $i \in B_t$ there exist multipliers $y_A \geq 0$ supported on proper ϕ-tight coalitions, $\eta \in \mathbb{R}$, $\rho_l \in \mathbb{R}$ for $l \in B_1 \cup \cdots \cup B_{t-1}$, $\sigma_c \geq 0$ for $c \in B_t$, and $\gamma_j \geq 0$ supported on the active coordinates $\phi_j = 0$, with, for every player j,*

$$-\sum_{A\ni j} y_A + \eta + [\,j \in B_{<t}\,]\rho_j - [\,j \in B_t\,]\sigma_j - [\,\phi_j = 0\,]\gamma_j \;=\; [\,j = i\,].$$

The mirror statement characterizes Lmax*, with the level sets taken in decreasing order, σ entering with a plus sign, the same negative active-coordinate term, and right-hand side $-[\,j = i\,]$.*

Proof sketch. Leximin is computed stagewise: with the earlier strata pinned, the level μ_t is the largest uniform lower bound on the unpinned players, and B_t collects the players whose maximum subject to that equals μ_t. $\phi = \mathrm{Lmin}(v)$ is equivalent to the family of stage conditions $\max\{x_i : x \in \mathrm{Core}(v),\ x_l = \phi_l\ (l \in B_{<t}),\ x_c \geq \mu_t\ (c \in B_{\geq t})\} = \phi_i$ for $i \in B_t$. Each is a linear program whose feasible set contains ϕ; the constraints active at ϕ are efficiency, the ϕ-tight coalitions, the pinned equalities, and the level bounds on B_t itself, together with $x_j \geq 0$ whenever $\phi_j = 0$. Dualizing over the tangent cone at ϕ and computing the polar cone gives exactly the displayed identity; complementary slackness confines the support of y to tight coalitions and eliminates the bilinear terms $y_A v(A)$, which is what makes the certificate linear in (v, ϕ, y) jointly. The converse direction is the same computation read backwards. The certificate was derived independently twice within the program (once per player and stage as displayed, once per payoff cut) and the two forms verified against each other in both directions on 30 games and 6 non-optimal core points; the mirror was re-derived from scratch and validated on 1050 games (5,250 ceiling certificates, 10,500 counting the floor form, no failures), with negative controls rejecting a false Lmax claim and the core midpoint 261/261 times each. □

With these in place the supremum closes. Recall that a *configuration* is a triple: the Lmin stratum composition, the Lmax stratum composition, and the table of overlaps between the two stratifications. At $n = 5$ there are 2961.

Theorem 33 (the $(4, 1)$ regime)**.** *Every singleton-normalized exact five-player game whose* Lmin *strata form the pattern* $(4, 1)$ *satisfies* $\rho(v) \leq \frac{1}{3}$*, and the bound is attained.*

Theorem 34 (five-player supremum)**.** *Every singleton-normalized exact five-player game satisfies* $\rho(v) \leq \frac{1}{2}$*. With Theorem 22,* $s(5) = \frac{1}{2}$*.*

Proof of both, by configuration-complete finite check. Relabelling players reduces every game to the canonical layout of its configuration, so it suffices to exclude $\rho > \frac{1}{2}$ (resp. $\rho > \frac{1}{3}$ on the $(4, 1)$ class) in each configuration. This was done by two pipelines whose Lmin-optimality components are independently formulated; no broader machinery-independence claim is made. Coverage is machine-audited: every one of the 2961 configurations lands in exactly one bucket, no run errored, and none was left feasible.

The first pipeline is structural: 1953 configurations are impossible by Lemma 30(a) alone (Lmax composition ending in a singleton); 800 are excluded by one game-free exact program each — levels, orderings, totals and the anchor conditions of Lemma 30 against $\rho > \frac{1}{2}$; 100 fall to a finite check with v free, both core memberships, the certificates of Theorem 24, the floor and ceiling stage-one pinnings (dualized as in Theorem 26), and the exact-cone facet inequalities as lazy cuts; and the remaining 108 fall to the same check with the complete multi-stage pinning ladders — every Lmin stratum pinned given the strata below it, every Lmax stratum pinned given the strata above it, the reduced-game caps encoded exactly as minima over preimages. All programs are infeasible in exact rational arithmetic with no indeterminate verdicts.

The second pipeline re-proves the same verdicts from a minimal hypothesis basis: singleton normalization and scale; core membership of both vectors; exactness imposed *existentially* (for each proper coalition S, an attaining core point w^S with $w^S(S) = v(S)$ — the definition of exactness, linear, with no appeal to the facet description); and Lmin/Lmax imposed through Lemma 32 and its mirror. All 2961 programs are again infeasible — so the theorems depend neither on Theorem 24 nor on the facet catalogue. The two pipelines agree configuration by configuration.

Soundness controls, both pipelines: dropping the pinning ladders (resp. certificates) restores feasibility wherever the configuration is possible at all; the hypothesis sets hold, in exact arithmetic, on 1050 profiled random exact games spanning 21 configurations — the first pipeline's 2874 stage certificates in one run, the second pipeline's 10,500 Farkas certificates in another, no failures in either; and on all eight configurations for which a game above some target ratio is known, the same programs run at that target return feasible models that *audit end-to-end* as genuine exact games in the assumed configuration — including the extremal configuration at ratio exactly $\frac{1}{2}$ and the $(4, 1)$ regime at exactly $\frac{1}{3}$. The configurations declared possible include every configuration any harvest has produced. Attainment in Theorem 33 is by explicit games (111 harvested at exactly $\frac{1}{3}$). All drivers, logs and the coverage audit are archived with the verifiers. □

The two-regime picture is therefore a theorem: Lmin strata $(4, 1)$ force $\rho \leq \frac{1}{3}$ while $(3, 2)$ allows $\frac{1}{2}$, both attained. One calibration deserves record: under the single-stage hypothesis set the final 108 configurations are all *satisfiable*, and every one of the 108 models audits as a genuine exact game — in the wrong configuration. The complete ladders are what make the encoding decisive rather than merely stronger; this is the fourth, and largest, instance of the certificate-not-realizable phenomenon in this program, and the reason every feasible model here is audited before it is believed.

6 Obstructions to an upper bound

6.1 The exact six-player bound

Conceptually, this is a theorem about a core polytope inside the normalized simplex. Singleton exactness makes that polytope touch every coordinate facet, through six witness allocations. If two canonical Lorenz selections were farther than 2/3, their coordinatewise comparison would have one of fifteen sign patterns. In each pattern, every infinitesimal exchange that improves one selection must be stopped by a tight coalition on which the exchange has negative sum; the opposite selection and the six facet witnesses impose seven necessary inequalities on that blocker. The finite part proves that no such blockers can coexist. Thus the computation is a fifteen-case convex-geometric exhaustion of necessary exchange obstructions, not an enumeration of games or core vertices.

The first nontrivial fixed dimension can now be closed without a generator normal form or a support-table census.

Proof architecture. There are three dependencies and only three. A human reduction maps any strict counterexample to one of fifteen sign-labelled systems of necessary exchange blockers. A finite linear-arithmetic lemma proves those systems infeasible, with independently checked proof objects. Finally, the explicit game $G(2, 3)$ supplies the matching lower bound. The earlier generator census, equality census, continuity arguments, and support tables are historical discovery routes and are not used in this proof.

After orbit reduction, the nineteen-schema formulation of the three large sign cases uses only the following hand-checkable exchange calculus. The original minimized labelled menus contain additional direction shapes; they are proof-checked in the fifteen-case formulation but are not being classified by this lemma. Naming the reduced templates separately isolates what is conceptual from the later blocker-coexistence certificates.

Lemma 35 (Three exchange templates)**.** *Let $y \in \mathbb{R}^6$ and let $h(N) = 0$.*

(i) For $h = -e_a + e_r$, h is Lmin-admissible exactly when $y_r < y_a$; it has 16 negative nonempty proper cuts. The Lmax condition is the same inequality $y_a > y_r$, read at the highest affected level.

(ii) For $h = -e_a - e_b + e_r + e_s$, h is Lmin-admissible exactly when $\min(y_r, y_s) < \min(y_a, y_b)$; it has 20 negative nonempty proper cuts. For Lmax, the condition is $\max(y_a, y_b) > \max(y_r, y_s)$.

(iii) For $h = e_p + 2e_{p'} - e_{m_1} - e_{m_2} - e_{m_3}$, h is Lmax-admissible exactly when $\max_j y_{m_j} > \max(y_p, y_{p'})$; it has 24 negative nonempty proper cuts. Negating h gives the Lmin dual.

Ties make the displayed strict activation fail; they introduce no additional case.

Proof. Admissibility asks whether the first affected payoff stratum contains a coefficient of the improving sign and none of the opposite sign. Reading the lowest affected stratum for Lmin and the highest for Lmax gives the stated minimum and maximum tests. For (i), negativity means $a \in S, r \notin S$, with the other four players free. For (ii), if d donors and u recipients belong to S, negativity is $d > u$; summing the cases $(d, u) = (1, 0), (2, 0), (2, 1)$ gives $8 + 4 + 8 = 20$. In (iii), negativity means $|S \cap \{m_1, m_2, m_3\}| > \mathbf{1}_{p \in S} + 2\mathbf{1}_{p' \in S}$; the unused sixth player is optional, and the three positive-side loads contribute $14 + 8 + 2 = 24$ cuts. Negation proves the dual statements. □

Theorem 36 (Six-player bound)**.** *Every singleton-normalized exact six-player game satisfies*

$$\| \operatorname{Lmin}(v) - \operatorname{Lmax}(v) \|_1 \leq \frac{2}{3} v(N).$$

The bound is attained by $G(2, 3)$, and hence $s(6) = 2/3$.

Proof. Scale $T = v(N)$ to one and write $x = \operatorname{Lmin}(v)$, $z = \operatorname{Lmax}(v)$. Put

$$P = \{i : x_i > z_i\}, \qquad M = \{i : z_i > x_i\}, \qquad E = \{i : x_i = z_i\}.$$

If $x \neq z$, both P and M are nonempty. Up to a common relabelling there are therefore fifteen triples (p, m, e) with $p, m \geq 1$ and $p + m + e = 6$, and

$$\|x - z\|_1 = 2 \sum_{i \in P} (x_i - z_i). \tag{28}$$

Exactness at each singleton supplies a core point q^i with $q^i_i = 0$. For a point y and direction h, define the first order-crossing time

$$\Delta(y, h) = \min \left\{ \frac{y_j - y_i}{h_i - h_j} : y_i < y_j,\ h_i > h_j \right\},$$

with the empty minimum equal to $+\infty$. Every $0 < \varepsilon < \Delta(y, h)$ preserves all strict orders between distinct old payoff strata.

Three elementary selection anchors will be used. First, z cannot have a unique top player i: along the core segment toward q^i, the direction $q^i - z$ is negative at that unique top because $q^i_i = 0 < z_i$. Taking $0 < \varepsilon < \min\{1, \Delta(z, q^i - z)\}$ lowers the first entry of the decreasing rearrangement. Second, some z-top player has $x_i \geq z_i$, since otherwise the direction $x - z$ is negative throughout the old top stratum and the same explicit bound lowers the first changed decreasing-lex entry. Dually, some x-bottom player has $z_i \leq x_i$, using the direction $z - x$. Convexity keeps all three segments in the core; the crossing bound handles tied strata without an informal limiting argument.

Now let $h(N) = 0$. On the Lmin side call h admissible if its lowest affected payoff stratum has a positive coefficient and no negative coefficient. On the Lmax side call it admissible if its

highest affected stratum has a negative coefficient and no positive coefficient. Such a direction strictly improves the relevant sorted lexicographic vector for every positive step below $\Delta(y,h)$. If all proper tight coalitions had nonnegative derivative, set

$$\varepsilon_* = \min\left(\{1, \Delta(y,h)\} \cup \left\{\frac{y(S)-v(S)}{-h(S)} : \varnothing \neq S \subsetneq N,\ y(S) > v(S),\ h(S) < 0\right\}\right).$$

Every finite member is positive and the set contains 1, so $\varepsilon_* > 0$. Any $0 < \varepsilon < \varepsilon_*$ preserves both the payoff-stratum order and every core inequality, while $h(N) = 0$ preserves efficiency. This would give a core improvement, a contradiction. Optimality therefore forces a proper tight coalition S with $h(S) < 0$. Core membership of the opposite selection and the six witnesses gives, respectively,

$$z(S) \geq x(S), \quad q^i(S) \geq x(S)\ (i \in N), \tag{29}$$

or

$$x(S) \geq z(S), \quad q^i(S) \geq z(S)\ (i \in N). \tag{30}$$

As a finite differential control, the pivot implementation of admissibility was compared with a separately written first-affected-stratum criterion on all 4,683 labelled weak payoff orders, all 1,610 bound-two directions, and both selections: 15,079,260 comparisons with no mismatch.

For each sign triple, impose nonnegativity and efficiency of x, z, q^i, the strict signs, (28) with left side $> 2/3$, the three anchors, and the necessary blocker disjunction (29) or (30) for a finite menu of admissible integer directions. Thirteen cases use only efficiency-preserving directions with coefficients in $\{-1, 0, 1\}$; $(2, 3, 1)$ and $(3, 3, 0)$ use coefficients in $[-2, 2]$. After deletion of redundant direction clauses the fifteen menu sizes are

$$1, 1, 1, 1, 0, 2, 5, 14, 8, 2, 52, 71, 2, 90, 2, \tag{31}$$

in lexicographic order of (p, m, e). Every resulting exact QF_LRA formula is infeasible. Z3 and cvc5 agree on every formula; cvc5 emits CPC with proof pedantry enabled and trust disabled, and Ethos checks all fifteen proofs as `correct`. The formulas are relaxations: they retain only necessary consequences of tightness, so their infeasibility excludes every exact-game counterexample. The source-to-formula regeneration, hashes, proof controls, and all retained directions are in the repository under `verifiers/s6/strict_exchange_*`.

The dependency structure is simpler than the discovery history. First, the game-theoretic lemmas above map any genuine counterexample into the finite relaxation. Second, the retained exchange menus give an exhaustive relay of admissible directions and necessary blockers. Third, exact rational identities close every arithmetic leaf. Solver searches were used to discover and minimize these objects but are not evidence for the last layer. Nine of the eleven localized packet rows have independently specified hand-checkable relays; the two largest have complete solver-free arithmetic conditional on frozen relay/use maps. External validation of the first layer and independent reconstruction of those two maps remain distinct from rerunning the solvers.

The two bound-two cases were additionally reconstructed by a separate encoder that imports none of the production encoder, lex predicate, blocker builder, direction generator, minimizer, or SMT serializer. Its base systems are satisfiable, its full systems are infeasible in both solvers, and its fresh trust-disabled CPC proofs check in Ethos. A third formulation imports neither encoder nor Z3: a standard-library generator writes SMT-LIB directly and expresses activation as strict separation of the positive and negative payoff supports. All 458 assertion comparisons with both prior formulations are equivalent; both third-formulation bases are satisfiable and both full systems are infeasible. These controls reduce correlated encoding risk but do not replace

external review of the human implication. At the orbit level, the three largest menus admit a further unguarded reduction: their 213 labeled directions and 44 sign-role orbits shrink to 136 directions in 19 orbits. By Lemma 35, the $(3, 2, 1)$ and $(4, 2, 0)$ cases then use only one-for-one and two-for-two exchanges; the balanced $(3, 3, 0)$ case adds one asymmetric coefficient-two shape and its reverse. Their activation side is completely described by the same 19 role-level order-statistic tests: compare the lowest recipient with the lowest donor for Lmin, or the highest donor with the highest recipient for Lmax. Their 2,512 eligible negative cuts use only three unoriented templates, with respectively 16, 20, and 24 cuts for one-for-one, two-for-two, and asymmetric coefficient-two exchange. The remaining finite content is therefore blocker coexistence, not lexicographic activation or cut enumeration. All three reduced formulas have separate trust-disabled CPC proofs checked by Ethos, and deleting any selected orbit makes the corresponding relaxation feasible. These controls reduce all three large cases using only the two cross-selection anchors; the no-unique-Lmax-top anchor is dispensable there. In $(4, 2, 0)$ they further reduce each blocker alternative to the opposite-allocation comparison and the four singleton-witness comparisons indexed by P; both M-indexed comparison families are dispensable, while each of the five retained sources is irredundant in the relaxation. Anchor localization is sharper: if the x-bottom and z-top anchors $a, b \in P$ coincided, then $x_m \geq x_a > z_a \geq z_m > x_m$ for $m \in M$. After relabelling the distinct anchors, a separate 18-real certificate uses nineteen directions in ten residual orbits and only the one exactness witness q^b. These controls reduce implementation correlation and make the finite anatomy readable. In that localized case two exchanges are automatic. Their negative blockers form a pair $\{b, 4\} \cup A, \{b, 5\} \cup B$ with $A, B \subseteq P \setminus \{b\}$. Of the 64 labelled pairs, 32 are base-infeasible and the other 32 form eleven classes under the residual symmetry; every class closes with a subset-minimal menu of at most seven remaining exchanges. Adding the two blocker inequalities also gives a short aggregate packet inequality by hand. This eleven-row decomposition has nine hand-checkable rows. One is a two-exchange blocker relay whose four terminal blocker pairs force the distance back to at most $1/3$. In the second packet orbit the single exchange $-e_1 - e_3 + e_4 + e_5$ is forced to improve Lmin, while every one of its twenty eligible blockers violates the opposite-allocation comparison. The third is a depth-one three-leaf star: the root activation, exact three-blocker list, and all child activations have structural inequalities. Its terminal Lmax leaf has sixteen possible complement cuts and its two terminal Lmin leaves reduce to four base cuts. Twenty explicit nonnegative weighted identities, checked by pure rational coefficient arithmetic without a solver, show that every such cut forces the distance to at most $1/3$. The fourth is a unique three-exchange chain. Its first twenty cuts reduce structurally to seven; six source-specific identities leave the unique first blocker, fifteen leave the unique second blocker, and sixteen exclude every terminal cut. The fifth is the depth-three class: its apparent large tree reduces to eight exact difference-cone cut censuses, sixteen source-specific blocker filters, and nine activation identities. Exact Fourier–Motzkin elimination independently checks the cut table, and pure rational arithmetic checks all 25 identities. The sixth class closes after the exhaustive split $x_3 \leq x_4$ versus $x_3 > x_4$. Its five-leaf relay reduces to ten exact difference-cone cut censuses and 71 certificate uses. Deduplication leaves 43 blocker identities, seven activation identities, and two terminal-anchor identities; the terminal leaves share the threshold contradiction $x_1 > 1/3$ versus $x_1 \leq 1/3$. A solver-free exact-rational checker reconstructs every premise at every use site and rejects coefficient, use-site, candidate, use-map, and strict-target mutations. The seventh class closes after the same root order split and one nested comparison. Its nine leaves reduce to eighteen exact cut censuses and 108 identities checked at 122 uses. All terminal leaves force the distance directly to at most $1/3$; 103 of its 105 blocker exclusions use only the single witness q^1, while two weak-branch exceptions use only the opposite selection. These first seven hand-checkable rows cover twenty-four labelled packets. The eighth class is the singleton packet orbit. After the universal root split, its fifteen-leaf relay reduces to twenty-one exact difference-cone censuses and 85

blocker exclusions. The disjunctive two-donor activation is equivalent to comparing the minimum donor level with the minimum recipient level, reducing its failure expansion to four conjunctions and 75 cases over the tree. Deduplication leaves 65 blocker and 44 activation identities checked at all 160 uses by a solver-free rational checker. The ninth class has an asymmetric packet sandwich and an eleven-leaf relay. The same minimum-level lemma and its maximum-level dual reduce its activation side to 72 cases; 21 exact cut censuses and 108 identities close all 132 uses. Thus the nine hand-checkable rows cover twenty-seven labelled packets. The largest remaining row has a symmetric packet satisfying $d_0 + d_1 \le \delta_4, \delta_5 \le d_2 + d_3$; the anchor chains $z_0 < x_0 \le x_4 < z_4 \le z_1$ and $x_0 \le x_4 < z_4 \le z_1 < x_1$ make both one-for-one exchanges automatic. For the Lmin exchange $e_0 - e_1$, the same sandwich reduces all sixteen eligible negative cuts to ten possible opposite blockers, uniformly at all ten occurrences in the frozen tree. Exact coefficient arithmetic checks the sandwich and cut partition without a solver or production import. Three orbit identities further exclude seven of the thirty formal pairs between the three possible preceding blockers and the ten packet candidates. The root split is hand-classified as well: its weak and strict twenty-cut censuses reduce to three and four blockers, respectively, using one seven-term witness identity beyond signs. Across all 90 direction nodes, 1,736 eligible cuts split exactly into 1,464 pure difference-cone exclusions, 34 packet/ancestor-context exclusions, 150 target-witness exclusions, and 88 relay branches. The 184 exclusions have sparse incidence over-covers and reduce to 107 exact rational identities. The 285 elementary activation failures split into 195 exact allocation contradictions and 90 unit-weight witness over-covers; 72 exact rational identities settle all of them. Solver-free checkers verify every identity and the exhaustive partitions. Thus no arithmetic leaf in the 51-terminal tree remains machine-only, conditional on its frozen relay/use maps. Its 24 labelled cut menus collapse to fifteen rich-player-parametric templates, and a path-local grammar generates all 88 edges and 51 leaves. Forty-one path-free activation identities cover 253 uses; only 32 uses require an ancestor comparison, in three relay families. The rich-player swap reduces the 72 identities to 47 affine motifs, twenty of them path-dependent. This is still not a tenth independently derived hand proof: the grammar and the 32 residual placements require external reconstruction. The other remaining row has an asymmetric packet sandwich $d_0 + d_1 \le \delta_4 \le d_3$ and $d_0 + d_1 + d_2 \le \delta_5 \le d_2 + d_3$; it reduces both twenty-cut roots to two blockers using signs alone. Across its 49 direction nodes there are 924 eligible negative-cut uses. Exact Fourier–Motzkin elimination removes 711 in the pure difference cone and leaves 213 candidates. The ten distinct candidate sets form seven direction menus plus one binary refinement: below the weak-root h_{11} blocker 124, the h_7, h_{10}, h_{15} menus each gain one explicit cut. Seven path-local transition rules then generate all 47 relay edges: they depend only on the weak/strict regime, the preceding h_{10} cut, and the same $h_{11} : 124$ bit. Eleven candidates are excluded by the already active packet/ancestor witness context, 155 by the target singleton witness, and 47 are exactly the relay branches. The eleven intermediate exclusions have nine solver-free rational identities. All 155 target-witness exclusions have incidence over-covers of support at most five; 150 use unit coefficients and five require coefficient two. This is a complete human explanation of the witness-side upper bounds, but the allocation-side strict reductions have also been exactified: 53 rational identities, checked solver-free at all 155 uses, show that the corresponding lower sums strictly exceed their incidence bounds. Of the uses, 118 need no path comparison, 33 use ancestors only, and four also use the target comparison. An independent exact-arithmetic checker verifies that the $11 + 155 + 47$ uses exhaust all 213 candidates, so every terminal no-blocker closure is arithmetic rather than solver-only. The coverage maps remain relay-derived. The same incidence geometry controls activation: anchor/order chains remove 84 of 124 elementary failure cases, four more are allocation-infeasible, and each of the other 36 has a packet/ancestor over-cover of support at most four. The activation incidence bounds are human; 26 solver-free rational identities prove the 36 allocation lower bounds, and the four direct cases share one more identity. Thus no class-6

arithmetic leaf remains machine-only, conditional on the frozen relay/use maps. Of the full 40-use activation map, 19 uses are path-free and the 21 path-dependent uses split into three relay families; the 27 identities reduce to 25 affine motifs after source-depth aliases are forgotten. Those maps remain encoding-shared and require external human-to-machine review. Both remaining rows also close after imposing the same root split $x_3 \le x_4$ versus $x_3 > x_4$. Indeed, that single split gives complete forced-exchange relay trees for all eleven rows; class 10 alone needs the nested x_1–x_2 split. The last two trees still rely on the parent encoding rather than separate exact hand certificates. This does not replace the human obligation to verify the counterexample-to-formula implication.

The single-witness filters have a common exact geometry. For any fixed coalition menu, the map $q \mapsto (q(C))_C$ sends $\{q \ge 0 : q(N) = 1, q_b = 0\}$ onto the convex hull of the non-anchor incidence columns. Hence a nonnegative weighted sum of witness comparisons is bounded exactly by its maximum non-anchor incidence load, attained at a unit vector. The six witnesses are otherwise independent, so after choosing one blocker for each active exchange this projection eliminates all witness coordinates exactly; their six anchor-specific support bounds collapse to the second-largest player incidence load. This support system has an exact finite reduction. For anchor i and a proposed maximizing non-anchor player k, partition the nonnegative weight orthant by the cones

$$K_{i,k} = \{\lambda \ge 0 : \lambda \cdot (a^k - a^\ell) \ge 0 \text{ for every } \ell \notin \{i,k\}\}.$$

On $K_{i,k}$ the support bound is linear, so its extreme rays give a complete rational inequality system. A ray with s positive coordinates requires $s-1$ active player-comparison hyperplanes; because there are only $n-2$, $s \le n-1$. Thus every six-player witness inequality uses at most five blocker requirements, and a five-coalition incidence circuit shows this bound is sharp. This constructive ray statement strengthens the corresponding pointwise Helly bound and gives an exact facet-enumeration algorithm. The remaining system is a finite cut-transversal problem coupled to this symmetric incidence support function and the opposite-selection signs. The class-5 crossing sums are sparse incidence over-covers of this form. This standard support-function identity supplies an LP/facet algorithm. Applied exhaustively to class 6, it gives the 155 target-witness supports and eleven context supports just described; these couple to 53 and nine exact allocation identities, respectively. The exact cut census identifies the remaining 47 candidates with the relay branches and settles every terminal no-blocker closure. This arithmetic exhaustion does not independently validate the relay/use maps or the counterexample-to-formula implication.

Thus the upper bound holds. The family formula gives $\rho(G(2,3)) = 2(2(3-1)-1)/(3(2+1)) = 2/3$, proving equality. □

6.2 No constant below 2

Theorem 18 already rules out one shape of answer. Any bound of the form $\|\operatorname{Lmin}(v) - \operatorname{Lmax}(v)\|_1 \le c \cdot v(N)$ with a constant $c < 2$, valid for all n, is false: choose n even with $2 - 8/n > c$. So the only possible clean upper bound is dimension-dependent, and the right target is the finite-n statement

$$s_2(n) = 2 - \frac{8n}{n^2 - (n \bmod 2)} \qquad (n \ge 6),$$

now proved for the full exact class at $n = 6$ by Theorem 36, and confirmed by complete two-generator enumeration at $n = 7, 8$, together with the question of whether $s(n) = s_2(n)$ in larger dimensions.

6.3 No bound linear in the leximin floor

A second shape is also excluded, and this one was a live candidate. The inequality

$$\|\operatorname{Lmin}(v) - \operatorname{Lmax}(v)\|_1 \;\le\; 4(n-4)\,\mu(v), \qquad \mu(v) = \min_i \operatorname{Lmin}(v)_i,$$

holds with no violations in 900 fresh random five-player games and is tight both on the $n = 5$ extremal class and on $G(3,4)$ at $n = 8$, where both sides equal 12. It fails from $n = 9$ on:

game	n	$\|\operatorname{Lmin} - \operatorname{Lmax}\|_1$	$4(n-4)\mu$
$G(4,4)$	9	88/5	16
$G(3,5)$	9	33/2	15
$G(4,5)$	10	24	96/5
$G(5,5)$	11	95/3	70/3

The reason is structural rather than accidental: in $G(p,k)$ one has $\mu = p/(p+1) < 1$, bounded, while the gap grows like $2T$. No bound of the form $C(n)\mu(v)$ can track the divergence, whatever C is.

6.4 Non-monotonicity under majorization

The most natural route to an upper bound on the two-generator class is to argue that a uniform h-block is without loss of generality — the extremal family $G(p,k)$ has one — and then to optimize over g alone using the floor formula of Conjecture 9. That route is closed.

Define a *smoothing step* as a Robin Hood transfer of one unit from a larger to a smaller h-coordinate, so that h moves down in the majorization order, toward uniform. The premise that a smoothing step never decreases ρ is false, and not only for small blocks:

$\lvert B\rvert$	2	3	4	5
smoothing steps tested	40	76	84	84
violations	18	27	11	21
rate	45%	36%	13%	25%

In particular it fails for $|B| \ge 4$, so the block-size threshold one might guess from the $n = 5$ anomaly does not exist. An explicit verified counterexample at $|B| = 4$, with $n = 7$ and $T = 10$, both allocations confirmed in the core:

$$g = (0,5,0,3,2,0,0), \qquad h = (1,0,2,0,0,5,2),$$

$$\operatorname{Lmin} = (1,2,1,1,1,3,1), \qquad \operatorname{Lmax} = \left(\tfrac{5}{6}, \tfrac{5}{2}, \tfrac{5}{6}, \tfrac{4}{3}, \tfrac{7}{6}, \tfrac{5}{2}, \tfrac{5}{6}\right), \qquad \rho = \tfrac{1}{5}.$$

Smoothing the h-block $(5,2) \mapsto (4,3)$, i.e. replacing h by $(1,0,2,0,0,4,3)$, gives

$$\operatorname{Lmin} = \operatorname{Lmax} = (1,2,1,1,1,2,2), \qquad \rho = 0.$$

So a single Robin Hood step inside one generator collapses the ratio from 1/5 to 0, and ρ is not monotone under majorization of h.

The two facts sit together uncomfortably and both are true: smoothing makes a game more symmetric and tends to *collapse* the two selections onto each other, yet the extremal family has a perfectly uniform h. The dependence is therefore genuinely non-monotone, and any upper-bound proof must handle general h directly rather than reducing to the uniform case.

7 What the results do and do not establish

Proved. Lemma 1 (lower-envelope generation and its converse); Lemma 3 (ceiling and overlap identity); Lemma 5 (support disjointness); Theorem 7 (two-generator core duality, with identity (1) valid without disjointness); Theorem 8 ($\mu \le \lambda$); Proposition 10 (the cover inequality (2) for balanced covers); the validity half of Observation 11; Theorem 14 with Lemma 15 (closed forms for $G(p,k)$, including membership for all $p \ge 1$, $k \ge 2$ and the exact list of tight coalitions); Corollary 16 and Remark 17 (the deficit identity, twice); Theorem 18 ($s(n) \ge 2 - 8n/(n^2 - (n \bmod 2))$ and $\sup_n s(n) = 2$); Lemma 21 (core closing); Proposition 23 (the covering obstruction); Theorem 22 ($\rho = 1/2$ at $n = 5$, by hand). The failure of the $4(n-4)\mu$ bound from $n = 9$, and the non-monotonicity of ρ under majorization of h, are proved by the explicit games displayed in Section 6: each is a finite exact-rational computation on a named game, and each is a refutation, so a single instance suffices.

Proved by exhaustive finite check. Theorem 24 ($\mu = \lambda$ and $M = \nu$ for every exact five-player game; 1142 symmetry classes, each infeasible in exact rational arithmetic), whose M-side Corollary 28 consumes; and Theorems 34 and 33 ($s(5) = 1/2$ and the $(4,1)$ bound $1/3$; 2961 configurations, two pipelines with independently formulated Lmin-optimality components, machine-audited coverage); and Theorem 36 ($s(6) = 2/3$; fifteen table-free sign-count relaxations, independently checked exact CPC proofs). Theorem 26 is now a *hand* result (its one-page proof replaced the original 10,974-program check, which remains archived as confirmation); Lemma 30 and Corollary 31 are hand results; Lemma 32 is a hand result derived independently twice within the program. The checks are proof dependencies; the corollary's remaining additional hypothesis (the ν-attaining triple) is open in general and noted under (4) below.

Enumerated, not proved. The entries of Table 2 are maxima over the stated integer ranges, hence complete only for rational shapes whose primitive integer representative has total in that range; they are lower bounds for $s_2(n)$, not its value. The strict failure of equality in (3) is established by the displayed six-player game, but the claim that this is the *smallest* such failure rests on the enumeration. The tightness rates for (3) and for Conjecture 9 are sample statistics over 586 random games. The structural findings about the $n = 5$ extremal class in Section 5.3 — unique solution pair up to permutation, $(3,2)$ and $(1,4)$ stratum compositions, $|\operatorname{supp}(\mathrm{Lmax})| = 4$, the refutation of the seesaw hypothesis — are computational, over 360 profiled games; of these only the two-regime split is now proved (Theorem 33), the rest remaining sample statistics. The local-optimality certificates from seeded c-space climbing at $n = 5, 6, 7$ are numerical searches over the full exact class; the unseeded variant is under-powered at $n \ge 6$ and certifies nothing.

Conjectured. Four problems remain open.

(1) **Conjecture 9**, equality in the leximin floor formula. Reduced by linear-programming duality to the fractional-cover inequality (2); proved for balanced covers (Proposition 10); the over-covering case is open in general. At five players Theorem 24 *proves* the equality, for every exact game and not only two-generator ones. Verified in 586/586 random two-generator games at $n = 4, \ldots, 7$.

(2) **An upper bound for $n \ge 7$.** Theorem 18 forces any true bound to be asymptotically 2, so every bound $c \cdot v(N)$ with $c < 2$ is false for large n. The right target is $s_2(n) = 2 - 8n/(n^2 - (n \bmod 2))$ for $n \ge 7$, confirmed by complete enumeration at $n = 7, 8$ only. The six-player case is now proved: $s(6) = s_2(6) = 2/3$ (Theorem 36). Section 6 closes

the majorization route and the $C(n)\mu$ route. $n = 5$ is settled (Theorem 34); larger fixed dimensions remain open.

(3) **Is the two-generator class extremal?** Adding a generator lowers the envelope and grows the core, so this is not automatic. All evidence is negative for a counterexample: zero improvements from about 200 third generators at each of $n = 5, 6, 7, 8$, and no improvement from seeded c-space climbing over the full exact class at $n = 5, 6, 7$. Open globally; settled affirmatively at $n = 5, 6$, where the two-generator lower bounds equal $s(5)$ and $s(6)$.

(4) **A structural explanation of the $n = 5$ versus $n \geq 6$ regime change.** The question now has a sharp two-part form, replacing the earlier "why five". *Part one, answered*: the starved family generalizing the extremal configuration (four players at $T/4$, a poor stratum of $n - 2$ absorbing the zeros) exists at *every* $n \geq 5$ with ratio exactly $(n - 4)/(n - 3)$ — now a theorem, by a six-row closed-form construction whose hand proof rests on two coalition values (the identity $T - (n - 3)\mu = 3T/4 = 3M$ lets one cap pin the poorer side from below and fix the ceiling from above, uniformly in n; archived with machine guards at $n = 5..9$). At $n = 5$ the member realizes the extremal configuration and solution pair of Theorem 22 at ratio $\frac{1}{2}$ — though as a *second, non-isomorphic* game (no relabelling aligns the two; the extremal configuration is attained by at least two essentially different games) — so the family is tight there; its value *exceeds* the two-generator bound at $n = 5$, *ties* it at $n = 6$, and undershoots thereafter (it saturates at 1 while s_2 climbs to 2, so it is dominated for large n and says nothing about the supremum beyond $n = 5$); the regime change is the handover between champion families. The crossover test is now resolved by Theorem 36: both families attain the common value $s(6) = 2/3$. *Part two, open*: the floor side of Theorem 24 reduces (archived, with hand lemmas) to a single inequality — for every ordered pair (p, q) of players, some core point satisfies $2y_p + y_q \geq 3\lambda$ — needed only in the regime $\min_{S \subsetneq N} f(S)/|S| \leq T/5$ (outside it the equality $\mu = \lambda$ holds for every game with a nonempty core), and sharp: the margin vanishes on real games. Any proof must engage the dual region in which the efficiency multiplier is negative (active on about 1% of sampled instances), cannot factor through contraction to fewer players, and cannot proceed by saturating the nested pair of caps (both routes are refuted in the archive); and it must consume *exactness*: a balanced non-exact five-player game exists, in the relevant regime, with $\mu = \frac{1}{18} < \frac{1}{12} = \lambda$ (archived and independently re-verified), so core membership alone provably does not suffice. Remaining residues: the M-side of Theorem 24 by hand, a hand distillation of the 2961-configuration check of Theorem 34 (archived progress: 95.5% of configurations fall to twelve hand lemmas), and Corollary 28's ν-attaining-triple hypothesis.

A reduced formulation of problem (2). There is a dimension-free, exactly checkable reduction that narrows the second problem without solving it. Fix one complete staged tight-coalition support $\mathcal{C}_x$ for $x = \mathrm{Lmin}(v)$, one sparse first-stage over-cover support $\mathcal{C}_z$ for $z = \mathrm{Lmax}(v)$, and a singleton exactness witness $w^r \in \mathrm{Core}(v)$ with $w^r_r = 0$. It is sufficient to retain

$$x(S) \leq z(S), w^r(S) \quad (S \in \mathcal{C}_x), \qquad z(S) \leq x(S), w^r(S) \quad (S \in \mathcal{C}_x \cap \mathcal{C}_z),$$

together with the common totals, nonnegativity, the two ordered strata and the actual hybrid sign cut. Indeed the actual triple (x, z, w^r) is feasible in this relaxation, so a lower bound on its hybrid-overlap objective is a lower bound for the game. On a common coalition the displayed rows reduce to $x(S) = z(S) \leq w^r(S)$: the sparse leximax-min certificate glues two rooted tangents only where it meets the complete leximin packet. The over-cover may be chosen with at most n coalitions, and affine Farkas duality gives a support-minimal target certificate with at most $3n - 3$ defining

rows. For a fixed labelled branch, validity, equality vertices and the relative-facet test are finite exact rational computations; projection to a facet of the full exact-game cone still requires the indecomposable semi-balanced test of Studený and Kratochvíl (2022).

There is a useful dimension-free reduction before any packet rows are used. Write $\mu = \min_i x_i > 0$, let $M = \max_i z_i$, and let H be the top block of z. The normalized first-stage over-cover gives a probability α supported on H with $\alpha \cdot x \geq M$, so some $r \in H$ satisfies $x_r \geq M$. At this ceiling root,

$$\min(x_r, z_r) = M, \qquad \min(x_i, z_i) \geq \min(\mu, z_i) \geq \frac{\mu}{M} z_i \quad (i \neq r).$$

Consequently

$$\sum_i \min(x_i, z_i) \geq M - \mu + \frac{T\mu}{M}. \tag{CRI}$$

In normalized variables $t = M/T$, $u = \mu/T$, this is $O/T \geq t - u + u/t$. If $c_n = n/\lfloor n^2/4 \rfloor$, the target is therefore automatic for $t \geq c_n$, and for $t < c_n$ it is automatic whenever

$$u \geq \frac{t(c_n - t)}{1 - t}.$$

Thus an unresolved positive-floor counterexample lies in the explicit wedge $1/n \leq t < c_n$ and $0 < u < t(c_n - t)/(1 - t)$. For $0 < \mu < M$, equality in (CRI) forces every nonroot $z_i \in \{0, M\}$ and every nonroot top player to have $x_i = \mu$, hence $T = qM$ for an integer q. This rigidity does not prove the target: on the two sharp seven-player orientations the right side of (CRI) is respectively $79T/144$ and $37T/72$, both below $7T/12$. Exact rational replay on eight archived controls and three equality patterns is included as an arithmetic audit; none of the eight hard controls is closed by (CRI) alone.

The singleton witness can now be eliminated exactly. Put $\mathcal{C}_\cap = \mathcal{C}_x \cap \mathcal{C}_z$. On this intersection the direct rows already give $x(S) = z(S)$, so the z-to-witness comparisons are duplicates. For nonnegative weights u on $\mathcal{C}_x$, let

$$p_i = \sum_{S \in \mathcal{C}_x : i \in S} u_S.$$

There exists a nonnegative total-T witness w with $w_r = 0$ dominating the required x- and z-coalition sums if and only if

$$\sum_{S \in \mathcal{C}_x} u_S x(S) \leq T \max_{i \neq r} p_i \tag{RHP}$$

for every such u. Necessity follows by weighting the witness rows; sufficiency is separation from the closed down-set generated by the rooted simplex. After normalizing $p_i \leq 1$ for $i \neq r$, it suffices to check the vertices of a rational rooted fractional packing polytope. Every vertex uses at most $n-1$ coalition weights, since its positive columns are independent on the $n - 1$ nonroot capacity rows. Thus the witness projection has a finite exact inequality description with a linear support bound. On the sharp family, the uniform extreme packing on the block-star hinges gives both inequalities used in the hand proof and recovers the bound with equality. These are relative rooted-hinge inequalities; only projected systems passing the published indecomposable semi-balanced test are exact-cone facets.

The packing has a direct rooted balance form. If $p_i = \sum_{S \ni i} u_S$, $p_i \leq 1$ off the root, and $\delta = p_r - 1$, then (RHP) is

$$\delta x_r \le \sum_{i \ne r} (1 - p_i) x_i. \tag{RDB}$$

When the packing uses only common hinges, the same inequality holds with z in place of x. For the master family the uniform block-star packing has $\delta = 1/(k-1)$, zero deficit on the block players, and unit deficit on the p outside players. The two versions of (RDB) are exactly the hand inequalities $X \le p(k-1)q$ and $M \le p(k-1)l$, giving a structural explanation of the sharp family handover. A universal proof cannot demand $\delta > 0$ on every nonspectral branch: one hostile control reaches its target from the direct common-hinge gluing rows before any projection cut is needed.

The block-star argument in fact needs only one common hinge. This gives a dimension-free mixed branch theorem. Let $N = \{r\} \dot\cup O \dot\cup B$, with $|O| = p \ge 1$, $|B| = k \ge 2$, choose $b_0 \in B$ and $j \in B \setminus \{b_0\}$, and put $C_b = \{r\} \cup (B \setminus \{b\})$. Suppose a labelled branch has

$$x_r = c, \qquad x_b = b \ (b \in B \setminus \{b_0\}), \qquad x_i = a \ (i \in O \cup \{b_0\}), \qquad 0 \le a \le b \le c,$$

while $z_r = z_b = A$ on $B \setminus \{b_0\}$, $z_{b_0} = B_0$, and $z(O) = L$. Assign B to x and $\{r\} \cup O$ to z in the hybrid cut. If a nonnegative total-T witness w, with $w_r = 0$, satisfies

$$x(C_b) \le w(C_b) \quad (b \in B), \qquad z(C_j) \le x(C_j), \tag{MBS}$$

then

$$O(x, z) \ge \frac{nT}{k(p+1)} \ge \frac{nT}{\lfloor n^2/4 \rfloor}. \tag{MBB}$$

Indeed, summing the k witness rows and subtracting $(k-1)x(N)$ gives $c \le p(k-1)a$. Put $P = p(k-1)$ and $D = k(p+1)$. The exact identity

$$(P-1)T - D(c-b) = [kP + p + 1](b-a) + n(Pa - c) \ge 0 \tag{MBI}$$

gives $c - b \le (P-1)T/D$. Since $T = kA + B_0 + L$ and the one common hinge gives $(k-1)A + B_0 \le c + (k-2)b + a$, the fixed-cut overlap obeys

$$O = A + L + (k-1)b + a \ge T - (c-b) = nT/D.$$

Finally $k + (p+1) = n$, so $D \le \lfloor n^2/4 \rfloor$. Equality at the universal target requires the balanced product, $b = a$, $c = p(k-1)a$, and equality in the displayed packet rows.

A second repair pattern covers tied ceiling loads without a strict over-cover leader. Let $N = \{r, u, j\} \dot\cup Q$, where $|Q| = m \ge 1$ and $n = m + 3$. Suppose nonnegative x, z, w have common total T, $w_r = 0$, and

$$x_r = c, \quad x_u = x_j = b, \quad x_q \le b \ (q \in Q), \qquad z_r = z_u = z_j = A, \quad z_q \le x_q \ (q \in Q), \qquad b \le A \le c.$$

If

$$x(\{r, u\}) \le w(\{r, u\}), \qquad x(\{r, j\}) \le w(\{r, j\}), \qquad z(\{r, j\}) \le x(\{r, j\}), \tag{TCS}$$

then

$$\frac{O(x, z)}{T} \ge \frac{n}{2(n-2)} \ge \frac{n}{\lfloor n^2/4 \rfloor}. \tag{TCB}$$

To see this, put $X_Q = x(Q)$. The two witness rows give $T - 2(c+b) = X_Q - c \ge 0$, the Lmin order gives $mb - c \ge 0$, and the reverse common pair gives $c + b - 2A \ge 0$. Since $O = T - 2A + 2b$, the full certificate is the identity

$$O - \frac{m+3}{2(m+1)} T = (c + b - 2A) + \frac{2(mb - c)}{m+1} + \frac{m-1}{2(m+1)} [T - 2(c+b)] \ge 0. \tag{TCI}$$

Finally $\lfloor n^2/4 \rfloor \geq 2(n-2)$ for $n \geq 4$. This theorem humanly closes two of the three positive-floor tied-repair templates in the six-player census, records 601[103] and 602[129]. The irregular record 602[82] does not have the same two-spoke form. The exact-cone row behind the theorem is

$$\mathbf{1}_{\{r,u\}} + \mathbf{1}_{\{r,j\}} + \sum_{q \in Q} \mathbf{1}_{\{r,q\}} - (m+1)e_r = \mathbf{1}_N. \tag{TCSF}$$

This is the $k = 2$, $p = m$ member of the dimension-free block-star family (GPL), so the theorem already proved that it is an indecomposable purely min-semi-balanced facet of the exact-game cone.

That last template nevertheless has a separate human certificate. On six players $\{r, u, g, j, l_1, l_2\}$, suppose

$$x = (c, b, a, a, a, a) \quad \text{on } (r, u, l_1, l_2, g, j), \qquad z_r = z_u = z_g = A, \quad z_j = D, \quad z_{l_1} = z_{l_2} = E,$$

with $0 \leq E \leq a \leq D$ and $a \leq b \leq A \leq c$. Let a nonnegative total-T witness satisfy $w_r = 0$ and

$$\begin{aligned} x(\{r,u,g,l_2\}) &\leq w(\{r,u,g,l_2\}), \\ x(\{r,u,j\}) &\leq w(\{r,u,j\}), \\ x(\{r,g,j\}) &\leq w(\{r,g,j\}), \\ z(\{r,u,j\}) &\leq x(\{r,u,j\}). \end{aligned} \tag{TP}$$

Summing the first three rows gives $3c+2b+5a \leq 2T$, hence $c \leq 3a$. Writing $h = c+b+a-2A-D \geq 0$, one obtains the exact identity

$$O - \frac{2}{3}T = h + \frac{a+b}{3} + \frac{2}{3}(3a - c) \geq 0. \tag{TPI}$$

This proves record 602[82], whose actual overlap is 7/9. Thus all three common-hinge repair templates isolated by the sequential census have human branch proofs. What remains open is the structural selection theorem forcing an arbitrary residual into the Lmin-root-price or repair alternatives; these branch theorems do not show that every residual has one of the displayed geometries.

The last repair also has an exact-cone explanation. Giving the three witness rows in (TP) weight 1/2 produces root excess 1/2, deficit one on l_1, and deficit 1/2 on l_2. Its rooted pair lift is

$$\mathbf{1}_{\{r,u,g,l_2\}} + \mathbf{1}_{\{r,u,j\}} + \mathbf{1}_{\{r,g,j\}} + \mathbf{1}_{\{r,l_2\}} + 2\,\mathbf{1}_{\{r,l_1\}} - 4e_r = 2\,\mathbf{1}_N. \tag{TPF}$$

The six support columns have rank six. For a candidate exceptional coalition with indicators $(e_r, e_u, e_1, e_2, e_g, e_j)$, the Studený–Kratochvíl exceptional-set equations have nonnegative support weights exactly when $L \leq \rho \leq U$, where

$$\begin{aligned} L = \max\{ &-e_u - e_g + e_j, -e_u - e_j + e_g, -e_g - e_j + e_u, \\ &-2e_2 + e_u + e_g - e_j, -e_1\}, \\ U = &(e_r - e_j - e_2 - e_1)/2. \end{aligned} \tag{TPE}$$

For binary indicators the feasible sets are precisely the six support coalitions and N. Hence no new proper exceptional coalition exists: (TPF) is indecomposable and defines a facet of the exact-game cone. An independently formulated exact-linear-arithmetic replay confirms all 56 unsupported proper candidates infeasible; the displayed interval argument, not the solver result, is the proof.

Only the scalar root estimate is special to the star. More generally, suppose a branch has one root at c, g players at b, and m players at a, with $n = 1 + g + m$, and its cross row gives

$O \geq T - (c - b)$. If a rooted pair lift has root excess $\delta > 0$, deficit weights supported on the a-level loss players, and total deficit weight Q, then (PFRDB) and nonnegative pair premiums give $c \leq Pa$, where $P = Q/\delta$. Putting $D_P = P + n - 1$ gives

$$(P-1)T - D_P(c-b) = [g(P-1) + D_P](b-a) + n(Pa-c) \geq 0.$$

Hence

$$O/T \geq n/(P+n-1), \tag{MFP}$$

which proves the universal target whenever $P \leq \lfloor n^2/4 \rfloor - n + 1$. The ratio Q/δ is therefore an effective *marked* facet price. For the block-star lift it is $p(k-1)$. Every one of the nine extreme mixed-projection cuts in the hostile control suite has price at most its dimensional threshold. Eight put every deficit on a nongaining player, but only five put every deficit on the first Lmin floor; the still-missing property is therefore multistage payoff/sign alignment, not a bound on the primitive coefficients.

The same calculation permits unstructured reserve players. Suppose only a subset U has the root–middle–low form, with $n_0 = 1 + g + m$, every player outside U is a nongainer, and the full gain set G consists of the g middle players and one low player. If a mixed row gives $z(G) - x(G) \leq c - b$, then the preceding identity applies to $T_U = x(U) \leq T$, while $O = T - [z(G) - x(G)]$. Therefore

$$O/T \geq n_0/(P+n_0-1). \tag{MFR}$$

If $T_U = 0$ the cross row gives $O = T$ directly; otherwise the order and $c \leq Pa$ force $P \geq 1$. This proves the n-player target whenever $P \leq n_0\lfloor n^2/4 \rfloor/n - n_0 + 1$. Extra nongaining strata are thus an overlap reserve. What remains in a general branch is not a coefficient bound, but simultaneous alignment: the rooted deficits must lie in the low loss stratum and one mixed row must pay the entire gain by $c - b$.

This misalignment has an exact scalar penalty. For an arbitrary rooted pair lift put

$$e_q = \max\left\{0, \frac{\sum_i q_i(x_i - a) - \sum_i q_i v(\{r,i\})}{\delta}\right\}.$$

Then (PFRDB) gives $c \leq Pa + e_q$, and the same identity yields

$$\frac{O}{T} \geq \frac{n}{P+n-1} - \frac{ne_q}{(P+n-1)T}. \tag{MPE}$$

With a reserve subset, replace the numerator by $n_0(T - e_q)$ and the denominator by $(P + n_0 - 1)T$. Thus the missing multistage theorem must pay a specific amount, not merely establish qualitative alignment. In the full n-player form, the denominator margin absorbs the defect whenever $e_q/T \leq 1 - (P+n-1)/\lfloor n^2/4 \rfloor$.

Thus $k-1$ hinges may be Lmin-only: one aligned Lmax-tight hinge is enough to couple the complete star to the target. Eliminating the outside witness coordinates produces the same rooted pair-spoke shear as (GPL), so the support-minimal full-cone factor is the previously classified indecomposable block-star facet despite the two-colour labelled proof. A new exact three-generator seven-player control realizes this branch away from the sharp game: its actual overlap is $11/12$, its maximum-overlap root is unique, and its rooted one-witness relaxation has exact optimum $7/12$, attained by the normalized $G(2,4)$ triple. Fourteen explicit staged identities and a hard-coded exact primal/dual certificate are included with the verifier.

One exactness row strengthens this projection at no dimension-dependent cost. Put $m_i = v(\{r,i\})$ for $i \neq r$ and $M_r = \sum_{i\neq r} m_i$. A core witness with $w_r = 0$ has $w_i \geq m_i$, hence $M_r \leq T$. Translating the rooted simplex by m shows that its support function in direction p is

$$\sum_{i\neq r} p_i m_i + (T - M_r) \max_{i\neq r} p_i.$$

Consequently the exact analogue of (RHP), with only these pair floors retained from the witness core, is

$$\delta x_r + \sum_{i\neq r} q_i v(\{r,i\}) \leq \sum_{i\neq r} q_i x_i, \tag{PFRDB}$$

with the same inequality for z when the packing uses common hinges. The new term is also the exact price of transporting the exceptional singleton: replacing $\{r\}$ by $\{r,i\}$ weakens the induced exact-game inequality by $\delta[v(\{r,i\}) - v(\{r\}) - v(\{i\})]$, namely $\delta v(\{r,i\})$ under singleton normalization. Thus the pair-floor projection retains precisely the premium lost by the canonical semi-balanced decomposition. It is sharp in the controls: seven of the eight nontrivial hostile/sharp cuts satisfy (PFRDB) with equality, while the neutral-endpoint cut has exact slack 9/9523.

Crucially, (PFRDB) has a different canonical exact-cone lift. If $Q = \sum_{i\neq r} q_i$, then

$$\sum_{S\in\mathcal{C}_x} u_S \mathbf{1}_S + \sum_{i\neq r} q_i \mathbf{1}_{\{r,i\}} - (\delta + Q) e_r = \mathbf{1}_N. \tag{RPL}$$

This is obtained from (RMS) below by the invertible column shear $e_i \mapsto e_r + e_i$. Hence a vertex packing gives a min-semi-balanced system of support at most n, with exceptional coalition $\{r\}$ when $\delta + Q > 0$. Its induced exact-game inequality on the tight Lmin coalitions is exactly (PFRDB), and the same inequality evaluates at Lmax on common hinges. By Studený and Kratochvíl (2022), these rooted pair lifts are full exact-cone facets exactly when indecomposable. They therefore form a distinguished, exactly enumerable facet class attached to a staged menu. Exact-LP and independent Z3 decomposition tests agree on the eight nontrivial controls: six pair lifts are facets, including both sharp seven-player orientations; the remaining two have respectively three and twenty decompositions.

For the full sharp family this facet statement is dimension-free. With $N = \{r\} \dot\cup O \dot\cup B$, $|O| = p$, $|B| = k$, and $C_b = \{r\} \cup (B \setminus \{b\})$, the lift is

$$\sum_{b\in B} \mathbf{1}_{C_b} + (k-1) \sum_{o\in O} \mathbf{1}_{\{r,o\}} - [p(k-1)+1] e_r = (k-1)\mathbf{1}_N. \tag{GPL}$$

It is indecomposable for every $p \geq 1, k \geq 2$. Indeed, normalizing a prospective new exceptional coalition E to coefficient -1, and writing $b = |E \cap B|$, $o = |E \cap O|$, gives $A = (k\rho + b)/(k-1)$ and $A + (p-1)\rho + o + d - \mathbf{1}_{r\in E} = 0$ with all displayed support coefficients nonnegative. Separating the cases $o < p$ and $o = p$ leaves only $\varnothing, N, \{r\}, \{r,o\}$ and the supported C_b; hence no new proper exceptional set exists. On $G(p,k)$, (GPL) is active at both selections and its evaluations are exactly the two hand inequalities producing equality in the hybrid-overlap bound.

There is now a direct coupling to the sparse Lmax certificate. Suppose a normalized top overcover supported on common hinges is $\sum_S \lambda_S \mathbf{1}_S = \eta \mathbf{1}_N + \alpha$, with $\sum_i \alpha_i = 1$. For a top root r, let $\theta_- := \max_{i\neq r} \alpha_i$ and $\theta := \eta + \theta_- > 0$. Then $u = \lambda/\theta$ is a rooted packing and

$$\delta = \frac{\alpha_r - \theta_-}{\theta}, \qquad q_i = \frac{\theta_- - \alpha_i}{\theta}. \tag{OCP}$$

Thus a strict deleted-load leader produces a positive-root pair lift, and its two evaluations are

$$(\alpha_r - \theta_-)y_r + \sum_{i \neq r}(\theta_- - \alpha_i)v(\{r,i\}) \leq \sum_{i \neq r}(\theta_- - \alpha_i)y_i, \tag{OCPL}$$

for both $y = x$ and $y = z$. This follows simply by dividing the over-cover loads $\eta + \alpha_i$ by $\eta + \theta_-$ and applying (RPL). In the bounded positive-top case, decomposing this packing into vertices decomposes (OCPL) into the corresponding rooted pair-lift facet candidates. For $G(p,k)$, $\alpha_r = 1/k$, $\theta_- = (k-1)/k^2$, the block deficits vanish and the outside deficits equal one after normalization, exactly recovering (GPL). Moreover the two rows have the same pair-lift slack. Equivalently,

$$\sum_{i \neq r}(\theta_- - \alpha_i)(x_i - z_i) = (\alpha_r - \theta_-)(x_r - z_r). \tag{OCD}$$

This follows from common tightness, which gives $\alpha \cdot x = \alpha \cdot z$, and equal totals. Thus a coupled pair-lift is active at one selection iff it is active at both; (OCD) is its selection-visible equality law.

There is also a useful quotient in the facet algorithm. If a rooted pair $\{r,i\}$ already has packing weight $t > 0$, replacing $u_{\{r,i\}}$ by $u_{\{r,i\}} - t$, q_i by $q_i + t$, and δ by $\delta - t$ leaves (RPL) unchanged: both $u_{\{r,i\}} + q_i$ and $\delta + \sum_j q_j$ are invariant. Hence all rooted-pair columns may be deleted before packing-vertex enumeration and restored only as deficit columns. This dimension-free gauge quotient removes duplicate encodings of the same exact-cone row.

This quotient does not preserve the marked price. With $K = \lfloor n^2/4 \rfloor - n + 1$, the same move sends

$$(\delta, Q, K\delta - Q) \longmapsto (\delta - t, Q + t, K\delta - Q - (K+1)t). \tag{RPGP}$$

Thus the unlabelled facet normal does not determine affordability. There is nevertheless a canonical best admissible marking once the Lmin-tight pairs are retained. For a rooted completed row whose positive nonpair columns are x-tight, let $a_S > 0$ be its positive coefficients. Split a rooted-pair coefficient as $a_{\{r,i\}} = t_i + q_i$, where t_i is packing mass, q_i is deficit mass, and require $t_i = 0$ unless $\{r,i\}$ is x-tight. Then

$$\delta(t) = \sum_{|S| \geq 3} a_S + \sum_i t_i - 1, \qquad Q(t) = \sum_i (a_{\{r,i\}} - t_i),$$

and consequently

$$K\delta(t) - Q(t) = K\delta(0) - Q(0) + (K+1)\sum_i t_i. \tag{RPM}$$

Every admissible t_i therefore takes its upper endpoint in the optimal marking. Affordability is determined exactly by the normalized facet normal together with its Lmin-tight pair labels; no particular dual-solver history is needed. More explicitly, let c be the positive exceptional-root coefficient and let $Q_{\text{bad}} = \sum_{i:\{r,i\} \text{ not } x\text{-tight}} a_{\{r,i\}}$. Root balance gives $c = \delta_* + Q_{\text{bad}}$ for the canonical best marking, so

$$K\delta_* - Q_{\text{bad}} = Kc - (K+1)Q_{\text{bad}}. \tag{RPM$'$}$$

Consequently the row is affordably markable exactly when $Q_{\text{bad}}/c \leq K/(K+1)$ and $\delta_* > 0$. This bad-spoke fraction is an exact finite labelled-facet test at fixed dimension.

It also has an exact hypergraph dual. Let G be the nonroot players with an Lmin-tight rooted pair, let B be its complement, and let A range over the offroot traces of the other root-containing Lmin-tight coalitions. For a fractional packing a of these traces, complete every unused player load by its rooted pair. With

$$C = K(|G| - 1) - |B|, \qquad w(A) = K(1 - |A \cap G|) + |A \cap B|,$$

the best marked margin is

$$C + \max_{a \geq 0} \left\{ \sum_A w(A) a_A : \sum_{A \ni i} a_A \leq 1 \ (i \neq r) \right\} \tag{LPC}$$

and equals the weighted-cover minimum

$$C + \min_{y \geq 0} \left\{ \sum_{i \neq r} y_i : \sum_{i \in A} y_i \geq w(A) \ (A \in \mathcal{H}_r) \right\}. \tag{LCD}$$

This follows by substituting the unused loads in (RPM) and applying ordinary fractional packing/cover duality. A vertex in packing-plus-slack standard form has at most $n-1$ positive columns, giving the earlier support bound and the same rooted semi-balanced/facet test. In particular, two Lmin-tight rooted pairs make the zero nonpair packing affordable, since its margin is at least

$$K - (n-3) = \lfloor n^2/4 \rfloor - 2n + 4 > 0 \qquad (n \geq 6). \tag{LPC2}$$

This closes the scalar price-selection problem in that branch. It does not by itself align the remaining deficits with the required payoff/loss strata; that labelled alignment, or an Lmax row paying its defect, remains necessary for the hybrid-overlap conclusion.

The one-pair branch has a complementary rigidity theorem. Suppose the unique Lmin-tight rooted-pair partner is j. Assigning unit mass to one nonpair trace A gives marked margin

$$K|G \setminus A| - |B \setminus A|. \tag{LPC3}$$

For $n \geq 6$, $K \geq n - 2 = |B|$. Thus if $j \notin A$, this margin is nonnegative and the root excess is one, so the row is affordable. If every nonpair trace contains j, its incidence constraint gives $\sum_A a_A \leq 1$, and (LPC) gives root excess zero for every feasible packing. Therefore a one-pair root is affordable if and only if some larger rooted tight coalition omits its partner. An unaffordable one-pair root has a *locked partner* in every rooted packing column. This explains why the direct arm must absorb that partner with the root, but does not yet select the rest of the direct block.

There is also a complete pair-free alternative. Put $m = n - 1$. Then

$$K = \lfloor (m-1)^2/4 \rfloor = \max_{2 \leq t \leq m} (t-1)(m-t), \qquad \lfloor n^2/4 \rfloor = K + m. \tag{LPC4}$$

Suppose $G = \varnothing$ and the nonpair traces have empty total intersection. Choose an inclusion-minimal empty-intersection subfamily $A_1, \ldots, A_t$. For each j, minimality supplies a distinct witness

$$v_j \in \bigcap_{k \neq j} A_k \setminus A_j.$$

Every selected trace contains the other $t-1$ witnesses, and every vertex occurs in at most $t-1$ selected traces. Hence assigning $a_{A_j} = 1/(t-1)$ is feasible, has root excess $1/(t-1) > 0$, and has marked packing value

$$\frac{\sum_j (K + |A_j|)}{t-1} \geq \frac{t(K+t-1)}{t-1} \geq K + m = \lfloor n^2/4 \rfloor. \tag{LPC5}$$

After multiplication by $t-1$, the final difference is $K - (t-1)(m-t) \geq 0$. Conversely, a common carrier in every trace forces $\sum_A a_A \leq 1$ and therefore nonpositive root excess. Thus, for $n \geq 6$, a

pair-free root is affordable if and only if its nonpair traces have empty total intersection. A pair-free price failure is exactly a rooted common carrier star.

Equality in the displayed construction forces $K = (t-1)(m-t)$ and $A_j = W \setminus \{v_j\}$ for $W = \{v_1, \dots, v_t\}$. Thus the canonical equality support is the balanced complement-of-singletons block-star geometry. Additional tight columns may still make the best marked margin positive, so this is an equality statement for the constructed packing rather than a complete game-equality theorem.

The canonical equality completion also has an exact-cone interpretation. Put $p = m - t \geq 1$, $U = V \setminus W$, $C_j = \{r\} \cup (W \setminus \{v_j\})$, and $P_u = \{r, u\}$. Then

$$\sum_j \mathbf{1}_{C_j} + (t-1) \sum_{u \in U} \mathbf{1}_{P_u} - [p(t-1)+1]\mathbf{1}_{\{r\}} = (t-1)\mathbf{1}_N. \tag{CSF}$$

This is a saturated indecomposable purely min-semi-balanced system for every $t \geq 2$ and $p \geq 1$, hence an exact-cone facet by the Studený–Kratochvíl criterion. Indeed, the positive offroot incidence matrix has diagonal blocks $J_t - I_t$ and I_p, so the displayed n columns have full rank. If E is another feasible exceptional set with common load γ, put $q = |E \cap W|$ and $s = |E \cap U|$. Its necessarily nonnegative support coefficients are

$$a_j = \frac{\gamma + q - (t-1)\mathbf{1}_{\{v_j \in E\}}}{t-1}, \quad b_u = \gamma + \mathbf{1}_{\{u \in E\}},$$

$$\lambda_r = \mathbf{1}_{\{r \in E\}} - s - \frac{q}{t-1} - \frac{1 + p(t-1)}{t-1}\gamma. \tag{CSX}$$

For $\gamma < 0$, these inequalities leave only $E = N$. For $\gamma \geq 0$, they leave only $E = \varnothing$, the root singleton, one displayed C_j, or one displayed P_u. This is precisely indecomposability. The condition $p \geq 1$ matters: the $p = 0, t = 3$ system is decomposable. Thus the balanced minimal empty-intersection packing selects a distinguished exact-cone facet, not merely a nonminimal Farkas row. We make no novelty claim for this facet family before comparison with the published semi-balanced catalogues.

Combining the two-pair bound, the locked-pair theorem, and the pair-free alternative completely classifies scalar affordability at every root. In the failure branch the remaining task is discrete: carriers chosen for the unaffordable roots must be converted into the upper nongainer block and gainer packet required by direct gluing.

Carriers have a useful intrinsic order. Let $x = \mathrm{Lmin}(v)$, let $x_r > 0$, and suppose c belongs to every x-tight coalition containing r. Then

$$x_c \geq x_r. \tag{CMO}$$

Indeed, $e_c - e_r$ is nonnegative on every active coalition constraint, and a sufficiently small positive movement preserves all slack constraints and nonnegativity. If $x_c < x_r$, the first changed entry of the increasing rearrangement lies in the x_c stratum and increases, while the compensating decrease occurs only in the strictly higher x_r stratum. This contradicts leximin optimality. Consequently carrier arrows point weakly upward in the Lmin order and every carrier cycle lies in one Lmin stratum.

This implication uses the full active tight menu; a carrier found only in one selected staged support need not define a feasible tangent. As finite evidence, all 554 full-tight-menu unaffordable maximum-overlap roots in the six-player archive have a full-menu carrier that is itself another maximum-overlap root. The narrower selected staged menus have 560 unaffordable roots and the same finite closure. Thus the archived carrier graph closes on the root block, but that closure is not yet a dimension-free theorem. Proving it, or finding its sharp failure, is the next selection bridge.

The mirror tangent argument removes the need to assume that closure. If $z_c < z_r$ and no z-tight coalition contains r while omitting c, then $e_c - e_r$ is feasible at z and strictly improves the decreasing lexicographic objective, contradicting $z = \mathrm{Lmax}(v)$.

Let $R = \arg\max_i \min(x_i, z_i)$, assume its coordinates are positive, and suppose every $r \in R$ is unaffordable on its full Lmin-tight menu. Then either

(i) there is a full Lmin carrier arrow $r \to c$, with $r \in R$ and $z_c < z_r$, and a z-tight separator D with $r \in D$, $c \notin D$; or

(ii) there is a directed carrier cycle $H \subset R$, both x and z are constant on H, every x-tight coalition meeting H contains all of H, and every member of H has the same gainer/nongainer sign.

Indeed, in the first case carrier monotonicity gives $x_c \geq x_r$; failure of c to maximize overlap therefore gives $z_c < \min(x_r, z_r) \leq z_r$, and the Lmax tangent argument supplies D. If every root has a carrier in R, choose one per root. The resulting finite functional digraph has a cycle; monotonicity forces its x coordinates to be equal, and the carrier implications force every tight coalition meeting the cycle to contain it wholesale. If z is not constant on that cycle, some carrier arrow descends in z and gives the first branch. Otherwise z is constant too; the common x, z levels make the sign uniform.

This carrier-cycle/separator alternative is a dimension-free coupling of the complete Lmin tight menu to an Lmax row. It does not yet prove the overlap bound, but both branches sharpen further. If D is z-tight, contains a root r, and omits its Lmin carrier c, then

$$x(D) > z(D), \qquad z(N \setminus D) - x(N \setminus D) > 0. \tag{SCP}$$

Core feasibility gives the weak first inequality, and equality would make D an x-tight coalition separating the carrier pair. Efficiency gives the second inequality. Thus complementation turns every separator into a strict direct-comparison payment. No dimension-free lower bound on its magnitude is known, so this does not alone repair a deficient direct cardinality base or a negative order residual.

There is also an exact contraction theorem for the cycle arm. Call H a tight-face module at y if every y-tight proper coalition contains all or none of H. A carrier cycle is a tight-face module at x, and

$$\{d : \operatorname{supp} d \subseteq H, d(H) = 0\}$$

lies in the two-sided lineality space of the active core tangent cone at x. If a z-tight coalition separates two cycle members, the same argument as (SCP) supplies a strict complement payment, even though their z coordinates are equal. Otherwise H is a tight-face module at both selections.

In the latter case contract H to a role of integer weight $h = |H|$, with coalitions expanded by replacing that role by all of H and role payoff equal to the sum on H. The quotient game is exact: projecting an original core point gives a quotient core point, and projecting an exactness witness for each expanded coalition attains its quotient worth. Rank the contracted role by its per-capita payoff, repeated h times. The projections of x and z are respectively weighted Lmin and weighted Lmax. Otherwise a better quotient point, distributed equally over H, would give a locally feasible improving direction in the original core; concavity of lowest-rank sums and convexity of highest-rank sums preserve the relevant lexicographic improvement on a sufficiently short segment. Total payoff and hybrid overlap are unchanged by this weighted contraction.

The contracted role has singleton worth $v(H)$ and need not be singleton-normalized. Subtracting this endowment can change the two lexicographic selections. Nevertheless the weighted endowed

formulation has an exact singleton-normalized block expansion. For role weights h_a, replace role a by a clone block B_a and set

$$w^{\uparrow}(S) = w(\{a : B_a \subseteq S\}). \tag{CBE}$$

Assume $w(\{a\}) = 0$ for every unit-weight role, as holds for the contraction above. Then $w^{\uparrow}$ is singleton-normalized and exact. Indeed block summation maps its core onto the role core. Given a role exactness witness q for the complete blocks of S, place the total q_a of every partially met block on a clone outside S; the resulting clone-core point attains $w^{\uparrow}(S)$. Block-permutation symmetry and uniqueness force both lex selections to be constant on every block, where they are exactly the weighted role selections. Total payoff and overlap are preserved.

The correct induction parameter is the number of roles, not the atomic weight. Let $\mathcal{P}_r$ denote the target inequality for every nonnegative exact weighted game on at most r roles, with positive integer weights of total $n \geq 6$ and zero singleton worth on each unit-weight role.

Proposition 37 (common-module induction)**.** *Assume $\mathcal{P}_{r-1}$. If an r-role instance has a nontrivial set H whose per-capita coordinates are constant at each selection and which is a tight-face module at both selections, then it satisfies the target inequality. Consequently a counterexample minimal in its number of roles has no such common module.*

Indeed contract H to one role of weight $\sum_{i\in H} h_i$. Exactness and nonnegativity descend, all remaining unit-role singleton worths stay zero, and the new role is permitted to have positive singleton worth because its weight is at least two. The quotient has at most $r-1$ roles and the same atomic weight n, while contraction preserves weighted Lmin, weighted Lmax, T, and O. Hence $\mathcal{P}_{r-1}$ gives precisely the original target. Theorems 40 and 43, together with weighted coincidence through three roles, establish the base $\mathcal{P}_4$.

Thus the contraction arm is closed in a strong-induction proof. It does not require a weighted affordability theorem after contraction: the entire smaller-role assertion is the induction hypothesis. What still needs a weighted formulation is the selection step at the current role count. In particular, whenever that step supplies a carrier cycle, a role-minimal counterexample must have an Lmax-tight coalition splitting it; otherwise the cycle is a common module and the proposition applies. Incidence-matrix contraction is the selection-face inverse of clone expansion and supplies the exact-cone interpretation of the discharged branch.

The induction base starts with complete coincidence through three roles. For three roles every proper coalition is a singleton or a singleton complement, so efficiency writes the role core as

$$K = \{q : q(R) = T, \ell_a \leq q_a \leq u_a\}, \qquad \ell_a = w(\{a\}), \quad u_a = T - w(R \setminus \{a\}). \tag{W3R1}$$

The same form covers one and two roles. For positive integer role weights h_a, the unique clipped weighted equalizer is

$$q^*_a = \mathrm{clip}_{[\ell_a,u_a]}(h_a\lambda), \qquad \sum_a q^*_a = T. \tag{W3R2}$$

It Lorenz-dominates every feasible vector after repeating q_a/h_a exactly h_a times. Indeed, whenever two unblocked roles have unequal per-capita payoffs, a sufficiently small transfer from the higher to the lower stays in the box and is a strict Robin Hood transfer. The only point admitting no such transfer satisfies exactly (W3R2). Hence weighted Lmin and weighted Lmax coincide. The two-role formula is $q^*_1 = \mathrm{clip}_{[\ell_1,u_1]}(h_1T/(h_1+h_2))$. Canonical block expansion preserves this coincidence and gives overlap T. An exact regression grid checks 170,244 feasible weighted boxes and 311,364 vertices through total atomic weight eight; the exchange argument, not that grid, is the proof. Four or more weighted roles remain open.

The same argument closes elementary four-role subcases. If all role weights are equal and all role singleton worths vanish, repetition preserves the ordinary four-coordinate lex orders, so the companion singleton-normalized four-player exact-partition theorem gives coincidence. The role-singleton hypothesis cannot be dropped: nonunit contracted roles may retain positive endowments under canonical expansion. If one role total is constant, projecting it away leaves the three-role core with coalition bounds

$$q(A) \geq \max\{w(A), w(A \cup \{r\}) - q_r\};$$

the weighted three-role equalizer remains Lorenz-greatest after adjoining the same fixed block to every repeated vector. Finally, if all three complementary pair-sum slabs are redundant, the four-role core is a box intersected with efficiency, and the clipped weighted water-filling proof applies without change. The same conclusion holds with one nonredundant complementary pair slab: the coordinate and pair intervals form a laminar constraint tree. Progressive weighted water filling from below and from above freezes the same nested nodes in reverse order, producing the unique point with no feasible Robin Hood transfer; the first-differing partial-sum exchange then makes it Lorenz-greatest. Hence any hard four-role quotient has four nonconstant role totals and at least two crossing nonredundant pair slabs. Equal weights are harmless only under the stronger role-level singleton normalization. This is a structural reduction, not yet an overlap estimate.

The crossing requirement can be imposed at each selector. If a canonical selector q has at most one active complementary pair-sum slab, retain efficiency, the coordinate intervals, and only the pair slabs active at q. This is the laminar relaxation above and has a Lorenz-greatest weighted water-filling point g. The selector q is Lorenz-undominated. If $g \neq q$, a short segment from q toward g strictly Lorenz-dominates q and remains feasible for every dropped pair inequality, since those inequalities are strict at q. This is impossible. Thus $q = g$ and it Lorenz-dominates the original core, forcing the two selectors to coincide. Consequently, when the selectors differ, each is on endpoints of at least two independent crossing pair slabs. Since there are only three pair partitions, their active edge systems share a partition. The unresolved case is now a transition between different two-flat edges (including an opposite endpoint on a shared partition), not a locally laminar selector.

This residue has a useful exact coordinate form. Set

$$y_0 = q_0 + q_1, \qquad y_1 = q_0 + q_2, \qquad y_2 = q_0 + q_3. \tag{FEC1}$$

The complementary pair-sum slabs become independent intervals in the three y-coordinates, and efficiency gives

$$\begin{aligned} 2q_0 &= y_0 + y_1 + y_2 - T, \quad 2q_1 = T + y_0 - y_1 - y_2, \\ 2q_2 &= T - y_0 + y_1 - y_2, \quad 2q_3 = T - y_0 - y_1 + y_2. \end{aligned} \tag{FEC2}$$

Hence the four coordinate intervals are four diagonal slabs cutting the pair-slab box, and the remaining cases move between different clipped box edges. If $\delta_j = y_j(x) - y_j(z)$, the inverse Hadamard transform in (FEC2) yields

$$\|x - z\|_1 = \max\left\{ \sum_{j=0}^{2} |\delta_j|, 2 \max_j |\delta_j| \right\}. \tag{FEC3}$$

Indeed, after ordering the absolute values as $a \geq b \geq c$, the four absolute Hadamard forms sum to $4a$ when $a \geq b + c$ and to $2(a + b + c)$ otherwise. Since $O = T - \|x - z\|_1/2$, the outstanding four-role theorem is now an atomic-weight-aware bound on this explicit edge-to-edge displacement. Equation (FEC3) is only an exact reformulation; it does not itself close that bound.

Up to relabelling pair coordinates, reversing their interval orientations, and interchanging the selectors, the first edge may be fixed at $y_0 = L_0, y_1 = L_1$. Besides the closed common-edge case, the second edge is one of

$$(y_0 = L_0, y_1 = U_1), \quad (y_0 = L_0, y_2 = L_2), \quad (y_0 = U_0, y_1 = U_1), \quad (y_0 = U_0, y_2 = L_2). \tag{FEC4}$$

The first two rows share one endpoint equality, lie in a common two- dimensional pair face, and satisfy $\|x - z\|_1 = 2\max\{|\delta_1|, |\delta_2|\}$. The last two are the genuinely face-changing opposite-endpoint cases. Additional active endpoints only specialize these four types or return to the common-two-flat theorem. This exhausts the minimal active-edge incidences but not the weighted payoff-stratum patterns within them.

The pure first row of (FEC4) is impossible when the third pair slab is strict at both selectors. On the common face $y_0 = L_0$, drop the y_1, y_2 slabs and call the Lorenz-greatest point of the resulting laminar fixed-pair relaxation g. At a selector x on $y_1 = L_1$, local Lorenz undominance forces $y_1(g) < L_1$: otherwise a short segment toward g stays inside the y_1 interval and inside the y_2 slab, which is strict at x. But at a selector z on $y_1 = U_1$, the inequality $y_1(g) < L_1 < U_1$ makes the same short segment feasible and strictly Lorenz-improving, a contradiction. Thus a surviving common-face transition is the hinged second row or has a third pair endpoint active, specializing a three-flat boundary case.

In fact a weighted dispersion argument closes the whole common-face residue. Fix a common endpoint, relabel its two sides as $0, 1$ and $2, 3$, and write

$$q_0 = s, \quad q_1 = W - s, \quad q_2 = t, \quad q_3 = T - W - t. \tag{FEC5}$$

For positive role weights h_i, the within-pair weighted-equality splits are

$$s_* = \frac{h_0 W}{h_0 + h_1}, \qquad t_* = \frac{h_2(T - W)}{h_2 + h_3}. \tag{FEC6}$$

Put $d_1 = |s - s_*|$ and $d_2 = |t - t_*|$. Moving either split toward its equality point is a Pigou–Dalton transfer between the corresponding clone blocks. Hence weakly decreasing both dispersions, one strictly, strictly Lorenz-improves the full repeated payoff vector.

Theorem 38 (weighted common-face rigidity)**.** *Let K be a compact four-role core with positive integer role weights. If two Lorenz-undominated points of K lie on the same endpoint of one complementary pair-sum slab, then either they coincide or they lie on the same endpoint of a second independent complementary pair-sum slab. Consequently distinct weighted Lmin and weighted Lmax cannot form either common-face row of (FEC4).*

Proof. Work in the displayed common face. Two undominated points lie in a common closed orthant about (s_*, t_*). Otherwise their segment crosses, say, $s = s_*$. At the crossing the first dispersion is zero and convexity of absolute value makes the second no larger than the larger endpoint dispersion. The crossing therefore Lorenz-dominates whichever endpoint has that larger second dispersion.

In a fixed orthant write

$$s = s_* + \epsilon_1 d_1, \qquad t = t_* + \epsilon_2 d_2, \qquad d_1, d_2 \geq 0. \tag{FEC7}$$

The singleton and complementary-singleton bounds give independent intervals for d_1, d_2. The two remaining pair coordinates have normals

$$(\epsilon_1, \epsilon_2), \qquad (\epsilon_1, -\epsilon_2). \tag{FEC8}$$

Exactly one normal has equal signs. Two distinct coordinatewise-minimal points are incomparable, so the lower convex boundary between them has negative slope. Coordinate boundaries are vertical or horizontal, and a boundary of the opposite-sign strip has positive slope; each contains at most one coordinatewise-minimal point. The only available negative-slope boundary is the unique inward endpoint of the equal-sign pair strip, where $d_1 + d_2$ is fixed. All such edges are parallel, so convexity makes them one collinear boundary rather than separate efficient pieces. Any two distinct Lorenz-undominated points share that second pair endpoint. □

Both canonical selections are Lorenz-undominated. The theorem therefore sends every common-one-endpoint incidence to Theorem 42. Only the opposite-endpoint, face-changing third and fourth rows of (FEC4) remain in the four-role quotient problem.

The opposite-endpoint residue disappears after connecting the two selectors inside the Lorenz frontier. Repeat q_i/h_i exactly h_i times and let $\Lambda_k(q)$ be the sum of the k smallest repeated coordinates, $1 \le k < n$. Each Λ_k is continuous, concave, and polyhedral: it is the minimum of the sums over all k clone coordinates.

Lemma 39 (connected positive-Lorenz continuation)**.** *Any two uniquely specified lexicographic extremes of* $(\Lambda_1, \dots, \Lambda_{n-1})$ *lie in one connected set of Lorenz-undominated core points.*

Proof. The lower-comprehensive Lorenz image

$$\{u : \exists q \in K,\ u_k \le \Lambda_k(q)\ (1 \le k < n)\}$$

is a convex polyhedron, because it is the projection of the common hypograph of the concave polyhedral Λ_k. Its finite normal fan gives strictly positive hierarchical vectors a^x, a^z for which the two lexicographic extremes uniquely maximize

$$F_a(q) = \sum_{k=1}^{n-1} a_k \Lambda_k(q). \tag{FEC9}$$

Put $a(t) = (1-t)a^x + ta^z$ and $A(t) = \arg\max_K F_{a(t)}$. Every $A(t)$ is nonempty, compact, and convex, because $F_{a(t)}$ is concave. Its graph is compact. If the graph were the disjoint union of two nonempty closed sets, their projections would be closed sets covering $[0,1]$, so they would meet at some t. They would then disconnect the convex fiber $A(t)$, a contradiction. Thus the graph and its projection to K are connected. Every projected point is Lorenz-undominated because every coefficient of $a(t)$ is positive. Finally, two distinct role allocations cannot be undominated with the same full Lorenz vector: their repeated payoff vectors are permutations, and their midpoint is strictly more equal unless they agree. Thus the canonical role point at either endpoint is unique. □

The earlier two-active-crossing argument applies to every Lorenz-undominated point. Unless the core has a Lorenz-greatest point (when the selectors coincide), every point in the connected continuation lies on endpoints of at least two pair slabs, hence in the one-skeleton of the three-dimensional pair-slab box.

Theorem 40 (weighted four-role edge confinement)**.** *In every compact four-role core with positive integer role weights, weighted Lmin and weighted Lmax either coincide or share endpoints of two independent complementary pair-sum slabs.*

Proof. If a pair slab has zero width, its endpoint is common to every point and Theorem 38 supplies the second endpoint. Otherwise use the connected continuation in the ordinary box one-skeleton. A connected subset meeting interiors of two different box edges must meet interiors of two

edges incident at some vertex. Points on those incident edges share exactly one endpoint equality, contradicting Theorem 38. Hence the continuation lies in one clipped box edge, and its endpoints share the two pair equalities defining that edge. □

Thus (FEC4) has no surviving edge-transition residue. The quantitative atomic theorem below still uses the common-two cost calculation and its finite exact audit; Theorem 40 is the structural bridge making that branch exhaustive.

This boundary is sharp. On four roles of weights $(3,1,1,1)$, the exact role game with worths in binary coalition order

$$(0,11,0,14,0,14,0,16,0,12,0,15,0,15,0,18) \tag{W4O1}$$

has weighted role-total selections

$$q^{\min}=(12,2,2,2), \qquad q^{\max}=(11,3,3,1). \tag{W4O2}$$

Its canonical six-clone expansion has

$$x=(4,4,4,2,2,2), \qquad z=(11/3,11/3,11/3,3,3,1), \qquad O/T=8/9. \tag{W4O3}$$

Thus weighted coincidence fails already at four roles, although this control lies far above the six-player overlap target $2/3$. Exact LP witnesses attain all sixteen role worths, their canonical placements attain all 64 clone worths, and an independent exact simplex recomputes both displayed traces.

There is a dimension-free four-role family behind this boundary. For $m \geq 2$, use role weights $(m,1,1,1)$ and define the game as the lower envelope of

$$(2m+4,0,0,0),(2m,2,2,0),(2m+2,0,2,0),(2m+2,2,0,0),(2m+3,0,0,1). \tag{W4F1}$$

The five common-total generators make the role game exact. Its weighted selections are

$$q^{\min}=(2m+1,1,1,1), \qquad q^{\max}=(2m,2,2,0). \tag{W4F2}$$

Indeed the complement of the first three roles caps role 3 at one. On the weighted-Lmin floor-one face, the two pair bounds $q_0+q_1, q_0+q_2 \geq 2m+2$ and efficiency force $q_0=2m+1$, $q_1=q_2=q_3=1$. For weighted Lmax, the singleton bound $q_0 \geq 2m$ makes the per-capita ceiling at least two. At ceiling two it pins $q_0=2m$; the same pair bounds and efficiency force $(q_1,q_2,q_3)=(2,2,0)$. Canonical expansion to $n=m+3$ players gives

$$\frac{O}{T}=\frac{m+1}{m+2}, \qquad \frac{\|x-z\|_1}{T}=\frac{2}{m+2}. \tag{W4F3}$$

This proves that four-role weighted coincidence fails in every atomic dimension, but the family tends toward full overlap and is not extremal for the main conjecture. The symbolic identities are audited through $m=40$.

The stronger proposed bridge "every normalized exact weighted four-role game has $O/T \geq 3/4$" is false. For weights $(m,1,1,1)$, $m \geq 5$, take the lower envelope of

$$(0,5,3,4),(1,3,3,5),(2,4,3,3),(5,0,7,0),(0,6,0,6). \tag{W4N1}$$

Its binary-order worth table is

$$(0,0,0,4,0,0,6,6,0,4,0,5,6,6,7,12), \tag{W4N2}$$

so all role singleton worths vanish, and attainment by the displayed common-total generators proves exactness. Direct staged elimination gives

$$q^{\min} = (5,1,5,1), \qquad q^{\max} = (3,3,3,3), \qquad O/T = 2/3. \tag{W4N3}$$

For Lmin, complementing $w(123) = 7$ gives $q_0 \le 5$, so the first floor is at most $5/m$, attained by the displayed candidate. On the face $q_0 = 5$, efficiency and $q_1 + q_2, q_2 + q_3 \ge 6$ give $q_3, q_1 \le 1$; the next floor therefore pins $(q_1, q_2, q_3) = (1, 5, 1)$. At $m = 5$ the equal levels merge. For Lmax, $q_1 + q_2 \ge 6$ forces ceiling at least three. At ceiling three it pins $q_1 = q_2 = 3$; then $q_2 + q_3 \ge 6$ pins $q_3 = 3$, and efficiency gives $q_0 = 3$. An independently encoded exact staged LP audits every pinning step. Canonical expansion gives singleton-normalized exact examples for every $n = m + 3 \ge 8$. They remain above the conjectured atomic target, but prove that four-role contraction cannot be closed by an atomic-size-independent $3/4$ estimate. A valid quotient theorem must retain atomic weight, largest-block weight, or labelled carrier information. The $m = 8$ verifier also constructs explicit attainment witnesses for all 2^{11} clone coalitions.

The common crossing face of this family has a complete atomic-weight-aware explanation. More generally, suppose both weighted selectors lie on

$$q_0 + q_1 = W, \qquad q_0 + q_3 = W, \qquad T = 2W. \tag{BWS1}$$

The face is the interval

$$q(t) = (t, W - t, t, W - t), \qquad \ell \le t \le u. \tag{BWS2}$$

Put $P = \max(h_0, h_2)$, $p = \min(h_0, h_2)$ and similarly $Q = \max(h_1, h_3)$, $q = \min(h_1, h_3)$. Weighted water filling from below and above gives, before common-interval clipping,

$$t_{\min} = \frac{WP}{P+Q}, \qquad t_{\max} = \frac{Wp}{p+q}. \tag{BWS3}$$

Clipping cannot increase their distance, and $O/T = 1 - |t_{\min} - t_{\max}|/W$. Orienting the difference gives

$$\frac{O}{T} \ge \frac{1}{P+Q} + \frac{1}{p+q} = \frac{n}{(P+Q)(p+q)} \ge \frac{4n}{n^2 - (n \bmod 2)}. \tag{BWS4}$$

The last step is the sharp integer product bound for two positive summands of total n. Thus the full conjectured quotient inequality holds on every balanced common seesaw face. Equality in the harmonic/product proof occurs only at atomic weights four and five, so this branch has strict target slack for $n \ge 6$. An exact audit covers all 91,390 positive weight vectors through $n = 40$ and 4,112,550 rational clipping intervals; the displayed argument is the proof.

Each flat in (BWS1) is a tight pair-coalition face. It is a complementary exact-partition equality only in the special case where the opposite slab endpoints coincide. Thus this bound is selection-order arithmetic on the intersection of two core faces, not a new full exact-game-cone facet. Its strict margin for $n \ge 6$ also excludes this branch from the six- and seven-player target-equality orientations, which must remain in the labelled rooted or mixed facet arms.

The same common-two-flat geometry admits a useful endowed reduction even without the balance assumption. Parameterize its core interval by

$$q(t) = (e_0 + t, e_1 + s - t, e_2 + t, e_3 + s - t), \qquad 0 \le t \le s, \quad T = E + 2s, \quad E = \sum_i e_i. \tag{EWS1}$$

The theorem below assumes $e_i \ge 0$ and the stronger face-normalization condition $e_i = 0$ whenever the role multiplicity h_i is one. Zero singleton worth of a unit role does not by itself imply this

condition: its minimum on this particular line may exceed its minimum on the full core. For $r \in \{0,2\}$ and $f \in \{1,3\}$ put

$$\kappa_{rf} = \frac{h_r(e_f + s) - h_f e_r}{h_r + h_f}. \tag{EWS2}$$

An envelope comparison of the two rising and two falling per-capita lines gives

$$t_{\min} = \operatorname{clip}_{[0,s]} \max_r \min_f \kappa_{rf}, \qquad t_{\max} = \operatorname{clip}_{[0,s]} \max_f \min_r \kappa_{rf}. \tag{EWS3}$$

The crossing geometry supplies an order restriction not valid for a generic 2×2 matrix. A positive gap cannot separate the two diagonals: at a parameter between them the four per-capita lines would form a strict cycle $R_0 > F_1 > R_2 > F_3 > R_0$. Hence the low entries share a row or column. After this exclusion, the two max–min values in (EWS3) are exactly the two middle clipped crossings. Writing them as $\kappa_2 \le \kappa_3$ gives

$$\{t_{\min}, t_{\max}\} = \{\kappa_2, \kappa_3\}, \qquad \frac{O}{T} = 1 - \frac{2(\kappa_3 - \kappa_2)}{T}. \tag{EWS4}$$

The remaining adjacent middle-gap inequality admits a closed atomic proof. After row/column relabelling and, for a column, reversing the parameter, let p be the low-row weight, r the other row weight, and $q \le b$ the column weights.

Theorem 41 (face-normalized endowed common-two-flat bound)**.** *Normalize a positive adjacent middle gap to one. The minimum possible total, over all nonnegative endowments with $e_i = 0$ on every unit-weight role, is*

$$S(p,r;q,b) = \begin{cases} 2 + p/q + (p+q+b)/r, & p > 1, \\ +\infty, & p = r = 1, \\ 2(2r+b+1)/(r-1), & p = 1 < r,\ q = 1, \\ (q+1)(2r+q+b)/[q(r-1)], & p = 1 < r,\ q > 1. \end{cases} \tag{EWS5}$$

Consequently the atomic bound below holds when the face is normalized in this way and its total atomic weight is $n \ge 6$:

$$\frac{O}{T} \ge \frac{4n}{n^2 - (n \bmod 2)}. \tag{EWS6}$$

Equality is possible only at $n = 6$ and $(p,r,q,b) = (2,2,1,1)$, up to the displayed symmetries.

Proof. Let $\ell, \ell+1$ be the scaled middle crossings and set $z_f = e_f + s$ for $f \in \{q,b\}$. Clipping gives $\ell \ge 0$ and $s \ge \ell + 1$. The crossing inequalities and total are

$$pz_f - fe_p \le (p+f)\ell, \qquad rz_f - fe_r \ge (r+f)(\ell+1), \qquad T = e_p + e_r + z_q + z_b. \tag{EWS7}$$

If $p > 1$, add both high-row inequalities and substitute the low-row q-inequality. This gives exactly

$$T \ge 2 + \frac{p}{q} + \frac{p+q+b}{r} + \frac{n}{r}(\ell + e_r). \tag{EWS8}$$

Equality is attained at $\ell = e_r = 0$ by $z_f = (r+f)/r$, $e_p = p(r+q)/(qr)$, $s = z_q$; the remaining endowment is $(b-q)/r$ on role b, which is nonunit whenever positive.

If $p = 1$, then $e_p = 0$. For $r = 1$ the two rows coincide. Otherwise (EWS7) implies the sharper $\ell \ge [r + q + qe_r]/[q(r-1)]$. Summing the high inequalities and substituting this bound gives the

last line of (EWS5), plus a nonnegative multiple of e_r, when $q > 1$. When $q = 1$, also $e_q = 0$ and $z_q = s$; the low and high q inequalities give $\ell \geq (r + 1 + e_r)/(r - 1)$. Adding the two high lower bounds to the total and substituting this inequality yields the third line plus a nonnegative multiple of e_r. Equality in both cases occurs at $e_r = 0$; the remaining explicit choices are recorded in the exact verifier.

Put $F = \lfloor n^2/4 \rfloor$ and $H_n = 2/(1-n/F) = 2+2n/(F-n)$. Since (EWS4) gives $O/T = 1-2d/T$, it remains to prove $S \geq H_n$. For $p > 1$, set $k = p + q + b = n - r$. Since $q \leq b$ and $p \geq 2$, Engel's inequality gives

$$S - 2 \geq \frac{4}{k-2} + \frac{k}{n-k} \geq \frac{(2+\sqrt{k})^2}{n-2} \geq \frac{16}{n-2} \geq \frac{2n}{F-n}. \tag{EWS9}$$

The last comparison is $16/(n-2) \geq 8/(n-4)$ for even n; for odd $n \geq 7$ it is equivalent to $n^2 - 6n - 2 \geq 0$. The $p = q = 1$ line gives $S - 2 = 2n/(r-1) \geq 2n/(n-4)$. Finally, if $p = 1, q > 1$, put $d = n - 2q - 2 \geq r - 1$. Dropping positive terms and applying Engel again gives

$$S - 2 \geq \frac{2}{q} + \frac{2(q+1)}{d} \geq \frac{2}{q} + \frac{6}{d} \geq \frac{(\sqrt{2}+\sqrt{6})^2}{q+d} > \frac{14}{n-4} > \frac{2n}{F-n}. \tag{EWS10}$$

Equality in the chain is therefore possible only in the first case at $n = 6, (p, r, q, b) = (2, 2, 1, 1)$. □

The equality role interval has endpoints $(3, 3/2, 0, 3/2)$ and $(9/2, 0, 3/2, 0)$. Its lower-envelope role game is exact; the complementary values $w(01) = w(03) = 9/2$, $w(23) = w(12) = 3/2$, and $w(R) = 6$ force its core onto the displayed line, and the zero singleton bounds cut out exactly that interval. Canonical expansion at weights $(2, 1, 2, 1)$ is singleton-normalized and has

$$\mathrm{Lmin} = (2, 2, 1/2, 1/2, 1/2, 1/2), \qquad \mathrm{Lmax} = (3/2, 3/2, 3/2, 0, 0, 3/2). \tag{EWS11}$$

Thus it is precisely six-player equality survivor 1408, the $(4, 2)$ versus $(2, 4)$ profile. At $n = 5$, the earlier $((1, 1, 1, 2), (0, 0, 0, 2), 1)$ row gives $3/4$, so the dimensional threshold is sharp. At $n = 7$ the sharp bound within this face-normalized branch is $5/8 > 7/12$, so neither seven-player equality orientation lies here. A common pair-slab endpoint is a tight-coalition face, not automatically a complementary exact-partition equality; the latter requires opposite slab endpoints to coincide. The theorem is therefore selection arithmetic on its stated face and is not a new full exact-cone facet. The verifier checks 1,950,536 oriented formula rows through $n = 60$ and independently solves the original six-variable unit-gap LP on all 4,844 positive weights through $n = 20$, recovering every formula optimum or infeasible equal-unit-row case. It also checks all 64 exactness minima and both staged selectors of (EWS11), eight independently formulated role controls, and the previous 2,484,224 bounded endowment rows. The displayed inequalities, not those finite checks, are the proof. A constant half-overlap shortcut is still false at the abstract seesaw level.

The face-normalization hypothesis cannot be inferred from zero role singletons. With weights $(3, 1, 1, 1)$, the lower envelope of

$$(0, 5, 3, 4), (0, 3, 5, 4), (2, 4, 3, 3), (5, 0, 7, 0), (0, 5, 0, 7)$$

has zero singleton worths and weighted selections $(9/2, 3/2, 7/2, 5/2)$ and $(4, 2, 3, 3)$. They share the upper endpoints of pair partitions $01|23$ and $03|12$, but the maximal nonnegative extension of that line has direction $(1, -1, 1, -1)$ and endowment $(1, 0, 0, 1)$. Thus unit-weight role 3 has positive effective face endowment. Canonical expansion and an independent six-player replay give overlap $11/12 > 2/3$. This is a scope control, not a target counterexample: the unrestricted common-two-flat proof must use the off-face exactness witness attaining the unit singleton to pay for its positive face endowment.

There is already a sharp payment when the positive unit role is the common hub. Write the two tight pair coalitions defining the line as $C = \{h, a\}$ and $D = \{h, b\}$, and let o be the fourth, opposite role. For any q on their common face, take an exactness witness y for the singleton $\{h\}$. Since $q(C) = w(C)$, $q(D) = w(D)$, $y_h = w(h)$, and $y_o \geq w(o)$,

$$T - y_o + w(h) = y(C) + y(D) \geq q(C) + q(D) = T - q_o + q_h,$$

and hence

$$q_h - w(h) \leq q_o - y_o \leq q_o - w(o). \tag{EWS20}$$

The hub and opposite coordinates move in the same direction on the line. Thus, if h has unit weight and zero singleton worth, evaluating at their common minimizing endpoint gives $e_h \leq e_o$. This common-face hub payment is a structural exactness lemma. It does not charge a unit arm: an arm occurs in only one of the two tight coalitions. The two arm orientations therefore remain open in the unrestricted common-two-flat branch.

The remaining common-face problem is finite in atomic weight. If all effective-endowment restrictions are dropped, the same projection as above has the exact unit-gap cost

$$S_0(p, r; q, b) = 2 + \frac{p}{q} + \frac{p + q + b}{r}, \qquad q \leq b. \tag{EWS21}$$

The lower bound is the calculation leading to (EWS8), and its displayed equality point remains nonnegative without the unit-endowment restrictions. For $p \geq 2$ the preceding integer argument already proves the atomic target. If $p = 1$, then

$$S_0 - 2 = \frac{1}{q} + \frac{n - r}{r}.$$

For $q = 1$ this is at least $n/(n-3)$. For $q \geq 2$, $b \geq q$ and Engel's inequality give

$$S_0 - 2 \geq \frac{1}{q} + \frac{2q + 1}{r} \geq \frac{(\sqrt{2} + \sqrt{2q+1})^2}{2q + r} > \frac{13}{n - 1}.$$

Direct comparison with $2n/(\lfloor n^2/4 \rfloor - n)$ proves the target for every $n \geq 10$. Thus off-face exactness is needed on common-two-flat faces only for $n = 6, 7, 8, 9$; the complete subtarget list, up to ordering the side weights, is

$$\begin{array}{c|l} 6 & (1,2;1,2), (1,3;1,1) \\ 7 & (1,3;1,2), (1,4;1,1) \\ 8 & (1,5;1,1) \\ 9 & (1,6;1,1). \end{array} \tag{EWS22}$$

These are relaxed weight rows, not claimed exact-game counterexamples.

One orientation of every displayed row is already excluded structurally. All have $p = q = 1$. If the unit low role p is the common hub, the hub payment gives $d = e_r - e_p \geq 0$. Comparing the low and high inequalities in the unit side column yields

$$e_p + \ell \geq \frac{r + 1 + d}{r - 1}.$$

Adding the two high crossing bounds to the total, and writing $A = 2r + b + 1$, then gives

$$T \geq \frac{A}{r}(e_p + \ell + 1) + \frac{r + b + 1}{r} d \geq \frac{2(2r + b + 1)}{r - 1}. \tag{EWS23}$$

This is the safe face-normalized cost. Hence only the orientations in which the unit low role is not the common hub remain in the finite common-face residue; it can be an arm or the role opposite the hub.

For the finite audit we use the standard fairness characterization without guessing selector strata. A point x of a compact convex weighted role core is weighted Lmin exactly when, for every y in the core,

$$\frac{y_i}{h_i} > \frac{x_i}{h_i} \quad \Longrightarrow \quad \exists j : \ \frac{x_j}{h_j} \le \frac{x_i}{h_i}, \quad \frac{y_j}{h_j} < \frac{x_j}{h_j}. \tag{EWS24}$$

Weighted Lmax satisfies the reversed min–max implication. A violation gives a strict repeated-payoff lexicographic improvement on a short segment; the converse follows by taking the direction to the lexicographic optimizer and its first improved level. Thus (EWS24), the interval–pair-slab core description, and the unit-singleton attainment witnesses form a quantified linear-real encoding with no assumed payoff stratum or active coordinate face. A separately formulated rational regression computes 236 staged selectors on 118 exact lower-envelope games and excludes a violating no-repair direction in all 944 role tests. This audits the formulation; the preceding convex argument is the proof of the characterization.

The finite audit is conclusive. For every row of (EWS22), it splits the 12 possible choices of two common pair-slab endpoint labels and the 16 coordinatewise overlap sign cells. All $6 \cdot 12 \cdot 16 = 1{,}152$ quantified linear-real sentences are unsatisfiable, with no unknown result. The encoding retains a core point attaining every zero unit-role singleton but relaxes attainment of all other coalition bounds, so its feasible set contains every exact game under consideration.

Theorem 42 (atomic common-two-flat quotient bound)**.** *Let an exact nonnegative weighted four-role game have positive integer role weights of total $n \ge 6$, with zero singleton worth on every unit-weight role. If weighted Lmin and weighted Lmax share endpoints of two independent complementary pair-sum slabs, then*

$$\frac{O}{T} \ge \frac{4n}{n^2 - (n \bmod 2)}.$$

For $n \ge 10$ this is the human unrestricted cost argument. At $n = 6, 7, 8, 9$, the same formula closes every weight row outside (EWS22), and the exact finite audit excludes strict violation on the six remaining rows. That is exactly the asserted weak inequality. Together with Theorem 38, it gives the following stronger structural branch closure.

Theorem 43 (atomic weighted four-role quotient bound)**.** *Let an exact nonnegative weighted four-role game have positive integer role weights of total $n \ge 6$, with zero singleton worth on every unit-weight role. Then*

$$\frac{O}{T} \ge \frac{4n}{n^2 - (n \bmod 2)}.$$

If the selectors coincide the claim is immediate. Otherwise Theorem 40 supplies two common pair endpoints and Theorem 42 applies. This completes the weighted four-role quotient objective. The finite component inherited from the common-two theorem is computer-assisted and has not received external specialist review; equality classification within its six finite rows is not supplied by the UNSAT audit alone.

On the 228 archived six-player rows where all maximum-overlap roots are unaffordable on their full menus, every row enters the nongainer-cycle arm. Of 420 cycles, 169 are separated by the Lmax face and 251 are common-face modules. At the row level, 163 expose at least one inductively contractible common module. The other 65 are role-irreducible with respect to these cycles, and every one has an Lmax-tight separator. This is the exact finite shadow of Proposition 37, not

evidence that the separator payments already close the target. Of the 65 role-irreducible rows, the cycle itself gives a tight target-reaching direct block on 62. In each of the remaining three, the unique splitting-row complement contains every gainer and is therefore outside the proper-packet hypothesis of the direct identity; replacing at least one gainer by one additional upper nongainer gives an Lmin-tight block with cardinality base 4/5. Every row has such a target-reaching tight direct block containing a cycle, but over all 228 rows the cycle itself works on only 216; the other 12 need one additional nongainer, and full carrier closure still does not select it. These exact finite counts falsify two tempting cycle-selection rules; they are not evidence for an unbounded-dimension inequality. Eighty-five of those rows have an Lmax-tight coalition that splits a carrier cycle, but complementing one such separator supplies a target-valid direct identity with nonnegative order residual on only one row. Thus strictness in (SCP) is qualitative currency: a general proof must aggregate separator payments, repair their order residuals, or use weighted contraction. Of the six fixed seven-player hostile and sharp controls, only the top-and-bottom-fixed branch has every full-menu maximum-overlap root unaffordable. Its roots form a nongainer common-face module and a separate direct packet reaches 7/12. Both sharp equality orientations instead remain on the affordable rooted-facet arm.

The full Lmax staged packet aggregates the separator currency. Let H be a nongainer carrier cycle with $x_i = a$, $z_i = b$, and $L = a - b \geq 0$ on H, and let Q be the cumulative Lmax upper set through its stratum. For each $i \in H$, sign-coherent pinning gives

$$\sum_S \lambda_S^i \mathbf{1}_S - \sum_{j \in Q} c_j^i e_j = k_i \mathbf{1}_N, \qquad \lambda^i, c^i \geq 0, \quad c_i^i > 0, \tag{ASB1}$$

with z-tight support. Evaluating at $d = x - z$ and summing gives

$$\sum_S \Lambda_S [x(S) - z(S)] = LC_H + \sum_{j \in Q \setminus H : d_j \geq 0} C_j d_j - \sum_{j \in Q : d_j < 0} C_j (z_j - x_j), \tag{ASB2}$$

where $C_H = \sum_{j \in H} C_j > 0$. Every bracket on the left is nonnegative, and every row splitting H is strictly positive by (SCP).

The negative terms have exactly the missing labels. If $j \in Q$ is a gainer, then $z_j \geq b$, while maximality of the overlap on H forces $x_j \leq b \leq a$. When $L > 0$, each debt edge from j to H is strict in the Lmin order and is available to the sign-coherent Lmin transport. It is not a high–low overlap rectangle, because

$$\max(x_j, z_j) = z_j \geq b = \min(x_i, z_i), \qquad i \in H.$$

Thus the canonical label-currency packing has no such interval. The split

$$z_j - x_j = (z_j - b) + (b - x_j). \tag{ASB3}$$

does not make either term independently countable overlap currency. Direct comparison–residual cancellation below avoids matching these pieces separately: it turns the whole direct arm into one normalized scalar budget $U_s - L_H$. The aggregate separator packet may still be needed to prove that budget for a selected nongainer block or to supply a marked-facet alternative. If $L = 0$, strict Lmin availability may also disappear. Circuit amplification makes unnormalized coefficient bounds meaningless. If the chosen Lmax coalitions never split H, their coalition-incidence part factors through the weighted contraction and canonical block expansion above. Proposition 37 discharges that branch. Thus, in a role-minimal counterexample, the aggregate packet may be restricted to

cycles with at least one splitting row. The player-unit normalization still requires weighted pinning coordinates in the current-role selection step; no general weighted affordability theorem is assumed.

The proper-packet condition itself is automatic away from the tied block. Suppose $L = a - b > 0$ and let j be any gainer. Maximality of the overlap on H gives $x_j \le b < a = x_i$ for $i \in H$. If every x-tight coalition containing i also contained j, carrier monotonicity would give $x_j \ge x_i$, a contradiction. Hence there is an x-tight coalition C_j containing i and omitting j; the module property forces $H \subset C_j$. Thus:

Lemma 44 (strict-loss proper-packet support)**.** *For a nongainer carrier cycle with $a - b > 0$, every gainer is omitted by some Lmin-tight coalition containing the whole cycle. In particular the gainer part of that coalition is proper and its comparison payment is nonnegative.*

The coalition may depend on the omitted gainer. The lemma does not prove the order-residual condition, and when $L = 0$ the strict Lmin-order contradiction disappears. The cardinality condition, however, is automatic.

Lemma 45 (carrier-cycle cardinality)**.** *Let p, g, h be the atomic multiplicities of the three sign blocks in the direct identity, with $n = p + g + h \ge 6$, $p, g \ge 1$, and let $h \ge 2$ because the nongainer block contains a carrier cycle. For every proper gainer packet of atomic multiplicity $0 \le s < g$,*

$$1 - \frac{gp}{(g+h)(p+g-s)} \ge \frac{4n}{n^2 - (n \bmod 2)}. \tag{CCC}$$

Equality is possible only at $(n, p, g, h, s) = (6, 2, 2, 2, 1)$.

Proof. The left side decreases with s, so take $s = g - 1$. At fixed n, g, the subtracted product

$$\frac{g}{g+h}\frac{p}{p+1}, \qquad p = n - g - h,$$

decreases with h, so take $h = 2$. With $g = n - p - 2$, the remaining base is $(n+p)/[(n-p)(p+1)]$. For even n, cross multiplication against $4/n$ leaves n times

$$4p^2 + (4 - 3n)p + n^2 - 4n,$$

whose discriminant is $-7n^2 + 40n + 16 < 0$ for $n \ge 8$; at $n = 6$ direct substitution of $1 \le p \le 3$ gives the unique zero $p = 2$. For odd n, the cross-product quadratic has leading coefficient $4n$ and discriminant

$$-(n+1)(7n^3 - 47n^2 + 9n - 1) < 0 \qquad (n \ge 7),$$

because the cubic is positive at seven and increasing thereafter. This proves (CCC) and its equality statement. □

Consequently the strict-loss carrier-cycle branch has only one remaining direct condition, and even that condition has an exact scalar reduction. Recall from the direct master identity that, for $d = p + g - s$ and $B = g/d$,

$$d(R_x + R_z) = g[U_s - \Delta(J)], \qquad z(H \cup J) - x(H \cup J) = \Delta(J) - L_H, \tag{CB1}$$

where $L_H = x(H) - z(H)$ and $\Delta(J) = z(J) - x(J)$. Therefore

$$B[z(H \cup J) - x(H \cup J)] + R_x + R_z = B(U_s - L_H). \tag{CB2}$$

The gainer labels have disappeared. Direct substitution in U_s also gives

$$U_{s+1} - U_s = \frac{z(G) - x(G)}{g} \geq 0. \tag{CB3}$$

Hence the best proper packet always has $s = g - 1$, and any such packet gives the same combined margin.

Lemma 46 (comparison–residual cancellation)**.** *Once a sign-compatible nongainer block H and nonempty complement P are fixed, the direct identity with any proper gainer packet reaches its best non-low-coordinate margin at $s = g - 1$. It closes that part of the target exactly when $U_{g-1} \geq L_H$; no separate comparison or residual-sign condition and no choice among equally sized gainer labels remains.*

Thus neither contraction, packet properness, target cardinality, nor labelled residual matching remains open on this arm. The remaining choice is the nongainer block H containing a carrier cycle. The full nongainer maximum-overlap root block satisfies the scalar budget on all 65 role-irreducible six-player archive rows; the cycle alone does so on 63. The following aggregate formula remains useful, but a new seven-player control shows that its proposed universal sign is false.

That candidate statement has a particularly simple aggregate form. Let M be the maximum overlap coordinate, let H be all nongainer roots, let P be the other nongainers, and put

$$U = \sum_{i \in P} (M - z_i), \qquad R = \sum_{j \in G} (M - x_j), \qquad D = z(G) - x(G).$$

Thus D is also the total nongainer loss. Substitution of $x(P) = pM - U + [x(P) - z(P)]$, $x(G) = gM - R$, and $z(H) = hM$ into $U_{g-1} - L_H$ gives the exact equivalence

$$U_{g-1} - L_H = \frac{g(g+h)U - gpR - [g + h(p+1)]D}{g(g+h)}. \tag{RBB}$$

Consider the lower envelope of the three singleton-normalized allocations

$$(6, 7, 1, 0, 2, 0, 1), \quad (7, 0, 0, 0, 0, 10, 0), \quad (0, 0, 0, 5, 4, 6, 2).$$

Its exact selections are

$$x = (3, 7/2, 1, 3, 2, 7/2, 1), \qquad z = (14/5, 16/5, 2/5, 13/5, 16/5, 16/5, 8/5).$$

The maximum-overlap roots are $H = \{1, 5\}$ at $M = 16/5$. They are mutually locked, both full rooted menus are unaffordable, and Lmax-tight coalitions 33 and 101 split the carrier cycle. Thus the block is role-irreducible, but with $P = \{0, 2, 3\}$ and $G = \{4, 6\}$ one has

$$U = 19/5, \quad R = 17/5, \quad D = 9/5, \qquad g(g+h)U - gpR - [g + h(p+1)]D = -8. \tag{RBB-CE}$$

Hence the full-root-block RBB conjecture is false. This does not threaten the overlap target: the actual ratio is $76/85 > 7/12$. The unique target-reaching tight direct row is coalition 59; its nongainer part $\{0, 1, 3, 5\}$ adjoins two labelled nongainers to the cycle, leaves $P = \{2\}$, and gives base $5/6$ and scalar margin $17/30$. It is selected by the primitive first-stage Lmin circuit

$$\mathbf{1}_{59} + e_2 + e_6 = \mathbf{1}_N, \tag{RBB-C}$$

where $2, 6$ are the first-floor players and 6 is the omitted gainer. The unit columns in (RBB-C) are stage-floor columns, not tight singleton game inequalities. Hence this is a relative selection-cell circuit, not evidence for a new full exact-game-cone facet. The archived exact verifier reconstructs all coalition minima, checks staged dual pinning relations, and independently rules out every staged improving direction in rational Z3 arithmetic. The two optimality encodings share the reconstructed core and tight menus, but not their multiplier/direction formulations.

There is nevertheless a dimension-free localization of the missing charge. Put $d = x - z$. For a gainer j, every positive-loss source i in a sign-coherent Lmin transport satisfies $x_i \le x_j < M$. A nongainer root has $x_i \ge M$, so every such source lies among the nonroot nongainers P. Moreover, with $u_i = M - z_i$ and $r_j = M - x_j$,

$$u_i = (M - x_i) + (x_i - z_i) \ge r_j + (x_i - z_i). \tag{RET}$$

Lemma 47 (root-exclusion transport)**.** *Every gainer's sign-coherent Lmin transport draws its positive loss currency entirely from nonroot nongainers. Each eligible source pays both its own loss and the gainer's root-distance reserve as in* (RET)*.*

The lemma identifies the correct source block, but the known transport theorem does not bound its joint column congestion. That congestion, or a mixed marked-facet alternative, is the quantitative gap.

The scalar objective on candidate supports has an additional exact simplification. Let $L = N \setminus G$ be the full nongainer set, choose $H \subset L$, put $P = L \setminus H$, and retain $p = |P|$, $g = |G|$, $h = |H|$. Substitution in the maximal-proper budget gives

$$U_{g-1} - L_H = \frac{n}{g+h} z(H) - x(L) + \frac{p - g + 1}{g} x(G) + \frac{(g-1)(g+h) - hp}{g(g+h)} z(G). \tag{CRS}$$

Lemma 48 (cardinality-ranked support)**.** *Within any family of admissible nongainer blocks of one fixed cardinality h, a block maximizes the direct scalar budget if and only if it maximizes $z(H)$.*

Every term in (CRS) except $nz(H)/(g+h)$ is fixed at cardinality h, and that coefficient is positive. Thus the staged Lmin packet decides which labelled supports are admissible, while Lmax ranks equal-size candidates by one ordinary coalition sum. What remains is an existence/threshold theorem across cardinalities. One cannot simply take the globally largest Lmax coordinates, since their block need not occur as the nongainer trace of a cycle-containing staged support.

There is a stronger cross-cardinality form. Put $T = x(N) = z(N)$, $D = z(G) - x(G)$, and $Q(P) = \sum_{i \in P} (T/n - z_i) = pT/n - z(P)$. Total balance reduces (CRS) to

$$U_{g-1} - L_H = \frac{pT - nz(P)}{n - p} - \frac{p+1}{g} D = \frac{n}{n-p} Q(P) - \frac{p+1}{g} D. \tag{EDQ}$$

Lemma 49 (excluded-block deficit quota)**.** *A staged-admissible nongainer support $H = L \setminus P$ closes the maximal-proper scalar budget if and only if*

$$Q(P) \ge \frac{(n-p)(p+1)}{ng} D. \tag{EDQ2}$$

On the carrier-cycle arm, (CCC) says that the cardinality base already reaches the target, so strict scalar closure is stronger than the target requires. Let $\tau_n = n/\lfloor n^2/4 \rfloor$ and retain the surplus

of the carrier cardinality base. Since the maximal-proper weight is $g/(p+1)$, (EDQ) gives the exact nonlow target reserve

$$\begin{aligned}\mathcal{R}_{\tau_n}(P) &= \frac{g}{p+1}(U_{g-1} - L_H) + \left(1 - \frac{gp}{(n-p)(p+1)} - \tau_n\right) T \\ &= (1-\tau_n)T - D - \frac{gn}{(n-p)(p+1)} z(P).\end{aligned} \tag{EDQ5}$$

Lemma 50 (target-reserve quota)**.** *If $\mathcal{R}_{\tau_n}(P) \geq 0$, then the maximal-proper direct identity reaches the target without requiring the strict scalar quota (EDQ2).*

Indeed, the first line is the normalized comparison–residual contribution plus cardinality slack; substitute $Q(P) = pT/n - z(P)$ to obtain the second. The omitted low-Lmax payment is nonnegative. Whenever the cardinality base is at least the target—in particular, under the carrier-cycle hypotheses of (CCC)—this criterion is no stronger than (EDQ2), and is strictly weaker when the base surplus is positive. Without that base hypothesis the two tests need not be ordered. The lemma still requires selection of an admissible support and the proved comparison premise.

This subsumes cardinality ranking and removes the opaque comparison across block sizes: the excluded nongainers must carry a specified Lmax deficit below the common mean. On a fixed labelled selection cell one can decide the direct carrier arm by scanning the complete staged Lmin support rows containing the cycle and testing (EDQ2) on their excluded blocks. This is a relative support-menu algorithm, not a full exact-game-cone facet or a polynomial full-cone enumeration theorem.

There is also an exact marginal law. If $P_{p+1} = P_p \dot\cup \{i\}$ and E_p is the quota surplus in (EDQ2), then

$$E_{p+1} - E_p = T/n - z_i - \frac{n-2p-2}{ng} D. \tag{EDQ4}$$

The required marginal deficit decreases linearly with p and becomes nonpositive after the midpoint. This explains why some six-player cells require an excluded pair even though neither available singleton closes, whereas adding above-mean Lmax mass destroys the larger excluded blocks in the seven-player control. The admissible support family need not itself be nested.

The rooted and direct currencies admit an exact common interface. Fix a root r and a root-containing fractional packing u supported on x-tight coalitions. Put

$$\delta = \sum_S u_S - 1, \qquad q_i = 1 - \sum_{S \ni i} u_S \quad (i \neq r), \tag{RCA1}$$

and assume the offroot loads are at most one. For any $A \subseteq N \setminus \{r\}$ set $P_S = A \setminus S$. Coordinate counting and core feasibility give

$$\sum_S u_S z(P_S) = \delta z(A) + \sum_{i \in A} q_i z_i,$$

$$\sum_S u_S |P_S| = \delta |A| + \sum_{i \in A} q_i, \tag{RCA2}$$

$$\sum_{i \neq r} q_i (x_i - z_i) \geq \delta (x_r - z_r). \tag{RCA3}$$

Indeed, player $i \in A$ is absent with packing mass $(1+\delta) - (1-q_i) = \delta + q_i$. Also $z(S) - x(S) \geq 0$ on every supported x-tight coalition; summing and using $x(N) = z(N)$ proves (RCA3). Thus these identities require no facet classification.

There is a fixed-cardinality reserve consequence. Suppose all supported rows contain the selected nongainer carrier, are proper on the gainer side, and have $|P_S| = p$. Proper x-tightness supplies the comparison premise, and with $c_p = gn/[(n-p)(p+1)]$,

$$\sum_S u_S \mathcal{R}_{\tau_n}(P_S) = (1+\delta)((1-\tau_n)T - D) - c_p \left(\delta z(A) + \sum_{i \in A} q_i z_i \right). \tag{RCA4}$$

Nonnegativity of the right side therefore selects a target-reaching row. This is a literal coupling of rooted deficit mass to excluded-block reserve, but it is not yet the universal handover: a marked price optimizer need not be supported on proper rows of one complement cardinality, and a negative average does not imply an affordable facet.

The cost of repairing an improper packing row can nevertheless be separated exactly into its two currencies. To avoid confusing tight-pair labels with payoff gainers, put $\Gamma_r = G$ and $B_r = B$ in the notation of (LPC), and write $\pi_r(S) = w(S \setminus \{r\})$ for the labelled-price coefficient of a rooted row. For any two rooted rows S, C,

$$\pi_r(S) - \pi_r(C) = K\big(|(C \setminus \{r\}) \cap \Gamma_r| - |(S \setminus \{r\}) \cap \Gamma_r|\big) + |(S \setminus \{r\}) \cap B_r| - |(C \setminus \{r\}) \cap B_r|. \tag{PXC1}$$

This is the exact labelled-price charge; it depends on tight rooted-pair labels, not merely on payoff strata. On the reserve side, let $P^+ = P \,\dot\cup\, \{i\}$, $p = |P| \le n-2$, and keep the selected gainer and nongainer strata fixed. Direct subtraction in (EDQ5) gives

$$\mathcal{R}_{\tau_n}(P^+) - \mathcal{R}_{\tau_n}(P) = \frac{gn\{(n-2p-2)z(P) - (n-p)(p+1)z_i\}}{(n-p)(p+1)(n-p-1)(p+2)}. \tag{PXC2}$$

Deleting payoff gainers from a coalition does not change its excluded nongainer block and hence costs no reserve; deleting an additional nongainer changes reserve by exactly (PXC2). Thus (PXC1)–(PXC2) reduce properization to an explicit price-versus-reserve charge, but do not prove that a charged replacement row is present in the tight support.

In particular, on the pair-free branch $\Gamma_r = \varnothing$ and $\pi_r(S) = K + |S \setminus \{r\}|$. Hence every deletion has labelled-price charge exactly equal to its incidence size. Suppose a co-singleton tight row $S = N \setminus \{j\}$ omits a nongainer j, has $\mathcal{R}_{\tau_n}(\{j\}) \ge 0$, and admits a carrier-preserving proper tight subset C. If $S \setminus C$ contains only payoff gainers, then C has the same reserve and closes the direct target. If it additionally contains exactly one nongainer i, it closes if and only if the right side of (PXC2), added to $\mathcal{R}_{\tau_n}(\{j\})$, is nonnegative. At $n = 6, p = 1$ the increment is

$$\frac{g}{10}\big(z_j - 5z_i\big). \tag{PXC3}$$

This conditional properization lemma is dimension-free before the final specialization. Its hypotheses make explicit the two facts still missing from a universal theorem: selection of the tight subset and payment of any additional omitted nongainer.

Exactness removes the apparent crossing obstruction completely. More generally, let S, C be tight at one core point x and suppose $S \cup C = N$. Put $I = S \cap C$. Core feasibility gives $v(S) + v(C) = v(N) + x(I) \ge v(N) + v(I)$. By exactness, however, there is a core point y with $y(I) = v(I)$, and then $v(S) + v(C) \le y(S) + y(C) = v(N) + v(I)$. Hence equality holds throughout:

$$v(S) + v(C) = v(N) + v(I), \qquad x(I) = v(I), \qquad z(I) - x(I) = [z(S) - x(S)] + [z(C) - x(C)] \ge 0. \tag{PXC4}$$

Thus tight sets covering the grand coalition intersect tightly in an exact game. In particular, intersecting any proper carrier row with a tight co-singleton preserves tightness, the carrier, and

properness, and automatically supplies the comparison premise. This elementary exact-game uncrossing fact is stated within the standard exact-cone framework (Csóka et al., 2011; Studený and Kratochvíl, 2022); no novelty claim is made for it. After this reduction the only remaining charge is the telescoped target-reserve change in (PXC2), not a comparison-slack defect.

That reserve has a useful ordered-cut lower certificate. Let $B \subseteq H$ be a retained block of b nongainers with $z_i = M$ for $i \in B$, put $K = H \setminus B$, suppose $z_i \le M^- \le M$ for every $i \in P$, and put $p = |P|$, $g = |G|$, and

$$c_p = \frac{gn}{(n-p)(p+1)}.$$

Since $D = T - z(H) - z(P) - x(G)$, the target reserve admits the exact rewrite

$$\mathcal{R}_\tau(P) = -\tau T + x(G) + z(H) + (1 - c_p)z(P). \tag{PXC5}$$

The retained block gives $z(H) = bM + z(K)$. If $c_p \le 1$, the final term is nonnegative; if $c_p > 1$, then $z(P) \le pM^-$. Hence

$$\mathcal{R}_\tau(P) \ge bM + z(K) + x(G) - \tau T - p(c_p - 1)_+ M^-. \tag{PXC6}$$

If M^- bounds every nongainer outside B, dropping $z(K) \ge 0$ gives the root-block-only corollary

$$\mathcal{R}_\tau(P) \ge bM + x(G) - \tau T - p(c_p - 1)_+ M^-. \tag{PXC7}$$

Relative to the still coarser substitution $M^- = M$, the next Lmax stratum supplies the exact additional credit

$$p(c_p - 1)_+(M - M^-). \tag{PXC8}$$

This lemma is dimension-free algebra, but nonnegativity of its right-hand side is a selection obligation, not a theorem for every carrier row. The importance of (PXC6)–(PXC8) is that they retain both sides of the ordered cut: the additional high-z nongainers kept in K and the low-z price of the excluded block P. This is the same colour of information exposed by the sparse Lmax pinning packet.

The complementary Lmin-side certificate comes from the comparison payment itself. Let I be the selected tight row, let $A = G \setminus I$ be its omitted gainers, put $\ell(P) = x(P) - z(P)$ and $\Delta(A) = z(A) - x(A)$. Total balance gives

$$\sigma_I := z(I) - x(I) = \ell(P) - \Delta(A) \ge 0. \tag{PXC9}$$

Substituting $z(P) = x(P) - \ell(P)$ in the target reserve therefore yields

$$\mathcal{R}_\tau(P) = (1 - \tau)T - D - c_p x(P) + c_p \Delta(A) + c_p \sigma_I, \tag{PXC10}$$

and in particular

$$\mathcal{R}_\tau(P) \ge (1 - \tau)T - D - c_p x(P) + c_p \Delta(A). \tag{PXC11}$$

Thus the same row has two independently useful lower certificates: (PXC6) prices its Lmax ordered cut, while (PXC11) prices the Lmin mass of its excluded block against the gain deliberately omitted by properness. More importantly, (PXC10) says that the entire difference between the latter certificate and the true reserve is the normalized Lmax comparison payment; there is no third residual currency. These are not independent proof systems—both are algebraic consequences of the same selected tight row—but they expose the two currencies that a genuine Lmin/Lmax gluing theorem must couple.

Their exact coupling is a minimum-residual identity. Take $M^- = \max_{i\in P} z_i$ in (PXC6), and define

$$E_I^x = c_p\sigma_I, \qquad E_I^z = \begin{cases} (1-c_p)z(P), & c_p \le 1, \\ (c_p-1)[pM^- - z(P)], & c_p \ge 1. \end{cases}$$

Then (PXC5), (PXC6), and (PXC10) give, exactly,

$$\max\{\text{PXC6}(I), \text{PXC11}(I)\} = \mathcal{R}_\tau(P) - \min\{E_I^z, E_I^x\}. \tag{PXC12}$$

Both errors are nonnegative. Consequently the selected-row max-certificate is equivalent to the single hybrid inequality $\mathcal{R}_\tau(P) \ge \min\{E_I^z, E_I^x\}$. This is stronger than mere nonnegativity of the reserve and is the precise load-bearing statement for a minimum-deletion exchange proof. It also prevents a proof from silently combining two error terms on a fictitious row; the outer existence quantifier is handled separately below.

For the corrected disjunction, however, the two optimizations themselves do decouple. Put $K_\tau = (1-\tau)T - D = O - \tau T$ and define the two row prices

$$Z_I = z(P) + (c_p - 1)_+ pM^-, \qquad X_I = c_p[z(P) + \sigma_I] = c_p[x(P) - \Delta(A)].$$

Direct substitution gives

$$\text{PXC6}(I) = K_\tau - Z_I, \qquad \text{PXC11}(I) = K_\tau - X_I. \tag{PXC13}$$

Consequently, for any nonempty admissible row family $\mathcal{F}$, if $Z_* := \min_{I\in\mathcal{F}} Z_I$ and $X_* := \min_{I\in\mathcal{F}} X_I$, then

$$\max_{I\in\mathcal{F}} \max\{\text{PXC6}(I), \text{PXC11}(I)\} = K_\tau - \min\{Z_*, X_*\}. \tag{PXC14}$$

Thus a row closes precisely when the cheaper of the independently optimized Lmax cut price and Lmin omission price is at most the global overlap surplus. One does not add certificates from different rows: a minimizer of the cheaper price supplies the actual selected row. This turns the sufficient max-certificate selection step into the scalar hybrid-price inequality $\min\{Z_*, X_*\} \le K_\tau$ and gives a direct interface to the two LP duals.

At fixed deletion cardinality these two prices are ordinary maximum-weight row problems. Let $L = N \setminus G$, let $\mathcal{F}_p$ be the carrier-containing proper x-tight rows with $|L \setminus I| = p$, and put $m_I = \max_{i\in L\setminus I} z_i$. Define

$$S_z(I) = z(I\cap L) - (c_p - 1)_+ pm_I, \qquad S_x(I) = x(I\cap L) - \Delta(I\cap G).$$

Then direct complementation gives

$$Z_* = z(L) - \max_{I\in\mathcal{F}_p} S_z(I), \qquad X_* = c_p\left[x(L) - D - max_{I\in\mathcal{F}_p} S_x(I)\right]. \tag{PXC15}$$

For $c_p \le 1$ the Lmax score is simply retained nongainer z-mass. For $c_p > 1$, stratifying by the highest excluded Lmax level makes the additional term constant on each stratum, so (PXC15) is still a finite sequence of linear maximum-weight row problems. This proves an exact algorithm on an explicit tight menu and exposes two support-function duals. Their extreme inequalities are relative selection-cell facets; promotion to facets of the full exact-game cone still requires the semi-balanced indecomposability test. Scanning every nonempty cardinality class gives

$$Z_*^{\text{all}} = \min_p Z_{*,p}, \qquad X_*^{\text{all}} = \min_p X_{*,p}. \tag{PXC16}$$

By (PXC14), some admissible carrier row has a nonnegative two-colour max-certificate if and only if $\min\{Z_*^{\mathrm{all}}, X_*^{\mathrm{all}}\} \le K_\tau$. Minimum deletion is a particularly rigid finite selector, but it is only a sufficient subfamily; the dimension-free proof should use the full scan and may close at a larger deletion size.

There is a dimension-free equality consequence. If $K_\tau = 0$, then every Z_I, X_I is nonnegative, and

$$\max\{\mathrm{PXC6}(I), \mathrm{PXC11}(I)\} \ge 0 \quad \Longleftrightarrow \quad z(P) = 0. \tag{PXC17}$$

Indeed either closing price must vanish, hence $z(P) = 0$. Conversely, nonnegativity of z makes $z(P) = 0$ imply $M^- = 0$, so $Z_I = 0$ and PXC6 is exactly zero. Thus equality selection on the carrier-direct arm reduces to excluding only zero-Lmax nongainers; no gainer-label enumeration remains. Combined with (CCC), this recovers the existing unique six-player carrier-direct equality parameters, while leaving the rooted seven-player equality arm separate.

There is an exact exchange law along every co-singleton-generated deletion chain. Let I be an admissible row with excluded nongainer block P of size p, let $j \in (I \cap L)$ lie outside the retained carrier, and suppose $N \setminus \{j\}$ is x-tight. Since $I \cup (N \setminus \{j\}) = N$, (PXC4) makes $J = I \setminus \{j\}$ tight; it remains admissible, has excluded block $P \cup \{j\}$, and has the same omitted-gainer block. Put

$$m = \max_{i \in P} z_i, \qquad m' = \max\{m, z_j\}, \qquad a_p = p(c_p - 1)_+.$$

Then direct subtraction in the two price formulas gives

$$\begin{aligned} X_J - X_I &= (c_{p+1} - c_p)[z(P) + \sigma_I] + c_{p+1} x_j, \\ Z_J - Z_I &= z_j + a_{p+1} m' - a_p m, \qquad \text{(PXC18)} \\ \frac{c_{p+1}}{c_p} - 1 &= \frac{2p + 2 - n}{(p+2)(n-p-1)}. \qquad \text{(PXC19)} \end{aligned}$$

Thus an Lmin-colour deletion cannot lower price once $2p + 2 \ge n$. An Lmax-colour deletion lowers price exactly when $z_j < a_p m - a_{p+1} m'$. Before the midpoint, these are explicit low-payoff thresholds; after the coefficient turns upward, nonnegativity prevents the corresponding Lmin improvement. The lemma localizes every improving step on an available tight-co-singleton chain. It does not assert that every two admissible rows lie on one such chain, or that the required co-singleton is tight.

The contrapositive gives the exact local obstruction needed by the separator arm. Choose a global X-price minimizer I_X and a global Z-price minimizer I_Z (they need not coincide). For every co-singleton deletion available from the relevant row, (PXC18) and optimality give the following, with p, P, m interpreted for the minimizer in the corresponding line:

$$\begin{aligned} x_j &\ge \frac{n - 2p - 2}{(n-p)(p+1)}[z(P) + \sigma_{I_X}] \qquad (2p + 2 < n), \\ z_j + a_{p+1} \max\{m, z_j\} &\ge a_p m. \qquad \text{(PXC20)} \end{aligned}$$

In particular, when $z_j \le m$ the second inequality is $z_j \ge (a_p - a_{p+1})m$. Repeated strict improvement terminates because each deletion increases p, so each colour has a finite co-singleton descent to a row satisfying (PXC20). Therefore simultaneous failure of the two global price tests produces two explicit exchange-stable rows and these positive player-price inequalities. What remains is to aggregate those payments by the staged Lmax separator, or to show that a shortage of available tight co-singletons exposes one of the rooted facet packets. This is a reduction of the missing theorem, not its proof.

The X-stability payment already lies in an elementary balanced class. Let R_X be the retained noncarrier nongainers whose co-singletons $C_j = N \setminus \{j\}$ are x-tight at an X-price minimizer, and assume $2p+2 < n$. Total balance and x-tightness give $z(C_j) - x(C_j) = x_j - z_j$. Summing the first line of (PXC20) therefore yields

$$z(R_X) + \sum_{j \in R_X} [z(C_j) - x(C_j)] \geq |R_X| \frac{n-2p-2}{(n-p)(p+1)} [z(P) + \sigma_{I_X}]. \tag{PXC21}$$

The underlying incidence packet is simply

$$\sum_{j \in R_X} \mathbf{1}_{C_j} + \sum_{j \in R_X} e_j = |R_X| \mathbf{1}_N.$$

Thus each required exchange payment splits exactly into retained Lmax mass and comparison slack. The aggregate is the decomposable sum of the balanced partition identities $\mathbf{1}_{N \setminus \{j\}} + e_j = \mathbf{1}_N$; it is not the distinct indecomposable complement-star facet family proved earlier, and it cannot itself be a new indecomposable exact-cone facet. Consequently any genuinely new facet needed here must enter through the remaining supply-or-sufficiency obstruction, not through (PXC21) itself. Inequality (PXC21) does not prove either that enough co-singletons are tight or that the available atoms have sufficient total weight to close the target.

For the direct-versus-separator theorem, the exact reserve price is weaker than either certificate price and should be optimized first. Define

$$Y_I = c_p z(P).$$

Then (EDQ5) and (PXC13) give

$$\mathcal{R}_\tau(P) = K_\tau - Y_I, \qquad Y_I \leq Z_I, \qquad Y_I \leq X_I. \tag{PXC22}$$

The second inequality uses $\sigma_I \geq 0$. For the first, if $c_p \leq 1$ then $Z_I = z(P) \geq c_p z(P)$; if $c_p \geq 1$, then $p m_I \geq z(P)$ gives $Z_I = z(P) + (c_p - 1) p m_I \geq c_p z(P)$. At fixed cardinality and over the full family, therefore,

$$\begin{aligned} Y_{*,p} &= c_p \left[z(L) - \max_{I \in \mathcal{F}_p} z(I \cap L) \right], \\ Y_*^{\mathrm{all}} &= \min_p Y_{*,p} \leq \min\{Z_*^{\mathrm{all}}, X_*^{\mathrm{all}}\}, \\ \text{a direct row closes} \quad &\Longleftrightarrow \quad Y_*^{\mathrm{all}} \leq K_\tau. \end{aligned} \tag{PXC23}$$

Thus failure of the direct price automatically implies failure of both certificate colours, but the converse need not hold. The actual hard branch for the direct-or-separator theorem is $Y_*^{\mathrm{all}} > K_\tau$; (PXC16) is a stronger sufficient certificate conjecture, not the necessary direct target. Formula (PXC23) is a pure Lmax support function evaluated on the complete Lmin-tight menu, giving the cleanest interface to the sparse Lmax dual.

The co-singleton exchange becomes correspondingly simpler:

$$Y_J - Y_I = (c_{p+1} - c_p) z(P) + c_{p+1} z_j. \tag{PXC24}$$

If I_Y is a global direct-price minimizer with $2p+2 < n$, every available deletion therefore satisfies, with R_Y denoting all such retained nongainers,

$$z_j \geq \frac{n-2p-2}{(n-p)(p+1)} z(P), \qquad z(R_Y) \geq |R_Y| \frac{n-2p-2}{(n-p)(p+1)} z(P). \tag{PXC25}$$

Unlike (PXC21), this payment is entirely retained Lmax mass and hence enters overlap directly. What remains is still a supply-or-sufficiency theorem: neither (PXC23) nor (PXC25) proves that enough co-singleton deletions exist or that their aggregate mass forces $Y_*^{\mathrm{all}} \le K_\tau$.

When the co-singletons do exist, their deletions are jointly compatible, not merely pairwise available. Let R be any set of retained noncarrier nongainers of an admissible row I such that every $N \setminus \{j\}$, $j \in R$, is x-tight. Repeatedly applying (PXC4) gives

$$I \setminus Q \text{ is } x\text{-tight and admissible for every } Q \subseteq R. \tag{PXC26}$$

Indeed, after any earlier deletions the current row still contains the next j, so its union with $N \setminus \{j\}$ is again N. Thus the available co-singletons generate a full Boolean deletion cube. If $z_{(1)} \le \cdots \le z_{(|R|)}$ are their Lmax payoffs, then for every $0 \le q \le |R|$,

$$\min_{\substack{Q \subseteq R \\ |Q|=q}} Y_{I\setminus Q} = c_{p+q}\left[z(P) + \sum_{t=1}^{q} z_{(t)}\right], \qquad \min_{Q \subseteq R} Y_{I\setminus Q} = \min_{0 \le q \le |R|} c_{p+q}\left[z(P) + \sum_{t=1}^{q} z_{(t)}\right]. \tag{PXC27}$$

For fixed q the coefficient is constant, so the exchange argument is just the order-statistic choice of the q cheapest players. This proves an exact one-dimensional direct-price algorithm on every co-singleton cube. It does not show that an arbitrary admissible row has a cube large enough to reach the coefficient midpoint. The six-player full-menu replay checks all 157 available cubes and 314 vertices/order-statistic layers; every archived cube has dimension one, so higher-dimensional coverage rests on the human uncrossing proof rather than that finite control.

The shortage side also has an exact geometric meaning. For each player put

$$u_j := T - v(N \setminus \{j\}).$$

Every core allocation y satisfies $y_j \le u_j$. Exactness at $N \setminus \{j\}$ supplies a core allocation w^j with $w^j(N \setminus \{j\}) = v(N \setminus \{j\})$, and therefore

$$u_j = \max_{y \in \mathrm{Core}(v)} y_j = w^j_j, \qquad N \setminus \{j\} \text{ is } x\text{-tight} \iff x_j = u_j. \tag{PXC28}$$

Consequently every unavailable co-singleton direction has the positive upper slack

$$\kappa_j := u_j - x_j = x(N \setminus \{j\}) - v(N \setminus \{j\}) > 0$$

and comes with an exact endpoint witness w^j. Thus cube shortage is not an unstructured support failure: it is a collection of coordinate-exposing exactness witnesses. Formula (PXC28) does not show that the separately chosen witnesses w^j can be aggregated with common normalized weights, nor that their supports form an indecomposable semi-balanced packet. That aggregation-or-separator step remains the missing theorem.

On the six-player full tight menus the exact replay checks (PXC28) in 497 retained noncarrier nongainer occurrences: 157 are upper-pinned and 340 have positive upper slack. The extracted staged supports give 7 pinned and 338 slack occurrences among 345 checks. These multiplicity counts are finite diagnostics; exactness proves (PXC28) in every dimension. Every one of the 340 full-menu shortage sets (and all 338 extracted-support sets) is a singleton, so exactness gives a common endpoint witness in every archived instance and no positive (PXC30) separator gap appears. Of the 340 full-menu slack occurrences, 300 lie in the Lmax top stratum and 40 lie below it. Thus this archive checks the one-direction endpoint arm but supplies no nontrivial finite test of multi-endpoint compatibility.

The endpoint witnesses admit a genuine aggregate alternative. For a nonempty set R of missing directions define the co-singleton slack $s_j(y) := y(N \setminus \{j\}) - v(N \setminus \{j\}) = u_j - y_j$ and

$$\epsilon_R := \min_{y \in \mathrm{Core}(v)} \max_{j \in R} s_j(y).$$

Compact minimax duality gives

$$\epsilon_R = \max_{\beta \in \Delta(R)} \min_{y \in \mathrm{Core}(v)} \sum_{j \in R} \beta_j s_j(y). \tag{PXC29}$$

In particular, $\epsilon_R = 0$ if and only if the endpoint faces $\{y : y_j = u_j\}$, $j \in R$, have a common core point. If their intersection is empty, then $\epsilon_R > 0$ and an optimizing β, together with core LP duality, supplies nonnegative coalition weights λ_S and a free scalar γ such that

$$\sum_S \lambda_S \mathbf{1}_S + \gamma \mathbf{1}_N = \sum_{j \in R} \beta_j \mathbf{1}_{N \setminus \{j\}}, \qquad \sum_S \lambda_S v(S) + \gamma T - \sum_{j \in R} \beta_j v(N \setminus \{j\}) = \epsilon_R > 0. \tag{PXC30}$$

For fixed β an extreme dual needs at most n positive coalition weights. Complementary slackness makes those coalitions tight at a core point minimizing the weighted endpoint slack. Thus missing cube directions either aggregate into one common endpoint witness or expose a sparse exact core separator with a positive normalized payment; separate exactness witnesses are no longer the endpoint of the argument.

This is not yet the full sparse-Lmax arm of the carrier theorem. The coalitions in (PXC30) are tight at the minimizing core point and need not be tight at z or belong to the chosen first-stage Lmax over-cover. The remaining gluing obligation is to transfer the positive payment to that over-cover, or to show that failure of transfer yields one of the already classified rooted or mixed-witness packets. Consequently (PXC29)–(PXC30) prove an endpoint-face supply-or-separator theorem, not the target overlap inequality.

On the common-witness arm there is an exact projection theorem parallel to the rooted mixed-witness projection. Suppose a nonempty proper R has a common endpoint witness, put $U_R = \sum_{j \in R} u_j$, and let $\mathcal{D}$ be any family of selection-tagged requirements $a_d \le w(S_d)$. For $\rho_d \ge 0$ write $p_i(\rho) = \sum_{d : i \in S_d} \rho_d$. A nonnegative total-$T$ vector w with $w_j = u_j$ for every $j \in R$ and satisfying all the requirements exists if and only if

$$\sum_d \rho_d a_d \le \sum_{j \in R} p_j(\rho) u_j + (T - U_R) \max_{i \notin R} p_i(\rho) \tag{PXC31}$$

for every $\rho \ge 0$. Necessity is weighted summation. Sufficiency follows by separating from the down-set generated by the endpoint-fixed simplex; the right side is exactly its support function.

After normalizing $p_i \le 1$ for $i \notin R$, extreme packings use at most $n - |R|$ positive requirement columns, apart from rays supported wholly on R. With $q_i = 1 - p_i$ off R and $\delta_j = p_j - 1$ on R, their incidence completion is

$$\sum_d \rho_d \mathbf{1}_{S_d} + \sum_{i \notin R} q_i e_i - \sum_{j \in R} \delta_j e_j = \mathbf{1}_N. \tag{PXC32}$$

For purely x-tagged requirements, (PXC31) is equivalently

$$\sum_{j \in R} \delta_j (x_j - u_j) \le \sum_{i \notin R} q_i x_i. \tag{PXC33}$$

This exposes the exact sign of the shortage currency: positive root excess δ_j is helped by $x_j - u_j = -\kappa_j$, whereas negative root excess must be paid. Equations (PXC31)–(PXC33) give a finite, linear-support algorithm for the common-endpoint mixed-packet projection. They do not prove that its branch optimum reaches the overlap target. Moreover the multi-exception incidence system (PXC32) must pass the full Studený–Kratochvíl indecomposability test before any factor is called an exact-cone facet.

The endpoint alternative couples directly to the sparse Lmax over-cover. Let (λ, η, α) satisfy (OC), so α is a probability on the Lmax top block and $\alpha \cdot y \geq M$ for every core point y. If a nonempty proper R has a common endpoint witness, set $a_R := \max_{i \notin R} \alpha_i$. Evaluating the support function of its endpoint-fixed simplex gives

$$M \leq \sum_{j \in R} \alpha_j u_j + (T - U_R) a_R = a_R T + \sum_{j \in R} (\alpha_j - a_R) u_j. \tag{PXC34}$$

In particular, if $\theta = \max_i \alpha_i$, then

$$M \leq \theta T - \sum_{j \in R} (\theta - \alpha_j) u_j. \tag{PXC35}$$

Thus fixing an endpoint outside the maximally loaded over-cover coordinates strictly reduces the crude dispersion ceiling $M \leq \theta T$.

The positive-separator arm has a priced version of the same inequality. Let β attain (PXC29), put $d_j := \alpha_j - a_R$, and suppose $d_j \geq 0$ on R. With

$$q_R := \min_{\substack{j \in R \\ \beta_j > 0}} \frac{d_j}{\beta_j},$$

choose a core point y^* minimizing the β-weighted endpoint slack. Writing $s_j = u_j - y^*_j$ and placing the remaining mass on a largest off-R over-cover coordinate yields

$$M \leq \alpha \cdot y^* \leq a_R T + \sum_{j \in R} d_j u_j - q_R \epsilon_R. \tag{PXC36}$$

Indeed, the same mass calculation before the final inequality gives $a_R T + \sum_j d_j u_j - \sum_j d_j s_j$, and $\sum_j d_j s_j \geq q_R \sum_j \beta_j s_j = q_R \epsilon_R$. This is a dimension-free coupling between the sparse Lmax dual and the endpoint-face Farkas separator. Its alignment hypothesis is essential: if some $d_j < 0$, endpoint slack at that coordinate can increase rather than pay the ceiling bound. The remaining carrier theorem must select an α-dominant shortage block, close directly, or turn this precise alignment failure into a priced rooted/mixed packet.

The six-player full-menu replay verifies (PXC34)–(PXC35) on all 340 common-endpoint shortage rows. The shortage coordinate is α-dominant on 276 rows; on the other 64, (PXC35) is a strict improvement over $M \leq \theta T$. The extracted-support counts are 276 and 62 among 338 rows. Only two of the 340 full-menu coupling bounds are tight; 338 have positive upper-bound slack, and (PXC34) coincides with the weaker (PXC35) on every archived row. This exact finite handover checks both sides of the alignment split, but also shows that the new ceiling bound alone is usually not decisive for the overlap target.

The positive-gap arm has a sharp exact seven-player sign control. Take the lower envelope of $(1, 0, 0, 0, 0, 0, 0)$ and $(0, 1, 0, 0, 0, 0, 0)$ and let $R = \{0, 1\}$. Its core is the two-player simplex with five null coordinates, the two endpoint faces are disjoint, and $\epsilon_R = 1/2$. Here $\alpha = \beta = (1/2, 1/2, 0, 0, 0, 0, 0)$ on their respective supports, $a_R = 0$, and $q_R = 1$, so (PXC36) is the equality

$M = 1/2$. An explicit (PXC30) dual has $\gamma = 1/2$ and weight $1/2$ on each of the five dummy singletons. This is a decomposable dummy-extension control for the signs and separator price, not a carrier hard case or a new facet.

A tempting quantitative corollary is false even on the six-player archive. Combining the upper ceiling (PXC35) with the ceiling-root interpolation bound (CRI), and minimizing (CRI) over the entire permitted interval $T/n \le M \le M_{\mathrm{PXC35}}$, reaches the target on only 28 of the 340 full-menu shortage rows. Using the actual ceiling reaches it on 165. Thus neither (PXC35) nor its strictness can replace the labelled direct-price or mixed-packet argument; further optimization of this scalar ceiling is a local maximum, not the remaining proof route.

There is nevertheless a canonical way to choose the endpoint shortage at the actual direct-price optimum. Let I globally minimize Y_I, let H be its retained carrier cycle, let $C = (I \cap L) \setminus H$, and assume $2p + 2 \le n$. Put

$$\rho_p := \frac{n - 2p - 2}{(n - p)(p + 1)}, \qquad R^- := \{j \in C : z_j < \rho_p z(P)\}, \qquad C^+ := C \setminus R^-.$$

Then (PXC24) and global optimality imply

$$j \in R^- \Longrightarrow N \setminus \{j\} \text{ is not } x\text{-tight.} \tag{PXC37}$$

Moreover, by the definition of C^+,

$$z(C^+) \ge |C^+|\rho_p z(P). \tag{PXC38}$$

Thus R^- feeds exactly into (PXC28)–(PXC36); if it is empty, every retained noncarrier already pays the direct exchange price. Writing $h = |H|$, using $z_i = M$ on H, and expanding the actual reserve gives

$$\mathcal{R}_\tau(P) \ge hM + x(G) - \tau T + z(R^-) + [1 - c_p + |C^+|\rho_p]z(P). \tag{PXC39}$$

This is the promised direct-price handover: a negative right side is the exact residual amount that the endpoint separator or a labelled mixed packet must pay. It does not disappear into a scalar ceiling.

The discarded amount is itself an exact nonnegative labelled currency:

$$\Lambda_I := z(C^+) - |C^+|\rho_p z(P) \ge 0, \qquad \mathcal{R}_\tau(P) = B_I + \Lambda_I, \tag{PXC40}$$

where B_I is the right side of (PXC39). Hence the remaining problem is not to invent another residual, but to certify this retained Lmax support-function excess by the sparse staged Lmax packet or a marked facet.

On the six-player archive every one of the 920 full-menu direct-optimal rows through the midpoint has $R^- = \varnothing$. The high-mass bound (PXC39) closes 682 rows and at least one optimum on 296 of the 415 cycles. The remaining 238 rows, spanning 119 cycles, require additional labelled currency. Thus the canonical endpoint arm is absent at the archived direct optima even though arbitrary nonoptimal rows have endpoint shortage. This prevents the endpoint theory from being mistaken for a complete selection proof and localizes the remaining finite obstruction to the labelled packet term.

All 238 negative-B_I rows have the single profile $(g, h, p, |C^+|) = (2, 2, 1, 1)$, and their ordered-cut error is zero, so (PXC40) is exactly the established PXC6/Lmax colour there. In 236 rows the retained currency player belongs to the first Lmax top stratum. The only two off-top rows are the same exact record and have $z(P) = 0$, so the direct price already vanishes. Thus every archived

positive-price residual is visible to the first-stage Lmax top support; a dimension-free proof of that visibility is still missing.

The relevant carrier closure is a preorder, not an arbitrarily chosen cycle. For positive x define

$$r \preceq_x c \quad \Longleftrightarrow \quad c \text{ belongs to every } x\text{-tight coalition containing } r. \tag{PXC41}$$

This relation is reflexive and transitive. Every strong equivalence class E is therefore a tight-face module at x: if a tight coalition contains one member, mutual reachability forces it to contain all of E. Carrier monotonicity in both directions also makes x constant on E. If z is constant on E—in particular, when E consists of Lmax-top nongainer roots—the established carrier-module theorem gives the exact alternative: an Lmax-tight coalition splits E and supplies strict comparison payment, or E contracts to one weighted role while preserving exactness, both selectors, total, and overlap.

Thus a top-stratum player contributing to Λ_I is automatically a maximum-overlap root, because it is a nongainer with $z_j = M$ and hence $x_j \geq M$. If it lies in the carrier component of H, (PXC41) enlarges the module before applying separator-or-contraction. This is dimension-free; whether the resulting separator is sufficiently priced remains the quantitative obligation.

The one-way case also has a canonical quotient normal form. Let E be a carrier strong component and let a root $j \notin E$ satisfy $j \preceq_x e$ for every $e \in E$, but not conversely. If no Lmax-tight coalition splits E, contract it to a role $*$ of weight $|E|$. The quotient retains

$$z_j = z_* = M, \qquad x_*/|E| \geq x_j \geq M, \qquad j \preceq_{\bar{x}} *. \tag{PXC42}$$

Here the last relation says that every quotient Lmin-tight coalition containing j contains $*$. Exactness, the weighted selectors, total, and overlap are preserved. Thus a one-way top-currency satellite either has an Lmax splitter, or reduces to a weighted locked-satellite quotient. This does not put that quotient under the four-role theorem unless at most three other roles remain.

This quotient has an immediate atomic-size cutoff. The role $*$ and its satellite contribute $W = |E| + 1$ atomic coordinates with overlap M, while nonnegativity and efficiency give $M \geq T/n$. Hence

$$\frac{O}{T} \geq \frac{W}{n} \geq \frac{4n}{n^2 - (n \bmod 2)} \tag{PXC43}$$

whenever $W \geq 4$ for even n, or $W \geq 5$ for odd n. Indeed the even target is $4/n$, while for odd n the required integer satisfies $W \geq 4n^2/(n^2-1) > 4$. Thus a surviving locked satellite has $|E| \leq 2$ in even dimension and $|E| \leq 3$ in odd dimension. The reduction does not bound the number of roles outside this small locked block.

Of the 236 archived top-currency occurrences, 222 join the carrier component. Ninety enlarged modules are also Lmax modules; their contraction leaves four weighted roles and is therefore closed by the weighted four-role theorem. The other 132 have an Lmax-tight splitter. In every one, some strict separator payment s satisfies $-B_I \leq s \leq \Lambda_I$, so $B_I + s \geq 0$ is an exact finite certificate. The remaining 14 occurrences, in seven records, are one-way carrier satellites: the currency root carries the two-cycle, but not conversely. These counts audit the handover and leave the one-way satellite as the next structural case; they do not prove a dimension-free lower bound on strict separator payment.

For the 14 archived satellites, four cycle modules have an Lmax splitter with payment in $[-B_I, \Lambda_I]$ and close directly. The other ten contract to the five-role weight profile $(2, 1, 1, 1, 1)$ satisfying (PXC42). This is the exact residual normal form, not a claim that weighted five-role locked satellites are already covered by Objective A.

The apparent five-role residue has a further dimension-free marked-row dichotomy. In the quotient let $*$ denote the contracted carrier component. Because $j \preceq_{\bar{x}} *$ is not reversible, some Lmin-tight coalition A contains $*$ and omits j. For every such row, and for every Lmax-tight row D which splits the locked block $\{*, j\}$, core feasibility gives

$$\delta_A := z(A) - x(A) \geq 0, \qquad \sigma_D := x(D) - z(D) \geq 0. \tag{PXC44}$$

The first inequality is strict unless A is a common hinge. The second is strict unless D is a common hinge; moreover it is necessarily strict when D contains j but omits $*$, since such a row cannot be Lmin-tight under $j \preceq_{\bar{x}} *$. Thus every one-way satellite exposes either positive comparison payment in one of the two colours or a labelled common hinge. This is an existence and sign theorem, not a lower bound on the payment and not by itself the target inequality.

There is, however, an immediate certificate on a labelled subarm. If a zero-payment marked row A consists only of Lmax-top players, then

$$\lambda_A = \frac{1}{|A|}, \qquad \eta = 0, \qquad \alpha_i = \frac{1}{|A|}\mathbf{1}_{\{i \in A\}} \tag{PXC45}$$

is a normalized Lmax top over-cover supported on a common hinge, with $\theta = 1/|A|$ and zero cross-surplus. This follows directly from $\lambda_A \mathbf{1}_A = \alpha$ and $x(A) = z(A) = |A|M$. The resulting spectral bound is exact certificate evidence; whether it reaches the target must still be checked unless a separate common-hinge theorem applies.

If the common-support arm fails, its deficiency dual makes the obstruction quantitative. Let $H = \arg\max z$, let (y, ϵ_*) be an optimum of (CSD-D), and let A be any common hinge. Dual feasibility and zero slack on A give $y(A) \leq 0$, while $y_i \geq \epsilon_*$ for $i \in H$. Therefore

$$-y(A \setminus H) \geq |A \cap H|\epsilon_*, \qquad y(H \setminus A) \geq |H \setminus A|\epsilon_*. \tag{PXC46}$$

Thus a common marked row which both meets and omits the top block has an exact alternative: either $\epsilon_* = 0$ and a common top over-cover exists, or its off-top labels absorb a negative dual price proportional to the number of retained top labels while every omitted top label carries positive price. This is a signed labelled separation witness at the common-support/ mixed-support interface. It is not yet a feasible polar of the full mixed-witness projection, a nonnegative mixed packet, a target-price bound, or an indecomposable exact-cone facet; converting its polar prices into one of those objects is the remaining gluing step.

The exact six-player replay checks (PXC46) on all 2,065 common rows in 374 records. Seventy-three records have positive common-support deficiency, and 21 common rows both meet and omit the top block there. A sharp price-transfer control is record 71 of the first archive: coalition 23 contains three of four top labels, $\epsilon_* = 1/6$, its off-top price is $-1/2 = -3\epsilon_*$, and the omitted top price is $1/6$. Only 11 of the 21 positive-transfer rows have any optimal over-cover row with the same top trace as the hinge. Record 17 of the second archive is a four-row control with support $(21, 26, 54, 60)$ and weights $(2/5, 1/5, 1/5, 1/5)$, none having the hinge's top trace. Hence a one-row nested shear is false even finitely; the facet conversion must retain the aggregate sparse over-cover.

The aggregate conversion is nevertheless canonical when it has positive root excess. Adjoin the tagged common hinge to the support of an optimal deficiency over-cover, choose a top root r, and maximize the root load p_r subject to all off-root loads $p_i \leq 1$. If $\delta = p_r - 1 > 0$, put $q_i = 1 - p_i$. Colour erasure followed by the proved rooted-pair shear gives

$$\sum_C \bar{\rho}_C \mathbf{1}_C + sum_{i \neq r} q_i \mathbf{1}_{\{r,i\}} - \left(\delta + \sum_{i \neq r} q_i\right) e_r = \mathbf{1}_N. \tag{PXC47}$$

This is a singleton-exceptional min-semi-balanced system supported on the full aggregate packet. Support minimality and the Studený–Kratochvíl indecomposability test determine whether it is a full exact-cone facet; no new criterion is being asserted.

Exactly 11 of the 21 positive-deficiency transfer rows have $\delta > 0$; the other ten are tied at zero for every top root. All eleven positive systems are support-minimal and indecomposable, independently replayed in exact Z3 linear arithmetic. They form three rooted support orbits with multiplicities $4, 3, 4$. Two orbits already occur in the five-orbit six-player residual atlas. The third has canonical support $(1, 3, 13, 49)$ and identity $\mathbf{1}_3 + \mathbf{1}_{13} + \mathbf{1}_{49} - 2e_0 = \mathbf{1}_N$: it is the saturated rooted star whose three off-root blocks have sizes $(1, 2, 2)$. Hence the positive aggregate arm lies entirely in the already proved saturated/clone-expanded rooted-star class. With $K = 4$, all eleven displayed markings also have positive margin $K\delta - Q$: their effective prices Q/δ are only 1 or 2. This settles finite scalar affordability, but not the payoff-stratum alignment required by the rooted price theorem. The tied ten, not facet novelty, are the remaining alignment cell. Every tied row is already closed by an earlier arm in the archive: eight have a direct target packet, three satisfy ceiling interpolation, and five have an affordable labelled root, with overlaps among those counts. Hence none is a new finite hard case. What remains universal is to prove that tied aggregate geometry forces an earlier arm, and to prove sufficient hybrid-overlap price for the positive aggregate facet; facet existence and scalar affordability alone do not supply that alignment.

The tied geometry has its own exact cover dual. For a top root r, let $\mathcal{D}$ be the aggregate tagged support and define

$$L_r = \max\left\{\sum_{C \ni r} \rho_C : \rho \geq 0, \sum_{C \ni i} \rho_C \leq 1 \ (i \neq r)\right\}.$$

Ordinary rational LP duality gives

$$L_r = \min\left\{\sum_{i \neq r} \beta_i : \beta \geq 0, \sum_{i \in C \setminus \{r\}} \beta_i \geq 1 \quad (C \in \mathcal{D},\, r \in C)\right\}. \tag{PXC48}$$

Thus positive rooted excess is exactly $L_r > 1$; the tied arm has a unit-cost fractional carrier cover for every top root occurring in the aggregate support. In fact the cover is automatically integral in the only sense needed here. If its total mass is one, every constraint has left side at most one and hence equality; every label with positive β_i therefore belongs to every rooted trace. Conversely, one common label gives the unit cover. Thus

$$L_r = 1 \quad \Longleftrightarrow \quad \bigcap_{C \in \mathcal{D} : r \in C} (C \setminus \{r\}) \neq \varnothing. \tag{PXC48$'$}$$

This is a dimension-free aggregate-carrier theorem. All 31 nonzero tied archive roots exhibit it. The qualification is load-bearing: the carrier belongs to every row of the selected aggregate support, not necessarily to every tight row of either full selection face, so common-module contraction does not yet follow. The next bridge is support reselection: either promote an aggregate carrier cycle to a full tight-face module, or use an escaping tight row as the positive rooted/separator packet.

Passing from one optimal deficiency support to the full Lmax-tight menu makes that reselection alternative exact. Fix a common hinge A and apply (PXC48) to the tagged family consisting of A and every proper z-tight coalition. Then at least one of the following holds:

(a) some $r \in A \cap H$ has $L_r > 1$, and an extreme optimizer gives a positive-root mixed packet with at most $n - 1$ tagged rows;

(b) a nontrivial z-carrier cycle $E \subseteq A \cap H$ is split by an x-tight row D, in which case $z(D) - x(D) > 0$;

(c) E is a tight-face module at both selections and contracts exactly to one weighted role.

We call this alternative (PXC49). Indeed, if every root in $A \cap H$ is tied, (PXC48′) supplies a carrier in every full z-tight row containing it and in A. Such a carrier is itself in H: otherwise the feasible tangent moving mass from the top root to the lower carrier lexicographically improves Lmax. Choosing one carrier per root gives a finite directed graph on $A \cap H$ and hence a cycle E. Every z-tight row meeting E contains all of it. If an x-tight row splits E, equality at z would contradict that module property, so its reverse comparison is strict. If no such row exists, E is also an x-module; Lmin optimality equalizes x on it, and the proved common-module contraction applies.

This is a dimension-free carrier/separator/contraction theorem at the marked hinge interface. It supplies structure and strict sign, not the missing lower bound on payment. In the ten formerly tied archive rows, full-menu reselection gives three positive packets, two strict Lmin splitters (the same record in two marked rows), and five common modules. Thus the structural trichotomy is exact on the finite residue; quantitative price/alignment is the remaining universal step.

The positive arm nevertheless has an exact quantitative conservation law. Delete every root-omitting row from a maximum-root-load optimizer (such rows do not increase the objective and only consume off-root capacity), and write p_i for its incidence loads, $q_i = 1 - p_i$ for $i \neq r$, and $\delta = p_r - 1 > 0$. If $\sigma_C = x(C) - z(C) \geq 0$ is the comparison payment of a supported z-tight row and $\Gamma = \sum_C \rho_C \sigma_C$, then the singleton completion evaluated at the two selections gives

$$\Gamma + \sum_{i \neq r} q_i (x_i - z_i) = \delta (x_r - z_r). \tag{PXC50}$$

This is just subtraction in $\sum_C \rho_C \mathbf{1}_C + \sum_{i \neq r} q_i e_i - \delta e_r = \mathbf{1}_N$; no facet or optimality assertion beyond the chosen packing is used.

For a maximum-overlap root put $G = \{i : z_i > x_i\}$, $D = \sum_{i \in G} (z_i - x_i)$, $\ell_i = x_i - z_i$ on the nongainers, and define the exact alignment defect

$$E = \sum_{i \in G} (1 - q_i)(z_i - x_i) + \sum_{\substack{i \notin G \\ i \neq r}} q_i \ell_i.$$

Then

$$\Gamma = D + \delta \ell_r - E, \qquad \max_{C : \rho_C > 0} \sigma_C \geq \frac{\Gamma}{1 + \delta}. \tag{PXC51}$$

The second inequality uses that every retained row contains r, so its total weight is $1 + \delta$. Thus $E \leq \delta \ell_r$ is an exact sufficient condition for the aggregate comparison charge to cover the total payoff gain. Perfect gainer coverage with no nongainer deficit gives $E = 0$.

This condition is not automatic, even after optimizing the alignment cost over the entire maximum-root-load face and then across all eligible roots. Among the 21 positive-deficiency marked rows in the exact six-player replay, six contain no eligible hinge-top maximum-overlap root, seven have tied full-menu root load, and eight have positive eligible root excess. At their exact lexicographic optima, only three of those eight have $E = 0$, and only four satisfy $\Gamma \geq D$. Hence positive root excess and optimal support reselection alone do not solve payoff-stratum alignment. The three positive-packet rows in the formerly tied (PXC49) residue are exactly the zero-defect controls: $\delta = 1$, $D = 5/4$, $\ell_r = 1/4$, $\Gamma = 3/2$, and the normalized payment floor is $3/4$. The

remaining universal obligation is now the scalar one: absorb $E - \delta\ell_r$ by an earlier direct/separator arm. Choosing a better packet on the same maximum-load face is an exactly refuted universal route.

Nor can the shortfall always be handed to the Lmin-only rooted projection. In the exact seven-player neutral control, the common hinges 10 and 26 both give an eligible full-menu packet with $\delta = 1$, but every positive row is common, so $\Gamma = 0$ and $E - \delta\ell_r = D = 711/19046$. The exact complete-Lmin one-witness relaxation at that root has bound

$$\frac{8}{15} < \frac{7}{12},$$

whereas adjoining the full coloured Lmax-tight menu raises the exact bound to $5/6$. This is a hostile necessity control: positive aggregate root excess does not imply either comparison-charge closure or an Lmin-only root certificate. The missing universal statement must genuinely use the coloured mixed-witness projection (or a separate direct separator), not just the uncoloured incidence support.

There is a useful exact correction to the comparison-charge viewpoint. For an arbitrary two-colour packet in (MCP4), put

$$M = \sum_{(\sigma,C)} \rho_{\sigma,C}\sigma(C), \qquad S = T - M \geq 0,$$

and let

$$A = \sum_{(x,C)} \rho_{x,C}[z(C) - x(C)], \qquad B = \sum_{(z,C)} \rho_{z,C}[x(C) - z(C)].$$

Tightness and core feasibility give $A, B \geq 0$. Evaluating the incidence completion at z and x, respectively, gives the two exact identities

$$\boxed{A = S + \delta z_r - \sum_{i \neq r} q_i z_i, \qquad B = S + \delta x_r - \sum_{i \neq r} q_i x_i.} \tag{PXC52}$$

Thus both colours share one and the same unused core-budget currency. The inequalities obtained by deleting S are valid but can lose the entire mechanism. In particular, on the neutral seven-player control the maximum root-load face has $\delta = 1$, $A = B = 0$, and $S = 125/19046$; maximizing comparison charge subject only to rooted feasibility still gives $A + B = 0$, now with $\delta = 0$ and $S = 53/107$. Nevertheless the full mixed one-witness relaxation reaches the target, as above. Hence no universal theorem depending only on actual-selection comparison charge can prove the coloured handover. The packet inequality $M \leq T$ must instead be evaluated at a hypothetical bad branch point, where it may be violated even though it has positive slack at the actual game. This is exactly the separation currency in the mixed-witness projection theorem.

The same exact optimization on ten maximum-overlap roots from six selected hard archive rows exposes a second tradeoff: every comparison-charge maximizer has zero root excess, although the corresponding maximum root loads range from $3/2$ to 2. Maximizing excess first and maximizing charge first are therefore distinct faces. These controls do not disprove the one-witness root-selection conjecture; they refute the proposed reduction of that conjecture to a scalar charge at the archived selections. The live dimension-free statement is now a bad-point separation theorem: some distinguished root's mixed packet cuts, together with the direct/stratum rows, must exclude every branch point below the hybrid target.

One target-reaching bad-point atom can now be proved in every dimension. Normalize only for display, so $x(N) = z(N) = T$. Suppose the gainer set is exactly $\{a, b\}$ and there is a distinct root r such that

$$x_r \geq x_a \geq x_b \geq x_j \quad (j \notin \{r, a, b\}), \qquad z_r = z_a \geq z_b. \tag{PXC53a}$$

Assume the two spokes $\{r, a\}, \{r, b\}$ are x-tight and the aligned spoke $\{r, a\}$ is also z-tight. Then

$$\boxed{O(x, z) \geq \frac{n}{2(n-2)} T \geq \frac{n}{\lfloor n^2/4 \rfloor} T \quad (n \geq 6).} \tag{PXC53}$$

Indeed, let w^r be a root-zero exactness witness. The two x-tight spokes give

$$2x_r + x_a + x_b \leq w_a^r + w_b^r \leq T. \tag{PXC53b}$$

The common aligned spoke gives $2z_a = z_r + z_a \leq x_r + x_a$. Hence the total gain is at most

$$(z_a - x_a) + (z_b - x_b) \leq x_r - x_b. \tag{PXC53c}$$

On the other hand, efficiency and the order in (PXC53a) give $T \leq x_r + x_a + (n-2)x_b$. Eliminating x_a with (PXC53b) yields $x_r \leq (n-3)x_b$, while $x_a \geq x_b$ in (PXC53b) yields $x_r + x_b \leq T/2$. Their intersection gives

$$x_r - x_b \leq \frac{n-4}{2(n-2)} T,$$

which proves the first inequality in (PXC53). The second is equivalent to $\lfloor n^2/4 \rfloor \geq 2(n-2)$, immediate separately for even and odd $n \geq 6$.

This theorem explains the neutral control by hand. There $(r, a, b) = (3, 0, 1)$, the spokes are coalitions 9 and 10, and coalition 9 is common. The resulting floor is $7/10$, exactly the independently computed one-witness branch optimum and strictly above $7/12$. The theorem also applies to 46 of the 374 archived six-player records (with 90 root/orientation certificates). Those counts audit scope only. The dimension-free proof closes the entire two-gainer/two-spoke alignment class; the remaining selection theorem must force this atom, its multi-gainer analogue, a direct separator, or a quotient contraction.

The spoke argument extends to any number of gainers while requiring only one Lmax alignment. Let $G = \{i : z_i > x_i\}$ have size $g \geq 2$, let $r \notin G$, and put $X = x(G)$ and $u = \min_{i \in G} x_i$. Suppose

$$x_r \geq x_i \geq x_j \quad (i \in G,\ j \notin G \cup \{r\}), \tag{PXC54a}$$

every spoke $\{r, i\}$, $i \in G$, is x-tight, and for some $a \in G$ maximizing z on G, the spoke $\{r, a\}$ is also z-tight and $z_r \geq z_a$. Then

$$\boxed{\frac{O(x, z)}{T} \geq \begin{cases} \dfrac{n}{2(n-2)}, & g = 2, \\ \dfrac{2(g-1)}{n-2}, & g \geq 3, \end{cases} \geq \frac{n}{\lfloor n^2/4 \rfloor}.} \tag{PXC54}$$

To prove it, sum the g x-tight spoke inequalities against a root-zero witness:

$$g x_r + X \leq T. \tag{PXC54b}$$

The one common spoke and the z order give

$$z(G) \le g z_a \le \frac{g}{2}(x_r + x_a).$$

Since $x_a \le x_r$, the total gain $D = z(G) - X$ satisfies $D \le g x_r - X$. Efficiency and (PXC54a), with $m = n - g - 1$, also give

$$T \le x_r + X + mu, \qquad X \ge gu.$$

Eliminating u, X, x_r from these inequalities and (PXC54b) gives

$$g x_r - X \le \frac{n - 2g}{n - 2} T. \tag{PXC54c}$$

(If $2g > n$, the hypotheses are consequently infeasible.) Thus $O = T - D \ge 2(g-1)T/(n-2)$. For $g \ge 3$ this dominates the target because $4\lfloor n^2/4 \rfloor \ge n(n-2)$. When $g = 2$, retaining the identity of the other gainer rather than replacing both by z_a gives the sharper (PXC53) argument.

This is the useful generalization: the Lmin side supplies the full star, but the Lmax side need align only one highest-gainer spoke. The exact archive scan finds 90 orientations across 46 six-player records and closes eight of the 16 positive-floor residuals; those finite matches all have $g = 2$. None of the seven fixed hostile/equality branches has the full star, so the theorem does not replace their already proved direct or rooted-facet arms. The next selection problem is correspondingly combinatorial: show that a role-minimal hard branch exposes such an aligned star, or use a missing spoke to construct the strict separator/marked-facet alternative.

No new full-cone facet class is hidden here. The incidence completion of the x-star is

$$\sum_{i \in G} \mathbf{1}_{\{r,i\}} + \sum_{j \notin G \cup \{r\}} e_j - (g-1) e_r = \mathbf{1}_N. \tag{PXC54d}$$

Applying the proved rooted-pair shear to every outside singleton replaces it by $\{r, j\}$ and changes the exceptional coefficient to $-(n-2)$. The result is the saturated rooted-star/complement-star system already classified as indecomposable. The one common-spoke comparison is a relative two-selection alignment row, not a new characteristic-function facet. Finally, both alternatives in (PXC54) are strict above the hybrid target for $n \ge 6$. Consequently this whole aligned-star branch is excluded from the equality classification; the balanced $G(p,k)$ orientations correctly remain in their separate block-star mechanism.

The complementary rooted co-star admits a sharper theorem without the equal-stratum assumptions of the earlier mixed block-star calculation. Keep $G = \{i : z_i > x_i\}$, $|G| = g \ge 2$, $r \notin G$, and suppose

$$x_r \ge x_i \ge x_j \quad (i \in G,\ j \notin G \cup \{r\}). \tag{PXC55a}$$

For $a \in G$ put $C_a = \{r\} \cup (G \setminus \{a\})$. Assume every C_a is x-tight and that, for some a maximizing z on G, C_a is also z-tight and $z_r \ge z_a$. Then

$$\boxed{O(x,z) \ge \frac{n}{g(n-g)} T \ge \frac{n}{\lfloor n^2/4 \rfloor} T.} \tag{PXC55}$$

Let $X = x(G)$, $u = \min_{i \in G} x_i$, and $m = n - g - 1$. Weight every x-tight co-star by $1/(g-1)$ and evaluate at a root-zero witness. This gives

$$\frac{g}{g-1}x_r + X \le T. \tag{PXC55b}$$

Efficiency and (PXC55a) give $T \le x_r + X + mu$ and $X \ge gu$. Subtracting (PXC55b) from the first inequality yields $x_r \le m(g-1)u$. Combining this with (PXC55b) and $X \ge gu$ gives

$$x_r - u \le \frac{g(n-g)-n}{g(n-g)}T. \tag{PXC55c}$$

The one common co-star gives

$$z_r + z(G \setminus \{a\}) = x_r + x(G \setminus \{a\}).$$

Since $z_a \le z_r$, total gain is at most $x_r - x_a \le x_r - u$. Subtracting (PXC55c) from T proves the first inequality; the second is the ordinary balanced-product bound $g(n-g) \le \lfloor n^2/4 \rfloor$.

This theorem absorbs the earlier equal-stratum mixed block-star bound at the selection level and exactly matches its facet system. The incidence completion is

$$\frac{1}{g-1}\sum_{a\in G}\mathbf{1}_{C_a} + \sum_{j\notin G\cup\{r\}} e_j - \frac{1}{g-1}e_r = \mathbf{1}_N. \tag{PXC55d}$$

After the rooted-pair shear its outside-pair coefficients are one, its hinge coefficients are $1/(g-1)$, and its exceptional coefficient is $-[m+1/(g-1)]$: precisely the already proved $G(p,k)$ rooted block-star facet with $(p,k) = (m,g)$. No new full-cone class is introduced.

Target equality inside this branch requires $g(n-g) = \lfloor n^2/4 \rfloor$ and equality throughout the scalar proof. In particular every nonroot x coordinate equals u and $x_r = m(g-1)u$, the normalized Lmin profile of the sharp family. The branch argument alone does not force the remaining Lmax coordinates; their stage equations do so in the two archived balanced seven-player orientations.

Among 374 archived six-player records, (PXC55) applies to 84 records and 186 root/omitted-gainer orientations, with gainer counts two, three, and four. More importantly, it covers all 16 positive-floor residual rows left after direct gluing, interpolation, and null descent. Among the seven fixed hostile/equality branches, it rejects all five hostile controls and accepts the sharp $G(2,4)$ and $G(3,3)$ orientations at respectively four and three balanced roots. These counts do not prove universal selection. They show that the former five-orbit finite residue is now explained by one dimension-free ordered co-star inequality in an existing facet class. The remaining all-n obligation is to force this co-star, the aligned star (PXC54), a direct/separator certificate, or quotient contraction.

There is no additional structural branch between a rooted packet and the proved carrier–separator–contraction alternative. Fix a positive ceiling root r, so $z_r = M = \max_i z_i$ and $x_r \ge M$, and assume r is pinned by at least one proper x-tight coalition. For every gainer a,

$$x_a < z_a \le M \le x_r. \tag{PXC56a}$$

Carrier monotonicity therefore implies that some x-tight coalition contains r and omits a: otherwise a would be a full-menu carrier of r, forcing $x_a \ge x_r$.

Let $\mathcal{T}_x(r)$ be the proper x-tight rows containing r, and define the rooted packing number

$$\nu_r = \max\left\{ \sum_{C\in\mathcal{T}_x(r)} \rho_C : \rho \ge 0,\ \sum_{C\ni i} \rho_C \le 1\ (i \ne r) \right\}. \tag{PXC56b}$$

Then exactly one of the following holds.

(i) $\nu_r > 1$. An extreme optimizer has at most $n-1$ positive rows and, with $q_i = 1 - \sum_{C \ni i} \rho_C$ and $\delta = \nu_r - 1 > 0$, gives

$$\sum_C \rho_C \mathbf{1}_C + \sum_{i \neq r} q_i e_i - \delta e_r = \mathbf{1}_N. \tag{PXC56c}$$

Thus the branch exposes a sparse singleton-exceptional rooted packet; its support-minimal factors are tested by the existing semi-balanced indecomposability criterion.

(ii) $\nu_r = 1$. Rooted packing/cover duality gives a nonnegative cover y of total one with $y(C) \geq 1$ for every $C \in \mathcal{T}_x(r)$. Since $y(N \setminus \{r\}) = 1$, every player in the positive support of y belongs to every row in $\mathcal{T}_x(r)$. Hence r has a full-menu common carrier. No such carrier is a gainer, by (PXC56a) and the preceding omission argument. The carrier-cycle/separator theorem then supplies either a strict Lmax separator payment or a common module and exact weighted contraction.

Indeed a nonempty rooted row can be scaled to root load one, so $\nu_r \geq 1$; the two cases are exhaustive. In the unit-cover case, $y(C) \geq y(N \setminus \{r\})$ forces equality and puts the support of y in the intersection of all rooted rows. This proves the asserted carrier.

Call this the ceiling-root gainer-exclusion trichotomy (PXC56). It is a dimension-free structural theorem, not the missing quantitative pricing statement. The exact six-player replay contains 703 ceiling roots: 149 have positive excess and 554 are tied, exactly the previously audited full-menu unaffordable roots. Every one of the 16 positive-floor residual roots has positive excess. On the seven fixed hostile/equality controls, the sole all-tied case is the already proved top-and-bottom common-module branch; all other relevant cases expose a positive packet. The universal gap is therefore narrowed to this statement: on the positive side of (PXC56), prove target pricing/alignment (with (PXC54)–(PXC55) as complete large classes), or factor the packet into a direct separator or one of the already classified rooted facets.

The first incidence split on that positive side is now dimension-free. Let G be the gainer set, $g = |G| \geq 2$, and let an x-tight proper row C contain the ceiling root r while omitting a gainer. Put

$$H = C \setminus G, \quad J = C \cap G, \quad h = |H|, \quad s = |J|, \quad p = n - g - h.$$

Then $h \geq 1$ because $r \in H$, and necessarily $p \geq 1$. Indeed, if C contained every nongainer, its loss between x and z would be the total gain D, while omission of a genuine gainer would make the gain inside J strictly smaller than D. This would give $z(C) < x(C) = v(C)$, contrary to $z \in \mathrm{Core}(v)$.

Suppose now that C omits at least two gainers, so $s \leq g - 2$. The cardinality base in the direct master identity is

$$b(n; g, h, p, s) = 1 - \frac{gp}{(g+h)(p+g-s)}. \tag{PXC57a}$$

For fixed n, g, the subtracted product decreases as h increases, while b decreases as s increases. Hence its minimum is at $h = 1$ and $s = g - 2$, where $p = n - g - 1$ and

$$b \geq \frac{n+g+1}{(g+1)(n-g+1)}. \tag{PXC57b}$$

For every $n \geq 8$ this is at least the hybrid target:

$$\boxed{b(n; g, h, p, s) \geq \frac{4n}{n^2 - (n \bmod 2)}.} \tag{PXC57}$$

To check the last comparison, write $n = 2k$. After multiplication by the positive denominator, the relevant quadratic is $2g^2 - 3kg + 2k^2 - 3k - 2$, whose continuous minimum is $(k-4)(7k+4)/8$. For $n = 2k+1$, comparison with $(2k+1)/[k(k+1)]$ has continuous minimum numerator

$$\frac{7(k-4)^4 + 102(k-4)^3 + 505(k-4)^2 + 898(k-4) + 239}{4(2k+1)},$$

which is strictly positive for $k \geq 4$. Equality in (PXC57) occurs only at $(n, g, h, p, s) = (8, 3, 1, 4, 1)$. The complete lower-dimensional failures are $(6, 2, 1, 3, 0)$, $(6, 3, 1, 2, 1)$, $(7, 2, 1, 4, 0)$, and $(7, 3, 1, 3, 1)$, with bases respectively $3/5, 5/8, 5/9, 11/20$.

The upper-budget qualification is real. At root 6 of the fixed eight-player primitive-coefficient control, the five x-tight multi-omission rows $70, 85, 90, 104, 204$ all satisfy (PXC57) and have zero lower comparison slack, but every residual budget is smaller than the loss on its nongainer block. Thus cardinality plus Lmin tightness cannot finish this arm by termwise positivity; the remaining labelled payment must come from Lmax colour, the low-block term, or an earlier separator.

Call (PXC57) the multi-omission direct-base gate. Because C is x-tight, core feasibility of z automatically supplies the lower comparison inequality: the gain on J is at least the loss on H. Thus from eight players onward a multi-omission blocker can fail the direct theorem only at its *upper* labelled comparison budget; it cannot fail because of cardinality. If no such blocker exists, the gainer omissions forced by (PXC56) select, one per gainer, rooted rows containing every other gainer. This is a decorated complete co-star: the remaining work is to remove or price its nongainer decorations and force one Lmax-aligned row. This dichotomy is a strict narrowing of the positive arm, not the missing selection/gluing theorem itself.

The remaining budget has a useful exact Lmax form. Let $D = z(G) - x(G)$ be the total gain and define the below-mean deficit of the omitted nongainer block

$$Q(P) := \sum_{i \in P} \left(\frac{T}{n} - z_i \right).$$

If R_s denotes the direct residual budget for a packet J of size s and $L_H = x(H) - z(H)$, then

$$\boxed{R_s - L_H = \frac{n}{n-p} Q(P) - \frac{p+g-s}{g} D.} \tag{PXC58}$$

Indeed total gain equals total nongainer loss, so $x(G) = z(G) - D$ and $x(P) = z(P) + D - L_H$. Substitution in the direct residual formula first gives

$$R_s - L_H = \frac{p}{g+h} z(G \cup H) - z(P) - \frac{p+g-s}{g} D.$$

Using $g + h = n - p$ and $z(G \cup H) = T - z(P)$ yields (PXC58). Consequently the exact combined comparison/residual term is nonnegative if and only if

$$Q(P) \geq \frac{(n-p)(p+g-s)}{ng} D. \tag{PXC58a}$$

For $s = g-1$ this is precisely the previously proved maximal-packet deficit-quota identity; (PXC58) shows that each additional omitted gainer raises the required quota by $(n-p)D/(ng)$. Hence multi-omission buys the cardinality gate (PXC57) at an exact Lmax price.

The primitive eight-player control fails (PXC58a) on all five rows. Its quota margins are

$$-\frac{48}{13}, \quad -\frac{40}{39}, \quad -\frac{124}{39}, \quad -\frac{308}{195}, \quad -\frac{4}{13}.$$

Conversely every displayed co-star row in both balanced seven-player equality orientations has quota margin zero. Their direct cardinality bases are below target, so this is equality compatibility rather than a direct proof of the sharp cases. The live Lmax-colour problem is now explicit: pay the negative part of (PXC58), or use quota failure to select a separator/contraction or a decorated co-star covered by (PXC55). Merely enlarging the uncoloured packing cannot change this scalar deficit.

The decorations in the co-star arm can nevertheless be priced rather than deleted. Put $q = n-g$ and $m = q-1$, and for every $a \in G$ suppose that

$$C_a = \{r\} \cup (G \setminus \{a\}) \cup K_a$$

is x-tight, where K_a is a set of nongainers outside $G \cup \{r\}$. Assume $q \geq 2$, the order condition (PXC55a), and the packing condition that every outside player belongs to at most $g-1$ of the K_a. Define

$$A := \frac{1}{g-1} \sum_{a \in G} x(K_a), \qquad \mu := \frac{(g-1)n}{g(n-g)}.$$

If, for one a maximizing z_a on G, the row C_a is also z-tight and $z_r \geq z_a$, then with $L_a := x(K_a) - z(K_a)$,

$$\boxed{O \geq \frac{n}{g(n-g)} T + \mu A - L_a.} \tag{PXC59}$$

Indeed, weight the decorated co-stars by $1/(g-1)$ and evaluate them at a root-zero exactness witness. The packing hypothesis bounds every nonroot witness load by one, and hence

$$\frac{g}{g-1} x_r + x(G) + A \leq T. \tag{PXC59a}$$

Efficiency and the order condition give $T \leq x_r + x(G) + m\,u$. Comparison with (PXC59a) yields

$$x_r \leq (g-1)(m\,u - A). \tag{PXC59b}$$

The common decorated row and $z_a \leq z_r$ give, for total gainer gain D,

$$D \leq x_r - x_a + L_a \leq x_r - u + L_a. \tag{PXC59c}$$

Set $k = 1 - n/[g(n-g)]$. Since $g, q \geq 2$ and $n \geq 6$, $k \geq 0$; use (PXC59a) as a lower bound on T, then (PXC59b) as an upper bound on x_r. Exact simplification gives

$$D \leq kT + L_a - \mu A,$$

which is equivalent to (PXC59). Therefore the hybrid target follows whenever

$$L_a \leq \mu A + n \left(\frac{1}{g(n-g)} - \frac{1}{\lfloor n^2/4 \rfloor} \right) T. \tag{PXC59d}$$

At balanced cardinality this says simply $L_a \leq \mu A$. Clean co-stars have $A = L_a = 0$ and recover (PXC55), including both sharp seven-player orientations. In the six-player archive the priced theorem applies to 90 of 374 records, six more than the clean theorem. It has 728 labelled certificates, of which 186 are clean, and all 728 reach the target. The six new records use genuine decorations; in a representative certificate $A > 0$ while the aligned decoration has $L_a = 0$. All

16 positive-floor residuals still close, every fixed hostile control is rejected, and the two sharp orientations retain respectively four and three clean certificates. These counts audit the signs and show that decoration pricing has nonvacuous scope. They do not select a packing-admissible family or force (PXC59d). Decorated supports may also factor into known rooted facets or expose a new support-minimal semi-balanced system; that facet classification remains an objective-C obligation.

The packing hypothesis in (PXC59) has an exact fractional alternative. Keep the no-multi-omission branch of (PXC57), and for any rooted packing ρ write

$$\nu = \sum_{C} \rho_C, \qquad \omega_a = \sum_{C \not\ni a} \rho_C \quad (a \in G).$$

Every supported row omits at most one gainer. Since the load of gainer a is $\nu - \omega_a \le 1$, one has $\omega_a \ge \nu - 1$, while $\sum_a \omega_a \le \nu$. Therefore

$$\boxed{1 \le \nu_r \le \frac{g}{g-1}.} \tag{PXC60}$$

At equality every inequality just used is tight: no positive-weight row omits zero gainers, and the total weight of rows omitting each a is exactly $1/(g-1)$. Hence a saturated optimizer is a *fractional decorated complete co-star*. Its gainer loads are one, its root load is $g/(g-1)$, and every outsider load is at most one. If

$$A_\rho := \sum_{C} \rho_C x(K_C),$$

the root-zero witness gives (PXC59a) with A_ρ in place of A. Consequently the proof of (PXC59) applies verbatim whenever a positive-weight row omitting a highest-z gainer is also z-tight and satisfies the same root comparison. Integral packing-admissible co-stars are only the special case assigning weight $1/(g-1)$ to one row of each omission type.

Saturation itself has an exact weighted Hall criterion. For each $a \in G$, let $\mathcal{F}_a$ be the rooted tight rows omitting a, and write $K(C)$ for the outsider decoration of C. Then a fractional decorated co-star exists if and only if, for every nonnegative outsider price vector β,

$$\boxed{\frac{1}{g-1} \sum_{a \in G} \min_{C \in \mathcal{F}_a} \beta(K(C)) \le \sum_{j \notin G \cup \{r\}} \beta_j.} \tag{PXC61}$$

Necessity follows by pricing the outsider-capacity inequalities and then minimizing separately within every omission type. For sufficiency, the product of the type simplices

$$\sum_{C \in \mathcal{F}_a} \rho_C = \frac{1}{g-1}, \qquad \rho_C \ge 0,$$

has a compact convex image in outsider-load space. If that image misses the downward unit box, strict separation supplies a nonnegative β violating (PXC61); conversely such a violation separates every typed mixture from the box. This is finite-dimensional separation, equivalently exact LP duality.

Thus failure of saturation is not an unspecified selection failure: after normalizing $\sum_j \beta_j = 1$, it has a strict rational weighted-Hall certificate. For every omission type with positive minimum, every row in that type meets $\operatorname{supp} \beta$; this is a labelled collective carrier obstruction. At $\nu_r = 1$ it collapses to the common-carrier arm of (PXC56). For $1 < \nu_r < g/(g-1)$, converting this weighted

obstruction into a strict Lmax separator payment or a target-reaching mixed rooted facet is the remaining all-dimensional arrow.

There is a direct tangent interpretation of a strict Hall certificate. Use the normalization $\sum_j \beta_j = 1$, put

$$m_a := \min_{C \in \mathcal{F}_a} \beta(K(C)),$$

and define a zero-sum direction y by

$$y_j = \beta_j \quad (j \notin G \cup \{r\}), \qquad y_a = m_a - 1 \quad (a \in G), \qquad y_r = g - 1 - \sum_a m_a. \tag{PXC62a}$$

A Hall violation gives $\sum_a m_a > g - 1$, so $y_r < 0$, while $0 \le m_a \le 1$ makes every gainer coordinate nonpositive. For a rooted row $C \in \mathcal{F}_a$,

$$y(C) = \beta(K(C)) - m_a \ge 0. \tag{PXC62b}$$

Order players increasingly by x_i, breaking every tie increasingly by y_i. Suppose every proper prefix has nonnegative y-sum and at least one has positive sum. This is exactly the condition that a sufficiently small move from x along y strictly Lorenz-improves x; it holds in particular when every positive-priced outsider is strictly below the root and every gainer with $m_a < 1$. If $y(S) \ge 0$ on every x-tight coalition, that small move stays in the core, contradicting Lmin optimality. Consequently some x-tight coalition D has $y(D) < 0$. By (PXC62b) and the no-multi-omission hypothesis, such a blocker is necessarily either root-free or root-containing with every gainer present:

$$\boxed{\; r \notin D \quad \text{or} \quad \{r\} \cup G \subseteq D. \;} \tag{PXC62}$$

This Lorenz–Hall-blocker lemma is a conventional tangent-cone consequence, not a claim that the blocker already pays the overlap target. Weak Lmin ties can prevent the exact prefix condition and remain a separate contraction/stage-pinning case.

On the ordered arm the prefix premise is in fact automatic away from boundary ties. Each omission type has a row by (PXC56) once the root is pinned by a proper Lmin-tight coalition, so $0 \le m_a \le 1$, and a Hall violation normalizes to $\sum_a (1 - m_a) = 1 - \epsilon < 1$. Orderedness places every outsider weakly below every gainer, so in the increasing order the running sum first accumulates nonnegative outsider prices, and once every outsider has been passed it stays at least

$$1 - \sum_{b \text{ processed}} (1 - m_b) \;\ge\; \epsilon \;>\; 0, \tag{PXC74}$$

because the subtracted terms are nonnegative and total $1 - \epsilon$. Hence if no positively priced outsider ties the minimum gainer and no gainer ties the root, every prefix is nonnegative and some prefix is positive, the direction is a strict Lorenz tangent, and Lmin optimality forces a root-free or all-gainer rooted blocker as in (PXC62). The prospective ordered-saturation theorem—every ordered no-multi-omission ceiling root is saturated, with no exception in the six-player census or among the designed seven-to-nine-player roots—is therefore equivalent, away from boundary ties, to excluding those two blocker shapes. Both shapes are genuinely tight coalitions in general, so the exclusion must use their interaction with the obstruction prices; this remains open.

On the saturated arm, even common z-tightness can be removed at an exact scalar price. Put $d_a = z_a - x_a$, let $d_{\max} = \max_{a \in G} d_a$, and write $\ell_r = x_r - z_r \ge 0$. For any positive-weight co-star row C of omission type a, define its Lmax slack $S_C = z(C) - x(C) \ge 0$. Direct expansion gives the row-independent identity

$$S_C + L_C = D - d_a - \ell_r, \qquad L_C = x(K_C) - z(K_C). \tag{PXC63a}$$

The comparison used in (PXC59) now gives

$$D \le x_r - x_a + L_C + S_C.$$

The fractional packet proof gives, before using any common row,

$$x_r - u \le \left(1 - \frac{n}{g(n-g)}\right) T - \mu A_\rho.$$

Choose a maximizing d_a and substitute (PXC63a). Then

$$\boxed{O \ge \frac{n}{g(n-g)} T + \mu A_\rho - (D - d_{\max} - \ell_r).} \tag{PXC63}$$

Thus an ordered saturated fractional co-star reaches the hybrid target without any common Lmax row whenever

$$D - d_{\max} - \ell_r \le \mu A_\rho + n \left(\frac{1}{g(n-g)} - \frac{1}{\lfloor n^2/4 \rfloor}\right) T. \tag{PXC63b}$$

The left side is nonnegative: by (PXC63a), it is $S_C + L_C$ for the maximum-gain omission type. It measures precisely the gain dispersed away from the strongest gainer after crediting the root loss. Equation (PXC63b), not common-row existence, is therefore the remaining scalar statement on the saturated arm.

The Lmax input is indispensable. At the allocation level, take

$$x = (9, 1, 1, 1, 1, 1, 1, 1), \qquad z = (7/2, 25/8, 25/8, 25/8, 25/8, 0, 0, 0).$$

With root zero and gainers $1, 2, 3, 4$, the four clean co-stars are x-tight by stipulation, their uniform packet saturates $T = 16$, the root is the Lmax ceiling, and every co-star has nonnegative comparison slack $7/8$. Nevertheless $O = 15/2 < 8$, the eight-player target, while (PXC63) gives only $57/8$. This is an exact relaxed allocation-level control, not an asserted Lmin/Lmax pair of any exact game. It proves that order, saturated packet geometry, efficiency, and the individual core-side co-star comparisons do not imply (PXC63b). Any universal proof must use genuine Lmax optimality—the sparse over-cover, staged Lmax pinning, or an equivalent property.

The first-stage Lmax payment can be isolated exactly. Under the ordered co-star hypotheses, every outsider j satisfies $z_j \le x_j \le u < z_a$ for each gainer a. Hence the Lmax top stratum $H_{\max}$ is contained in $\{r\} \cup G$, and r is its only nongainer. Let α be the probability on $H_{\max}$ supplied by a normalized sparse first-stage Lmax over-cover, so $\alpha \cdot x \ge M = z_r$. Then

$$\boxed{\alpha_r \ell_r - \sum_{a \in G \cap H_{\max}} \alpha_a d_a = \alpha \cdot x - M \ge 0.} \tag{PXC64}$$

This is genuine normalized Lmax currency, but it is weighted currency. It does not by itself control the unweighted dispersed-gain term in (PXC63b). Indeed all 89 ordered saturated archive roots give positive nonroot top mass, but only two have positive surplus in (PXC64); the other 87 are tight. Thus top-anchor surplus is not the missing universal reserve.

The labels of the sparse over-cover are much sharper in the finite census. At 88 of the 89 ordered saturated roots, its support contains a positive-weight fractional co-star row. The sole exception is `harvest6_601.json`, record 71, root 4: its one-row sparse Lmax support is coalition $26 = \{1, 3, 4\}$, which omits two gainers, and its (PXC63b) margin is $19/12$. This is discovery evidence for a

common-co-star-or-Lmax- multi-omission alternative, not a universal theorem. It shows that the next proof must retain the labelled sparse support rather than only α.

The complete staged Lmax packet gives a second exact description of the same missing reserve. Choose $a^* \in G$ with $d_{a^*} = d_{\max}$. For every $a \neq a^*$, normalize a sign-coherent Lmax pinning relation for a, and let c_a be the coefficient of the root-loss column. Under the ordered hypotheses the root is the only positive-loss source in the cumulative Lmax upper set: outsiders satisfy $z_j \leq x_j < z_a$, and gainers are not loss sources. Evaluation at $x - z$ therefore gives

$$d_a \leq c_a \ell_r.$$

Put $\sigma_a = c_a \ell_r - d_a \geq 0$ and $\kappa = \sum_{a \neq a^*} c_a$. Then, without an inequality,

$$D - d_{\max} - \ell_r = (\kappa - 1)\ell_r - \sum_{a \neq a^*} \sigma_a. \tag{PXC65a}$$

Substitution into (PXC63) gives the mixed Lmin–Lmax bound

$$\boxed{O \geq \frac{n}{g(n-g)} T + \mu A_\rho - (\kappa - 1)\ell_r + \sum_{a \neq a^*} \sigma_a.} \tag{PXC65}$$

Thus coefficient congestion is not itself the invariant obstruction: increasing a root price also creates the evaluated transport surplus that must be retained with it. The target follows exactly when

$$(\kappa - 1)\ell_r \leq \mu A_\rho + n \left(\frac{1}{g(n-g)} - \frac{1}{\lfloor n^2/4 \rfloor} \right) T + \sum_{a \neq a^*} \sigma_a. \tag{PXC65b}$$

This is an exact coupling of the fractional Lmin packet to complete staged Lmax pinning, not a proof of (PXC65b): by (PXC65a), the latter is equivalent to (PXC63b). In the exact six-player audit the canonical minimum-root-price relations have $\kappa \in [1, 3]$. The tempting unit-congestion claim holds on only 77/89 ordered saturated roots, and the raw congestion inequality closes only 79/89, with minimum margin $-25/12$. After retaining the transport surpluses, all 89 close and the minimum margin is zero. The next load-bearing task is therefore to factor or aggregate the labelled σ_a certificates; no scalar bound on κ alone can supply the missing theorem.

The surplus in (PXC65) has a labelled nonnegative decomposition. Write the normalized sign-coherent Lmax relation for gainer a as

$$\sum_C \lambda_C^a \mathbf{1}_C - \sum_{j \in Q_a} p_j^a e_j = c^a \mathbf{1}_N + e_a, \qquad \lambda^a, p^a \geq 0, \tag{PXC66a}$$

where Q_a is the cumulative Lmax upper set and $p_r^a = c_a$. Put $R_C = x(C) - z(C) \geq 0$ on its z-tight coalition support. Orderedness again makes r the only loss coordinate in Q_a; every other nonzero term is a gainer b with $z_b \geq z_a$. Dotting (PXC66a) with $x - z$ gives

$$\boxed{\sigma_a = c_a \ell_r - d_a = \sum_C \lambda_C^a R_C + \sum_{b \in G \cap Q_a} p_b^a d_b.} \tag{PXC66}$$

Thus every unit of transport surplus is labelled either by a genuine reverse-payment coalition or by another gainer in the same or a higher Lmax stratum. In the latter terms the directed edge $a \to b$ weakly increases z; consequently every directed cycle lies within one tied Lmax stratum. If all R_C on a transport component vanish, all its coalition columns are common x–z tight columns. A terminal

tied component is therefore either split by a common tight row, supplying the common-hinge arm, or is unsplit by the full common menu and is a common module eligible for the established contraction theorem. This is a qualitative handover, not yet the quantitative assertion that the selected hinge or contracted quotient pays (PXC65b); root-paid terminal leaves must also be aggregated without reintroducing the false scalar-congestion claim.

The exact minimum-root-price audit records the complete decomposition rather than only c_a. Of the ten ordered saturated roots where raw congestion fails, three use a positive reverse-payment row: two have a root-free row and one has a rooted multi-omission row. The other seven have pure cross-gainer surplus on common coalition support. Nine of the ten also have a fractional co-star row in the canonical sparse top over-cover; the sole exception is the rooted multi-omission case. Restricting to the four positive-floor residual roots leaves two root-free reverse-payment cases and two pure-pin cases, and all four have a sparse common co-star. These are finite support classifications and depend on the canonical exact LP solutions; they motivate, but do not prove, the dimension-free common-co-star/separator/module alternative.

There is a coefficient-only way to aggregate the transport bundles without double-counting their imported gains. Fix a maximum-gain gainer a^* and put

$$B_{ab} = \mathbf{1}_{\{a=b\}} + p_b^a \qquad (a, b \in G).$$

For one normalized relation per gainer, define the fractional bundle-cover number

$$\begin{aligned} \tau(a^*) &= \min\left\{ \sum_{a\in G} c_a t_a : \sum_{a\in G} B_{ab} t_a \geq 1 \ (b \neq a^*),\ t \geq 0 \right\} \\ &= \max\left\{ \sum_{b\neq a^*} y_b : \sum_{b\neq a^*} B_{ab} y_b \leq c_a \ (a \in G),\ y \geq 0 \right\}. \end{aligned} \tag{PXC67a}$$

The equality is ordinary finite LP duality. Dropping only the nonnegative reverse-coalition term in (PXC66) gives $c_a \ell_r \geq \sum_b B_{ab} d_b$. Multiply by any feasible t_a and use the cover constraints. Then

$$\sum_{b\neq a^*} d_b \leq \tau(a^*)\ell_r.$$

Consequently the saturated co-star bound strengthens to the conditional mixed-packet theorem

$$\boxed{O \geq \frac{n}{g(n-g)} T + \mu A_\rho - \big(\tau(a^*) - 1\big)\ell_r.} \tag{PXC67}$$

In particular the hybrid target follows whenever

$$\big(\tau(a^*) - 1\big)\ell_r \leq \mu A_\rho + n\left(\frac{1}{g(n-g)} - \frac{1}{\lfloor n^2/4 \rfloor}\right) T. \tag{PXC67b}$$

Unlike (PXC65b), this is a genuinely stronger certificate condition rather than an algebraic restatement of (PXC63b): it replaces the actual dispersed gain by an independently checkable fractional cover of normalized Lmax relations. It is also exactly the sought sparse Lmax over-cover dual to the fractional Lmin co-star packet.

The exact six-player audit solves both sides of (PXC67a). It proves (PXC67b) at all 89 ordered saturated roots, with minimum margin zero. At 88 roots the optimal aggregate root price is at most one. The sole value 4/3 occurs at `harvest6_602.json`, record 22, root 2; that record is

already direct-closed, has $A_\rho = 1/4$, and retains margin $7/30$ even after all reverse-coalition credit is discarded. For the actual-gain-weighted cover, 88/89 optima cover demand exactly; the same exceptional record overcovers by $1/10$. These facts are exact bounded evidence, not a proof of the dimension-free budget (PXC67b). The new load-bearing problem is to prove that the bundle-cover dual prices are paid by decoration/cardinality reserve, or else force a separator, common-hinge, or contraction arm.

The dual in (PXC67a) already forces that labelled alternative whenever its price exceeds one. Extend an optimal dual by $y_{a^*} = 0$, so $\sum_{b \in G} y_b = \tau$, and define the zero-sum player price

$$w_r = -\tau, \qquad w_b = y_b \ (b \in G), \qquad w_j = 0 \ (j \notin G \cup \{r\}). \tag{PXC68a}$$

Dot (PXC66a) with w. The dual constraint gives, for every normalized gainer relation,

$$\boxed{\sum_C \lambda_C^a w(C) = \sum_{b \neq a^*} B_{ab} y_b - c_a \tau \leq c_a(1 - \tau).} \tag{PXC68}$$

Because $d_a > 0$ forces $c_a > 0$, $\tau > 1$ makes this average strictly negative. Hence every relation support contains a proper z-tight row C with $w(C) < 0$. A root-free row has $w(C) \geq 0$, while a rooted row has

$$w(C) = -\sum_{b \notin C} y_b.$$

Thus the selected row is rooted and omits positive dual price. This is a dimension-free aggregate Lmax separator theorem. If its reverse slack R_C is positive, it supplies genuine strict comparison payment. If $R_C = 0$, it is a common x–z tight row; according as it omits several positive-price gainers, exactly one, or is unsplit across a tied component, it enters the rooted multi-omission, common co-star/common-hinge, or common- module arm. What remains is quantitative: prove that one of those selected rows pays the excess $(\tau - 1)\ell_r$ not already covered by decoration and cardinality reserve.

The only audited root with $\tau > 1$ is the decorated direct-closed control above. Its exact dual is $(y_4, y_5) = (1/3, 1)$, and every normalized relation exposes a negative rooted omission row. The returned rows $28, 45, 46$ all have zero reverse slack and are common co-stars; the decoration reserve already proves (PXC67b). This is a sign and label calibration of (PXC68), not evidence that the same quantitative payment holds in every dimension.

The staged relations behind (PXC67a) are only one certificate basis. The same evaluation argument accepts any nonnegative combination of full-menu Lmax-tight rows. Call a representation

$$\sum_C L_C \mathbf{1}_C = C_0 \mathbf{1}_N + \sum_{i \in N} q_i e_i, \qquad L_C \geq 0, \tag{PXC69a}$$

with C ranging over the proper z-tight coalitions, C_0 free, $q_b \geq 1$ for every gainer $b \neq a^*$, $q_{a^*} \geq 0$, $q_j \leq 0$ at every outsider with $x_j > z_j$, and q_j free at tied outsiders, a *menu cover* at (r, a^*), and let $\tau_{\mathrm{menu}}(a^*)$ be the minimum of q_r over all menu covers. Aggregating any admissible staged relation system with optimal (PXC67a) weights yields a menu cover whose root coefficient is its cover price: coverage gives $q_b \geq 1$, orderedness keeps strict outsiders out of the transport columns, and tied outsiders may enter with either sign. Hence $\tau_{\mathrm{menu}}(a^*) \leq \tau(a^*)$ for every admissible system, and τ_{menu} is the exact value of the best cover-type certificate. Dotting (PXC69a) with $x - z$ kills the ambient term, pays each used row its nonnegative reverse slack, and keeps every outsider term nonpositive, so $\sum_{b \neq a^*} d_b \leq \tau_{\mathrm{menu}}(a^*)\ell_r$ and the saturated co-star theorem strengthens to

$$\boxed{O \geq \frac{n}{g(n-g)} T + \mu A_\rho - \big(\tau_{\mathrm{menu}}(a^*) - 1\big)\ell_r.} \tag{PXC69}$$

Ordinary LP duality computes the same number as a root-exchange rate:

$$\tau_{\text{menu}}(a^*) = \max\Big\{\sum_{b\neq a^*} y_b\Big\} \tag{PXC69b}$$

over all y with $y(C) \le 0$ on every proper z-tight coalition, $y(N) = 0$, $y_r = -1$, $y \ge 0$ on the gainers, $y \le 0$ on strict outsiders, and $y = 0$ on tied outsiders. Negated, an optimal y is a core tangent direction at the Lmax point that raises the root by exactly one unit, weakly lowers every gainer, weakly raises every strict outsider, and freezes the tied outsiders. Thus τ_{menu} is the exact root-versus-gainer exchange rate of the Lmax tangent cone in the selector orthant. The direction $(x - z)/\ell_r$ is feasible for (PXC69b)—and $\ell_r > 0$ on this arm, since the staged pinning inequality $d_a \le c_a\ell_r$ of (PXC65) forbids a zero root loss once a gainer exists—so

$$\tau_{\text{menu}}(a^*) \ge \frac{D - d_{\max}}{\ell_r}.$$

No cover-type certificate prices below the actual dispersed gain, and equality states that the selector pair itself is an extremal exchange. The target therefore follows from (PXC69) whenever

$$\big(\tau_{\text{menu}}(a^*) - 1\big)\ell_r \le \mu A_\rho + n\left(\frac{1}{g(n-g)} - \frac{1}{\lfloor n^2/4\rfloor}\right)T, \tag{PXC69c}$$

and in the sharper unit-exchange form whenever $\tau_{\text{menu}} \le 1$. The unit-exchange bound forces $D - d_{\max} \le \ell_r$, equivalently, through efficiency, the outsider-loss bound $\sum_j (x_j - z_j) \le d_{\max}$ over the outsiders; the two statements coincide exactly when the menu cover is exchange-exact.

The exact audit solves both programs at all 89 ordered saturated archive roots and both sharp seven-player orientations. Every optimum is exchange-exact and equal to one: $\tau_{\text{menu}} = (D - d_{\max})/\ell_r = 1$, so the outsider loss equals the maximum gain exactly at every audited root, and (PXC69c) closes all 89 roots and all 16 positive-floor residuals with minimum margin zero. Every optimal cover burns no reverse slack, so its support consists of common x–z tight rows only. The extremal exchange direction is the literal selector direction $-(x - z)/\ell_r$ at 25 roots and a genuinely different tied optimum at the other 64. Eleven roots admit a strictly cheaper menu cover than the canonical staged bundle cover, with maximum saving one; in particular the sole $\tau = 4/3$ exception of (PXC67) prices at $\tau_{\text{menu}} = 1$ through the single common row $52 = \{2, 4, 5\}$ and needs no decoration reserve. Both sharp orientations sit at $\tau_{\text{menu}} = 1$ with zero margin, so the unit-exchange bound is equality-compatible. The eight-player relaxed control has outsider loss 3 and maximum gain 17/8, violating the unit-exchange bound at the allocation level; genuine Lmax optimality remains indispensable. These are exact finite calculations. The uniform equality of outsider loss and maximum gain is a striking archived regularity, possibly a harvest artifact, and is not asserted as a theorem. At this stage the working problem was the unit root-exchange bound $\tau_{\text{menu}} \le 1$, or its reserve-padded form (PXC69c); both were later refuted, as the sequel and the saturated-arm summary record.

The archived optima localize that bound to one row. Call an Lmax-tight coalition C an *aligned co-star row* at (r, a^*) if $r \in C$, $G \setminus \{a^*\} \subseteq C$, $a^* \notin C$, and C contains no strict outsider. Gainers in C contribute $-d_b$ to the selector difference, the root contributes ℓ_r, and tied outsiders contribute nothing, so

$$x(C) - z(C) = \ell_r - (D - d_{a^*}). \tag{PXC70a}$$

Core feasibility of x on the z-tight row makes the left side nonnegative. One aligned co-star row therefore already pays the entire saturated arm:

$$\boxed{D - d_{\max} \le \ell_r, \qquad \tau_{\text{menu}}(a^*) \le 1 \text{ via the menu cover } L_C = 1,} \tag{PXC70}$$

and (PXC69) reaches the target with the whole decoration and cardinality reserve to spare. Dually, if the Lmin packet contains an x-tight type-a^* co-star row all of whose decorations are tied, core feasibility of z on that row reverses the sign in (PXC70a) and gives $D - d_{\max} \geq \ell_r$. A common aligned row therefore forces the exact equality $D - d_{\max} = \ell_r$, equivalently outsider loss equal to maximum gain.

The exact audit searches the complete menu at all 89 ordered saturated archive roots and both sharp orientations. Every root has an aligned co-star row, and every root has one that is moreover common and lies in the Lmin tight support. This explains the uniform unit exchange rate: the archived equality is exactly the co-existence of both memberships in one row. In the sharp orientations every gainer has maximum gain and each of the g aligned rows is common and Lmin-tight, matching the shared-hinge geometry of $G(p, k)$. The four clean co-star rows of the eight-player relaxed control all carry Lmax slack $7/8$, so none is aligned there and the penalty is positive. At this stage the working frontier was one selection statement—an Lmax-tight aligned co-star row at some maximum-gain gainer, with the reserve-padded (PXC69c) as fallback; the next subsection refutes the universal existence of both.

Half of that localized statement is in fact unconditional, and the other half is false. For any gainer a and any Lmin-tight rooted co-star row C omitting exactly a, core feasibility of the Lmax selector on the x-tight row gives, by the coordinate expansion of (PXC70a),

$$\boxed{z(C) - x(C) = (D - d_a) - \ell_r - \sum_{j \in C \cap O} (x_j - z_j) \;\geq\; 0,} \tag{PXC71}$$

so $D - d_a$ exceeds ℓ_r by at least the strict decoration loss of the row. Saturation supplies a positive-weight row of every omission type, so on the entire ordered saturated arm

$$D - d_{\max} \;\geq\; \ell_r, \qquad \tau_{\text{menu}}(a^*) \;\geq\; \frac{D - d_{\max}}{\ell_r} \;\geq\; 1.$$

The menu-relaxed root price is never below one, and the following are equivalent at an ordered saturated root: the unit-exchange bound $D - d_{\max} \leq \ell_r$; the exact equality $D - d_{\max} = \ell_r$; existence of an aligned co-star row; commonness, with tied decorations, of every positive-weight maximum-gain packet row; and $\tau_{\text{menu}} = 1$. Equality collapses (PXC71) on every maximum-type row simultaneously, which explains the archive: all 202 positive-weight maximum-gain packet rows across the 89 ordered saturated roots are common aligned rows, and the uniform unit exchange rate has a trivially proved half.

The open half fails at seven players. An exact design search over singleton-normalized lower envelopes of integer generators with a common total produced two permanent controls, both with exactly certified selectors. The lower envelope of $(2, 0, 6, 2, 0, 2, 0)$ and $(0, 4, 0, 0, 4, 0, 4)$ has $x = (4/3, 4/3, 4, 4/3, 4/3, 4/3, 4/3)$ and $z = (4/5, 12/5, 12/5, 4/5, 12/5, 4/5, 12/5)$. Root 2 is ordered and saturated with gainers $\{1, 4, 6\}$, $D - d_{\max} = 32/15 > 8/5 = \ell_r$, no aligned co-star row, and $\tau_{\text{menu}} = 4/3$; the cover is still exchange-exact, and the decoration reserve pays (PXC69c) with margin $9/5$. The lower envelope of $(0, 7, 2, 0, 1, 2, 0)$, $(4, 0, 0, 5, 0, 0, 3)$, and $(0, 11, 1, 0, 0, 0, 0)$ is sharper: root 1 is ordered and saturated with $x = (2, 4, 1, 2, 1, 1, 1)$, $z = (14/5, 14/5, 3/5, 14/5, 3/5, 3/5, 9/5)$, reverse penalty $D - d_{\max} - \ell_r = 2/5$, and $\tau_{\text{menu}} = 2$, strictly above the demand ratio $4/3$: menu exchange-exactness fails, and the menu-cover condition (PXC69c) fails with margin $-37/60$. Both controls satisfy the raw dispersed-gain reserve (PXC63b), with margins $9/5$ and $11/60$, and both lie far above the hybrid target ($O = 44/5$ and $48/5$ against 7). Hence the six-player uniform equality, universal aligned rows, and menu exchange-exactness are small-dimension artifacts. Among the saturated-arm certificates only the reserve-padded dispersed-gain condition (PXC63b) survived this round of

archived and designed controls; it too fell to the extended search below, and both witnesses join the mandatory hostile battery for any future saturated-arm lemma.

That reinstatement is also too optimistic. First, (PXC71) pins the common rows exactly: for a common co-star row of omission type a, the vanishing left side gives

$$\boxed{\sum_{j\in C\cap O}(x_j - z_j) = (D - d_a) - \ell_r.} \tag{PXC72}$$

The strict decoration loss of every common single-omission row is therefore not a selectable quantity but a scalar invariant of its type. Consequently the common-row bounds of (PXC55) and (PXC59) close a type-a row exactly when μA_ρ plus the cardinality reserve covers $D - d_a - \ell_r$: every packet-side closure on this arm—(PXC55), (PXC59), (PXC63b), (PXC67b), (PXC69c)—is governed by the single scalar family $D - d_a - \ell_r$ against the single reserve μA_ρ + slack.

An extended exact design search then refutes that last reserve condition as well; its seeded drivers and complete progress logs are archived in `verifiers/divergence/search_archive`. Over 4,624 singleton-normalized integer lower envelopes with certified selectors (2,034 divergent pairs, 4,269 candidate ceiling roots, 109 ordered saturated roots), every ordered no-multi-omission root was saturated, 99 of the 109 roots satisfy alignment and equality, and seven violate (PXC63b). Three pinned controls summarize the boundary. The mutation control, the lower envelope of $(2,0,7,2,0,1,0)$ and $(0,4,0,0,4,0,4)$, has $x = (4/3, 4/3, 13/3, 4/3, 4/3, 1, 4/3)$ and the same Lmax point as the first control: its root 2 has (PXC63b) margin $-1/5$, (PXC69c) margin $-29/45$, no aligned row, and all three of its common co-star rows carry the pinned loss $1/5$, so (PXC55) and the canonical selection-free evaluation of (PXC59) fail as well. It is a genuine positive-floor residual—neither direct, ceiling, nor null-preclosed—whose overlap $44/5$ still clears the target 7: no selection-free packet certificate of this arm closes it. The chosen-family evaluation of (PXC59) does close it, as (PXC76) below records. The clean control, the lower envelope of $(0,3,0,0,0,3,6)$ and $(1,0,1,2,8,0,0)$, violates (PXC63b) by $5/7$ with no decorations, zero cardinality reserve, and no Lmax-tight top-omission row at all; it is direct-closed. The boundary control, the lower envelope of $(0,7,2,0,1,2,0)$, $(3,0,0,5,1,0,3)$, and $(0,11,1,0,0,0,0)$, sits at exact margin zero in both (PXC63b) and (PXC69c) with $\tau_{\text{menu}} = 1$ and an aligned row, so any corrected reserve statement must be equality-compatible with it. All three retain the hybrid target. Hence (PXC63b) is not universal either, and no selection-free packet reserve can be: closing the arm requires selection—of a chosen decorated family (PXC76), of a rooted Lmax multi-omission row, separator, hinge over-cover, or mixed facet in the (PXC66)–(PXC68) handover, or of the spectral certificate—and at the time of this refutation the mutation control was the sharpest outstanding residual.

The coloured certificate that closes the mutation control is already in the program: the sharpened spectral-anchor bound of the sparse first-stage Lmax over-cover reaches $127T/180$ there, above the target $7T/12$, with dispersion parameter $\theta = 2/5$. This is not an accident of one instance. Write the packet margin for the reserve condition (PXC63b) and the spectral margin for the sharpened anchor bound minus the target. On every audited and designed ordered saturated root the two are complementary:

$$\boxed{\max\big(\text{packet margin},\ \text{spectral margin}\big)\ \geq\ 0.} \tag{PXC73}$$

On the six-player archive the reserve closes all 89 roots while the spectral bound fails on 20 of them, with margin down to $-2T/21$; every one of those failures is a two-gainer root paid exactly by the normalized cardinality slack $T/12$. On the designed side the roles reverse: in the two-generator family with uniform gainer weights, heavy root, and integer outsider weights, the packet violation grows with the dimension (from $-1/5$ at $n = 7$ through $-3/7$ at $n = 8$ and -1 at $n = 9$ to $-4/3$

at $n = 11$ on the audited grid) while the spectral margin rescues every instance with slack growing even as a fraction of the total ($7/45$ against $-2/45$ at $n = 11$); the doubled-light equality member persists at $n = 11$ and $n = 12$, while its even-dimensional $n = 10$ sibling leaves the equality surface. The boundary is two-sided: the audited grid also contains packet-tight instances whose spectral margin is strictly negative (margins 0 against $-1/12$ and $-1/6$), and hill-climbing toward combined failure from those mirror seeds again plateaus at zero packet margin, even while the spectral deficit deepens. Both failure directions are walled at the boundary. The two margins coincide exactly on part of the closing region (25 of the 89 archive games, and on part of the audited family grid); since both margins share the true overlap, that coincidence is equivalent to the co-star bound's unexplained slack equalling the spectral top-selection and label currency total, and it is a partial-regime phenomenon, not a law. Hill-climbing the combined failure margin from five independent seeds plateaus at the exact double boundary: both margins vanish together, and every audited double-boundary instance—including the all-unit and doubled-light members of the parametric family and both plateau witnesses—has overlap exactly equal to the hybrid target and realizes the scaled sharp-family selector profile, one heavy Lmin player at $T - (n - 1)$ over a unit floor against the two-stratum Lmax point, in game realizations outside $G(p, k)$ itself. One plateau witness is a single unit move away from the mutation control: the equality surface of the conjecture and the total failure of the packet arm are adjacent in design space. These are exact finite calculations and a design-search law, not a theorem. The spectral-or-reserve alternative (PXC73) remains the governing empirical law of the arm, with the sharp profile as its equality boundary; its open half is proved below on the regimes of (PXC75) and (PXC78), and the consolidated statement is the selection conjecture (PXC79) of the saturated-arm summary.

That open half is a theorem under the Lmax geometry the designed violations actually exhibit. Call an ordered saturated ceiling root *two-stratum* if $z_r = z_a = M$ for every gainer and $z_j = \zeta$ for every outsider, with $M > \zeta \geq 0$, and *co-star fed* if every decorated co-star $\{r\} \cup (G \setminus \{a\}) \cup \{j\}$ is Lmax-tight. Put $\Delta = gk - g - 1$, where $k = n - 1 - g \geq 2$ is the outsider count. On the menu of a co-star fed root the uniform weights $\lambda_0 = 1/(g\Delta)$ on the gk decorated co-stars satisfy the first-stage over-cover system exactly, with ambient level $\eta = g\lambda_0$ and top weights $\alpha_a = (gk - k - g)\lambda_0$ and $\alpha_r = g(k - 1)\lambda_0$ summing to one; hence the canonical dispersion obeys

$$\theta \;\leq\; \frac{k-1}{\Delta}, \tag{PXC75a}$$

with equality whenever no other menu row helps. Now let $f = \min x$. Every outsider is a nongainer, so its Lmin value is at least its Lmax value ζ; orderedness lifts every gainer above every outsider, and the ceiling root has $x_r \geq M$; hence $f \geq \zeta$. Because every non-top Lmax value equals ζ and every Lmin value is at least ζ, the residual of every candidate top block is exactly $k\zeta$. Assume every nonroot Lmin value is at most M. An eligible block P satisfies $\sum_{i \in P}(x_i - f) \geq (M + s^* - f)/\theta \geq (M - f)\Delta/(k-1)$ by (PXC75a) and the nonnegative over-cover surplus; a rooted block then pays at least $M + k\zeta$ plus the eligible remainder, and a rootless block pays at least the same plus $x_r - M \geq 0$. In both cases

$$\boxed{B \;\geq\; \frac{(M-f)\Delta}{k-1} + (g+1)f + k\zeta + M - x_r \;\geq\; \frac{(M-\zeta)\Delta}{k-1} + n\zeta + M - x_r,} \tag{PXC75}$$

the second step from $f \geq \zeta$ and the positive f-coefficient $(g + 1) - \Delta/(k - 1) = k/(k - 1)$. In particular a two-stratum co-star fed root reaches the hybrid target whenever

$$x_r \;\leq\; M + \frac{(M-\zeta)\Delta}{k-1} + n\zeta - \frac{nT}{\lfloor n^2/4 \rfloor}. \tag{PXC75b}$$

This is the coloured closure the packet arm cannot supply, and it is dimension-free: no packing, decoration, or reserve enters.

The uniform outsider stratum is moreover only a normal form. If the outsiders split into several Lmax strata with values $\zeta_j < M$ and minimum $\zeta_{\min}$, then the non-top Lmax multiset is fixed, so for every candidate block the residual is at least $z(O) - \sum_j (\zeta_j - f)_+$, a block- and pairing-independent bound, and the identical argument gives

$$B \;\geq\; \frac{(M-f)\Delta}{k-1} + (g+1)f + z(O) - \sum_j (\zeta_j - f)_+ + M - x_r \;\geq\; \frac{(M-\zeta_{\min})\Delta}{k-1} + n\zeta_{\min} + M - x_r,$$

so (PXC75b) holds verbatim with $\zeta_{\min}$ in place of ζ. The exact audit checks the hypotheses—uniform top stratum, Lmax-tight decorated co-stars, nonroot Lmin values at most M—instance by instance on a fixed family grid: 37 roots qualify (27 with uniform outsiders), and the chain closes all 37, including all 16 packet-reserve violations in the grid and a battery of nine named violation controls with minimum margin 7/10; the weaker (PXC75b) closes 26. The mutation control closes at 23/3 against target 7, and the three-stratum $(5,5,5;8;(4,2,1))$ member is now covered. The canonical dispersion attains (PXC75a) with equality on 32 of the 37 roots; the strict cases exploit rootless extra menu rows, which only sharpen the bound. The theorem is genuinely complementary to the archive: none of the 89 archived ordered saturated roots satisfies its hypotheses (19 have split tops and 70 lack a fed co-star), yet all are packet-closed, so on the archive the two regimes are disjoint and the packet side suffices. Packet failure does not force the hypotheses: the clean undecorated control has uneven gainer weights, hence a split top, and its designed relabellings recur in random search as packet violations outside the theorem, each still spectral-closed with margin 2. What remains of the general spectral-or-reserve conjecture is therefore the split-top branch, unfed co-stars, and the archived two-gainer spectral failures.

The audit discipline then reverses the necessity conclusion once more. The decorated family in (PXC59) is a selectable object: each omission type may use any x-tight decorated co-star, subject only to the outsider load cap, and by (PXC72) the aligned common row's decoration loss is pinned at $D - d_{a^*} - \ell_r$ independently of its decorations. The strongest packet-side certificate on the arm is therefore the exact finite maximum

$$\boxed{\max_{\{K_a\}} \Big(\mu A + \text{slack}\Big) - \big(D - d_{a^*} - \ell_r\big),} \tag{PXC76}$$

over admissible families with one aligned common row at a z-maximal gainer—a pure Lmin-side packing optimization. Enumerating all admissible families exactly closes every named uniform-family packet-reserve violation: the mutation control closes at margin 29/30, through the family $K_1 = K_4 = \varnothing$, $K_6 = \{5\}$ and its enlargements, and the named battery closes with margins between 29/30 and 27/10. Only the split-top clean control leaves the certificate's hypotheses—it has no aligned common row at all—and it is direct-closed. Hence no audited root forces the spectral arm: the refutations of (PXC70)–(PXC72) establish the failure of the selection-free packet certificates, not of the packet colour itself, and the mutation control is no longer an outstanding residual. The selection problem in its consolidated five-mechanism form is stated as (PXC79) in the saturated-arm summary. The disjunction (PXC73) and the two-stratum theorem stand as stated, with the spectral arm a proved alternative rather than a forced necessity, and the split-top clean family remains the one violation class outside every packet-side hypothesis.

Two further exact facts sharpen that remaining class. Across every audited packet violation—the sixteen grid violations of the two-stratum audit and the ten archived search generator classes—top-uniformity and co-star feeding coincide: eighteen violations are top-uniform and fed, eight

are split-top and unfed, and no mixed class occurs, so the two (PXC75) hypotheses stand or fall together and the remaining branch of (PXC73) is one, not two. And the split-top branch now has an exact structural atlas (PXC77): the heavy gainer joins the root at the ceiling; the Lmax-tight menu is built from double co-stars containing the root, the heavy gainer, all sub-top gainers but one, and one outsider; the canonical over-cover concentrates its largest weight on the heavy gainer, with dispersion $4/7$ and weights $(3/7, 4/7)$ at the canonical control; both spectral currencies vanish, so the sharpened bound equals the actual-top bound; and the bound is exactly tight, $B = O = 9$ against target 7, at the control. On the audited split-family grid of 31 split-top roots the sharpened bound closes every instance, including both packet violations, and an archived extended grid of 75 split-top roots closes as well; the double co-stars are Lmax-tight at 11 of the 31, and the bound is tight at 14. Any split-top closure theorem must therefore be sharp, and its over-cover is heavy-gainer weighted rather than root weighted. These are exact finite structure facts; the single-heavy geometry now has its theorem. Suppose the top stratum is exactly $\{r, h\}$, the $s \geq 2$ sub-top gainers share one level and the $k \geq 2$ outsiders sit below it, every double co-star and the heavy-block row $\{h\} \cup O$ are Lmax-tight, and every nonroot Lmin value is at most M. Weighting the double co-stars uniformly by $\lambda_0 = 1/(sk + k - s)$ and the heavy-block row by $\nu = ((s-1)k - s)\lambda_0 \geq 0$ satisfies the over-cover system exactly, with ambient level $(s-1)k\lambda_0$ and top weights $\alpha_r = k\lambda_0$, $\alpha_h = s(k-1)\lambda_0$, so the canonical dispersion obeys

$$\theta \;\leq\; \frac{s(k-1)}{sk + k - s}, \tag{PXC78a}$$

the heavy weight dominating the root weight. With $f = \min x$, the non-top Lmax multiset bounds every residual by $sf + \sum_j \min(\zeta_j, f)$, since f lies below the sub-top level, and the eligibility self-bounding argument of (PXC75) gives

$$\boxed{B \;\geq\; \frac{(M-f)(sk+k-s)}{s(k-1)} + (s+2)f + \sum_j \min(\zeta_j, f) + M - x_r.} \tag{PXC78}$$

On the audited grid the eleven class-A roots—exactly those carrying both menu hypotheses—satisfy (PXC78a) with the canonical dispersion attaining the bound, and the chain closes all eleven, including both packet-reserve violations, with the canonical control at $31/4$ against target 7. The class-B and class-C sub-geometries of the atlas remain outside the theorem, but they carry no packet violation: all twenty audited such roots are reserve-closed. Every audited packet violation therefore lies in a regime with a proved closure theorem—(PXC75) for the uniform top, (PXC78) for the single-heavy split top, with the chosen-family certificate (PXC76) closing the uniform side independently—so on all audited data the outstanding universal problem is selection alone: prove that an ordered saturated ceiling root always verifiably enters one of the proved regimes or the direct arm. The two named hostile controls moreover carry three independent proved closures apiece: the exact mixed-witness projection of their labelled branches returns branch minima $11/15$ and $16/25$ against the target ratio $7/12$, each equal to its one-witness minimum after a single cut.

Saturated-arm summary

The preceding development corrected itself repeatedly; this summary states what survives, in one place and with full hypotheses. Throughout, r is an ordered saturated ceiling root of an exact singleton-normalized selector pair $x \neq z$: $z_r = M = \max_i z_i$, $x_r \geq M$, r is not a gainer, the gainer set G has $g \geq 2$ members, no Lmin-tight rooted row omits two gainers, the rooted packing is saturated at $g/(g-1)$, and $x_r \geq x_a \geq x_j$ for gainers a and outsiders j; there are $k = n - 1 - g$ outsiders, $d_a = z_a - x_a$, $D = \sum_a d_a$, $\ell_r = x_r - z_r$, and $f = \min_i x_i$.

Four dimension-free theorems stand. *Reverse inequality* (PXC71): core feasibility of z on any Lmin-tight co-star row omitting exactly a gives $z(C) - x(C) = (D - d_a) - \ell_r - \text{strictloss}(C) \geq 0$; saturation supplies rows of every type, so $D - d_{\max} \geq \ell_r$; together with the staged pinning inequality of (PXC65) the root loss is strictly positive, and the menu-relaxed root price is never below one; moreover the unit bound, the exact equality, aligned-row existence, commonness with tied decorations of every maximum-gain packet row, and $\tau_{\text{menu}} = 1$ are all equivalent. *Loss invariant* (PXC72): a common co-star row of type a has strict decoration loss exactly $D - d_a - \ell_r$, so every packet-side reserve condition is one scalar family against the reserve $\mu A_\rho + \text{slack}$. *Prefix lemma* (PXC74): a weighted-Hall violation has total gainer deficit below one, so away from boundary ties its tangent is strictly Lorenz-improving and Lmin optimality forces a root-free or all-gainer rooted blocker. *Spectral closures* (PXC75, PXC78): if the Lmax top stratum is $\{r\} \cup G$ and every decorated co-star is Lmax-tight, then $\theta \leq (k-1)/(gk-g-1)$ and $B \geq (M-f)(gk-g-1)/(k-1) + (g+1)f + z(O) - \sum_j (\zeta_j - f)_+ + M - x_r$; if instead the top stratum is exactly $\{r, h\}$ with $s \geq 2$ sub-top gainers on one level, $k \geq 2$ outsiders below it, and both the double co-stars and the heavy-block row $\{h\} \cup O$ Lmax-tight, then $\theta \leq s(k-1)/(sk+k-s)$ and $B \geq (M-f)(sk+k-s)/(s(k-1)) + (s+2)f + \sum_j \min(\zeta_j, f) + M - x_r$; both under the harmless normalization that every nonroot Lmin value is at most M.

Six closure mechanisms are proved, each with instance-checkable hypotheses:

mechanism	hypotheses	audited closures
packet reserve (PXC63b)	none beyond the arm	all 89 archive roots
direct packet	direct-arm criterion	18 archive roots; clean control
chosen family (PXC76)	aligned common row at a z-maximal gainer	every audited uniform violation
spectral chain (PXC75)	uniform top, fed co-stars	37/37 audited, 16 violations
spectral chain (PXC78)	single-heavy top, double-fed, heavy block	11/11 audited, 2 violations
branch projection	complete tight supports	both named control branches

The open problem of the arm is the selection conjecture: every ordered saturated ceiling root is closed by a mechanism whose hypotheses it verifiably satisfies,

$$\boxed{\text{reserve} \ \vee \ \text{direct} \ \vee \ (\text{PXC75 hypotheses}) \ \vee \ (\text{PXC78 hypotheses}) \ \vee \ (\text{chosen family}).} \tag{PXC79}$$

The reserve disjunct is essential: none of the 89 archived roots satisfies a spectral hypothesis set, yet all are reserve-closed. The evidence is exact and complete over the audited landscape: every audited packet violation is either uniform-top and co-star fed (18, closed by (PXC75) and independently by (PXC76)) or single-heavy split-top and unfed (8, all class-A, closed by (PXC78), with the clean control also direct-closed), and no mixed class occurs; the feeding implication "violation and uniform top force fed co-stars" held at 56 of 56 within two- and three-generator envelopes but fails in wider ones, as recorded below; the class-B/C split-top sub-geometries carry no violation at all. One step of the feeding implication is a theorem: evaluating the (PXC71) identity on a clean co-star or on the full top block shows that at a packet violation neither is ever Lmax-tight—z-tightness of a clean type-a row forces $D - d_a \leq \ell_r$, hence a maximum-gain and the penalty nonpositive, and z-tightness of $\{r\} \cup G$ forces $D \leq \ell_r$—so every Lmax-tight type row at a violation is decorated; the implication asks that all single-decoration rows be among them. A scope caution stands, in corrected form: every singleton-normalized exact game is the lower envelope of its core witnesses, which are nonnegative, so monotonicity is automatic and is no restriction at all; the genuine narrowness of the designed evidence is the generator count—two or three generators against up to exponentially many witnesses in general—and the implication has not been probed against envelopes of more generators. The boundary data any proof of (PXC79) must respect are the plateau witnesses (both certificate margins and the target margin vanish together, on the scaled sharp selector profile), the

archived two-gainer spectral failures (paid exactly by the normalized cardinality slack $T/12$), and the zero-margin boundary control.

Historically, this summary replaces a chain of six sharper universal statements—the unit root-exchange bound, universal aligned rows, menu exchange-exactness, the reserve conditions (PXC69c) and (PXC63b) as universal closures, and the necessity of coloured selection—each refuted in turn by exact designed controls that remain pinned in the verifiers and the protected-results ledger; the six-player census uniformities that suggested them are small-dimension artifacts. The surviving structure is the two-regime complementarity above.

The wide probe promptly exercised this summary. Every singleton-normalized exact game is the lower envelope of its nonnegative core witnesses, so the designed evidence's real narrowness was generator count; seeded envelopes of four to eight generators— perturbed copies of the violation shapes—immediately refute the feeding implication (185 counterexamples in the first archived run: uniform-top violations whose co-stars of one omission type are all Lmax-slack) and then (PXC79) as first stated: fourteen roots, all with the selector profile $x = (9/2, 5/4, \dots, 5/4)$ against the familiar two-stratum Lmax point, fail the reserve, the direct arm, the fully fed hypotheses, the (PXC78) hypotheses, and the chosen-family certificate simultaneously, the last because no common aligned row exists at all. The repair costs nothing. Suppose, in the uniform-top setting of (PXC75), only that the decorated co-stars of at least two omission types are Lmax-tight, say of the types in $F \subseteq G$ with $|F| \geq 2$. Weighting those $|F|k$ rows uniformly satisfies the over-cover system with ambient level $|F|\lambda_0$ and top weights $|F|(k-1)\lambda_0$ on the root and every unfed gainer and $(|F|k - k - |F|)\lambda_0 \geq 0$ on each fed gainer, and the normalizing denominator collapses by the identity $(k-1)(g-|F|+1) + |F|k - k - |F| = g(k-1) - 1$, independently of $|F|$:

$$\boxed{\theta \;\leq\; \frac{k-1}{g(k-1)-1} \quad \text{for every } |F| \geq 2,} \tag{PXC80}$$

the same dispersion bound as the fully fed case, so the entire (PXC75) chain holds verbatim under the weaker hypothesis. The chain closes the escape profile at $327/40$ and the archived feeding counterexample at $239/30$, both against target 7. Accordingly the selection conjecture (PXC79) is amended in place: its fully fed disjunct is replaced by the partially fed hypotheses of (PXC80), and no audited root escapes the amended disjunction.

Even the partially fed hypothesis is only a sufficient shape. Let $E \subseteq G \times O$ be the *tightness graph* of decorated pairs, $(a, j) \in E$ when $C_{a,j}$ is Lmax-tight. Any over-cover supported on such rows gives every outsider the common ambient level η and satisfies the normalization exactly, and writing L for the total row weight these force $L = k\eta$ and $(g+1)(L-\eta) = 1 + L$, hence $\eta = 1/(g(k-1)-1)$ and

$$\theta_{\text{cover}} \;=\; L - \eta \;=\; \frac{k-1}{g(k-1)-1}$$

for *every* pair-supported cover: the shared dispersion of (PXC75) and (PXC80) is not a coincidence but the unique value on this support class. Existence is therefore a pure transportation question with demand one per outsider, supply exactly $k-1$ per used type, and arcs E; Gale's criterion over column subsets collapses, because a proper subset has at most $k-1$ columns and needs only one reachable type, while the full column set needs two:

$$\begin{gathered} \text{the cover exists} \\ \iff \text{every outsider lies in} \\ \text{some tight pair} \\ \text{and at least two omission types appear} \end{gathered} \tag{PXC81}$$

This is strictly weaker than two fully fed types, the characterization is verified by exhaustion over all bipartite shapes with $g, k \le 4$ in the pinned verifier, and the (PXC75) chain holds verbatim under it. The selection conjecture is amended a second time: its feeding disjunct is now the pair-graph condition (PXC81), and the audited pair graphs already exercise the extra freedom, with partial types occurring alongside full ones. The condition itself is not universal at violations: three grid instances—uniform-top violations whose only tight pairs decorate the single heavy outsider—fail it and are closed by the direct arm and the chosen family, with chosen-family margins 219/70 and 8/3. It is the disjunction, not any single disjunct, that carries the conjecture. An extended family grid ($g \le 5$, $a \le 8$, $n \le 9$) sharpens this: of 508 uniform-top violations, 89 fail the condition, eleven of those escape the direct arm, and three—the first such on the two-generator family—escape the chosen family as well, carrying the rooted-triple menu of the wide-probe escapes; all three close under the rooted-menu chain (PXC83) with surplus 1/2 and chain 131/6 against target 56/3, and one is pinned exactly. The same grid refutes the floor law: an uncovered outsider can sit strictly above the Lmin floor.

The same support-class analysis closes a door on the split top. In the single-heavy split-top geometry of (PXC78), consider covers supported on the fed double co-stars—the rows of t of the s sub-top omitted types—together with the heavy-block row. An unfed sub-top type u lies in every fed row and not in the heavy block, so the ambient equality gives u the full row weight L, the same as the root's coverage; hence $\alpha_r = 0$, each fed type's weight is $L - \eta = 0$, so $L = \eta = 0$ and the normalization $\alpha_r + \alpha_h = 1$ fails:

$$\begin{gathered}\text{no pair-supported cover exists}\\ \text{on a split top with an unfed sub-top type;}\\ \text{feasible iff } t = s \text{ and } (s-1)k \ge s\end{gathered} \tag{PXC82}$$

The dichotomy is verified by exact linear programming over $2 \le s \le 5$, $2 \le k \le 4$, $1 \le t \le s$ in the pinned split-top verifier. The collapsing denominator of (PXC80) is therefore special to the uniform top, where unfed types sit in the top stratum and absorb their over-coverage as top weight; any class-B split-top closure must leave the pair support entirely.

The probe then delivered a third refutation, deeper than the first two. An eight-generator envelope on eight players produces a uniform-top packet violation, with the familiar escape profile $x = (11/2, 3/2, \dots, 3/2)$ against the two-stratum Lmax point $z = (8/3^{\times 5}, 8/9^{\times 3})$, whose Lmax-tight menu contains *no decorated co-star at all*: the tight rows are exactly the five rooted triples $\{r, a, b\}$ with a, b gainers. The pair graph is empty, so reserve, the direct arm, the pair-graph condition, the split-top hypotheses, and the chosen family all fail simultaneously— the selection conjecture as amended is refuted a third time, now because the menu itself leaves the decorated-co-star class. The spectral certificate is untouched: the canonical over-cover on the triple menu has dispersion $\theta = 1/3$ with ambient level zero, the sharpened bound closes at 34/3 against target 8, and the spectral-or-reserve disjunction (PXC73) survives every refutation round to date. The witness is pinned exactly, menu included, in the partially fed verifier. The sixth mechanism exists, and its proof is three lines, because over-cover feasibility on a rooted menu is trivial. Suppose the top stratum is uniform, every nonroot Lmin value is at most M, and some rooted undecorated row $\{r\} \cup P$ with $\emptyset \ne P \subsetneq G$, $|P| = p$, is Lmax-tight. Weighting that row alone satisfies the over-cover system with ambient level zero and $\alpha_r = \alpha_a = 1/(p+1)$ on the row members, so

$$\theta \;\le\; \frac{1}{p+1}, \tag{PXC83a}$$

and if the row can be chosen with nonnegative surplus $x_r - M \ge \sum_{a \in P}(M - x_a)$, the (PXC75) chain—which uses only the dispersion, the nonnegative surplus, the residual bound, and the caps,

never the support—holds verbatim:

$$\boxed{B \;\geq\; (p+1)(M-f) + (g+1)f + z(O) - \sum_j (\zeta_j - f)_+ + M - x_r.} \tag{PXC83}$$

At a packet violation no clean row is tight, so $p \leq g-2$ there and the mechanism is disjoint in support from the pair-graph regime. The chain closes the pinned escape exactly: the surplus is $1/2$ on any tight triple, the dispersion bound $1/3$ is attained, and the chain pays $65/6$ against target 8, under the certificate value $34/3$. The selection conjecture is amended a third time, its mechanism list now six: reserve, direct, the pair-graph condition, the split-top hypotheses, the chosen family, and the rooted-menu hypotheses of (PXC83); the completed design search around the escape closes every rooted menu (4,666 of 4,666), and the completed wide probes produce four further escapes of the five pre-(PXC83) mechanisms—all rooted-triple menus, all closed by the (PXC83) chain out of sample (surpluses $1/2$ and $5/4$, chains $65/6$ and $149/12$). The probes also produce the first ordered-but-unsaturated ceiling roots ever observed (seven, rooted packing value $5/4$ against the saturation value $4/3$ at $g=4$, every one target-safe with overlap $50/3$ against 10; one pinned): the ordered geometry does not force saturation, and the saturated-arm theorems are scope-bounded by the saturation hypothesis.

The singleton-cover principle behind (PXC83) is in fact stratum-agnostic, and it is what closes class B. Let T be the full Lmax top stratum. Any Lmax-tight row R with $r \in R \subseteq T$ is by itself an over-cover with ambient level zero and uniform weights, so

$$\theta \;\leq\; \frac{1}{|R|}, \tag{PXC84a}$$

and with the surplus hypothesis $\sum_{i\in R} x_i \;\geq\; |R|M$ and the nonroot Lmin cap, the (PXC75) argument—dispersion, nonnegative surplus, the block- and pairing-independent residual bound, block size $|T|$—gives the master chain

$$\boxed{B \;\geq\; |R|(M-f) + |T|f + z(N\setminus T) - \sum_{i\in N\setminus T} (z_i - f)_+ + M - x_r,} \tag{PXC84}$$

of which the (PXC75), (PXC78), and (PXC83) chains are the uniform, split-top, and rooted-menu instances. On the split top this reads: if the top pair $\{r,h\}$ is Lmax-tight, then $\theta \leq 1/2$ and, with $x_r + x_h \geq 2M$, the chain closes without any feeding or heavy-block hypothesis—exactly the class-B geometry (PXC82) proved unreachable by pair-supported covers. The design search that found class B confirms the mechanism: every observed split-top violation with a slack heavy-block row has the top pair Lmax-tight (41 of 41), the first such instance is pinned exactly (dispersion $1/2$, surplus 1, chain $29/2$ against target $35/3$ under certificate $31/2$), and the class-B closure conjecture—at a single-heavy split-top violation the class-A hypotheses hold or the top pair is tight—joins the amended selection disjunction.

One universal structural fact about the Lmax point itself emerged from these hunts and is worth recording. At the Lmax point, whenever $z_i > z_j$ there is a z-tight coalition containing i and avoiding j: otherwise the transfer of ε from i to j stays in the core—only coalitions containing i without j lose, and all are slack—and lexicographically improves the sorted vector, since the entry at i's position falls while everything above it is unchanged. (PXC85) The mirror holds at the Lmin point with the roles reversed. Between players at a common value no separation is forced—transfers between equal tops worsen the egalitarian objective in both directions—and the pinned menus realize this freedom: the triple-menu escape has no tight row containing a gainer and

avoiding the root. The lemma bounds every future mechanism from below: the tight menu must separate every strictly ordered pair, so a violation's menu always carries at least the separating family.

There is now a human explanation for the class-A/class-B handover on the continuous family through the pinned class-B control. Replace only its fourth generator by

$$q^3(t) = (0, 9 - t, 6, 5 + t, 0, 0, 0), \qquad 0 \le t \le 1,$$

leaving the other five generators displayed in the exact verifier. The lower selector is constant, $x = (6, 4, 2, 2, 2, 2, 2)$, while the upper selector is

$$z(t) = \begin{cases} ((9-t)/2, (9-t)/2, 4+2t, 4+2t, 1-t, 1-t, 1-t), & 0 \le t \le 1/7, \\ (31/7, 31/7, 30/7, 30/7, 6/7, 6/7, 6/7), & 1/7 \le t \le 1. \end{cases} \tag{PXC86a}$$

This is not a sampled formula. Three fixed lower rows, after complementation, successively pin the five coordinates at level 2, the coordinate at level 4, and the remaining coordinate at level 6. On the first upper branch $v(\{r, h\}) = 9 - t$ pins both top coordinates; summing the six double-co-star inequalities pins the two sub-top coordinates, and efficiency then equalizes the outsiders. On the second branch put $R = y_r + y_h$, $A = y_a + y_b$, and $O = \sum_j y_j$. The heavy-block inequality and the same six double-co-star inequalities give

$$O \ge 7 - \max_i y_i, \qquad 4R + A \ge 44, \qquad 4R + A = 3R + (R + A) \le 13 + 7 \max_i y_i,$$

so $\max_i y_i \ge 31/7$; equality pins R, then A, then O, giving the second line of (PXC86a). Full affine-envelope checking verifies core membership on the entire two parameter intervals. Consequently the two relevant row gaps are exactly

$$2M - v(\{r, h\}) = (t - 1/7)_+, \qquad z(\{h\} \cup O) - v(\{h\} \cup O) = \tfrac{1}{2}(1 - 7t)_+. \tag{PXC86}$$

The packet margin is $-(1 + t)$ before the breakpoint and $-8/7$ after it. Thus (PXC84) closes the pair-tight class-B side, (PXC78) closes the heavy-block-tight class-A side, and both apply at the breakpoint: the two active sets meet with no local gap. This continuous family theorem explains why direct perturbations of the control hand over between the two mechanisms; it does not prove the universal class-B closure conjecture.

That conjecture is in fact false, and its failure exposes a third split-top support rather than a failure of the overlap bound. Let P be the nonnegative total-20 polytope on players $0, \dots, 6$ defined by

$$\begin{gathered} y(0134) \ge 14, \quad y(0125) \ge 14, \quad y(16) \ge 37/7, \\ y(012) \ge 12, \quad y(013) \ge 12, \quad y(023) \ge 10, \\ y(0126) \ge 13, \quad y(0136) \ge 13, \end{gathered} \tag{PXC87a}$$

and put $v(S) = \min_{y \in P} y(S)$. This is an exact game by construction, and every singleton minimum is zero. Its selectors are

$$x = (6, 4, 2, 2, 2, 2, 2), \qquad z = \tfrac{1}{7}(31, 31, 23, 23, 13, 13, 6). \tag{PXC87b}$$

There is a short proof independent of selector software. Complements of the tight lower rows 012 and 013 bound the minimum coordinate by two; equality pins players $2, \dots, 6$ at two, and efficiency together with 023 then pins $(x_0, x_1) = (6, 4)$. For the upper selector, adding the first three inequalities of (PXC87a) gives $20 + y_0 + 2y_1 \ge 233/7$, so $\max_i y_i \ge 31/7$. Equality pins

$y_0 = y_1 = 31/7$ and all three rows; 16 pins $y_6 = 6/7$, the last two rows pin $y_2 = y_3 = 23/7$, and efficiency pins $y_4 = y_5 = 13/7$.

The unique ordered saturated root has packet margin $-1/7$, while the top-pair and heavy-block gaps are respectively $6/7$ and $26/7$, and the six double-co-star gaps are $(2, 0, 0, 2, 0, 0)$. Thus neither (PXC84) nor the full (PXC78) menu applies, refuting the proposed "class A or tight top pair" selection. The game is far from a counterexample to the main bound: $O(x, z) = 17$ against target $35/3$.

What replaces the failed dichotomy on this control is a *split-matching cover*. The three tight rows $0134, 0125, 16$, each with weight $1/3$, have ambient load $1/3$ and top weights $(\alpha_0, \alpha_1) = (1/3, 2/3)$, hence

$$\theta \le \tfrac{2}{3}. \tag{PXC87}$$

The incidence construction is dimension-free: for k outsiders, use one tight double co-star of each omission type on distinct outsiders and a tight heavy–outsider pair on every remaining outsider. Uniform weight $1/k$ gives $\eta = 1/k$, top weights $(1/k, (k-1)/k)$, and $\theta \le (k-1)/k$. This proves a seventh support mechanism, not its universal selection or a universal target-closing chain. The control is already closed by the direct packet arm, so it does not prove that a seventh closure arm is necessary. Exact-LP replay and an independently formulated 80-vertex rational replay are pinned in the supplement.

The split-matching arm is nevertheless genuinely nonredundant. In (PXC87a), replace the first three lower bounds by

$$y(0134) \ge 89/7, \qquad y(0125) \ge 97/7, \qquad y(16) \ge 41/7,$$

replace the last two by $y(0126), y(0136) \ge 97/7$, and leave the three Lmin-pinning rows unchanged. The same proofs give

$$x = (6, 4, 2, 2, 2, 2, 2), \qquad z = \tfrac{1}{7}(29, 29, 27, 27, 4, 12, 12). \tag{PXC88}$$

Here the packet margin is again $-1/7$, but exhaustive expansion of the direct master criterion finds no target-reaching direct packet. The top-pair and heavy-block gaps are $4/7$ and $16/7$, while the six double-co-star gaps are $(5/7, 0, 0, 13/7, 0, 0)$. Thus the pair, full (PXC78), and direct arms all fail, whereas the same three split-matching rows remain tight and carry the (PXC87) cover. The overlap is $113/7$, exceeding the target $35/3$ by $94/21$. A second verifier enumerates the 76 core vertices with a local rational eliminator and independently expands every direct-packet choice. This proves nonredundancy of the split-matching arm, not universal selection of it.

As a first bounded attack on selection, at the selector profile in (PXC88) there are 65 proper rows compatible with the comparison sign and 158 positive normalized top over-covers supported on exactly three rows. Among the 18 covers whose rows all have size at least three and whose two top weights are positive, exact minimal lifts leave only six ordered saturated roots with singleton normalization. All six land on packet margin zero and make the top pair tight; the other twelve fail singleton normalization. Hence this complete three-row/profile cell contains no escape. The four-row census has 714 positive-row-weight covers, of which 336 have both top coefficients positive and 168 are all-large-row lift candidates. Of their exact minimal lifts, 156 fail singleton normalization and twelve are valid, but none of the twelve is a packet violation. Thus the fixed-profile escape requires at least five rows. At five rows, collapsing duplicate incidence columns and the swaps of the two sub-top players and the two equal outsiders leaves 720 positive column covers, 924 labelled covers, and 252 orbits. Of the representative lifts, 250 fail singleton normalization; the remaining two are genuine packet violations, but both make the top pair tight. Hence no all-large-row fixed-profile support escape of size at most five exists. A zero-ambient cover on a two-player top can use

only rows inside that top; singleton normalization excludes the singleton rows, so it would force the top pair tight. Thus pair slack forces positive ambient load. With both top coefficients positive, the six equality rows then bound an extreme support by five rows. A zero top coefficient contributes one additional active equation and can permit six rows. The remaining fixed-profile cells therefore contain two-player rows or six-row one-sided top covers. This is a finite incidence-and-lift theorem only: different selector profiles and different Lmin certificates remain live.

The apparent one-sided-top exception can in fact be removed. Suppose the Lmax top is exactly $T = \{r, h\}$. Lexicographic optimality excludes every core-tangent direction with $d_r, d_h \le 0$ and at least one strict inequality: for sufficiently small steps the strict gap below T is preserved and the sorted ceiling vector improves. The polyhedral Gordan–Stiemke alternative therefore gives a normalized top over-cover with $\alpha_r, \alpha_h > 0$. At a single-heavy split packet, $x_r > M > x_h$, and the selector-payment inequality

$$\alpha_r(x_r - M) \ge \alpha_h(M - x_h)$$

rules out $\alpha_r = 0$ for every normalized top over-cover. Decompose the strictly positive cover into extreme covers. Recession rays have zero top weight, so some extreme cover has $\alpha_h > 0$, and the payment inequality also gives $\alpha_r > 0$. If the top pair is slack, its ambient load is positive by the preceding zero-ambient argument. After eliminating the two top weights, the normalized incidence system has $n - 1$ equality rows; an extreme cover with positive ambient variable consequently uses at most $n - 2$ tight coalitions. (PXC89)

For the seven-player coordinates in (PXC88), (PXC89) makes the incidence census complete through the extreme-support bound rather than merely all-large. Allowing every nonsingleton comparison-compatible row gives, at support sizes one through five, respectively $1, 6, 66, 312, 2576$ strictly-positive-top labelled covers and $1, 2, 20, 96, 719$ exact symmetry orbits. At sizes one and two every orbit is a valid packet-boundary lift with the top pair tight. At sizes three and four, respectively 16 of 20 and 92 of 96 representative minimal lifts fail singleton normalization; every remaining lift retains either the top-pair or split-matching support. At size five, 717 of 719 representatives fail singleton normalization and the other two are genuine packet violations, both with the top pair tight relative to their recomputed selectors. These are exact counts and a useful support atlas, but not a fixed-profile completion theorem: the census recomputes each minimal lift's selectors and does not assert that they remain the prescribed pair (x, z) from (PXC88). Singleton-normalization failure is completion-monotone, but a mechanism tight at a changed minimal-lift selector need not be tight at the prescribed z. The failure is realized explicitly in (PXC92a).

Those other profiles do refute the corrected four-arm classification. Let P be the nonnegative total-20 polytope cut out by the three Lmin rows $y(012), y(013) \ge 12$, $y(023) \ge 10$, and

$$y(0124), y(0134) \ge 94/7, \qquad y(0125), y(0136) \ge 95/7, \qquad y(1456) \ge 55/7, \tag{PXC90a}$$

and define $v(S) = \min_{y \in P} y(S)$. Then v is exact and singleton-normalized, with

$$x = (6, 4, 2, 2, 2, 2, 2), \qquad z = \tfrac{1}{7}(29, 29, 28, 28, 8, 9, 9). \tag{PXC90b}$$

The lower selector has the same short proof as (PXC87): the complement of 012 bounds the floor by two, equality and 013 pin players $2, \dots, 6$, and 023 pins the remaining ordered pair. For the upper selector, weight the five rows in (PXC90a), in their displayed order, by $(1, 1, 2, 2, 1)/7$. Their incidence load is

$$\frac{3}{7}\mathbf{1}_N + \frac{3}{7}e_0 + \frac{4}{7}e_1, \qquad \text{and their weighted lower bound is } 89/7.$$

Thus every core point has maximum at least $29/7$. Equality pins players $0, 1$ at $29/7$ and all five positive-weight rows; those equations with efficiency uniquely give the displayed z.

The unique ceiling root is ordered and saturated, with one heavy top gainer, two equal sub-top gainers, and packet margin $-2/7$. Nevertheless the top pair has gap $29/28$, the full (PXC78) menu fails because two double co-stars have gap $11/7$, no (PXC87) split matching is tight, and exhaustive expansion of the direct criterion returns false. Hence

$$\boxed{\text{top pair } \vee \text{ full (PXC78) } \vee \text{ (PXC87) matching } \vee \text{ direct packet}} \tag{PXC90}$$

is false even for seven-player singleton-normalized exact games. The paying top dual is a new forked support: both omission types meet outsider 4, one continues to outsider 5, the other to outsider 6, and the tight heavy block closes the outsider loads. It has weights $(1,1,2,2,1)/7$, ambient load $3/7$, and top weights $(3/7, 4/7)$. This refutes support selection, not the main overlap bound: $O(x,z) = 111/7$, exceeding $35/3$ by $88/21$. An independent verifier enumerates all 84 core vertices, reconstructs every coalition minimum, checks both human selector certificates, and independently expands the direct and matching predicates.

The exact boundary between the preceding theorem and its finite census can now be stated without reference to a menu shape. Let $\mathcal{C}_x$ be a complete staged Lmin certificate at x, let $\mathcal{R}$ support a normalized z-tight top over-cover, and let $P_{\mathcal{R}}$ be the polytope obtained from efficiency and nonnegativity by retaining the rows in $\mathcal{C}_x \cup \mathcal{R}$ at their displayed x- and z-values. Write $w_{\mathcal{R}}(Q) = \min_{y \in P_{\mathcal{R}}} y(Q)$ for its exact lower-envelope game. Two facts are profile-independent. First, the incidence identity and equal totals give the exact selector payment

$$\sum_{S \in \mathcal{R}} \lambda_S \big(x(S) - z(S)\big) = \alpha_r(x_r - M) - \alpha_h(M - x_h) \ \geq 0. \tag{PXC91a}$$

This is the equality behind (CB5), rather than a new inequality. Second, if P is any completion with $z \in P \subseteq P_{\mathcal{R}}$ and lower-envelope game v, then

$$w_{\mathcal{R}}(Q) \leq v(Q), \qquad w_{\mathcal{R}}(Q) = z(Q) \Longrightarrow v(Q) = z(Q). \tag{PXC91b}$$

Consequently a positive singleton minimum in the minimal lift rules out every singleton-normalized completion, while any row already tight at z remains tight in every completion retaining z. This is the monotone minimal-lift lemma used by the census.

What is fixed in that census is now explicit. The comparison-compatible row set (65 rows and 52 nonsingleton incidence columns), the swaps $2 \leftrightarrow 3$ and $5 \leftrightarrow 6$, the three staged Lmin rows, the minimal-lift selector calculations, and all labelled-cover and orbit counts belong to the particular pair (PXC88). Only comparison compatibility, (PXC91a)–(PXC91b), strict positivity of the two top weights, positive ambient load under pair slack, and the $n-2$ extreme-support bound survive a change of profile. In particular, "nonviolating" is not itself monotone under completion. No completion lift follows unless the witnessing closing rows are also verified tight at the prescribed z; the incidence census did not perform that fidelity check.

The tempting profile-independent bridge is false. Namely, one might ask whether every residual root has a strictly positive extreme top cover $\mathcal{R}$, $|\mathcal{R}| \leq n-2$, for which

$$\mathrm{Lmin}(w_{\mathcal{R}}) = x, \qquad \mathrm{Lmax}(w_{\mathcal{R}}) = z. \tag{PXC91c}$$

Here $\mathcal{C}_x$ is fixed complete staged-certificate data for the root; (PXC91c) is the unrestricted question over residual roots, not yet restricted to the later $s = 2$, $k \geq 2$, reserve/direct-failing cell. The game (PXC88) is already an exact counterexample. Its only proper nonsingleton rows tight at z are

$$0134, \quad 0125, \quad 16, \quad 0126, \quad 0136.$$

Writing their cover weights as (a, b, c, d, e) and the common outsider load as η, the outsider equations are

$$a = \eta, \quad b = \eta, \quad b + d = \eta, \quad a + e = \eta, \quad c + d + e = \eta.$$

Thus $d = e = 0$, $a = b = c = \eta$, and normalization gives $\eta = 1/3$. Consequently the unique normalized cover with both top coefficients positive has positive support $\mathcal{R} = \{0134, 0125, 16\}$ with weights $1/3$. Retaining this cover and the complete staged Lmin rows $012, 013, 023$ preserves x but changes the upper selector to

$$z' = \tfrac{1}{7}(29, 29, 26, 26, 5, 13, 12) \neq z. \tag{PXC92a}$$

These are all proper nonsingleton rows tight at x, and an exact replay of all seven proper subsets shows that each subset changes Lmin; hence no alternative complete retained row pool repairs this example. Indeed the cover equality fixes the first ceiling and gives $y_6 = 12/7$, $y_2 + y_5 = 39/7$, and $y_3 + y_4 = 31/7$; the retained Lmin rows then pin the second ceiling at $26/7$. The two zero-cover-weight rows 0126 and 0136 separately raise the two second-stratum coordinates to $27/7$ in the full game. Each is indispensable, and retaining both restores z. Hence an extreme top cover can certify the first Lmax stage while losing later stages. PXC90 remains the positive calibration: equality in its top-cover bound also pins $y_0 = y_1 = 29/7$; those coordinate equalities, the five row equalities, and efficiency have rank seven and uniquely pin its displayed z.

There is an exact replacement for the failed bridge. For a compact retained-row polytope P_H and $p \in P_H$, let $T_H(p)$ be its active tangent cone: $\mathbf{1}_N \cdot d = 0$, $d(S) \geq 0$ for every retained row active at p, and $d_i \geq 0$ whenever $p_i = 0$. Write the payoff strata of z in descending order as $B_1, \ldots, B_m$, and $E_{k-1} = B_1 \cup \cdots \cup B_{k-1}$. Then

$$\mathrm{Lmax}(P_H) = z \quad \Longleftrightarrow \quad \forall k \; \forall i \in B_k : \; \nexists d \in T_H(z) : d|_{E_{k-1}} = 0, d_j \leq 0 \; (j \in B_k), d_i < 0 \tag{PXC92b}$$

The Lmin statement uses ascending strata and excludes directions with $d|_{E_{k-1}} = 0$, $d_j \geq 0$ on B_k, and a designated $d_i > 0$. This is just the staged lexicographic optimality condition, but it exposes the missing hypothesis in (PXC91c). For example, $z' - z = (0, 0, -1, -1, 1, 1, 0)/7$ is an admissible second-stage direction in the PXC88 minimal lift.

Equivalently, each forbidden direction in (PXC92b) is excluded by a Farkas identity using active retained-row normals, the previously fixed coordinate normals, and sign normals on the current block. A complete staged Lmin certificate therefore preserves x automatically. Rows of $\mathcal{R}$ and all auxiliaries below are tight at the original z; comparison compatibility $x(S) \geq z(S)$ keeps x feasible, while $z(C) \geq x(C)$ on the x-tight rows keeps z feasible. The nonautomatic requirement is a complete collection of the later Lmax identities.

Assume $\mathrm{Top}(z) = \{r, h\}$ and begin with a normalized z-tight top cover $\mathcal{R}$ having $\alpha_r, \alpha_h > 0$. For each non-top player, choose one basic Farkas identity for its z-stage using the full original z-tight menu, and let $\mathcal{A}$ be the union of its coalition supports outside $\mathcal{C}_x \cup \mathcal{R}$. Retaining $\mathcal{A}$ preserves every later ceiling stage and hence both selectors. After quotienting by efficiency and previously fixed coordinate equalities, conic Carathéodory in a space of dimension at most n gives a certificate using at most n coalition normals (sign and nonnegativity normals only lower that count). Taking the union over the $n - 2$ non-top players gives the safe bound

$$|\mathcal{A}| \leq n(n - 2) \tag{PXC92c}$$

on additional rows when the top stratum has two players. Repeated rows can make the actual set much smaller: PXC88 needs exactly two and PXC90 needs none. Thus sparse cover selection has been replaced by a rigorous bounded auxiliary-row theorem.

There is a genuine six-player obstruction to compressing the PXC88 failure. Write the roles as r, h, a, b, j, k, where $\{r, h\}$ is the top stratum, $z_a = z_b = q < M$, and j, k are the two strict outsiders. Work in the ordered saturated strict single-heavy cell: the game is exact and singleton-normalized, x, z are efficient, $x_r > M > x_h$, the gainers are exactly h, a, b, the outsiders are nongainers, and the stated Lmin order holds. Suppose the following two rows are z-tight. The two-row split matching

$$R_j = \{r, h, b, j\}, \qquad R_k = \{r, h, a, k\}$$

has weights $1/2$, ambient load $1/2$, and top weights $(1/2, 1/2)$. The only opposite-omission rows on the same two outsiders are

$$A_j = \{r, h, a, j\}, \qquad B_k = \{r, h, b, k\}.$$

But A_j, B_k, again with weights $1/2$, form a second normalized positive top cover. Using the other assignment merely reproduces R_j, R_k. Consequently the two later-stage repair rows cannot both be forced to have zero coefficient in every positive cover, as $0126, 0136$ are in PXC88.

The two natural four-role alternatives are also impossible. Put $d_i = z_i - x_i$ on the three gainers, $L_j = x_j - z_j$, $L_k = x_k - z_k$, and $\ell = x_r - M$. Efficiency gives $d_h + d_a + d_b = \ell + L_j + L_k$. Hence

$$z(\{h, a, b, j\}) - x(\{h, a, b, j\}) = \ell + L_k > 0, \qquad z(\{h, a, b, k\}) - x(\{h, a, b, k\}) = \ell + L_j > 0,$$

so neither row can be z-tight while x remains in the core. On the other hand the split cover implies $y_r + y_h \geq 2M$ at every core point. If, say, $\{r, a, b, j\}$ were z-tight, singleton exactness would supply a core point with $y_r = 0$, and then

$$y(\{r, a, b, j\}) \leq T - 2M = 2q + z_j + z_k < M + 2q + z_j = z(\{r, a, b, j\}),$$

a contradiction; the $\{r, a, b, k\}$ case is symmetric.

In fact the literal two-row split cover admits a complete no-go theorem. On the face where the earlier Lmax stratum is pinned, a tangent direction u has $u_r = u_h = 0$. Put $t = u_a$ and $s = u_b$. Tightness of R_j, R_k forces

$$(u_a, u_b, u_j, u_k) = (t, s, -s, -t).$$

Every later z-tight row therefore cuts this two-dimensional tangent cone by one of the eight nonzero forms $\pm t,\ \pm s,\ \pm(t - s),\ \pm(t + s)$. Any family excluding $(t, s) \in \mathbb{R}^2_- \setminus \{0\}$ has a closing subfamily of at most two. After comparison compatibility with x, the canonical a-blocker types and their b-symmetric counterparts give the following pair table; the one-row $t + s$ variants are handled separately:

	bk	*hbk*	*rb*	*rab*	*rbk*	*rabk*	*rhb*	*rhbk*
aj	*I*	*I*	*F*	*X*	*I*	*F*	*X*	*I*
haj	*I*	*I*	*X*	*X*	*I*	*X*	*X*	*I*
ra	*F*	*X*	*F*	*X*	*F*	*F*	*X*	*F*
rab	*X*	*X*	*X*	*X*	*X*	*X*	*X*	*X*
raj	*I*	*I*	*F*	*X*	*I*	*F*	*X*	*I*
rabj	*F*	*X*	*F*	*X*	*F*	*F*	*X*	*F*
rha	*X*	*X*	*X*	*X*	*X*	*X*	*X*	*X*
rhaj	*I*	*I*	*F*	*X*	*I*	*F*	*X*	*I*

(PXC93)

Here I means that a common negative ray remains, X means that cover compatibility and the necessary reserve-failure inequality are inconsistent, and F means scalar feasibility but failure at the exactness witness with $y_r = 0$. The load-bearing identities are short. For example, compatibility

of rha gives $L_j + L_k \le d_b$, hence $d_h + d_a = \ell + L_j + L_k - d_b \le \ell$, contradicting reserve failure; compatibility of rab gives $d_a + d_b \le \ell$. In an F cell the split cover first gives $y_h \ge 2M$. For aj, rb, the two blocker rows then give

$$T \ge 3M + 2q + z_j > 2M + 2q + z_j + z_k = T,$$

and a row $rabj$ already gives the same contradiction by itself. The swap $a, b \leftrightarrow b, a,\ j, k \leftrightarrow k, j$ supplies the genuine symmetric cases; no monotonicity in top membership is used. An exact rational verifier classifies every displayed cell first without and then with the root witness, reproducing the table. Independently of that canonicalization it enumerates all 44 proper nonsingleton labelled coalition masks with nonzero tangent form and all 275 closing labelled one- or two-row families, including the four one-row $t + s$ variants; none survives, while the weaker base necessary system is feasible. It also exhausts all 192 closing subfamilies of the eight-form set and verifies that no minimal closing family needs more than two rows.

Thus, in the ordered saturated exact cell just fixed, reserve failure and the root singleton witness prevent the later z-tight rows from closing the second-stage Lmax tangent cone. A six-player game cannot realize the PXC88 failure through the literal two-row split cover at all. The unused third outsider is genuinely load-bearing. This excludes a whole cover mechanism, not the other nonsplit extreme-cover incidences through four rows.

The direct-packet route can be written as an explicit two-outsider inequality. Let $D = d_h + d_a + d_b$, $d_{\max} = \max(d_h, d_a, d_b)$, $E = D - d_{\max}$, and let A_ρ be the saturated decorated co-star load. Since $n = 6$ and $g = 3$, the cardinality term in (PXC63b) vanishes, so reserve failure is exactly

$$E - \ell > \frac{4}{3} A_\rho. \tag{PXC93a}$$

For $p \in \{j, k\}$ let u denote the other outsider, take $P = \{p\}$, $H = \{r, u\}$, and omit a maximum-gain player from the direct packet. The master identity has cardinality base $7/10 > 2/3$ and its exact admissible interval is

$$\ell + L_u \le E \le B_p := \frac{-x(G) - 3x_p - (M + z_u) + \frac{8}{5}(T - z_p)}{3}. \tag{PXC93b}$$

It would suffice to prove that (PXC93a) forces (PXC93b) for one of the two outsiders. This statement is not a consequence of allocation arithmetic alone. The following rank argument bypasses that implication and the false generic premise that every later blocker can be pivoted into a positive cover.

Assume first that the top pair is slack. By (PXC89), choose a normalized z-tight extreme top cover with positive ambient load η, positive top coefficients, and support $\mathcal{R} = \{S_1, \dots, S_m\}$, $m \le 4$. Let $A \in \{0, 1\}^{4 \times m}$ be its incidence matrix on the four non-top players a, b, j, k. The outsider equations are $A\lambda = \eta \mathbf{1}_4$. Extremality implies

$$\operatorname{rank} A = m. \tag{PXC94a}$$

Indeed, since $\eta > 0$, the column $-\mathbf{1}_4$ lies in the span of A. If $\operatorname{rank} A < m$, the outsider-balance solution cone has dimension at least two, and the one remaining normalization equation cannot isolate the strictly positive ray (λ, η).

This rank statement turns support four into selector fidelity. If $m = 4$ and y is any retained-row feasible point with $\max_i y_i \le M$, the cover inequality and efficiency give

$$\eta T + \alpha_r y_r + \alpha_h y_h \ge \eta T + M.$$

Because $\alpha_r, \alpha_h > 0$, their sum is one, and both top coordinates are at most M, equality is forced: $y_r = y_h = M$ and every positive cover row is tight. Thus $A^T(y_{\{a,b,j,k\}} - z_{\{a,b,j,k\}}) = 0$; (PXC94a) gives $y = z$. The cover therefore makes z the unique first-ceiling minimizer, while a complete staged $\mathcal{C}_x$ preserves x as in (PXC92b). The original singleton witnesses remain feasible in this relaxed polytope, so its lower envelope remains singleton-normalized.

The smaller supports cannot occur in the reserve-failing cell. Support one is impossible when $\eta > 0$. At support two the outsider columns are a nontrivial complementary partition, both rows contain r, h, and both weights are $1/2$. Singleton-gainer and $ab \mid jk$ partitions contradict comparison compatibility and (PXC93a); the two cross partitions are exactly the split faces excluded by (PXC93). The only new boundary is $rhj \mid rhabk$ (or its symmetric copy), where compatibility can force $L_j = 0$. Its second-stage tangent has precisely the six nonzero forms $\pm t, \pm s, \pm(t+s)$. An exact rational necessary-system enumeration finds 40 closing form families, each with a closing subfamily of at most two, and all 381 labelled closing one- or two-row families become infeasible when the blocker and root-zero witness are added to the full displayed necessary system.

At support three the balanced rank-three 0–1 hypergraphs have six colored placements of two complementary unweighted trace types,

$$\{a, b, jk\},\ \{a, j, bk\},\ \{ab, j, k\}, \quad \text{and their complements.} \tag{PXC94b}$$

The first and third shapes and all complement shapes are excluded by top positivity, comparison compatibility, efficiency, and the pair-sum consequence of (PXC93a). Of the 37 colored support-three orbits, 34 already fail that scalar system. Exactly three decorations of the delicate second trace survive it; the root-zero witness immediately excludes the versions $(ra, hj, rhbk)$ and $(ra, rhj, rhbk)$. The only unresolved equality boundary has rows ra, rhj, hbk and forces $d_a = \ell$, $L_j = 0$, $L_k = d_h + d_b$. Lmax optimality would then require a tight row containing b and omitting k. All 15 labelled possibilities are infeasible after adjoining that blocker to the full necessary system. Possible z-tight zero-valued outsider singletons account for 41 of the 807 labelled covers in 14 colored orbits: six size-three covers in two orbits and 35 size-four covers in 12 orbits. Every such orbit is already infeasible under the weaker scalar reserve system.

Consequently, in the ordered saturated strict single-heavy six-player cell, reserve failure and top-pair slack force every positive extreme top cover to have support four, and hence force a selector-faithful minimal lift:

$$\exists \mathcal{R}, \quad |\mathcal{R}| = 4, \qquad \mathrm{Lmin}(w_{\mathcal{R}}) = x, \quad \mathrm{Lmax}(w_{\mathcal{R}}) = z. \tag{PXC94c}$$

The complete incidence audit enumerates 807 labelled positive extreme covers in 259 role-preserving colored orbits. The sole zero-ambient cover is the tight top pair; removing it leaves 806 covers in 258 orbits. Every positive-ambient cover has the full outsider-incidence column rank in (PXC94a). Independent exact boundary verifiers retain feasible base systems, so the zero-survivor conclusions are not artifacts of an inconsistent relaxation.

The top-pair-tight case must be separated. It is the unique $\eta = 0$ cover $\{r, h\}$ and does not by itself preserve later Lmax stages. Nevertheless the target residual cell is empty. After the top block is pinned, enumerate the inclusion-minimal non-top incidence-trace families that close the second Lmax stage. There are 11, of size at most three. Decorating them by all admissible r, h memberships gives 445 labelled families and 5,340 branches after choosing a maximum gainer and one of the two strict failure sides of each interval (PXC93b). A continuous exact-rational necessary system retains role order, efficiency, comparison compatibility, reserve failure, both direct-interval failures, and a root-zero exactness witness. It has 642 feasible branches before the witness and none after it:

$$\{r, h\} \text{ tight}, \quad \text{reserve failure}, \quad \text{both (PXC93b) packets fail} \quad \Longrightarrow \quad z \neq \mathrm{Lmax}\,. \tag{PXC95}$$

This is a finite incidence theorem over continuous rational variables, not a bounded payoff search; the feasible pre-witness branches and a separate fail-open base prevent vacuous infeasibility. Hence every exact singleton-normalized game in the stated ordered cell with a tight top pair is already closed by the reserve or one of the two direct packets. PXC84 gives an independent closure of the same face for the hybrid-overlap inequality.

Together, (PXC94) and (PXC95) close the six-player reserve/direct-failing top-cover frontier. They do not reinstate the unrestricted conjecture (PXC91c), which fails at seven players by (PXC92a), nor do they prove pure selector fidelity for tight-top-pair games outside this residual cell.

The sharp dimension statement for the restricted frozen-domain conjecture is therefore now complete. The single-heavy $s = 2$ arm has three gainers—the heavy player and two sub-top players—and, by definition here, at least two outsiders; hence $n \geq 6$. At $n = 6$, reserve and direct failure rule out a tight top pair by (PXC95), while top-pair slack invokes (PXC94c). Thus every game in the frozen domain satisfies the faithful sparse-lift assertion at six players. The exact game (PXC88) satisfies the same hypotheses at seven players and fails it by (PXC92a). Consequently

$$\boxed{n = 6 : \text{ restricted (PXC91c) holds;} \qquad n = 7 : \text{ restricted (PXC91c) can fail.}} \tag{PXC96}$$

Here "restricted" means precisely the ordered, saturated, single-heavy $s = 2$, $k \geq 2$, reserve/direct-failing domain fixed above; this is not the unrestricted assertion (PXC91c). PXC88 is the smallest exact counterexample in that restricted domain. Its unique positive extreme cover omits exactly the two indispensable later-stage rows $0126, 0136$; (PXC92c) is the corrected lift theorem, with at most $n(n-2)$ additional rows in general and exactly those two rows in PXC88.

For the remaining global minimum-dimension question, the quantifiers are fixed as follows. An unrestricted instance is a tuple $(v, x, z, \mathcal{C}_x)$ on $4 \leq n \leq 6$ players such that v is exact and singleton-normalized, $x = \mathrm{Lmin}(v) \neq z = \mathrm{Lmax}(v)$, $\mathrm{Top}(z) = \{r, h\}$, and $\mathcal{C}_x$ is any fixed collection of original x-tight nonempty proper-coalition rows whose retained polytope (with efficiency and nonnegativity) has Lmin equal to x. Collections are identified when they define the same retained polytope, so repetitions and rows already supplied by efficiency or nonnegativity do not create distinct certificates. Thus the strong question is universal over every such fixed $\mathcal{C}_x$; the weaker question permits choosing $\mathcal{C}_x$ after (v, x, z) is known. A qualifying top cover means an extreme normalized identity

$$\sum_{S \in \mathcal{R}} \lambda_S \mathbf{1}_S = \eta \mathbf{1}_N + \alpha_r e_r + \alpha_h e_h, \qquad \lambda_S, \alpha_r, \alpha_h > 0, \quad \eta \geq 0, \quad \alpha_r + \alpha_h = 1, \tag{PXC97a}$$

using original z-tight proper-coalition rows and $|\mathcal{R}| \leq n - 2$. Ambient positivity is not assumed: when the top pair is tight, $\mathcal{R} = \{rh\}$ is the canonical $\eta = 0$ cover. The strong unrestricted assertion in dimension n is

$$\forall (v, x, z, \mathcal{C}_x)\ \exists \mathcal{R} : \qquad \mathrm{Lmin}(w_{\mathcal{R}}) = x, \quad \mathrm{Lmax}(w_{\mathcal{R}}) = z. \tag{PXC97b}$$

No orderedness, saturation, gain/loss pattern, reserve or direct condition, named support shape, or later-stratum tie pattern is part of (PXC97b). The global minimum failure dimension is the least n for which (PXC97b) fails; PXC88 previously proved only that it is at most seven.

Three unrestricted reductions are immediate. First, there are no instances of (PXC97b) for $n \leq 4$, because the small-game coincidence theorem gives Lmin = Lmax. Second, suppose a qualifying cover has $\eta > 0$ and its non-top incidence matrix $A_{\mathcal{R}}$ has full row rank $n - 2$. At any point with ceiling at most M, equality in the cover bound pins both top coordinates to M and

every positive-weight cover row to equality. Hence $A_{\mathcal{R}}^{\mathsf{T}}(y_{N\setminus\{r,h\}} - z_{N\setminus\{r,h\}}) = 0$, so full row rank pins $y = z$. Such a cover is faithful through every later stratum, independently of all role geometry. Third, if $\eta = 0$, positivity and the outsider load equations force every support row to avoid all non-top players. Singleton normalization excludes the top singletons from the z-tight menu, leaving only the tight pair $\{r, h\}$, with weight $1/2$. Thus

$$\boxed{\operatorname{rank} A_{\mathcal{R}} = n - 2 \Longrightarrow \mathcal{R} \text{ is faithful;} \qquad \exists \text{ a qualifying } \eta = 0 \text{ cover} \Longleftrightarrow z(rh) = 2M.} \tag{PXC97c}$$

At five players Theorem 24 gives $M = \nu$, and an attaining coalition lies inside the Lmax top stratum by Lemma 30(a). With exactly two top players it must be $\{r, h\}$; consequently every five-player instance of (PXC97b) is in the pair-tight, $\eta = 0$ cell. The five-player statement is therefore precisely whether the common-tight rows of every fixed $\mathcal{C}_x$, together with the top pair and active nonnegativity, supply all PXC92b polars on the remaining three coordinates.

The analogous six-player question is already false, however, and it fails in both the pair-tight and positive-ambient cells at once. On players $(r, h, a, b, j, k) = (0, 1, 2, 3, 4, 5)$ let P be the nonnegative total-17 polytope

$$y(rha) \geq 11, \qquad y(rhb) \geq 11, \qquad y(rj) \geq 6, \qquad y(rhjk) \geq 11, \tag{PXC98a}$$

and define $v(S) = \min_{y \in P} y(S)$. The polytope is nonempty and compact, every minimum is attained at a point of $P \subseteq \operatorname{Core}(v)$, and hence v is exact. The points $(0, 11, 0, 0, 6, 0)$ and $(17, 0, 0, 0, 0, 0)$ give a zero witness for r and for each other player respectively, so v is singleton-normalized. Exact staged optimization gives

$$x = \operatorname{Lmin}(v) = (9/2, 9/2, 2, 2, 2, 2), \qquad z = \operatorname{Lmax}(v) = (4, 4, 3, 3, 2, 1). \tag{PXC98b}$$

Indeed the complements of the first two rows in (PXC98a) have sums at most 6, so the Lmin floor is at most 2; equality pins a, b, j, k at 2 and then equalizes r, h at $9/2$. For Lmax, $v(rh) = 8$ forces the first ceiling to be at least 4. At equality $r = h = 4$; the first two rows force $a, b \geq 3$, efficiency pins $a = b = 3$ and $j + k = 3$, and the last two rows then give $j = 2, k = 1$ by the later ceiling stages.

The complete proper x-tight menu is exactly $\{rha, rhb\}$, and neither the empty row-set nor either one-row subset has Lmin equal to x. Thus $\mathcal{C}_x = \{rha, rhb\}$ is the unique complete certificate even under the weaker existential-$\mathcal{C}_x$ reading. The complete proper z-tight menu is

$$\{rh, rha, rhb, rj, rhjk\}.$$

Writing its cover weights as (A, B, C, D, E) in that order, the four outsider load equations and normalization reduce to

$$B = C = E = \eta, \qquad D = 0, \qquad 2A + 4\eta = 1. \tag{PXC98c}$$

Consequently the normalized cover polytope is a segment and has exactly two positive extreme supports of size at most four:

$$\{rh\} : \ (A, \eta) = (1/2, 0), \qquad \{rha, rhb, rhjk\} : \ (B, C, E, \eta) = (1/4, 1/4, 1/4, 1/4),$$

both with $(\alpha_r, \alpha_h) = (1/2, 1/2)$. The later blocker rj has coefficient zero in every normalized cover. Retaining either extreme cover with the unique $\mathcal{C}_x$ preserves Lmin and singleton normalization but changes the upper selector to

$$\operatorname{Lmax}(w_{\mathcal{R}}) = (4, 4, 3, 3, 3/2, 3/2) \neq z, \qquad d = (0, 0, 0, 0, -1/2, 1/2). \tag{PXC98d}$$

The direction d fixes every earlier z-stratum and is invisible to all retained rows, while $d(rj) = -1/2$. Thus (PXC98) refutes both the strong fixed-$\mathcal{C}_x$ and weak existential-$\mathcal{C}_x$ versions of (PXC97b) at $n = 6$. A standalone verifier exhausts all four x-tight subsets and every positive extreme top-cover support through size four, and an independent verifier imports no project LP, selector, cover, or game implementation. Their reproduction commands are

```
cd verifiers/divergence/search_archive
python3 unrestricted_pxc91c_exact_search.py \
  --output unrestricted_pxc91c_exact_search.json
cd ..
python3 verify_unrestricted_pxc91c_pxc98_independent.py \
  --output unrestricted_pxc91c_pxc98_independent.json
```

Therefore the unrestricted minimum failure dimension is at most six. Since $n \le 4$ is impossible, the five-player pair-tight residual problem is the only remaining lower-dimensional case. As finite evidence on that last cell, an exact audit of the existing five-player bank finds 165 distinct exact games, of which 36 have a two-player Lmax top stratum. It exhausts every x-tight row subset and all 37 complete certificates and 146 qualifying extreme covers across those games, with no strong or weak $\mathcal{C}_x$ failure. This is an exhaustive audit of the pinned bank, not of the five-player exact cone.

The remaining five-player statement has a compact exact polar form. Put $U = \{a, b, c\}$, let $\mathcal{C}_0 = \{C \in \mathcal{C}_x : x(C) = z(C)\}$, let $Z_0 = \{u \in U : z_u = 0\}$, write $q_C = \mathbf{1}_{C \cap U}$, and define

$$K = \operatorname{cone}\big(\{q_C : C \in \mathcal{C}_0\} \cup \{e_u : u \in Z_0\}\big).$$

After the pair row pins $r = h = M$, the PXC92b tangent criterion and Farkas' lemma give the following necessary-and-sufficient list, modulo relabelling:

residual strata of z	fidelity condition
$\{a, b, c\}$	automatic,
$\{a\} > \{b, c\}$	$e_a \in \operatorname{span}\{\mathbf{1}_U\} + K$,
$\{a, b\} > \{c\}$	$\exists \theta \in (0, 1) :\ \theta e_a + (1 - \theta) e_b \in \operatorname{span}\{\mathbf{1}_U\} + K$,
$\{a\} > \{b\} > \{c\}$	$e_a \in \operatorname{span}\{\mathbf{1}_U\} + K, \quad e_b \in \operatorname{span}\{\mathbf{1}_U, e_a\} + K$.

(PXC99)

The bank contains one important warning against a simpler argument: an x-tight row can be slack at z. In that case a smaller common-tight subcollection is already complete, but this common-subcertificate property need not be assumed in the proof below.

The finite exact-cone route closes the last case. Normalize $v(N) = 1$. For each residual order in (PXC99), enumerate every weak order of x and introduce a Boolean selector for every nonempty proper coalition row. The selected rows are required to be x-tight, and the full staged Farkas identities require them to certify $x = \text{Lmin}$ in their retained polytope. The mirror identities over all original z-tight rows require $z = \text{Lmax}$. Finally impose one normalized forbidden direction from the applicable line of (PXC99), including every active zero-coordinate normal. Exactness is encoded by the complete 280-inequality rational description of the five-player exact cone. Thus a feasible cell is exactly a failure of the canonical pair lift for some fixed complete $\mathcal{C}_x$, and is a PXC97b counterexample candidate whose other extreme covers would still require exhaustion; infeasibility of all cells is universal over the fixed certificate.

For a transparent exhaustive indexing, relabel players so that the weak order of x is increasing. Its tie blocks are then one of the 16 compositions of five. Independently choose the two top players of z in $\binom{5}{2} = 10$ ways, order the other three in $3! = 6$ ways, and choose their tie blocks from the

four compositions of three. This gives

$$16\binom{5}{2}3!\,4 = 3840 \tag{PXC99a}$$

cells. (Tied outsider orders create harmless duplicate cells.) Exact QF–LRA returns unsatisfiable in all 3840, with no unknown or error.

Two algebraically separate encodings give the same census. The primary driver uses the raw PXC92b polar/Farkas identity at every stage, with free coefficients on earlier pinned coordinates. The independent driver instead uses normalized nonnegative incidence identities: the designated current coefficient is scaled to one, other current coefficients are nonnegative, and coefficients on earlier pinned strata remain free. Both include active nonnegativity normals at every zero coordinate, quantify an arbitrary Boolean subset of all 30 proper rows for the fixed $\mathcal{C}_x$, and test every forbidden later-stage tangent. In a representative strict residual cell, separately deleting either singleton normalization or the exact-game cone changes UNSAT to SAT with an exact rational model. Thus the controls fail open rather than reporting UNSAT by construction.

The two census artifacts pin both drivers, the exact-cone catalogue, cell keys, solver version, and verdict counts. Consequently, for every five-player unrestricted instance (PXC97b) and every fixed complete staged Lmin certificate $\mathcal{C}_x$, the canonical tight-pair cover $\mathcal{R} = \{rh\}$ preserves both selectors. Hence unrestricted PXC91c holds through five players. Together with the fixed- and existential-$\mathcal{C}_x$ failure in PXC98, this gives the sharp threshold

$$\boxed{n \le 5:\ \text{(PXC97b) holds;} \qquad n = 6:\ \text{both formulations can fail.}} \tag{PXC100}$$

In particular, the global minimum failure dimension is exactly 6. This is a computer-assisted theorem: the reduction to the pair cover and the polar criterion are conventional, while the last finite implication uses the exact five-player cone and the 3840 rational cells. Reproduction is

```
python3 \
  verifiers/divergence/search_archive/unrestricted_pxc91c_n5_all_cells.py \
  --workers 4 \
  --cell-timeout-ms 120000 \
  --output \
  verifiers/divergence/search_archive/unrestricted_pxc91c_n5_all_cells.json
python3 \
  verifiers/divergence/search_archive/unrestricted_pxc91c_n5_all_cells.py \
  --workers 4 \
  --cell-timeout-ms 120000 \
  --cell-driver \
  verifiers/divergence/search_archive/verify_unrestricted_pxc91c_n5_cell_ni.py \
  --encoding-label independent-NI \
  --output \
  verifiers/divergence/search_archive/unrestricted_pxc91c_n5_all_cells_ni.json
```

The six-player witness is not isolated and is not a null-player padding artifact. It belongs to a rational chamber that persists in every larger dimension. Let $k \ge 2$, put

$$N = \{r, h, a, b\} \,\dot\cup\, O, \qquad O = \{j, o_1, \dots, o_{k-1}\},$$

and choose rational parameters

$$M > q > J > K > 0, \qquad q + (k-1)K \ge M, \qquad (k+2)q + (k-1)K > (2k+1)J. \tag{PXC101a}$$

This rational chamber has nonempty interior for every $k \geq 2$. Indeed, the rational point

$$(M, q, J, K) = (15/4, 3, 2, 1)$$

satisfies every displayed inequality strictly. Set

$$\begin{aligned} T &= 2M + 2q + J + (k-1)K, \\ L &= 2M + q, \qquad B = M + J, \qquad C = 2M + J + (k-1)K, \end{aligned}$$

and let $P_k(M, q, J, K)$ be the nonnegative total-T polytope

$$y(rha) \geq L, \qquad y(rhb) \geq L, \qquad y(rj) \geq B, \qquad y(rhO) \geq C. \tag{PXC101b}$$

As before define $v(S) = \min_{y \in P_k} y(S)$. This produces a positive-dimensional rational family of exact singleton-normalized games. Indeed P_k is nonempty because it contains the allocation displayed below. Compactness and the lower-envelope construction give exactness. For singleton normalization, the all-root allocation Te_r zeros every player other than r. A zero witness for r assigns

$$y_h = L, \qquad y_j = B, \qquad \sum_{\ell=1}^{k-1} y_{o\ell} = T - L - B = q + (k-1)K - M \geq 0,$$

and zero to r, a, b. It satisfies the fourth row because its full mass outside a, b is $T > C$.

Put

$$u = \frac{q + J + (k-1)K}{k+1}, \qquad p = \frac{T - (k+2)u}{2}.$$

Then exact staged optimization gives

$$\begin{aligned} \mathrm{Lmin}(v) &= (p, p, u, u, u^k), \\ \mathrm{Lmax}(v) &= (M, M, q, q, J, K^{k-1}). \end{aligned} \tag{PXC101c}$$

Here and below exponents denote repetition. For the lower selector, if $R = y_r + y_h$ and $W = y(O)$, the first two rows give

$$2L \leq 2R + y_a + y_b = T + R - W.$$

If every coordinate is at least m, then $R \leq T - (k+2)m$ and $W \geq km$, whence

$$m \leq \frac{T - L}{k+1} = u.$$

Equality pins a, b and all k outsiders at u, and the next stage equalizes r, h at p. The inequalities in (PXC101a) give $p > u$. They also give

$$\begin{aligned} 2(k+1)(p-u) &= 2M(k+1) + (k-2)q - 3J - 3(k-1)K > 0, \\ 2(k+1)(p+u-B) &= (k+2)q + (k-1)K - (2k+1)J > 0, \\ 2p + ku - C &= 2(q-u) > 0, \end{aligned}$$

so the blocker and the all-outsider row are slack at $\mathrm{Lmin}(v)$. The only proper x-tight rows are therefore rha, rhb. A one-row retention pins the complementary $k+1$ players at u and then equalizes the remaining triple at $L/3 \neq (p, p, u)$. Consequently

$$\mathcal{C}_x = \{rha, rhb\}$$

is the unique complete proper-row certificate.

For the upper selector, the first, second and fourth rows imply

$$2L \le T + R - W \le T + 2R - C,$$

and hence $R \ge 2M$. At ceiling M, this pins $r = h = M$; the first two rows then pin $a = b = q$, while efficiency and the fourth row pin $y(O) = J + (k-1)K$. The blocker gives $y_j \ge J$. Since $q > J > K$, the remaining stages pin $j = J$ and $o_1 = \cdots = o_{k-1} = K$, proving (PXC101c).

The normal-cone calculation is dimension-free. At x only the first two displayed rows are active. At z, solving $\mathbf{1}_S = \beta\mathbf{1} + \sum_{t=1}^{4} \lambda_t \mathbf{1}_{C_t}$, $\lambda_t \ge 0$, in the six coordinate types r, h, a, b, j, o means that the six coordinates of $\mathbf{1}_S$ are

$$\begin{aligned}&\beta + \lambda_1 + \lambda_2 + \lambda_3 + \lambda_4, \quad \beta + \lambda_1 + \lambda_2 + \lambda_4, \quad \beta + \lambda_1, \quad \beta + \lambda_2,\\ &\beta + \lambda_3 + \lambda_4, \quad \beta + \lambda_4.\end{aligned}$$

Requiring each to be zero or one gives exactly

$$\{rh, rha, rhb, rj, rhO\} \tag{PXC101d}$$

as the proper z-tight menu. If their normalized cover weights are (A_0, A_1, A_2, D, E), respectively, outsider balance and normalization give

$$A_1 = A_2 = E = \eta, \qquad D = 0, \qquad 2A_0 + 4\eta = 1. \tag{PXC101e}$$

Thus the cover polytope is the same segment in every dimension. Its two extreme supports are

$$\{rh\}: \ A_0 = \tfrac{1}{2}, \ \eta = 0, \qquad \{rha, rhb, rhO\}: \ A_1 = A_2 = E = \eta = \tfrac{1}{4}.$$

The blocker rj has coefficient zero in every normalized top cover. For either extreme cover, the retained upper selector is

$$z' = (M, M, q, q, \bar{K}^k), \qquad \bar{K} = \frac{J + (k-1)K}{k}, \tag{PXC101f}$$

which differs from z because $J > K$. The common forbidden direction is

$$d = z' - z = (0, 0, 0, 0, -(k-1)(J-K)/k, ((J-K)/k)^{k-1}).$$

It vanishes on rh, rha, rhb, rhO, but $d(rj) = -(k-1)(J-K)/k < 0$. Since every complete $\mathcal{C}_x$ equals the displayed pair and both extreme covers fail, both certificate quantifiers in (PXC97b) fail throughout this chamber.

The integer specialization

$$(M, q, J, K) = (4, 3, 2, 1), \qquad T = k + 15 = n + 11 \tag{PXC102a}$$

satisfies (PXC101a) for every $k \ge 2$. It gives the explicit all-dimension selector pair

$$\begin{aligned}x &= \left(\left(\frac{10k+7}{2(k+1)}\right)^2, \left(\frac{k+4}{k+1}\right)^{k+2}\right),\\ z &= (4, 4, 3, 3, 2, 1^{k-1}), \qquad z' = (4, 4, 3, 3, ((k+1)/k)^k).\end{aligned} \tag{PXC102b}$$

Consequently unrestricted sparse fidelity fails for every $n \geq 6$, by a genuine n-player exact game rather than by adjoining null players. Indeed

$$v(rha) - v(rh) = v(rhb) - v(rh) = q, \qquad v(rj) - v(r) = M + J,$$

and for every $o \in O$,

$$v(rhO) - v\big(rh(O \setminus \{o\})\big) = J + (k-1)K > 0.$$

Either top player has grand-complement marginal T. Thus every displayed player has a strictly positive marginal in an explicit coalition.

There is also an exact general formulation of the mechanism. Fix an instance, a complete $\mathcal{C}_x$, and let $\mathfrak{R}$ be the finite set of qualifying extreme top covers. Write

$$\mathcal{U}_z = \bigcup_{\mathcal{R} \in \mathfrak{R}} \mathcal{R}$$

for the cover-eligible z-tight rows. At a later Lmax stage, suppose d has zero total, vanishes on the top and earlier strata, is nonpositive on the current stratum with one strict entry, and respects the common-tight $\mathcal{C}_x$ rows and active nonnegativity normals. Then

$$\boxed{d \text{ is tangent to every } (\mathcal{C}_x \cup \mathcal{R}) \text{ lift} \iff d(S) = 0 \quad \text{for every } S \in \mathcal{U}_z.} \tag{PXC103}$$

Indeed tangency gives $d(S) \geq 0$ on a cover support, whereas dotting its normalized cover identity with d gives $\sum_{S \in \mathcal{R}} \lambda_S d(S) = 0$. Positive support weights force each term to vanish. The converse is immediate. Because the full game has upper selector z, any such common forbidden direction must be cut by some original z-tight row $Q \notin \mathcal{U}_z$ with $d(Q) < 0$; that row has coefficient zero in every top cover. This is a necessary-and-sufficient characterization of the *common-direction* forced-zero-blocker mechanism. It must not be enlarged without proof to all universal-cover failures: different covers may in principle leave different forbidden directions.

That warning is sharp. There is a second, genuinely cover-dependent six-player mechanism. Let $N = \{r, h, a, b, j, k\}$, impose nonnegativity and efficiency $y(N) = 23$, and define v as the lower envelope of

$$\begin{aligned} y(rha) \geq 16, \quad y(rhb) \geq 16, \quad y(rhj) \geq 13, \\ y(rhaj) \geq 18, \quad y(rhk) \geq 13, \quad y(rhak) \geq 18. \end{aligned} \tag{PXC104a}$$

Then the game is exact and singleton-normalized and

$$x = \big((41/6)^2, (7/3)^4\big), \qquad z = (6, 6, 5, 4, 1, 1). \tag{PXC104b}$$

Indeed, writing $R = y_r + y_h$, the first two rows give $2R + y_a + y_b \geq 32$, hence $R - (y_j + y_k) \geq 9$. If every coordinate is at least m, then $9 \leq 23 - 6m$, so $m \leq 7/3$. Equality pins all four outsiders at $7/3$, after which the next Lmin stage equalizes r, h at $41/6$. The other four displayed rows are strict there. Thus the only proper x-tight rows, and the unique complete proper-row certificate, are

$$\mathcal{C}_x = \{rha, rhb\}. \tag{PXC104c}$$

For the upper selector, either of the two identities below gives $R \geq 12$, so a ceiling-six point has $r = h = 6$. The five remaining tight rows give $b \geq 4$, $j, k \geq 1$, $a + j \geq 6$, and $a + k \geq 6$. Their minimum outsider mass is already $5 + 4 + 1 + 1 = 11$, the efficient residual, which pins z. Exactness follows from compact lower-envelope attainment. For singleton normalization, use $(0, 12, 5, 4, 1, 1)$ and its r, h swap for the two top-player witnesses; the all-root allocation supplies every outsider witness.

The proper z-tight menu is exactly

$$\{rh, rhb, rhj, rhaj, rhk, rhak\}.$$

Put weights (A, B, J, P, K, Q) on these rows. Outsider balance and normalization reduce exactly to

$$B = \eta, \quad J = \eta - P, \quad K = P, \quad Q = \eta - P, \quad A + 2\eta = \tfrac{1}{2}, \quad 0 \le P \le \eta \le \tfrac{1}{4}. \tag{PXC104d}$$

Hence the normalized-cover polytope is a triangle with precisely three extreme covers:

$$\{rh\} : \tfrac{1}{2}; \qquad \{rhb, rhj, rhak\} : \left(\tfrac{1}{4}\right)^3; \qquad \{rhb, rhaj, rhk\} : \left(\tfrac{1}{4}\right)^3. \tag{PXC104e}$$

With the unique $\mathcal{C}_x$, their reduced Lmax selectors are respectively

$$(6, 6, 4, 4, 3/2, 3/2), \qquad (6, 6, 4, 4, 1, 2), \qquad (6, 6, 4, 4, 2, 1), \tag{PXC104f}$$

so every qualifying extreme cover fails while Lmin and singleton normalization are preserved.

This failure has no PXC103 common direction. After the top block is fixed, the union of the three cover supports imposes $d_b = d_j = d_k = 0$ and $d_a + d_j = d_a + d_k = 0$, hence $d = 0$. Instead the second cover leaves $d^- = (0, 0, -1, 0, 0, 1)$, cut by $rhaj$; that blocker belongs positively to the third cover. The third leaves $d^+ = (0, 0, -1, 0, 1, 0)$, cut by $rhak$; that blocker belongs positively to the second cover. The pair-cover descent is cut by both. Thus there is no load-bearing row whose coefficient is zero in every cover: the obstruction is a two-cycle of *cross-cover blockers*. In particular, the proposed extension of (PXC103) from common-direction failures to all universal-cover failures is false.

The phenomenon is not isolated. For every rational parameter $4 < A < 11/2$, let v_A be the lower envelope on $N = \{r, h, a, b, j, k\}$ with total $18 + A$ of

$$\begin{aligned} y(rha) &\ge 16, & y(rhb) &\ge 16, \\ y(rhj) &\ge 13, & y(rhk) &\ge 13, \\ y(rhaj) &\ge 13 + A, & y(rhak) &\ge 13 + A. \end{aligned} \tag{PXC105a}$$

Then v_A is exact and singleton-normalized, and its selectors are

$$x_A = \left(\left(\frac{46 - A}{6}\right)^2, \left(\frac{A + 2}{3}\right)^4\right), \qquad z_A = (6, 6, A, 4, 1, 1). \tag{PXC105b}$$

The only proper x_A-tight rows are rha, rhb, so $\mathcal{C}_{x_A} = \{rha, rhb\}$ is the unique complete proper-row certificate. The proper z_A-tight menu and the three extreme covers are exactly those in (PXC104d)–(PXC104e). With this unique certificate, the three reduced upper selectors are

$$\left(6, 6, 4, 4, \frac{A - 2}{2}, \frac{A - 2}{2}\right), \qquad (6, 6, 4, 4, 1, A - 3), \qquad (6, 6, 4, 4, A - 3, 1). \tag{PXC105c}$$

Thus every rational A in the open interval is a universal-cover fidelity failure of cross-cover type, and (PXC104) is the member $A = 5$.

Here is a direct proof. A coalition omitting at least one of r, h has value zero: put the entire total on an omitted top player. For $U \subseteq \{a, b, j, k\}$, direct elimination of the four outsider bits gives

$$\begin{aligned} v_A(rh \cup U) = \max\big(&\{12\} \cup \{16 : a \in U \text{ or } b \in U\} \\ &\cup \{13 : j \in U \text{ or } k \in U\} \cup \{13 + A : a \in U,\ U \cap \{j, k\} \neq \varnothing\}\big). \end{aligned} \tag{PXC105d}$$

with the efficiency value at N. The lower bounds are the displayed rows and either three-row cover identity; the reverse inequalities follow from the following explicit endpoint minimizers. Coordinates are ordered (r, h, a, b, j, k); the two point columns correspond to $A = 4$ and $A = 11/2$, respectively.

U	w_U^4	$w_U^{11/2}$
$\varnothing$	$(12, 0, 4, 4, 1, 1)$	$(12, 0, 11/2, 4, 1, 1)$
a	$(16, 0, 0, 4, 1, 1)$	$(16, 0, 0, 5/2, 5/2, 5/2)$
b	$(16, 0, 6, 0, 0, 0)$	$(16, 0, 15/2, 0, 0, 0)$
ab	$(16, 0, 0, 0, 5, 1)$	$(16, 0, 0, 0, 5, 5/2)$
j	$(13, 0, 6, 3, 0, 0)$	$(13, 0, 15/2, 3, 0, 0)$
aj	$(17, 0, 0, 5, 0, 0)$	$(37/2, 0, 0, 5, 0, 0)$
bj	$(16, 0, 6, 0, 0, 0)$	$(16, 0, 15/2, 0, 0, 0)$
abj	$(17, 0, 0, 0, 0, 5)$	$(37/2, 0, 0, 0, 0, 5)$
k	$(13, 0, 6, 3, 0, 0)$	$(13, 0, 15/2, 3, 0, 0)$
ak	$(17, 0, 0, 5, 0, 0)$	$(37/2, 0, 0, 5, 0, 0)$
bk	$(16, 0, 6, 0, 0, 0)$	$(16, 0, 15/2, 0, 0, 0)$
abk	$(17, 0, 0, 0, 5, 0)$	$(37/2, 0, 0, 0, 5, 0)$
jk	$(13, 0, 6, 3, 0, 0)$	$(13, 0, 15/2, 3, 0, 0)$
ajk	$(17, 0, 0, 5, 0, 0)$	$(37/2, 0, 0, 5, 0, 0)$
bjk	$(16, 0, 6, 0, 0, 0)$	$(16, 0, 15/2, 0, 0, 0)$
$abjk$	$(22, 0, 0, 0, 0, 0)$	$(47/2, 0, 0, 0, 0, 0)$

Each w_U^e is feasible at its endpoint and attains the value in (PXC105d). Both the constraint right sides and the applicable branch of the maximum in (PXC105d) are affine on the closed interval, so the affine interpolation of $w_U^4, w_U^{11/2}$ is a feasible attaining allocation for every intermediate A. This proves the reverse inequality and hence exactness without an implicit attainment claim. The allocations $(0, 12, A, 4, 1, 1)$, its top swap, and the all-root allocation prove singleton normalization. No player is null: adjoining a or b to rh raises the value by 4, adjoining j or k raises it by 1, and adjoining either top player to its complement raises the value from zero to $18 + A$.

Writing $R = y_r + y_h$, the first two rows of (PXC105a) give $R - (y_j + y_k) \geq 14 - A$. If all coordinates are at least m, then $R - (y_j + y_k) \leq 18 + A - 6m$, whence $m \leq (A+2)/3$. Equality pins all four outsiders at this value and the next stage equalizes the tops at $(46 - A)/6$. The two long cross rows have slack $(11 - 2A)/3 > 0$; formula (PXC105d) therefore leaves only rha, rhb tight. With only one of them retained, the last three coordinates outside that row are pinned first at $(A + 2)/3$, while the three coordinates in the row equalize at $16/3$; hence neither proper subset is complete.

Adding the three inequalities in either positive-ambient cover gives $2R + (18 + A) \geq 42 + A$, so $R \geq 12$. A ceiling-six point therefore has $r = h = 6$. The remaining rows give $b \geq 4$, $j, k \geq 1$, $a+j \geq A+1$, and $a+k \geq A+1$. Their lower bounds sum to the residual $A+6$, pinning z_A. The cover triangle is unchanged because its incidence equations do not depend on A, and the three reduced selectors in (PXC105c) follow by the same equalization. Since $A > 4$, none equals z_A. The two cover-specific forbidden directions are positive multiples of $(0, 0, -1, 0, 0, 1)$ and $(0, 0, -1, 0, 1, 0)$; the opposite cross row blocks each one. Their cover-union kernel is zero for every A, completing the positive-dimensional cross-cover theorem. At $A = 4$ the a, b strata merge, while at $A = 11/2$ the long cross rows become x_A-tight; the open endpoints are therefore load-bearing, not cosmetic.

The cross-cover construction also lifts genuinely in dimension. Let J, K be disjoint nonempty blocks of sizes p, q, put $m = p + q$, and take a rational A satisfying

$$4 < A < 6, \qquad p(A-1) > (m+1)(A-4), \qquad q(A-1) > (m+1)(A-4). \tag{PXC106a}$$

On $N = \{r, h, a, b\} \dot\cup J \dot\cup K$, with total $16 + A + m$, define the lower-envelope game by

$$\begin{aligned} y(rha), y(rhb) &\geq 16, \\ y(rhJ) \geq 12 + p, &\quad y(rhaJ) \geq 12 + A + p, \\ y(rhK) \geq 12 + q, &\quad y(rhaK) \geq 12 + A + q. \end{aligned} \tag{PXC106b}$$

Writing

$$u = \frac{A+m}{m+1}, \qquad t = \frac{16-u}{2},$$

the selectors are

$$x = (t, t, u^{m+2}), \qquad z = (6, 6, A, 4, 1^m). \tag{PXC106c}$$

The unique complete proper-row certificate remains $\{rha, rhb\}$. The proper z-tight menu is

$$\{rh, rhb, rhJ, rhaJ, rhK, rhaK\},$$

and its three extreme covers are

$$\{rh\}, \qquad \{rhb, rhJ, rhaK\}, \qquad \{rhb, rhaJ, rhK\}, \tag{PXC106d}$$

with the same weights as in (PXC104e). Their reduced upper selectors are

$$\begin{aligned} &(6, 6, 4, 4, c^m), && c = 1 + \frac{A-4}{m}, \\ &(6, 6, 4, 4, 1^p, c_K^q), && c_K = 1 + \frac{A-4}{q}, \\ &(6, 6, 4, 4, c_J^p, 1^q), && c_J = 1 + \frac{A-4}{p}. \end{aligned} \tag{PXC106e}$$

Consequently every member fails for every qualifying extreme cover while preserving exactly two Lmax-top players, the unique complete $\mathcal{C}_x$, exactness, and singleton normalization.

For completeness, the proof is dimension-free. As in (PXC105d), every coalition omitting a top player has value zero. For a coalition $rh \cup U$, its value is the maximum of 12 and the bounds among

$$\begin{gathered} 16 \quad (a \in U \text{ or } b \in U), \qquad 12 + p \quad (J \subseteq U), \qquad 12 + A + p \quad (a \in U,\ J \subseteq U), \\ 12 + q \quad (K \subseteq U), \qquad 12 + A + q \quad (a \in U,\ K \subseteq U), \end{gathered}$$

with the efficiency value at N. Each lower bound is displayed or follows from a positive cover. The reverse inequality has the following uniform attainment construction. For a proper trace U, let L be the displayed maximum, put $y_r = L, y_h = 0$, and initially put zero on every selected non-top coordinate. If $a \notin U$, put $y_a = A$; if $b \notin U$, put $y_b = (16 - L)_+$. For a block $B \in \{J, K\}$ of size s that is not fully contained in U, choose $o_B \in B \setminus U$ and put

$$y_{o_B} = \max\{0,\ 12 + s - L,\ 12 + A + s - L - y_a\},$$

leaving the other coordinates of B at zero. A fully selected block receives zero. These assignments satisfy all six rows. Their used mass W is at most $16 + A + m$: if $a \notin U$, substituting $y_a = A$ gives the explicit expression

$$W = L + A + \mathbf{1}_{b \notin U}(16 - L)_+ + \sum_{\substack{B \in \{J,K\} \\ B \not\subseteq U}} (12 + |B| - L)_+,$$

whose eight choices of whether b, J, K are fully selected are each at most the total; if $a \in U$, then $L \geq 16$, and the cases of zero, one, or two fully selected blocks give respectively

$$W = 8 + 2A + m, \qquad W \leq 12 + A + \max\{p, q\}, \qquad W \leq 12 + A + \max\{p, q\},$$

all at most $16+A+m$ (the first uses $A<6$). Put the unused mass $16+A+m-W$ on any omitted non-top coordinate. It does not enter the coalition objective and can only increase the row left sides. Thus the constructed point attains L; the grand coalition is attained by the all-root point. This proves the asserted lower envelope and tight menus without an implicit projection claim. The allocations $(0,12,A,4,1^m)$, its top swap, and the all-root allocation prove singleton normalization.

The two short rows give

$$R-y(J)-y(K)\geq 16-A-m.$$

If all $m+4$ coordinates are at least s, the reverse side is at most $16+A+m-(2+2m)s$, so $s\leq u$. Equality pins every non-top coordinate at u, then equalizes the tops at t. The two long rows have respective slack

$$\frac{p(A-1)}{m+1}-(A-4),\qquad \frac{q(A-1)}{m+1}-(A-4),$$

which is positive precisely by (PXC106a); the shorter block rows are also strict. A one-row certificate instead equalizes its three members at $16/3$ after pinning its $m+1$-player complement at u, proving uniqueness of the complete pair.

Either positive cover in (PXC106d) partitions all non-top players among three rows. Adding its bounds gives $2R+(16+A+m)\geq 40+A+m$, hence $R\geq 12$. At ceiling six the top players are fixed, and the remaining rows force

$$y_b=4,\qquad y(J)=p,\qquad y(K)=q,\qquad y_a=A;$$

block equalization gives z. The same calculation after row deletion gives (PXC106e). The exact cover equations again form the triangle (PXC104d), now because each block member has the same incidence column. The cover union forces $d_a=d_b=d(J)=d(K)=0$. Internal block exchanges remain algebraically, but cannot be a staged Lmax descent on a tied block: all its entries would be nonpositive and their block sum zero. Thus there is no common forbidden direction, while $rhaJ$ and $rhaK$ cross-block the two cover-specific descents exactly as before.

Finally, for every $n\geq 6$, take

$$p=1,\qquad q=n-5,\qquad A=4+\frac{1}{n-3}. \tag{PXC106f}$$

Both inequalities in (PXC106a) are then strict. Every added player has payoff one and is non-null: for $i\in J$, $v(rhJ)-v(rh(J\setminus\{i\}))=p>0$, and the analogous K-marginal is $q>0$. The a,b marginals over rh equal 4, while either top player's grand-complement marginal is $16+A+m$. Thus this is not null-player padding. This proves a genuine cross-cover counterexample for every dimension $n\geq 6$.

There is an exact general obstruction alternative behind these two families. Fix an instance and complete $\mathcal{C}_x$, let $\mathfrak{R}$ be its finite set of qualifying extreme top covers, and let $\mathcal{U}_z=\bigcup_{\mathcal{R}\in\mathfrak{R}}\mathcal{R}$. For each cover $\mathcal{R}$, write $\mathcal{D}_{\mathcal{R}}$ for the union of all nonzero homogeneous later-stage directions satisfying the corresponding PXC92b tangent and descent conditions for the $\mathcal{C}_x\cup\mathcal{R}$ lift. Then

$$\text{every extreme-cover lift fails}\quad\Longleftrightarrow\quad \mathcal{D}_{\mathcal{R}}\neq\varnothing\quad\text{for every }\mathcal{R}\in\mathfrak{R}. \tag{PXC107a}$$

Whenever (PXC107a) holds, there is an exact dichotomy.

First, if $\bigcap_{\mathcal{R}\in\mathfrak{R}}\mathcal{D}_{\mathcal{R}}\neq\varnothing$, PXC103 applies: a common direction vanishes on every row in $\mathcal{U}_z$, and full-game Lmax optimality supplies a blocker outside that union. Second, if there is no common direction, every $d\in\mathcal{D}_{\mathcal{R}}$ satisfies

$$d(Q)<0\qquad\text{for some } Q\in\mathcal{U}_z. \tag{PXC107b}$$

Indeed, if instead $d(Q) \geq 0$ on the entire union, dotting d with the normalized identity of every cover $\mathcal{R}'$ gives $\sum_{Q\in\mathcal{R}'} \lambda_Q d(Q) = 0$. Strictly positive weights force $d(Q) = 0$ on every cover support, so d belongs to every $\mathcal{D}_{\mathcal{R}'}$, contradicting the absence of a common direction.

Consequently form a directed graph on $\mathfrak{R}$: choose $d_{\mathcal{R}} \in \mathcal{D}_{\mathcal{R}}$, choose a row $Q \in \mathcal{U}_z$ with $d_{\mathcal{R}}(Q) < 0$, and draw $\mathcal{R} \to \mathcal{R}'$ for a cover $\mathcal{R}' \ni Q$. Every vertex has an outgoing edge and no self-loop, so the graph contains a directed cycle of length at least two:

$$\mathcal{R}_1 \longrightarrow \mathcal{R}_2 \longrightarrow \cdots \longrightarrow \mathcal{R}_s \longrightarrow \mathcal{R}_1, \qquad s \geq 2. \tag{PXC107c}$$

The proof uses no enumeration. Positive weights and the normalized cover identity force $d(S) = 0$ on every row $S \in \mathcal{R}$. Since the full game has upper selector z, a forbidden lift direction must violate some original active coalition row. In the common case every such blocker is outside the cover union; otherwise (PXC107b) supplies the displayed outgoing edge and finiteness gives the cycle. Thus PXC98 realizes the common forced-zero branch and PXC104–PXC106 realize the cross-cover-cycle branch. There is no third cover-local forced-zero case: the cover identities collapse it to the common branch.

The preceding graph can be made invariant; no choices of directions or blockers are needed in its definition. Define the *cover–blocker incidence digraph* $\Gamma(v, x, z, \mathcal{C}_x)$ to have vertex set $\mathfrak{R}$ and the edge

$$\mathcal{R} \longrightarrow \mathcal{R}' \quad\Longleftrightarrow\quad \exists\, d \in \mathcal{D}_{\mathcal{R}}\ \exists\, Q \in \mathcal{R}' \quad d(Q) < 0. \tag{PXC108a}$$

This relation is determined by the retained-row tangent cones, not by a selected witness. It has no loop because positive cover weights force $d(S) = 0$ for every $S \in \mathcal{R}$. Under universal-cover failure and absence of a common direction, the collapse argument in (PXC107b) says that every vertex has positive outdegree. Hence every shortest directed cycle in Γ has length at least two and is chordless: any additional directed edge between nonconsecutive cycle vertices, or a reverse edge between consecutive vertices on a cycle of length greater than two, closes a shorter directed cycle. Thus the invariant obstruction core is

$$\begin{cases} (d, Q), & d \in \bigcap_{\mathcal{R}} \mathcal{D}_{\mathcal{R}},\ Q \notin \mathcal{U}_z,\ d(Q) < 0; \\ (\mathcal{R}_1, Q_1, d_1; \ldots; \mathcal{R}_s, Q_s, d_s), & \mathcal{R}_1 \to \cdots \to \mathcal{R}_s \to \mathcal{R}_1 \text{ a shortest cycle}, \end{cases} \tag{PXC108b}$$

where $Q_i \in \mathcal{R}_{i+1}$ and $d_i(Q_i) < 0$. The digraph and the set of its shortest cycles are canonical, although a graph with symmetry need not have a unique shortest cycle. Deleting any edge from the second certificate destroys that displayed cyclic certificate; this is the precise deletion-minimal meaning used here. It is not a claim that arbitrary removal of an extreme cover from the original cover polytope preserves the same quantified problem.

There are useful rank and support bounds. Let $T = \mathrm{Top}(z) = \{r, h\}$, $U = N \setminus T$, $u = n - 2$, and let $A_{\mathcal{R}}$ be the U-by-$|\mathcal{R}|$ incidence matrix of a positive-ambient extreme cover. Its outsider equations are

$$A_{\mathcal{R}}\lambda = \eta \mathbf{1}_U.$$

Since $\eta > 0$, $\mathbf{1}_U$ lies in the column span of $A_{\mathcal{R}}$. If $A_{\mathcal{R}}$ had column rank below $|\mathcal{R}|$, adjoining the one normalization equation could not isolate an extreme point with all displayed weights positive. Therefore

$$\operatorname{rank} A_{\mathcal{R}} = |\mathcal{R}| \leq u. \tag{PXC108c}$$

If equality holds, fixing the two top coordinates and using equality in the positive cover identity gives $A_{\mathcal{R}}^{\mathsf{T}} d_U = 0$, hence $d_U = 0$. Such a cover is faithful at every later Lmax stratum. Consequently a failing positive-ambient cover has support at most $n - 3$. When $\eta = 0$, nonnegative incidence

and positive row weights force every support row to avoid U; singleton normalization excludes a positive-payoff top singleton, leaving only the tight top pair rh, with support one. If M_z proper z-tight rows are cover-eligible, the crude finite bounds are therefore

$$2 \le s \le |\mathfrak{R}| \le 1 + \sum_{j=1}^{n-3} \binom{M_z}{j}$$

on a universal-failure cycle, with the initial 1 needed only for the possible zero-ambient pair cover.

The word circuit can now be used precisely, but only at the correct level (Björner et al., 1999). For a cover $\mathcal{R}$, append to each row column its number of top players, append the ambient column $(-\mathbf{1}_U, -2)$ when $\eta > 0$, and put $b = (0_U, 1)$. The active normalized-cover equations are

$$B_{\mathcal{R}}\theta = b, \qquad \theta > 0.$$

At an extreme cover the active columns are independent. Because every coefficient in the displayed representation of b is nonzero, the columns of $B_{\mathcal{R}}$ together with $-b$ form a minimally dependent signed set: an oriented-matroid circuit containing the distinguished normalization target. By contrast, an edge direction d supplies a covector sign pattern on the contracted coalition-normal arrangement. Accordingly (PXC108b) is a circuit–covector incidence cycle; in general the whole cycle is not one larger oriented-matroid circuit. Nor does the argument alone make it an indecomposable semi-balanced system or a facet of the exact-game cone: those are game-value relations, whereas Γ also depends on $\mathcal{C}_x$, the payoff strata, and staged selector signs. Any such stronger identification requires a separate projection/facet theorem.

This distinction also separates the construction from the three closest coalitional precedents. Reijnierse–Derks exact covers characterize coalition collections through their zero core; they do not carry ordered payoff strata or a pair of selected core points (Reijnierse and Derks, 1998). Minimal exact balanced collections allow the single exceptional sign needed to test exactness and have their own three-type generation theory (Lohmann et al., 2012); the augmented positive top-cover circuit above is not identified with one of those types. Finally, indecomposable semi-balanced systems correspond to facets of the full exact-game cone (Studený and Kratochvíl, 2022), whereas the edge relation in (PXC108a) changes when $\mathcal{C}_x$ or the Lmax stratum order changes while the underlying game-value cone does not. The standard ingredients are therefore prior art. The candidate new statement is only their selector-specific coupling in the invariant cover–blocker graph and its forced common/cyclic alternative; that novelty assessment remains subject to specialist review.

This gives an exact recognition algorithm. Enumerate the qualifying extreme covers; for every cover, later payoff stratum, and designated player solve the rational PXC92b tangent feasibility problem. A cover with no feasible problem is faithful. If all fail, impose every cover-union row as an equality and solve the same finite list to detect the common branch. Otherwise, for each ordered cover pair and target-support row, add the normalized inequality $d(Q) \le -1$; feasibility creates exactly the edge in (PXC108a). Breadth-first search returns a shortest chordless cycle and the LP solutions give checkable rational certificates. With $c = |\mathfrak{R}|$, this uses $O(c^2 M_z n)$ rational LPs after cover enumeration and it tests every later stratum, not only the first residual ceiling.

On the displayed calibrations, the primary verifier exhausts the complete-$\mathcal{C}_x$ and extreme-cover quantifiers. A solver-free rational active-set implementation independently reconstructs the games, selectors, cover systems, reduced lifts, and common kernels:

```
cd verifiers/divergence
python3 verify_pxc104_cover_dependent_counterexample.py \
  --output pxc104_cover_dependent_counterexample.json
```

```
python3 verify_pxc104_cover_dependent_\
counterexample_independent.py \
  --output pxc104_cover_dependent_\
counterexample_independent.json
```

The affine family has a separate primary audit of all sixty-four endpoint coalition programs and five interior rational anchors, and an independently formulated active-set replay:

```
python3 verify_pxc105_cross_cover_family.py \
  --output pxc105_cross_cover_family.json
python3 verify_pxc105_cross_cover_family_\
independent.py \
  --output pxc105_cross_cover_family_\
independent.json
```

The block lift is checked by exact full-game replays at dimensions six, seven, and eight and by a standalone dimension-free incidence audit:

```
python3 verify_pxc106_cross_cover_dimension_lift.py \
  --output pxc106_cross_cover_dimension_lift.json
python3 verify_pxc106_cross_cover_dimension_\
lift_independent.py \
  --output pxc106_cross_cover_dimension_\
lift_independent.json
```

The obstruction alternative has a standalone exact calibration on PXC98, PXC104, and the faithful full-rank PXC90 control:

```
python3 verify_pxc107_obstruction_alternative.py \
  --output pxc107_obstruction_alternative.json
python3 verify_pxc108_invariant_cover_blocker_circuits.py \
  --output pxc108_invariant_cover_blocker_circuits.json
python3 verify_pxc108_invariant_cover_blocker_circuits_independent.py \
  --output pxc108_invariant_cover_blocker_circuits_independent.json
```

The exact-cone projection can now be stated precisely. Fix n and a finite *selector signature* Σ: ordered equality strata for x, z, their active proper core rows, a fixed complete $\mathcal{C}_x$, the active zero-coordinate normals, the qualifying extreme-cover supports, and an edge/nonedge pattern for (PXC108a). In the variables (v, x, z), let

$$\begin{aligned}\mathcal{K}_\Sigma = \{(v,x,z) : v \in \mathcal{E}_n, \quad & v(i) = 0,\\ & x(N) = z(N) = v(N), \quad x(S), z(S) \geq v(S),\\ & x(S) = v(S)\ (S \in A_x), \quad z(S) = v(S)\ (S \in A_z),\\ & \text{the prescribed inactive-row and stratum inequalities are strict}\}.\end{aligned} \tag{PXC109a}$$

Here $\mathcal{E}_n$ is the exact-game cone, equivalently the cone cut out by the semi-balanced inequalities of Studený and Kratochvíl (2022). The PXC92b tangent test depends only on the incidence normals in $A_x, A_z, \mathcal{C}_x$, the zero set, and the ordered strata. Thus an admissible signature records its rational Farkas certificates once and for all; the same is true of every cover vertex and every edge or nonedge

in (PXC108a). It follows that $\mathcal{K}_\Sigma$ is a relatively open rational polyhedral cone. Linear projection gives

$$\mathcal{F}_n = \bigcup_{\substack{\Sigma:\ \mathfrak{R}_\Sigma \neq \varnothing \\ \mathcal{D}_\mathcal{R} \neq \varnothing\ \forall \mathcal{R}}} \pi_v(\mathcal{K}_\Sigma), \tag{PXC109b}$$

where $\mathcal{F}_n$ is the fixed-n, exactly-two-top, fixed-$\mathcal{C}_x$ failure locus after retaining $\mathcal{C}_x$ as part of the signature. Its projection to games records existence of a bad fixed certificate. The alternative quantifier requiring failure for every complete $\mathcal{C}_x$ is also a finite Boolean combination of these cells and, after distributing that finite combination, a finite union of relatively open rational polyhedra. The union in (PXC109b) is finite, and every member is the projection of a relatively open rational polyhedral cone. This is an explicit selector-augmented cone description: one obtains inequalities by eliminating x, z and the finitely many rational Farkas variables. It is not a claim that $\mathcal{F}_n$ is convex, closed, or a face of $\mathcal{E}_n$. Indeed the strict payoff order and descent signs normally put a PXC108 type in the relative interior of a selector cell rather than on one of its walls.

There is nevertheless a useful dimension-free facet bridge. Let $T \subsetneq N$, and let the nonempty blocks $B_1, \dots, B_m$, $m \geq 2$, partition $N \setminus T$. Then

$$\sum_{i=1}^{m} \mathbf{1}_{T \cup B_i} - (m-1)\mathbf{1}_T = \mathbf{1}_N. \tag{PXC110a}$$

The supported columns are independent: looking on each block first forces the coefficient of its unique row, and then the coefficient of T, to vanish in any homogeneous dependence. Thus (PXC110a) is support-minimal, and its empty-set coefficient is zero. It is in fact an indecomposable min-semi-balanced system. To see this directly, suppose a new exceptional coalition E has coefficient -1, let c be the common load, and put $e_T, e_i \in \{0, 1\}$ for its membership on T, B_i. Membership must be constant on every block, and the nonnegative old coefficients are forced to be

$$u_i = c + e_i, \qquad u_T = -(m-1)c + e_T - \sum_i e_i. \tag{PXC110b}$$

If some block is outside E, then $c \geq 0$, and (PXC110b) leaves only $E = T$ or $E = T \cup B_i$, besides the empty set. If every block is in E, it leaves only $E = N$. Hence no unsupported nonempty proper exceptional coalition decomposes (PXC110a), and the Studený–Kratochvíl criterion makes its inequality

$$\sum_i v(T \cup B_i) - (m-1)v(T) \leq v(N) \tag{PXC110c}$$

a facet of the exact-game cone.

This facet appears naturally when a two-top normalized cover has equal top coefficients. If $T = \{r, h\}$, $\alpha_r = \alpha_h = \alpha > 0$, and the positive support partitions the outsiders into the traces B_i, the cover identity can be scaled to (PXC110a). For unequal top coefficients no single coalition contained in T can absorb both exceptional coordinates. Thus equal-top *pair shear* is a sufficient exact-cone bridge, not a consequence of cover extremality alone.

The displayed families now have a sharp exact-cone interpretation. For (PXC101), the positive cover and the tight pair give the partition-star facet

$$v(rha) + v(rhb) + v(rhO) - 2v(rh) \leq v(N). \tag{PXC111a}$$

The forced-zero row rj is external to that facet support: the obstruction is a facet–blocker flag, not the facet by itself. At the generic six-player point $(M, q, J, K) = (15/4, 3, 2, 1)$, the complete published six-player catalogue has 552 tight facet inequalities whose normals have rank 54. Thus the game lies on a nine-dimensional minimal exact-cone face; the selector obstruction is much finer than merely lying on one facet.

For (PXC104)–(PXC105), the two positive covers give the two facets

$$\begin{aligned} v(rhb) + v(rhj) + v(rhak) - 2v(rh) &\le v(N), \\ v(rhb) + v(rhaj) + v(rhk) - 2v(rh) &\le v(N). \end{aligned} \tag{PXC111b}$$

The row $rhaj$ belongs only to the second support and cuts the descent of the first; $rhak$ belongs only to the first and cuts the descent of the second. Hence the PXC108 two-cycle is exactly the circuit–covector incidence graph of a minimal union of two indecomposable exact-cone facets. All three rational anchors $A = 17/4, 5, 21/4$ have the same 835 tight catalogue facets, of normal rank 53, so their minimal exact face has dimension ten. The conic PXC105 family has dimension two inside that face; the failure interval is a selector-cell interior, not a new facet.

The same statement is genuinely dimension-free for (PXC106). Its two block partitions are

$$\{b\} \mid J \mid (\{a\} \cup K), \qquad \{b\} \mid (\{a\} \cup J) \mid K.$$

Applying (PXC110) to each proves that the all-n cross-cover obstruction is a two-facet dependency cycle for every nonempty J, K. This is stronger than the finite PXC104 catalogue lookup, but the partition-star facets themselves are an instance of the established semi-balanced theory; the new candidate contribution is their staged-selector coupling.

The unrestricted proposed correspondence is false. In (PXC88), the unique positive cover has rows $0134, 0125, 16$, top coefficients $(1/3, 2/3)$, and indispensable later blockers $0126, 0136$. Form the seven-by-six matrix whose columns are the five displayed coalition incidence vectors and $-\mathbf{1}_N$. Its minor on player rows $0, \dots, 5$ has determinant -1. Hence the columns are independent: no nonempty subset of the natural cover–blocker support has even a constant incidence dependence, let alone an indecomposable semi-balanced one. PXC88 is therefore an exact counterexample to the assertion that every inclusion-minimal PXC108 obstruction projects to one exact-cone facet or a union of facets supported on its obstruction rows. It is the smallest certified unequal-top counterexample to this projection claim in the present record; no six-player minimality theorem is asserted.

The strongest correct replacement is consequently three-level. Every fixed combinatorial type has the rational polyhedral projection (PXC109). Every equal-top partition cover has the canonical facet lift (PXC110). Cycles made from such covers are facet-dependency cycles, as in (PXC104)–(PXC106); general PXC108 obstructions remain circuit–covector decorations of selector cells and need not contain an exact-cone facet. This also explains why oriented-matroid programming is the right local language while directed-cycle terminology alone does not prove a new matroid circuit or exact-cone facet.

The exact catalogue/rank audit and a standalone solver-free incidence replay are:

```
python3 verify_pxc109_exact_cone_projection.py \
  --output pxc109_exact_cone_projection.json
python3 verify_pxc109_exact_cone_projection_independent.py \
  --output pxc109_exact_cone_projection_independent.json
```

The first exhausts all 7,006 labelled six-player facet normals and checks the stated ranks. The second imports no project LP, selector, cover, catalogue, semi-balanced, or PXC108 code. Both retain PXC101, PXC104, and the full-rank faithful PXC90 calibration; the PXC88 rank test is the fail-open negative control for the overstrong facet correspondence.

Thus (PXC101) classifies the minimal five-row tight-menu type containing PXC98: two indispensable Lmin rows, two extreme cover supports, and one transverse later blocker outside the cover-eligible union. It does not claim that every six-player failure has this incidence type; (PXC104) gives a different six-row type with three extreme covers. The primary verifier recomputes all coalition minima, selectors, certificates and covers at $n = 6, 7, 8$, while a separate rational active-set implementation replays the six-player anchor and the dimension-free identities without importing project LP or selector code:

```
python3 verifiers/divergence/verify_pxc101_forced_zero_family.py \
  --output verifiers/divergence/pxc101_forced_zero_family.json
python3 verifiers/divergence/verify_pxc101_forced_zero_family_independent.py \
  --output verifiers/divergence/pxc101_forced_zero_family_independent.json
```

The finite evidence is now definition-correct but remains bounded. The two pinned six-player archives contain 89 ordered saturated roots, including 35 with three gainers and two outsiders, but none in the strict single-heavy two-player-top cell; every one of the 35 satisfies the reserve. A symmetry-reduced two-block integer grid through total 11 plus 1,000 seeded multigenerator envelopes gives 1,975 exact games, 592 divergent games, 104 three-gainer/two-outsider roots, and four in the strict single-heavy cell; all 104 satisfy the reserve. These are an archive replay and an exhaustive bounded grid plus a random probe, not an all-game theorem. Theorems (PXC94)–(PXC95), rather than these counts, close the six-player reserve/direct-failing top-cover frontier. Accordingly the displayed seven-player example is the first failure in that ordered saturated program, not merely the smallest found there. This conditional dimension threshold must not be advertised as a global six-player selector theorem outside the stated cell.

The exact six-player census has 78 rooted instances with a multi-omission row, 98 saturated no-multi instances, 27 intermediate instances, and 500 unit instances. Every one of the 27 intermediate instances is already closed by the direct arm. The exact Hall replay constructs 527 strict dual price certificates. Seven satisfy the Lorenz-tangent premise of (PXC62), and each has an x-tight rooted blocker containing every gainer; none requires the root-free alternative. More strikingly, all 89 ordered no-multi instances are saturated: the census contains no ordered intermediate or unit root. All 16 positive-floor residual roots are ordered and saturated, and close under the selection-free fractional (PXC63) bound. In fact all 89 ordered saturated archive roots satisfy (PXC63b). Both sharp seven-player orientations are saturated and retain the required equality compatibility. These exact facts motivate, but do not prove, the prospective ordered-saturation theorem. The weighted Hall alternative is the rigorous fallback: either prove order forces (PXC61), or price a violating β by the established separator/mixed-facet machinery.

The exact archive audit now closes eight of the ten quotient occurrences at the price interface: a positive value among the δ_A and σ_D lies in $[-B_I, \Lambda_I]$. Four use the marked Lmin reverse payment and four use an Lmax locked-block splitter. The remaining two occurrences are the same record; their only marked row is coalition $22 = \{1, 2, 4\}$, which is tight for both selections and has zero comparison payment. Its data are

$$x = (7/2, 17/4, 7/2, 4, 17/4, 7/2), \qquad z = (3, 4, 4, 4, 4, 4), \qquad T = 23, \qquad O = 22.$$

Hence the finite residue has collapsed from an unrestricted weighted five-role case to one common-hinge normal form. That last record closes by the direct marked-hinge route (PXC45): coalition

$22 = \{1, 2, 4\}$ contains only Lmax-top players, so it gives a normalized common over-cover with $\theta = 1/3$ and zero cross-surplus. The resulting exact spectral overlap bound is $43/2$, well above the six-player target $2T/3 = 46/3$. Thus all ten archived quotient occurrences close, eight by priced mixed-colour payments and two duplicate occurrences by one common-hinge/top-over-cover certificate. A universal theorem still has to show that the positive arm is sufficiently priced, or that a zero marked payment forces a target-reaching common-support packet; the latter implication does not follow from (PXC44) alone.

The distinction is already active at six players. On the full tight menu, minimum deletion is globally hybrid-price optimal on 386 of the 415 strict cycles, while a larger deletion size strictly lowers $\min\{Z_I, X_I\}$ on the other 29; every strict handover is from $p = 1$ to $p = 2$. Global optima occur at $p = 1$ on 351 cycles and at $p = 2$ on 110, with both sizes optimal on some cycles. On the extracted staged support the minimum class appears optimal on all 415 only because those improving rows are absent. Thus restricting the proof to the extracted basis or to minimum deletion would create a genuine local maximum; (PXC16) is the correct basis-independent target. The exchange audit removes the residual mystery in these 29 cases: every handover is generated by an x-tight co-singleton, and the newly deleted player has $x_j = z_j = 0$. Here $n = 6$, $g = 2$, and $p = 1$, so $c_1 = 6/5$, $c_2 = 1$, $a_1 = 1/5$, and $a_2 = 0$; (PXC18) strictly lowers both positive prices. All 157 available co-singleton exchange pairs in the full menus satisfy (PXC18) exactly. These are finite archive counts, not an all-dimensions supply theorem for cheap co-singletons.

The exact six-player replay finds 493 distinct selected-support blocks over 420 nongainer carrier cycles; 477 meet the quota, and every cycle has at least one. The smallest closing block has $p = 1$ on 325 cycles and $p = 2$ on 95. In the seven-player control the unique selected repair has $P = \{2\}$, $Q(P) = 71/35$, required deficit $54/35$, quota margin $17/35$, and hence budget margin $17/30$. These are finite controls, not existence proofs.

A stricter optimal-face audit makes the residual combinatorics explicit. On the 233 archived six-player records where every maximum-overlap root is unaffordable, maximize root excess over the entire optimal labelled-price face. All 521 root optima have root excess zero. Nevertheless every root has a target-reserve row in its complete staged support. The returned price-optimal packing itself contains one at 444 roots, while some packing on the full optimal price face contains one at 448. The remaining 73 are genuine price-face obstructions, spread over 33 records. Every one is in the pair-free root branch; its common carriers are exactly the other members of the maximum-overlap root block, the root optimum has labelled facet margin -1, and its returned optimizer is one unit of a co-singleton row. There are 52 rooted instances with a two-player root block and 21 with a three-player block. Root selector loss is strict in 71 and tied in two. Thus the existing strict-loss proper-packet lemma already supplies a carrier-preserving proper tight row in 71 cases; what it does not ensure is that this row also omits the player excluded by the co-singleton optimizer. In the archive 71 rooted instances have no crossing proper carrier row at all. The other two have exactly one crossing row, whose co-singleton intersection is tight by (PXC4). Every one of the 73 instances has at least two carrier-row intersections whose final reserve is nonnegative. This is a finite compatibility fact, not the missing dimension-free reserve theorem. More sharply, applying the root-only corollary (PXC7) with B equal to the complete maximum-overlap nongainer root block certifies all 150 such intersections: 142 lower bounds are positive and eight are zero. Thus all 73 rooted obstructions close without inserting their exact reserve into the test. The coarser $M^- = M$ estimate fails on ten rows and leaves two roots uncovered; in the latter record the omitted nongainer actually has $z_i = 0$, whereas $M = 5$, and the true reserve is $10/3$. The apparent residue was therefore loss of Lmax stratum information, not a new exceptional packet. These counts remain an archive theorem only; a universal proof must show that the staged support or the separator/facet alternative selects an ordered cut satisfying the right side of (PXC6). The fixed

seven-player full-root-block control proves that the distinction matters: (PXC7) is $-59/60$ on its selected repair, while the retained-extra term in (PXC6) restores a positive bound $53/12$ and the exact reserve is $289/60$. Thus root mass plus one global next-stratum ceiling is not a universal substitute for the selected support. The comparison certificate (PXC11) is nonnegative on 148 of the 150 intersections and still closes at least one row for every one of the 73 roots. Exactly those 148 intersections have $\sigma_I = 0$, so (PXC11) is their exact reserve; the other two have $\sigma_I = 1/2$ and are closed by (PXC7). On the selected seven-player repair (PXC11) again equals the exact reserve $289/60$. This complementary coverage is evidence for a two-colour max-certificate, not a proof that one of the two bounds is nonnegative in every dimension.

The broader strict-cycle audit exposes a more rigid selector than either certificate alone. Over the 415 positive-loss nongainer carrier cycles in the same two six-player archives, the full Lmin-tight menu contains 1,025 admissible carrier rows. Restrict, separately for each cycle, to the rows minimizing the number p of excluded nongainers. There are 830 such minimum-deletion rows, and every one has both the ordered-cut bound (PXC6) and the omitted-gainer bound (PXC11) nonnegative. The minimum is $p = 1$ on 380 cycles and $p = 2$ on 35. One extracted complete staged support gives the same conclusion on its 862 admissible and 829 minimum rows. The full-menu test shows that the observation is not an artifact of that basis extraction. This is not the tautology of choosing a row with largest reserve: minimum deletion is defined using only the support incidence.

The price decomposition explains the six-player coincidence. Every one of the 830 full-menu minimum rows has $E_I^z = 0$; 828 also have $E_I^x = 0$. On each of the two cycles containing a positive-comparison minimum row, another minimum row has $E_I^x = 0$. Hence $Z_* = X_*$ on all 415 cycles, and both equal-price routes close; the two exceptional cycles illustrate why the minima in (PXC14) may legitimately be attained on different rows. At seven players the new control has instead $X_* = 35/8 < Z_* = 5$: the Lmin colour takes over exactly when the ordered-cut coarsening becomes lossy. This is a finite structural handover, not a proof that a zero-error minimum row exists in every dimension.

The restriction matters. Six full-menu nonminimal $p = 2$ rows have actual target reserve $-1/3$ (and both lower certificates equal $-1/3$), while a minimum-deletion row closes the same cycle; four of those six rows occur in the extracted support. Thus "every proper carrier row closes" is false even in the archive. In the fixed seven-player hostile family, the sole strict carrier-cycle case has four admissible rows; all four are minimum-deletion rows and both certificates close. These exact finite facts initially suggested simultaneous closure, but that strengthening is false. Let v be the lower envelope of the four common-total allocations

$$(1,5,0,7,0,0,0),\ (2,0,3,1,2,3,2),\ (0,2,0,0,5,4,2),\ (0,1,5,6,0,1,0).$$

This is a nonnegative singleton-normalized exact seven-player game. Exact stage certificates give

$$x = (7/4, 7/4, 7/4, 17/8, 7/4, 17/8, 7/4), \qquad z = (5/3, 5/3, 5/3, 2, 2, 2, 2).$$

Its maximum-overlap roots $3, 5$ are both unaffordable and form a strict carrier cycle. The full Lmin-tight menu consists of rows 56 and 104; both are minimum-deletion rows with $P = \{0, 1, 2\}$. On each row the actual reserve and (PXC11) are $13/24$, whereas (PXC6) is $-1/12$. Thus the two errors in (PXC12) are $E_I^x = 0$ and $E_I^z = 5/8$: the two colours are genuinely disjunctive. The game is safely above the target: $O = 25/2 > 91/12$.

The corrected sufficient statement is that some minimum-deletion proper carrier row has

$$\max\{\text{PXC6}, \text{PXC11}\} \geq 0.$$

The necessary direct target is the still weaker actual-price test (PXC23). If it fails, the sparse Lmax separator dual must yield the priced rooted-facet alternative. Proving that disjunction requires an

exchange or Farkas argument for the weighted row scores; the counts and the new control do not prove it. They are reproduced by the archived deterministic verifier and JSON artifact.

Of the closest properizations, 25 delete one payoff gainer at incidence distance and price loss one, while 48 delete one payoff gainer and one nongainer at distance and price loss two. The four additional zero-price-loss distance-two cases are optimizer-selection swaps, so counting them as support-exchange obstructions would be a quantifier error. Consequently "some price-optimal marked packing already contains a closer" is false even in the six-player archive. The surviving missing lemma is a charged properization theorem for this rigid residual class, not an identity between two independently optimized certificates. These counts are exact finite evidence, not an all-six-player census or a dimension-free exchange classification.

The strict quota itself is not universal, even when the excluded block is a first Lmin-floor singleton. Consider the singleton-normalized exact lower envelope of

$$(6,7,1,0,1,0,1), \quad (7,0,0,0,0,9,0), \quad (0,0,0,5,3,6,2). \tag{QFH1}$$

Its two exact staged selectors are

$$x = (3,7/2,1,3,1,7/2,1), \qquad z = (3,3,0,2,3,3,2). \tag{QFH2}$$

The tied maximum-overlap roots are $\{0,1,5\}$. Roots 1 and 5 form an unaffordable pair-free carrier cycle, with labelled margins $-1,-1$; root 0 is affordable with margin 1. Every one of the six Lmin-tight rows containing the cycle misses (EDQ2). The least miss uses $P=\{2\}$, a first-floor singleton, and has quota margin $-2/7$ and scalar budget $-1/3$. The older interval-hitting filter therefore finds no packet. But the same row has base $5/6$, so (EDQ5) retains cardinality reserve 4 and gives $\mathcal{R}_{7/12}(\{2\}) = 4 - 1/3 = 11/3$: it proves the target exactly despite failure of the strict quota.

The same cell also exposes the facet arm. At root 0 the extreme marked packing puts weight $1/2$ on each of coalitions $59, 73, 81$. Its three active offroot equations are

$$a+b=1, \qquad a+c=1, \qquad b+c=1,$$

so the packing is an extreme half-integral vertex. Its root excess is $1/2$, its offroot deficits are $(1/2,1,0,0,1/2,0)$, and the rooted pair lift is

$$\tfrac{1}{2}\mathbf{1}_{59} + \tfrac{1}{2}\mathbf{1}_{73} + \tfrac{1}{2}\mathbf{1}_{81} + \tfrac{1}{2}\mathbf{1}_{3} + \mathbf{1}_{5} + \tfrac{1}{2}\mathbf{1}_{33} - \tfrac{5}{2}e_0 = \mathbf{1}_N. \tag{QFH3}$$

The support has rank seven and no proper exceptional-set decomposition; hence (QFH3) is an indecomposable exact-game-cone facet by the rooted pair-lift theorem and the Studený–Kratochvíl criterion. It is tight at (QFH1). Its canonical rooted support $\{1,3,5,9,49,81,103\}$ is distinct from the five six-player residual-atlas orbits, but the facet mechanism is the already proved extreme-packing pair lift, not a new full-cone class. No literature-novelty claim is made. Here $Q/\delta = 4$. At the first Lmin floor $a=1$, the weighted floor distance is $5/2$ while the rooted-pair premium is 3, so the multistage penalty (MPE1) is zero. The proved rooted price theorem therefore gives the independent target-reaching bound

$$\frac{O}{T} \geq \frac{7}{4+7-1} = \frac{7}{10} > \frac{7}{12}. \tag{QFH3a}$$

Finally, $D = z(\{4,6\}) - x(\{4,6\}) = 3$ and

$$O = T - D = 16 - 3 = 13 > \frac{7}{12}16. \tag{QFH4}$$

Thus a negative scalar quota is only failure of that stronger sufficient decomposition; it is not negative target evidence. The exact example is a clean boundary where strict quota fails, while both target reserve and a positively priced indecomposable rooted facet close. It does not prove universal selection of either the target-reserve row or the facet.

The root-exclusion transport now meets the scalar target in the correct currency. With $\alpha = M - T/n$, every eligible source satisfies

$$T/n - z_i = (M - z_i) - \alpha \geq r_j + (x_i - z_i) - \alpha. \tag{EDQ3}$$

The remaining obstruction is precisely the mean-to-root premium α together with repeated source congestion. Those charges must be controlled by the sparse Lmax separator or a marked-facet alternative; neither (RET) nor (EDQ3) alone proves selection of a quota-meeting support.

The two balanced seven-player equality orientations are compatible with this reduction in the sharp way. Their complete staged supports have quota-closing rows only at scalar equality with $h = 1$; the corresponding direct bases are $7/15$ for $G(2,4)$ and $7/16$ for $G(3,3)$, both below $7/12$. Hence (EDQ) does not remove the rooted equality arm: the carrier-cycle cardinality lemma requires $h \geq 2$, and the block-star facet remains necessary.

The same reduction classifies equality on the carrier-direct arm. By (CCC), its cardinality base is strictly above the target except at $(n, p, g, h, s) = (6, 2, 2, 2, 1)$. In that cell the remaining two terms are the nonnegative low-Lmax payment and $[g/(p+1)](U_{g-1} - L_H)$. Equality forces $z(P) = 0$ and quota equality, and (EDQ2) then gives $D = T/3$.

Corollary 51 (carrier-direct equality)**.** *A target-reaching direct packet whose nongainer support contains a carrier cycle is strict except possibly for $n = 6$ and $(p, g, h) = (2, 2, 2)$. Equality there requires $z(P) = 0$ and $z(G) - x(G) = T/3$.*

This classifies only the direct carrier arm. Rooted facets may attain the target in seven players, as the two balanced orientations do.

Every Lmin-tight coalition C meeting a carrier cycle contains the whole cycle. Writing $C = H_C \dot\cup J_C$ for its nongainers and gainers, core feasibility gives $z(C) - x(C) \geq 0$. Hence it is automatically a valid comparison coalition in (TDG). Equations (CB1)–(CB3) show that tightness and the separate residual sign are stronger than necessary: after choosing H_C, one may take any $g - 1$ gainers and test only $U_{g-1} \geq L_{H_C}$. By (CCC) its cardinality base already reaches the target. All 228 all-unaffordable six-player archive rows have such a tight support coalition; 216 use the carrier cycle as the nongainer block and twelve require one additional nongainer. The seven-player control above requires two. These are finite controls, not a proof. The current dimension-free missing lemma is therefore selected tight-support congestion, or a separator/marked-facet alternative; the refuted full-root-block sign may not be used.

For the coefficient-two certificate below, the normalized packing has $u_{24} = u_{44} = u_{79} = 1/2$ at root 3. Deleting the rooted pair 24 leaves the midpoint of the vertices supported on 44 and 79. Its row is the half-sum of

$$\mathbf{1}_{44} + \mathbf{1}_9 + \mathbf{1}_{10} + \mathbf{1}_{24} + \mathbf{1}_{72} - 4e_3 = \mathbf{1}_N, \qquad \mathbf{1}_{79} + \mathbf{1}_{24} + \mathbf{1}_{40} - 2e_3 = \mathbf{1}_N.$$

Both systems are indecomposable, active exact-cone facets (support sizes six and four). Exact-LP and independently formulated Z3 exceptional-set tests both return zero decompositions. Thus this strict leader yields two genuine facet currencies rather than only a nonminimal aggregate row.

There are two distinct root optimizations here. Minimizing the deleted load $\max_{i \neq r} \alpha_i$ gives the best root-deletion ceiling, whereas finding a positive pair lift requires maximizing $\alpha_r - \max_{i \neq r} \alpha_i$ after restricting the over-cover to common hinges. They need not select the same certificate. In

the coefficient-two seven-player control, a minimum-deleted-load certificate is tied, but the common certificate

$$\eta = \frac{1}{3}, \qquad \alpha = (0, 0, \frac{1}{3}, \frac{2}{3}, 0, 0, 0), \qquad \lambda_{24} = \lambda_{44} = \lambda_{79} = \frac{1}{3}$$

has strict leader gap $1/3$ and produces an active pair lift with $\delta = 1/2$. Exact rational hostile controls also separate the remaining premises: the primitive coefficient-seven and both sharp-family cases admit strict common leaders; the neutral rectangle admits common over-covers but no strict leader; and the optimized label-free, actual-top-fixed, and top-and-bottom-fixed controls admit no normalized top over-cover supported entirely on common hinges. These finite exclusions do not prove the hybrid bound, but show that its missing gluing lemma must contain strict-common, tied-common, and noncommon-support branches; optimizing deleted load alone cannot supply that disjunction.

The noncommon branch has an exact deficiency dual. For the proper z-tight family $\mathcal{C}_z$, top block H, and $d_S = x(S) - v(S) \geq 0$, let

$$\epsilon_* := \min \Big\{ \sum_{S \in \mathcal{C}_z} d_S \lambda_S : \sum_S \lambda_S \mathbf{1}_S = \eta \mathbf{1}_N + \alpha,\ \alpha(H) = 1,\ \mathrm{supp}(\alpha) \subseteq H,\ \lambda, \eta, \alpha \geq 0 \Big\}. \quad \text{(CSD-P)}$$

Ordinary rational LP duality gives

$$\epsilon_* = \max\{t : y(S) \leq d_S\ (S \in \mathcal{C}_z),\ y(N) \geq 0,\ y_i \geq t\ (i \in H)\}. \quad \text{(CSD-D)}$$

Since every primal summand is nonnegative, $\epsilon_* = 0$ exactly when a normalized top over-cover supported on common hinges exists. Thus failure of common support has a canonical nonnegative value and an exact player-price separation witness.

More precisely, take any normalized top over-cover and any top root r. Put $b_i = \theta_- - \alpha_i$ off the root, $a_r = \alpha_r - \theta_-$, $\epsilon = \alpha \cdot x - M$, and

$$\kappa_r = \theta_- T - M - \sum_{i \neq r} b_i v(\{r, i\}). \quad \text{(NSA)}$$

The rooted pair lift shows $\kappa_r \geq 0$. Direct evaluation at the two selections gives

$$\sum_{i \neq r} b_i z_i - a_r z_r - \sum_{i \neq r} b_i v(\{r, i\}) = \kappa_r, \quad \sum_{i \neq r} b_i x_i - a_r x_r - \sum_{i \neq r} b_i v(\{r, i\}) = \kappa_r - \epsilon. \quad \text{(NSAxz)}$$

Hence the noncommon support consumes pair-lift exact-cone slack exactly one-for-one, and the same rooted pair-floor inequality is valid at Lmin iff $\kappa_r \geq \epsilon$. Equivalently,

$$\sum_{i \neq r} b_i (x_i - z_i) = a_r (x_r - z_r) - \epsilon. \quad \text{(NSAD)}$$

An unrestricted "some top root absorbs the surplus" lemma is false. Null-padding a certified five-player exact control gives

$$x = (43/280, 43/280, 47/280, 21/80, 21/80, 0), \quad z = (1/10, 29/140, 31/140, 33/140, 33/140, 0).$$

Here $\epsilon_* = 3/112$, while both top roots have $\kappa_r = 0$. The overlap is $25/28 > 2/3$, so this does not challenge the hybrid bound; null-player descent returns it to the proved five-player theorem. Any absorption lemma must therefore be restricted to the irreducible positive-Lmin-floor branch. In the 374 archived six-player controls, the independently returned minimum-dispersion certificate is

also minimum-surplus in every record, 73 deficiencies are positive, and some top root absorbs the surplus in every record. These are finite falsification checks, not a theorem.

Finally, the strongest possible single capped spectral anchor is itself an LP. With $L = \lambda/\theta$, $e = \eta/\theta$, and $u = \alpha/\theta$, maximize

$$R_* = \sum_{i \in H} u_i(x_i - \mu), \qquad \sum_S L_S \mathbf{1}_S = e\mathbf{1}_N + u, \quad 0 \le u \le 1, \quad L, e \ge 0. \tag{BSA}$$

Perspective normalization proves $R_* = \max(M + \epsilon - \mu)/\theta$ over all top over-covers. Even this optimum gives only $605/36 < 35/2$, $39/2 < 21$, and $47/2 < 49/2$ on the optimized label-free, actual-top-fixed, and top-and-bottom-fixed controls. Thus the noncommon cases cannot be closed by choosing a better single scalar spectral anchor: multiple rooted or staged rows must genuinely be glued.

For comparison, the singleton-slack lift puts $q_i = 1 - p_i$ off the root. Every rooted packing satisfies the incidence identity

$$\sum_{S \in \mathcal{C}_x} u_S \mathbf{1}_S + \sum_{i \ne r} q_i e_i - \delta e_r = \mathbf{1}_N. \tag{RMS}$$

If u is a vertex, the positive columns in the standard-form system $A_{N\setminus\{r\}}u + q = \mathbf{1}$ are linearly independent; adjoining e_r preserves independence. Therefore, when $\delta > 0$, the nonzero support in (RMS) is a purely min-semi-balanced system with exceptional coalition $\{r\}$ in the precise terminology of Studený and Kratochvíl (2022, Lemma 2(c)). When $\delta \le 0$ the completion is min-balanced. For $n \ge 3$ the positive-root case defines a full exact-cone facet exactly when it passes their separate indecomposability test. Thus rooted-packing vertex enumeration gives an exact algorithm for singleton-exceptional facet candidates exposed by a fixed Lmin menu, but not an automatic facet correspondence. In fact all eight nontrivial cuts in the present hostile and sharp controls are min-semi-balanced but decomposable; an independently formulated Z3 check agrees with the exact-LP decomposition counts.

In fact the negative facet conclusion is dimension-free. If $x_r > 0$, no vertex completion (RMS) is indecomposable. For $\delta > 0$, (RDB) forces $q_i > 0$ for some $i \ne r$, so the completion contains $\{r\}$ and $\{i\}$. Its column independence excludes $\{r, i\}$, but $e_r + e_i - \mathbf{1}_{\{r,i\}} = 0$ makes $\{r, i\}$ exceptional in the enlarged system, which is a decomposition in the sense of Studený and Kratochvíl (2022, Definition 7). For $\delta \le 0$ the completion is min-balanced and likewise does not define a facet of the exact cone when $n \ge 3$. Therefore the rooted construction is a relative inequality and candidate-enumeration theorem in its singleton-slack form. The rooted-pair shear (RPL) is the required different lift and gives the exact facet class; the remaining task is to prove that sparse Lmax gluing selects enough such facet charge, not to enumerate more singleton-slack completions.

This reduction passes exact rational tests at all maximum-overlap roots of the 374 archived six-player pairs (703 rooted programs), as well as six hostile fixed branches and both sharp seven-player orientations. Those are falsification controls, not a proof for any unenumerated domain. An exact six-player negative control drops from $3/4$ to $5/8$ when the common-hinge rows are removed entirely, so a leximin-packet-only theorem would be false. The open dimension-free step is to prove that some maximum-overlap root and some sparse over-cover always make this common-hinge relaxation attain at least $n/\lfloor n^2/4 \rfloor$, or to produce an exact rational branch that refutes that statement.

As an exact control of (RHP), an independently formulated cutting-plane verifier separates over the fractional packing polytope and reproduces the direct three-generator minima at every maximum-overlap root of six hostile branches, including the neutral-endpoint control, and two sharp orientations. The neutral control needs one two-coalition cut and returns $7/10$; the sharp

cases require one packing cut, supported on four and three hinges respectively, and return 7/12 exactly. This checks the elimination arithmetic, not the universal existence of a target-proving packing.

One sharper scalar bridge suggested by this reduction is now known to be false. Minimize $\theta_+ = \max_i \alpha_i$ over normalized first-stage ceiling over-covers $\sum_S \lambda_S \mathbf{1}_S = \eta \mathbf{1}_N + \alpha$, and put

$$a = \min\{1/\theta_+, \lfloor n/2 \rfloor\}, \qquad u = a\alpha.$$

Then $0 \le u \le 1$, $u(N) = a$, and scaling the incidence identity gives

$$a \sum_S \lambda_S \mathbf{1}_S + (\mathbf{1}_N - u) = (1 + a\eta)\mathbf{1}_N.$$

If $\epsilon = \sum_S \lambda_S [x(S) - z(S)] \ge 0$, evaluation at the two selections gives $u \cdot x = a(M + \epsilon)$ and $u \cdot z = aM$, with complementary totals on $\mathbf{1}_N - u$. The proposed conjecture was

$$\frac{O(x,z)}{T} \;\ge\; \frac{1}{a} + \frac{1}{n-a}. \tag{CEB}$$

Condition (CEB) would imply the target because $a(n-a) \le \lfloor n^2/4 \rfloor$, is exact on both balanced orientations of the displayed family, and has the correct five/six boundary. Exact archived checks are falsification evidence only: (CEB) passes all 374 stored six-player pairs and the named seven-player controls, while failing on 12 of 124 distinct stored five-player pairs. Those controls missed an exact seven-player counterexample. For

$$g = (12/17, 4/17, 0, 1/17, 0, 0, 0), \qquad h = (0, 0, 6/23, 0, 6/23, 10/23, 1/23),$$

put $v(S) = \min\{g(S), h(S)\}$. The generators have common total one, are core points, attain every coalition worth between them, and have coordinatewise minimum zero; hence v is singleton-normalized and exact. Exact staged certificates give

$$x = (227/782, 48/391, 48/391, 1/34, 48/391, 13/46, 1/34), \quad z = (91, 12, 90, 12, 90, 91, 5)/391.$$

Exact primal and dual incidence witnesses give $\theta_+ = 3/5$ and $a = 5/3$, but

$$O(x,z) = \frac{613}{782} < \frac{63}{80} = \frac{1}{a} + \frac{1}{7-a}$$

with margin $-113/31280$. Thus (CEB) is refuted. The example does not challenge the main target: $613/782 - 7/12 = 941/4692 > 0$.

The cap in (CEB) is not an exogenous parameter. If $\mathcal{T}_z$ is the family of proper coalitions tight at z and H is its top payoff block, ordinary rational LP duality gives the human formula

$$\theta_+ = \max t: \quad y(S) \ge 0 \ (S \in \mathcal{T}_z), \quad y(N) \le 0, \quad t + y_i \le q_i \ (i \in H), \quad q \ge 0, \quad \sum_{i \in H} q_i \le 1. \tag{DD}$$

Formula (DD) remains valid: θ_+ depends only on the labelled top block and its complete tight-coalition incidence hypergraph; adding a tight coalition can only lower it. This explains why the fixed-cap one-witness relaxation is too strong on an exact six-player boundary: its minimizer lowers overlap only after θ_+ moves from 1/2 to 1/3. The seven-player example shows that retaining dispersion optimality still does not prove (CEB). Its positive hinge surplus instead makes the already

proved sharpened spectral floor exceed 7/12. The live route is therefore a spectral-or-common-hinge disjunction at the weaker universal target, not a reciprocal-mass theorem.

The label loss in that disjunction has an exact layer-cake form. Fix the actual ceiling block H, put $R = N \setminus H$ and $m = |R|$, and define $X_t = \{i \in R : x_i \geq t\}$ and $Z_t = \{i \in R : z_i \geq t\}$. If B_H is the spectral expression with H fixed and the residual spectra oppositely paired, then

$$O - B_H = \int_0^\infty \left(|X_t \cap Z_t| - \max\{0, |X_t| + |Z_t| - m\}\right) dt. \tag{LC}$$

Indeed, the first term is the layer-cake formula for residual overlap, while the second is the Fréchet lower intersection bound attained simultaneously by opposite sorting. The integrand is a nonnegative integer. On a fixed pair of ordered payoff partitions, (LC) is therefore a finite nonnegative linear form in the consecutive merged stratum gaps. If B_* denotes the spectral minimum over admissible ceiling assignments, then

$$O = B_* + (B_H - B_*) + (O - B_H),$$

separating top-selection currency from residual label currency. This is a proved identity, not the missing inequality: the remaining common-hinge problem is to show that the two currencies cover $\max\{0, nT/\lfloor n^2/4 \rfloor - B_*\}$. It does, exactly, identify a linear stratum-gap target for the rooted branch-circuit and relative-facet program without conflating those relative circuits with facets of the full exact-game cone.

There is also a direct bridge to the staged transports. If A_t denotes the players in R high in both selections and D_t those low in both, then the integrand in (LC) is exactly $\min\{|A_t|, |D_t|\}$. Put $a_i = \min\{x_i, z_i\}$ and $b_i = \max\{x_i, z_i\}$. Whenever $b_j < a_i$, the interval $(b_j, a_i]$ is a unit high–low rectangle. Moreover $x_j \leq b_j < a_i \leq x_i$, so the complete strict Lmin support theorem makes the pair (i, j) staged-available. This includes neutral players and is strictly larger than the opposite-sign carrier family used initially.

Sort the residual a-values downward and b-values upward and pair equal ranks. At threshold t, the active rank pairs form the intersection of an initial segment of size $|A_t|$ and one of size $|D_t|$; their number is therefore exactly $\min\{|A_t|, |D_t|\}$. Integrating gives the human identity

$$\sum_{k=1}^{m} (a_{[k]} - b_{[k]})_+ = O - B_H. \tag{CS}$$

Thus canonical weight-one rectangles saturate every label-currency layer. This corrects the omission of neutral endpoints and proves that the complete carrier packing is an identity, not the missing inequality.

For fixed payoff strata this packing problem has an exact combinatorial min–max theorem. Split the threshold line into the merged coordinate intervals J_ℓ, let $c_\ell = \min\{|A_t|, |D_t|\}$ there, and let $A_{\ell e}$ indicate whether carrier interval e contains J_ℓ. Then

$$\max\{\textstyle\sum_e |I_e| \rho_e : A\rho \leq c,\ \rho \geq 0\} = \min\{\textstyle\sum_\ell c_\ell p_\ell : \sum_{\ell : J_\ell \subset I_e} p_\ell \geq |I_e|,\ p \geq 0\}.$$

The columns of A have the consecutive-ones property, so A is totally unimodular. Consequently an integral-capacity instance has an integral optimal packing with support at most the number of threshold layers, and the primal and dual are exactly computable in polynomial time for rational data. The feasible cover $p_\ell = |J_\ell|$ recovers the label-currency upper bound. Writing $r_k = \sum_{\ell \leq k} p_\ell - \beta_k$ turns every carrier-cover constraint for $(\beta_s, \beta_t]$ into $r_t \geq r_s$ and nonnegativity into $r_k - r_{k-1} \geq -|J_k|$. Hence an extreme cover has a rooted spanning-tree certificate in the graph of tight path

and carrier arcs. This is a fixed-stratum theorem for restricted rectangle families and an exact interface to rooted incidence packets; it is not by itself a facet theorem for the exact-game cone. For the complete family, (CS) already gives the optimum $O - B_H$, so requiring it to pay the post-spectral deficit is exactly the original hybrid-overlap conjecture. A deliberately restricted, transport-oriented family might still give a nontrivial sufficient condition, but staged availability alone neither selects it nor proves its adequacy.

In fact, normalized circuit coefficients do not impose finite supplies. Write either sign-coherent relation as

$$\sum_S \lambda_S \mathbf{1}_S + \sigma \left(e_i + \sum_{j \in P} a_j e_j \right) = c\mathbf{1}_N, \qquad \lambda, a \geq 0,$$

where $\sigma = 1$ for Lmin and -1 for Lmax. If a normalized relation has $a_j > 0$, then for $K \geq 1$ multiplying all other coefficients by K and replacing a_i by $Ka_i + K - 1$ gives another normalized relation with carrier coefficient Ka_j. Convexly combine witnesses before amplifying to handle several carriers simultaneously. Hence the projection of the joint certificate/supply system is exactly the nonnegative orthant on those carrier edges which occur with positive coefficient in some relation. Coefficient magnitude disappears; only rooted-circuit support remains.

That support is actually complete. If $x_j < x_i$, then j was pinned at a strictly earlier Lmin stage. Adding a positive multiple of j's normalized relation to i's preserves the forced coefficient of e_i and gives e_j positive auxiliary coefficient. The sign-reversed Lmax argument applies whenever $z_j > z_i$. Hence every strict order-compatible residual pair is available; in particular every rectangle in (CS) is available through Lmin. Equality within a stratum gives only a zero-length interval.

The joint LP dual reaches the same conclusion. Dotting its staged price vector with a normalized relation forces the root price, and then every available carrier's supply price, to vanish. What remains is precisely the threshold-cover dual. On the complete family its optimum is the physical layer-width cover, while (CS) is the matching primal. Consequently neither circuit coefficient magnitudes, support patterns, nor threshold-tree prices can furnish the missing exact-cone facet mechanism.

In the two hostile controls where the spectral and top-selection terms do not reach the target, this complete-carrier system has short exact witnesses. Respectively,

$$\mathbf{1}_{\{0,2,4\}} + \mathbf{1}_{\{0,1,2,3,6\}} + \mathbf{1}_{\{4,6\}} + 2e_5 + e_1 + e_3 = 2 \cdot \mathbf{1}_N$$

supplies $(1/2, 8]$. The old opposite-sign packing and corrected full rectangle packing both have optimum 15, versus deficit 3/2. Similarly,

$$\mathbf{1}_{\{1,3,4,5,6\}} + e_2 + e_0 = \mathbf{1}_N$$

supplies $(1/2, 19/2]$. Its old opposite-sign packing optimum is 9, versus deficit 1, while the corrected full rectangle optimum is the full label currency 33/2. Primal packing, threshold-cover dual, and exact staged-availability witnesses agree. These historical opposite-sign checks are retained as regressions. The complete clipped family saturates label currency universally by (CS), but that fact does not prove the hybrid-overlap inequality.

The neutral-endpoint correction has an exact seven-player stress test. Its only residual gainer and loser satisfy $x_1 = x_2$ and $z_1 > z_2$, so the old opposite-sign family contains no strict interval. Yet its label currency is 15167/38092, exactly recovered by canonical rectangles of lengths 8081/38092 and 7086/38092 using neutral low endpoints. Fourteen explicit staged identities and a separately

formulated exact staged replay certify the selections. The example does not threaten the main bound: its overlap exceeds 7/12 by 43349/114276.

There is nevertheless an exact fixed-dimension decision procedure. Refine the exact-game cone into rational polyhedral cells fixing the active staged bases, payoff orders, ceiling block, spectral branch, and endpoint comparisons. On each cell both selections and every term in the hybrid-overlap inequality are rational linear in the game. Thus the existence of a violating game reduces to finitely many exact rational linear programs. Rooted cover trees remain useful only for deliberately restricted rectangle subfamilies; the complete family is already exhausted by (CS). This procedure may be exponential or worse, and a cellwise boundary row is only a candidate full-cone facet: it must still glue across cells and lift to an indecomposable semi-balanced system.

7.1 Multistage direct gluing

The direct-labelled branch has a dimension-free master identity. Partition $N = P \dot\cup G \dot\cup H$, with sizes $p, g, h \ge 1$, and put $m = p + g$, $q = g + h$, and $n = p + g + h$. Let P, H be nongainers and G the gainers, so

$$z_i \le x_i \quad (i \in P \cup H), \qquad x_i \le z_i \quad (i \in G).$$

Put $\ell = z(P)/p$ and $b = z(G \cup H)/q$. Let $J \subset G$, $s = |J| < g$, and assume $x(H \cup J) \le z(H \cup J)$. This comparison is automatic when $H \cup J$ is tight at x, because z lies in the core. Put

$$B = \frac{g}{m - s}, \qquad R_x = x(J) + (1 - B)x(G \setminus J) - Bx(P),$$
$$R_z = (1 - B)[z(H) - hb] - B[z(J) - sb].$$

Then

$$O - \left(1 - \frac{gp}{q(m - s)}\right) T = \frac{gpn}{q(m - s)}\ell + B[z(H \cup J) - x(H \cup J)] + R_x + R_z. \tag{TDG}$$

Indeed, expand the comparison term and use $x(H) = T - x(P) - x(G)$. Direct cancellation gives

$$B[z(H \cup J) - x(H \cup J)] + R_x + R_z = x(G) + z(H) - BT - (1 - B)hb + Bsb.$$

Substitution of $b = (T - p\ell)/q$, $B = g/(m - s)$, and $q = g + h$ yields (TDG).

Consequently the weaker condition $R_x + R_z \ge 0$ already implies

$$\frac{O}{T} \ge 1 - \frac{gp}{(g + h)(p + g - s)}. \tag{TDB}$$

Equality requires $\ell = 0$, a tight comparison, and zero residual sum. The right side of (TDB) reaches $n/\lfloor n^2/4 \rfloor$ exactly when

$$\lfloor n^2/4 \rfloor gp \le [\lfloor n^2/4 \rfloor - n](g + h)(p + g - s). \tag{TDT}$$

This criterion is one-dimensional. Put $\delta_i = z_i - x_i$ on G, $L_H = x(H) - z(H)$, and

$$U_s = \frac{(p - s)x(G) - gx(P) + (p - s)[z(H) - hb] + gsb}{g}.$$

Then

$$z(H \cup J) - x(H \cup J) = \delta(J) - L_H, \qquad (m - s)(R_x + R_z) = g[U_s - \delta(J)]. \tag{GPI}$$

Thus a cardinality-s packet closes the direct branch whenever $L_H \le \delta(J) \le U_s$. The arbitrary multistratum dependence has collapsed to an explicit interval-hitting problem on the gainer weights.

A simple sufficient condition for the residual signs is $s \le p$, $\max x(P) \le \min x(G)$, $\max z(G) \le \min z(H)$, with J the s lowest-z players of G. Then $0 < B \le 1$ and

$$R_x \ge [s + (1-B)(g-s) - Bp]\,\frac{x(P)}{p} = 0.$$

Moreover $\operatorname{avg} z(H) \ge b \ge \operatorname{avg} z(G)$ and $\operatorname{avg} z(J) \le \operatorname{avg} z(G)$, so $R_z \ge 0$.

Consequently, let L be the nongainers, choose a nonempty proper block $H \subset L$, put $P = L \setminus H$, and take J to be the s lowest-z gainers. If

$$s \le |P|, \quad \max x(P) \le \min x(G), \quad \max z(G) \le \min z(H), \quad H \cup J \text{ is Lmin-tight}, \tag{UBD1}$$

and

$$\lfloor n^2/4 \rfloor |G||P| \le [\lfloor n^2/4 \rfloor - n](|G| + |H|)(|P| + |G| - s), \tag{UBD2}$$

then the target hybrid-overlap inequality follows. Tightness and core feasibility pay the comparison term, (UBD1) pays both ordinal residuals, and (UBD2) is precisely the target cardinality test. This is a conditional upper-block handover; it does not assert that such a block and packet always exist.

The exact criterion is stronger. It recovers both archived hostile direct branches: the actual-top-fixed branch has base $7/10$, $R_x = 15/4$, and $R_z = 111/20$; the top-and-bottom-fixed branch has base $2/3$, $R_x = 1$, and $R_z = 11$, despite $s = 3 > p = 1$. Thus no two-stratum assumption on z is needed. Their gain packets are respectively $1/2 \in [1/2, 129/10]$ and $3/2 \in [3/2, 15/2]$; the fresh control below has the singleton interval $[1/2, 1/2]$. The remaining issue is a selection lemma: either a complete staged packet must supply a target-reaching cardinality-s gainer packet in $[L_H, U_s]$, or every interval miss must force an affordable mixed-colour rooted packet.

An exact finite atlas exhausts the sign-compatible partitions and proper gainer packets on the seven fixed hostile selections. The five nonsharp obstructions have direct bases $3/5$, $3/4$, $2/3$, $5/8$, and $5/9$ (the last at $n = 8$). The first four use Lmin-tight comparisons; the last has direct slack $5/39$. Neither sharp $G(2,4)$ nor sharp $G(3,3)$ has a target-reaching direct packet, and both are closed at equality by the rooted block-star theorem. The residual-sum criterion is essential: the $3/4$ certificate has $(R_x, R_z) = (15, -15/2)$, and the $5/9$ certificate has $(R_x, R_z) = (-8/9, 116/117)$. This is branch-level falsification evidence, not a universal dichotomy theorem.

A broader exact census refutes the tempting direct-only strengthening. Of 374 archived six-player pairs, 296 admit a target-reaching direct packet and 282 admit one with an Lmin-tight comparison. The other 78 have no direct packet, and 76 of those are strictly above $2/3$. The first strict control is

$$x = (1,2,1,1,1,1), \qquad z = (3/2, 3/2, 3/2, 5/6, 5/6, 5/6), \qquad O/T = 6/7.$$

The separately encoded common-hinge/root audit closes all 374 archived pairs with minimum exactly $2/3$. Thus the surviving candidate is genuinely disjunctive: rooted packets are needed even away from equality. These are finite exact censuses, not a six-player domain proof.

The converse finite handover has a sharper hostile boundary. On 233 six-player archive rows every maximum-overlap root has an unaffordable canonical labelled row, and every one has a target-reaching direct packet. Taking H to be exactly the entire maximum-overlap root block works with an Lmin-tight comparison on 228 rows, but fails on five. Each failure is repaired by adjoining one highest-x remaining nongainer to H: the remaining P is a singleton and a one-gainer packet gives an Lmin-tight comparison, $R_x = 0$, $R_z > 0$, and base $4/5$. Hence the fixed-root-block converse is

false even on this finite control set. The surviving candidate must select an upper nongainer block containing the roots, as in (UBD1)–(UBD2). These counts are exact-rational archive evidence, not a six-player domain classification.

The locked-pair theorem identifies four of the five failures: in those four rows the two maximum-overlap roots are each other's unique locked partner. The fifth failure is pair-free. Consequently, after the automatic $|G| \geq 2$ branch and the one-pair characterization, the genuinely unstructured price-failure case is the pair-free rooted hypergraph.

The ceiling-root interpolation theorem closes 28 of the 78 direct failures. Of the 50 remaining rows, null-player descent removes all 34 zero-floor cases. The genuinely live archive region comprises 14 strict positive-floor cases in the interpolation wedge and the two equality cases. Every direct failure has a ceiling root. Hence the unresolved rooted-packet lemma is confined to these 16 rows on the archive, without making that finite decomposition a general theorem.

The sequential support audit gives a sharper candidate. Every one of the 50 pre-descent residual rows has a common coalition between one complete staged Lmin support and one sparse Lmax top over-cover, and its ceiling root is the unique maximum-overlap root. Common-support sizes are one on 45 rows, two on three, and three on two. Every one of the 16 positive-floor residuals has a common coalition containing that root, and the corresponding rooted relaxation reaches 2/3. Deleting every Lmax row still leaves a target-reaching complete-Lmin rooted packet on 13 of the 16. Only three rows require the common-hinge repair, and all three have a tied over-cover root load; 15 of the 16 have a tied load in total. All 37 rows with empty common support were already closed by direct gluing or ceiling interpolation. The missing statement is thus a *root-price/common-repair lemma*: after null-player descent, failure of the first two arms must force either a target-priced complete Lmin rooted packet or a common root hinge whose reverse comparison repairs the deficit. The repair branch must cover tied roots. The audit verifies this only on the finite archive.

The exact-cone structure of that finite residual is unexpectedly rigid. For each of the 13 rows closed by the complete Lmin packet alone, we extract one exact optimal one-witness Farkas dual, normalize its witness component to a rooted packing, and apply the pair shear (RPL). Every resulting system has full support rank and passes the Studený–Kratochvíl indecomposability test: 13 facets and no decomposable row. Up to relabeling there are only five support orbits, with multiplicities $3, 4, 3, 2, 1$. The two-spoke and three-packet repair facets belong to the first and fourth of those same orbits. Hence all 16 live positive-floor archive rows are carried by only five indecomposable rooted pair-lift support orbits, with combined multiplicities $5, 4, 3, 3, 1$. An independently formulated exact-LRA replay excludes all 282 unsupported proper exceptional candidates for the five representatives.

This is a finite facet atlas, not a six-player domain classification: the archive is not exhaustive and the extraction fixes one optimal dual per row. It does sharpen the missing statement to a *marked-facet selection conjecture*: every positive-floor residual should expose an indecomposable rooted pair lift whose Lmin-tight labels give a target-priced canonical marking, with a common reverse comparison needed only when the Lmin-only marked price is insufficient. The finite gauge audit makes the distinction necessary: all 13 extracted Lmin-only packings are affordable, whereas the canonical rooted-pair deletion leaves only six affordable; restoring every eligible Lmin-tight pair recovers all 13. The full labelled packing/cover optimization goes one step further: all 16 positive-floor residual rows have an affordable canonical marked row. The three common-hinge repairs are therefore alignment failures, not scalar price failures. On the fixed seven-player hostile/sharp controls, each case either has a direct packet or an affordable labelled row at a maximum-overlap root. Both statements are finite audits, not a universal selection theorem.

The rooted status of the sharp rows does have a dimension-free explanation. If every Lmax coordinate is positive, then $\ell = z(P)/p > 0$ for every nonempty P. Equation (TDG) is therefore

strict whenever its comparison and residual-sum payments are nonnegative. Hence an equality game $O = c_n T$ with strictly positive Lmax cannot have a direct packet whose base reaches c_n. Balanced $G(p,k)$ has $k(p+1) = \lfloor n^2/4 \rfloor$, overlap exactly $c_n T$, and strictly positive Lmax coordinates $p(k-1)/k$ and $1/k$. Its rooted indecomposable block-star system is thus structurally necessary within this direct-or-rooted framework.

The handover can be classified on every $G(p,k)$. In (TDG), take $P = O$, $G = B$, $H = \{r\}$, and $J = B \setminus \{b\}$. The comparison coalition is the common block-star hinge. If

$$d = \frac{p(k-1)-1}{k(p+1)},$$

then each gainer has gain pd, while the root loss and the packet gain both equal $p(k-1)d$. Thus the comparison is tight and both residuals vanish. The direct and rooted bounds are respectively

$$\frac{n}{(p+1)(k+1)} \quad \text{and} \quad \frac{n}{k(p+1)}. \tag{GPH}$$

The direct bound reaches c_n exactly when $(p+1)(k+1) \le \lfloor n^2/4 \rfloor$. The rooted block-star bound always reaches it because $k + (p+1) = n$ implies $k(p+1) \le \lfloor n^2/4 \rfloor$. The one-unit denominator shift is the mechanism handover: $G(3,2)$ has direct base $1/2$ but rooted base $3/4 > 2/3$, while $G(2,4)$ and $G(3,3)$ have direct bases $7/15$ and $7/16$ but rooted base $7/12$ at equality. In facet-price notation the denominators are

$$D_{\text{direct}} = (p+1)(k+1), \qquad D_{\text{root}} = p(k-1) + n - 1 = k(p+1),$$

and differ by exactly $p+1$. The rooted pair-spoke shear is the indecomposable semi-balanced atom realizing the smaller denominator, whereas the direct identity retains the positive low-Lmax payment on the outside block. This is the exact-cone reason the common hinge alone loses the necessary factor near balance.

When both selections have the corresponding two-block form, the residuals vanish. A fresh exact seven-player control with $(p,g,h,s) = (3,2,2,1)$ has this specialization, and (TDG) becomes

$$O - \frac{5}{8}T = \frac{21}{8}z_0 + \frac{1}{2}[z(74) - x(74)],$$

where coalition $74 = \{1,3,6\}$ is tight at both selections. Its actual overlap is $11T/12$, while every one-witness branch relaxation has exact minimum $5/8 > 7/12$. The hand identity, not that finite control, is the dimension-free evidence.

7.2 Mixed-support rooted packets

The common-hinge projection is not the largest inequality obtainable from one singleton witness. Let v be singleton-normalized and exact, let x, z be core allocations, and fix $r \in N$. Exactness supplies a core allocation w^r with $w^r_r = 0$. Consider tagged coalitions (σ, S), where $\sigma \in \{x, z\}$ and S is tight at the tagged selection. If $\rho_{\sigma,S} \ge 0$, $q \ge 0$, and

$$\sum_{\sigma,S} \rho_{\sigma,S} \mathbf{1}_S + q = \xi \mathbf{1}_N + \tau e_r, \tag{MCP}$$

then

$$\sum_{\sigma,S} \rho_{\sigma,S} \sigma(S) \le \xi T. \tag{MCB}$$

Indeed, dot (MCP) with w^r. Its right side is ξT, whereas $w^r(S) \geq v(S) = \sigma(S)$ for every tagged tight coalition and $q \cdot w^r \geq 0$. This proves (MCB). Unlike common-hinge transfer, the packet may contain an x-only coalition and a different z-only coalition; their common exactness witness supplies the gluing.

These inequalities exactly eliminate the witness. Write $p_i(\rho) = \sum_{\sigma,S:i\in S} \rho_{\sigma,S}$. There exists a vector $w \geq 0$ with $w(N) = T$, $w_r = 0$, and $\sigma(S) \leq w(S)$ for every tagged coalition if and only if

$$\sum_{\sigma,S} \rho_{\sigma,S}\sigma(S) \leq T \max_{i\neq r} p_i(\rho) \qquad (\rho \geq 0). \tag{MWP}$$

Necessity is the weighted witness argument. For sufficiency, the tagged requirement vectors dominated by coalition sums of the rooted simplex form a closed convex down-set. A separating normal may therefore be chosen nonnegative, and the support function of the rooted simplex in that direction is precisely the right side of (MWP). Separation would contradict (MWP).

After normalizing every nonroot incidence load to at most one, the tagged packing polytope has $n-1$ capacity rows. Hence each vertex has at most $n-1$ positive tagged coalitions. Two tagged copies of the same incidence column cannot both occur at a vertex. If p_i denotes the resulting load, $q_i = 1 - p_i$ off the root, and $\delta = p_r - 1$, its completion is

$$\sum_{\sigma,S} \rho_{\sigma,S}\mathbf{1}_S + \sum_{i\neq r} q_i e_i - \delta e_r = \mathbf{1}_N. \tag{MCC}$$

Thus the colours disappear under projection to characteristic-function coordinates. Vertex enumeration, the rooted-pair shear, and the published indecomposability test give an exact finite algorithm for the entire one-witness projection and its corresponding min-semi-balanced facet candidates. The new content is their relative two-selection evaluation, not a claim that every candidate is a full-cone facet or that some root always proves the desired target.

The optimized seven-player spectral obstruction supplies the first genuinely noncommon packet. At root $r = 3$,

$$\mathbf{1}_{26} + \mathbf{1}_{57} + \mathbf{1}_{75} + 2e_2 + e_5 + e_6 = 2\mathbf{1}_N + e_3, \tag{MC7}$$

where coalitions $26, 57$ are tight at x, coalition 75 is tight at z, and only 57 is common. Therefore

$$x(26) + x(57) + z(75) \leq 2T. \tag{MC8}$$

Together with direct selection-comparison, stratum, nonnegativity, and efficiency rows, this packet gives an exact Farkas identity with branch minimum $3/4$. That identity now has a short human slack form. Normalize $T = 1$ and write $w = w^3$. On the displayed branch,

$$\begin{aligned} O - \frac{3}{4} = {} & z_2 + \frac{1}{2}w_2 + \frac{1}{4}w_5 + \frac{1}{4}w_6 \\ & + \frac{1}{4}[w(26) - x(26)] + \frac{1}{4}[w(57) - x(57)] \\ & + \frac{5}{4}[x(52) - z(52)] + \frac{1}{2}[x(57) - z(57)] + \frac{1}{4}[w(75) - z(75)] \\ & + \frac{3}{4}(x_0 - x_4) + \frac{1}{2}(z_4 - z_5) + (z_5 - z_1). \end{aligned} \tag{MCS}$$

Every term is nonnegative by witness/core comparisons, nonnegativity, or the fixed payoff orders. This is a hand proof of the branch bound, not merely a solver dual. At the actual control every coalition-comparison term in (MCS) vanishes; its strict margin $119/1080$ comes entirely from z_2 and

three payoff-stratum gaps. Thus the mixed packet fixes the baseline, while ordered label currency pays the multistage defect.

The old common-hinge one-witness projection at this root reaches only 25/34. The arbitrary-Lmax direct theorem now supplies a simpler target-closing bound 3/5 on this fixed selection; it also directly closes the coefficient-two and primitive coefficient-seven controls at 5/8 and 5/9. Thus (MCS) remains a stronger human rooted certificate and a test of the mixed-colour mechanism, but is no longer required to close those fixed controls. Exactly the two sharp positive-Lmax equality orientations lack a target-reaching direct packet and use the rooted block-star arm.

An exact rational verifier checks (MC7) coordinatewise and checks every coefficient in all three fixed-branch Farkas identities. These are proof objects for the displayed branches, not an all-dimension proof. The remaining question is whether failure of the sharpened spectral bound forces some distinguished root's direct labelled rows plus exact mixed-witness projection to prove the target.

The cutting-plane projection also finds a second mixed packet at the same root,

$$\tfrac{1}{2}(\mathbf{1}_{26}+\mathbf{1}_{57}+\mathbf{1}_{15})+\tfrac{1}{2}e_2+\tfrac{1}{2}e_5+e_6=\mathbf{1}_N+\tfrac{1}{2}e_3, \tag{MC9}$$

where $26, 57$ are tagged by x and 15 by z. This one cut raises the witness-free master from 25/34 to the exact one-witness minimum 3/4. Its singleton completion is purely min-semi-balanced but decomposable. The rooted-pair shear instead gives

$$\tfrac{1}{2}(\mathbf{1}_{12}+\mathbf{1}_{15}+\mathbf{1}_{26}+\mathbf{1}_{40}+\mathbf{1}_{57})+\mathbf{1}_{72}-\tfrac{5}{2}e_3=\mathbf{1}_N. \tag{MCF}$$

The seven columns in (MCF) have rank seven, the published exceptional-set test finds no decomposition, and its exact-game inequality is active on the hostile control. Hence (MCF) defines a facet of the seven-player exact-game cone. An independently formulated Z3 exact-LRA replay agrees with the exact LP decomposition count. Across eight fixed hostile and sharp controls, the projected cutting-plane and direct-witness minima agree at every maximum-overlap root; nine cuts are generated, seven rooted-pair shears pass the facet test, and (MC9) is the unique cut using both colours. These are finite controls and do not prove existence of a target-valid root in every dimension.

There is first a complete rectangular description for the positive-load singleton-exceptional class. Consider an indecomposable normalized identity

$$\sum_{j=1}^{k}\lambda_j\mathbf{1}_{S_j}-ce_r=\mathbf{1}_N,\qquad \lambda_j,c>0. \tag{RR1}$$

Every S_j contains r. Indeed, if S omits r and $S\cup\{r\}$ is an unsupported proper coalition, then $\mathbf{1}_S+e_r-\mathbf{1}_{S\cup\{r\}}=0$ is an exceptional-set decomposition. If the union is supported, the same three columns form a forbidden semi-balanced subsystem. If the union is N, the first two columns already form a balanced subsystem. The degenerate cases in which these columns are the whole support cannot satisfy the positive-load identity (RR1).

Write $S_j=\{r\}\cup A_j$ and let A be the $(n-1)\times k$ nonroot incidence matrix. Minimality is equivalent to $\operatorname{rank}A=k$, and the identity is

$$A\lambda=\mathbf{1},\qquad c=\mathbf{1}^{\mathsf{T}}\lambda-1. \tag{RR2}$$

Thus $k\le n-1$. For a proposed exceptional coalition with root indicator β and nonroot incidence b, the Studený–Kratochvíl decomposition equations reduce exactly to

$$A\mu - b = \gamma \mathbf{1}, \qquad \mu \geq 0, \qquad \beta + \gamma - \mathbf{1}^{\mathsf{T}}\mu \geq 0. \tag{RR3}$$

Consequently (RR1) is a facet exactly when (RR3) is infeasible for every unsupported proper nonempty (β, b). This gives an exact enumeration algorithm for the entire positive-load singleton-exceptional class at fixed n: enumerate full-column-rank zero-one matrices up to row and column permutations, retain the positive solutions of (RR2), and test all Boolean vectors by (RR3). Passing systems give primitive dual-cone extreme rays after clearing denominators. This is a specialization of the published criterion, with a support bound and a concrete matrix enumeration, not a new general characterization of the exact-game cone.

For the projected pair lifts this is also a decisive support filter. If an extreme rooted packing uses any positive coalition omitting its exceptional root, its pair lift is decomposable and cannot define a full exact-cone facet. Irreducible facet selection may therefore be restricted in advance to root-containing Lmin-tight coalitions, together with the restored rooted pair columns. The remaining hybrid theorem must show that some ceiling root exposes a member of this smaller menu whose canonical Lmin-labelled marking is sufficiently priced, or that an earlier direct/interpolation branch closes the target. The gauge quotient classifies facet rays but cannot test this marked price.

There is a closed inverse formula for the saturated rooted-star subclass containing (MCF) and the sharp-family facets. Remove a root r and let A be a nonsingular $(n-1) \times (n-1)$ zero-one matrix whose columns are the incidences of sets $A_j \subset N \setminus \{r\}$. Put

$$\lambda = A^{-1}\mathbf{1}, \qquad c = \mathbf{1}^{\mathsf{T}}\lambda - 1. \tag{RS1}$$

If $\lambda > 0$ and $c > 0$, then

$$\sum_j \lambda_j \mathbf{1}_{\{r\} \cup A_j} - c e_r = \mathbf{1}_N \tag{RS2}$$

is a purely min-semi-balanced system with n independent columns. It is the completion of a saturated rooted-packing vertex and attains the $n-1$ positive-coalition support bound; rooted-pair shears enter the class by treating their pair columns as completed-menu columns. Its exact-game inequality is a facet if and only if the following finite inverse-matrix test passes. For every proper nonempty coalition B outside the support, let $\beta = 1_{\{r \in B\}}$, let b be its incidence off the root, and put $u = A^{-1}b$. Then require

$$\max_j \frac{-u_j}{\lambda_j} > \frac{\beta - \mathbf{1}^{\mathsf{T}}u}{c}. \tag{RS3}$$

Indeed, a decomposition through B has nonnegative coefficients μ_0, μ_j and common load γ. Its off-root coordinates force $\mu = u + \gamma\lambda$, so γ is at least the left side of (RS3). Its root coordinate gives $\mu_0 = \beta - \mathbf{1}^{\mathsf{T}}u - \gamma c$, so γ is at most the right side. Thus a decomposition exists exactly when the reverse weak inequality holds. The facet assertion is then precisely the indecomposability theorem of Studený and Kratochvíl (2022). This specializes, rather than replaces, their general criterion.

Formula (RS3) gives an exact enumeration algorithm: enumerate nonsingular zero-one matrices up to row and column permutations, retain (RS1), and test all Boolean (β, b). Passing matrices give primitive dual-cone extreme rays after clearing denominators. Equality in the game inequality is $\sum_j \lambda_j v(\{r\} \cup A_j) = v(N)$.

This saturated class has a dimension-free unsaturated closure. Start from a purely min-semi-balanced rooted-star system and replace every nonroot role q by a nonempty block B_q of clones, expanding each positive coalition by taking the union of the blocks belonging to its roles. The

incidence identity is preserved coordinatewise. Minimality is preserved because any relation among expanded columns restricts, on one representative per block, to a relation among the base columns.

Moreover, this clone expansion is indecomposable if and only if its base is. Indeed, two players in the same clone block have identical rows in the positive-support incidence matrix. Subtracting their coordinates in the Studený–Kratochvíl exceptional-set equations forces a proposed exceptional coalition to contain either the whole block or none of it. Every expanded decomposition therefore collapses to a base decomposition, and every base decomposition expands blockwise. The published indecomposability criterion now implies:

Theorem 52 (nonroot clone expansion)**.** *Every nonroot clone expansion of a saturated rooted-star exact-cone facet is again an exact-cone facet.*

For fixed n this gives an exact enumeration algorithm for the enlarged class: enumerate the saturated bases using (RS1)–(RS3), then enumerate positive clone-block size compositions summing to $n-1$, modulo base automorphisms. This is complete for the clone-expanded rooted-star class, not for all exact-cone facets.

All five support orbits in the six-player positive-residual atlas belong to this class. Three are already saturated. The other two each compress one clone pair to a four-role saturated base, with determinants 1 and -2. Thus one dimension-free facet construction explains the finite atlas, but does not yet prove that every residual branch exposes a member with enough hybrid-overlap price.

The clone-expanded subclass is not exhaustive. The published catalogue of Studený and Kratochvíl (2022) shows that all 23 singleton-exceptional types through six players are clone-saturated, so the first failure is sharply at seven players. For $m \geq 4$ and $p \geq 0$, take

$$N = \{r\} \dot\cup E(K_m) \dot\cup O, \qquad |O| = p, \qquad S_v = \{r\} \cup \{e \in E(K_m) : v \in e\}.$$

Then the complete-graph rooted-pair identity is

$$\sum_{v \in K_m} \mathbf{1}_{S_v} + 2 \sum_{o \in O} \mathbf{1}_{\{r,o\}} - (m - 2 + 2p) e_r = 2\mathbf{1}_N. \tag{KG1}$$

Theorem 53 (complete-graph rooted pair facets)**.** *For every $m \geq 4$ and $p \geq 0$, (KG1) is an indecomposable purely min-semi-balanced system and defines a facet of the exact-game cone. It is not a clone expansion of a saturated rooted-star system. The first member has $(m, p) = (4, 0)$ and seven players.*

The incidence identity is immediate: each edge lies in its two endpoint stars and each outside player in its rooted pair. The unoriented vertex-edge incidence matrix of K_m has rank m because it contains an odd cycle; the pair and root columns raise the support rank to $m + p + 1$.

For the exceptional-set test, write its edge, outside, and root indicators as b_{ij}, o_a, β. The old-support coefficients obey

$$\mu_i + \mu_j - b_{ij} = \gamma, \qquad c_a = \gamma + o_a, \qquad d = \beta + \gamma - \sum_i \mu_i - \sum_a c_a \geq 0. \tag{KG2}$$

An elementary additive-binary graph lemma says that a zero-one edge labelling $b_{ij} = \mu_i + \mu_j - \gamma$ of K_m, $m \geq 4$, is only empty, complete, a star, or the complement of a star. If all potentials agree the labelling is constant. Otherwise choose $\mu_j > \mu_k$. For every other i, the binary difference $b_{ij} - b_{ik} = \mu_j - \mu_k$ is the same nonzero value, hence it is one and $(b_{ij}, b_{ik}) = (1, 0)$. All remaining potentials agree. Two remaining vertices exist, and their mutual edge bit and b_{jk} sum to one, giving

exactly the star and co-star alternatives. With a common gauge t, the four coefficient patterns are respectively

$$(\mu, \gamma) = (t\mathbf{1}, 2t), \quad (t\mathbf{1} + e_v, 2t), \quad (t\mathbf{1} - e_v, 2t - 1), \quad (t\mathbf{1}, 2t - 1).$$

Put $o = \sum_a o_a$ and $D = m + 2p - 2$. In the same four cases, (KG2) gives $d = \beta - o - Dt$, $\beta - o - 1 - Dt$, $\beta + p - o - Dt$, and $\beta + p - o - 1 - Dt$. Nonnegativity of the other coefficients requires $t \geq 0, 0, 1$, and either $t \geq 1/2$ or all outside bits one, respectively. The only feasible candidates are the empty coalition, the supported root, one supported rooted pair, one supported vertex star, or N. Hence there is no new proper exceptional coalition. Finally the nonroot star rows are the $\binom{m}{2}$ distinct two-subset vectors. Since $\binom{m}{2} > m$, no clone compression is saturated.

In rooted-packing coordinates the star weights are $1/2$, the pair weights are one, and

$$\delta = m/2 - 1, \qquad Q = p, \qquad P = Q/\delta = 2p/(m-2). \tag{KG3}$$

The $p = 0$ member belongs to the zero-root/null-floor mechanism when evaluated through a root-zero Lmin witness. Members with $p > 0$ are affordable positive-floor facet candidates, but the missing selection theorem is not implied by their existence. Exact enumeration on twelve parameter controls and an independent Z3 replay of all 127 seven-player exceptional candidates serve only as audits of the displayed proof.

For (MCF), $|\det A| = 2$, the weights are $(1/2, 1/2, 1/2, 1/2, 1/2, 1)$, $c = 5/2$, and every unsupported Boolean coalition has strict margin at least $2/5$ in (RS3). For $G(p, k)$, $A = I_p \oplus (J_k - I_k)$, giving pair weights one, hinge weights $1/(k-1)$, and $c = p + 1/(k-1)$; hence the new mixed facet and every sharp-family facet lie in the same saturated rooted-star class. Exact inverse arithmetic agrees with the LP decomposition counts for all four saturated control systems. The matrix theorem is dimension-free, but it still does not prove that every labelled selection branch exposes one of these facets with enough hybrid-overlap charge.

Not claimed. Exactness theory, lower-envelope exactification, and the characterizations of the egalitarian selections are prior art (Arin and Iñarra, 2001; Arin et al., 2003, 2008; Biswas et al., 1999; Dietzenbacher, 2021; Dietzenbacher and Yanovskaya, 2020; Llerena and Mauri, 2017; Csóka et al., 2011; Frank and Murota, 2022; Kroupa and Studený, 2019; Studený and Kratochvíl, 2022; Gong et al., 2022; Gong and Dietzenbacher, 2026; Lohmann et al., 2012; Reijnierse and Derks, 1998); the four-player rigidity is the companion paper's classification (Flanagan, 2026) combined with the Lorenz implication of Llerena and Mauri (2017). The candidate new content is the two-generator core duality and its one-sided floor certificate, the family $G(p, k)$ with its closed forms and deficit identity, the resulting evaluation $\sup_n s(n) = 2$, the identification of the $n = 5$ anomaly with a proof by hand, the one-coalition certificate theorem and the pinned extremal configuration at five players, the evaluation $s(5) = \frac{1}{2}$ with the $(4, 1)$-regime bound $\frac{1}{3}$, and the two closed routes of Section 6, together with the direct exactness- exchange evaluation $s(6) = \frac{2}{3}$, the singleton-root rectangular facet classification, and the complete-graph rooted-pair facet family. Novelty remains conditional on specialist review.

Terminology is also prior-art sensitive. The phrase *exact cover* already has a cooperative-game meaning for coalition collections (Reijnierse and Derks, 1998), and minimal exact balanced collections have their own three-type classification (Lohmann et al., 2012). The *normalized top cover* used in (PXC89)–(PXC107) is the particular positive dual identity displayed there; no identification with those earlier notions, and no novelty claim based on the word *cover*, is intended.

Reproducibility note

Every theorem not identified below is proved by hand and depends on no program. Theorems 24, 33, 34, and 36 depend on exhaustive finite exact-rational checks, described in their proofs; those checks are proof dependencies, archived in the public repository https://github.com/mfflanag1/counterexample-hunt on branch `codex/fixed-n-handover-facets`, with the theorem-to-driver map in `coopgames/verifiers/README.md` (for Theorems 33 and 34, both pipelines, the coverage audit, and the full control suite; for Theorem 36, the fifteen source-regenerated QF_LRA formulas, Z3/cvc5 records, trust-disabled CPC proofs, and Ethos checks under `s6/strict_exchange_*`, plus the independently formulated bound-two reconstruction under `s6/independent_bound_two/` and the standard-library direct SMT-LIB reconstruction under `s6/third_bound_two/`). The literal 72-assertion human-to-machine map is audited under `s6/strict_exchange_formula_interface.json`; a separately written semantic reconstruction is frozen as `s6/strict_exchange_semantic_interface.json`. These are transcription controls and do not replace external review of the human implication. The orbit anatomy under `s6/strict_exchange_large_case_templates.json` is explanatory; the later 19-orbit formulas under `s6/orbit_core_certificates/` are separately proof-checked reduced alternatives for the three large cases. The anchor-localized $(4, 2, 0)$ alternative under `s6/p4_anchor_core_certificates/` is separately proof-checked and uses only one singleton exactness witness. Its explanatory eleven-row forced-blocker decomposition is regenerated by `s6/audit_strict_exchange_p4_packet_classes.py`; its forced-exchange trees are regenerated by `s6/audit_strict_exchange_p4_relay_trees.py`. Theorem 26 is proved by hand; its superseded finite check and the hand proof's own symbolic guards are archived alongside. The remaining programs are bug audits, not proof dependencies. Exact rational programs in the same repository on branch `codex/hybrid-overlap-facets` archive the refuted effective-bipartition bridge, its exact seven-player counterexample, the dispersion primal/dual witnesses, and the threshold-label currency audit; the latter also replays canonical saturation, exact interval-packing/cover duals, and rooted potential trees. Separate exact replays archive staged availability for the clipped rectangles and the neutral-endpoint seven-player stress test; an independently formulated cutting-plane replay audits the rooted-hinge witness projection. A separate tagged cutting-plane replay audits the exact mixed-witness projection and rooted-pair facet filter, and an exact-LRA Z3 formulation independently checks every reported decomposition count. The complete-graph rooted-pair driver audits twelve finite parameter controls, and a separate exact-LRA encoding checks the first seven-player member; the graph case analysis is the proof. The marked-price gauge audit exactly replays the 13 staged six-player residual packings and their canonical tight-pair markings; equation (RPM), not that finite audit, is the marking lemma's proof. None is a proof dependency for the open all-dimension target. Other programs independently reconstruct all displayed games and coalition values, the closed forms of Theorem 14 at twelve parameter pairs, the full coalition-inequality check and the tight-set prediction of Lemma 15 for every (p, k) with $1 \le p \le 7$, $2 \le k \le 7$ and $n \le 13$, and the four-way agreement on the game of Theorem 22. They are bug audits, not proof dependencies. All arithmetic is exact rational; no floating-point value appears in any stated result. Reported optima come from a from-scratch exact-rational two-phase simplex with Bland's rule, from a staged exact solver over linear real arithmetic, or from exact vertex enumeration, with at least two methods agreeing on every displayed allocation; one library simplex implementation cycles on these degenerate programs and was not used for any reported number.

References

Björner A, Las Vergnas M, Sturmfels B, White N, Ziegler GM (1999) Oriented Matroids, 2nd ed. Cambridge University Press, Cambridge. https://doi.org/10.1017/CBO9780511586507

Arin J, Iñarra E (2001) Egalitarian solutions in the core. International Journal of Game Theory 30:187–193. https://doi.org/10.1007/s001820100073

Arin J, Kuipers J, Vermeulen D (2003) Some characterizations of egalitarian solutions on classes of TU-games. Mathematical Social Sciences 46:327–345. https://doi.org/10.1016/S0165-4896(03)00051-9

Arin J, Kuipers J, Vermeulen D (2008) An axiomatic approach to egalitarianism in TU-games. International Journal of Game Theory 37:565–580. https://doi.org/10.1007/s00182-008-0133-6

Biswas AK, Parthasarathy T, Potters JAM, Voorneveld M (1999) Large cores and exactness. Games and Economic Behavior 28:1–12. https://doi.org/10.1006/game.1998.0686

Csóka P, Herings PJJ, Kóczy LÁ (2011) Balancedness conditions for exact games. Mathematical Methods of Operations Research 74:41–52. https://doi.org/10.1007/s00186-011-0353-z

Fujishige S (1980) Lexicographically optimal base of a polymatroid with respect to a weight vector. Mathematics of Operations Research 5:186–196.

Frank A, Murota K (2022) Decreasing minimization on base-polyhedra: relation between discrete and continuous cases. Preprint, arXiv:2202.06561.

Gangam RR, Garg N, Shahkar P, Vazirani VV (2024) Equitable core imputations for max-flow, MST and b-matching games. Preprint, arXiv:2403.06037.

Izquierdo JM, Montes J, Rafels C (2024) Population Lorenz-monotonic allocation schemes for TU-games. Social Choice and Welfare 63:417–436. https://doi.org/10.1007/s00355-024-01538-x

Gong D, Dietzenbacher B (2026) Core bound reduced games and consistency. International Journal of Game Theory 55:13. https://doi.org/10.1007/s00182-026-00986-2

Gong D, Dietzenbacher B, Peters H (2022) Reduced two-bound core games. Mathematical Methods of Operations Research 96:447–457. https://doi.org/10.1007/s00186-022-00800-1

Gong D, Dietzenbacher B, Peters H (2024) One-bound core games. International Journal of Game Theory 53:859–878. https://doi.org/10.1007/s00182-024-00889-0

Klijn F, Vermeulen D, Hamers H, Solymosi T, Tijs S, Villar JP (2003) Neighbor games and the leximax solution. Mathematical Methods of Operations Research 58:191–208.

Kroupa T, Studený M (2019) Facets of the cone of totally balanced games. Mathematical Methods of Operations Research 90:271–300. https://doi.org/10.1007/s00186-019-00672-y

Lohmann E, Borm P, Herings PJJ (2012) Minimal exact balancedness. Mathematical Social Sciences 64:127–135. https://doi.org/10.1016/j.mathsocsci.2012.01.002

Llerena F, Mauri L (2016) Reduced games and egalitarian solutions. International Journal of Game Theory 45:1053–1069.

Vazirani VV (2024) Fair core imputations for the assignment game: new solution concepts and efficient algorithms. Preprint, arXiv:2402.11437.

Dutta B, Ray D (1989) A concept of egalitarianism under participation constraints. Econometrica 57:615–635.

Dietzenbacher B (2021) Monotonicity and egalitarianism. Games and Economic Behavior 127:194–205. https://doi.org/10.1016/j.geb.2021.03.006

Dietzenbacher B, Yanovskaya E (2020) Antiduality in exact partition games. Mathematical Social Sciences 108:116–121. https://doi.org/10.1016/j.mathsocsci.2020.04.006

Dietzenbacher B, Doğan E (2025) Population monotonicity and egalitarianism. International Journal of Game Theory 54, article 42. https://doi.org/10.1007/s00182-025-00966-y

Flanagan M (2026) Balanced transport at the egalitarian core: a complete boundary phase diagram. Preprint; companion paper.

Llerena F, Mauri L (2017) On the existence of the Dutta–Ray's egalitarian solution. Mathematical Social Sciences 89:92–99. https://doi.org/10.1016/j.mathsocsci.2017.07.002

Reijnierse H, Derks J (1998) Note on the core of a collection of coalitions. International Journal of Game Theory 27:451–459. https://doi.org/10.1007/s001820050084

Studený M, Kratochvíl V (2022) Facets of the cone of exact games. Mathematical Methods of Operations Research 95:35–80. https://doi.org/10.1007/s00186-022-00770-4

Schoot Uiterkamp MHH (2024) A characterization of simultaneous optimization, majorization, and (bi-)submodular polyhedra. Mathematics of Operations Research 50:252–276. https://doi.org/10.1287/moor.2023.0054